

I/O and Redirection

Standard I/O (1)

Normally each Unix process has three streams opened when it starts; one for input, one for output, and one for diagnostic or error messages.

- ◆ Standard Input (`stdin`)
 - default place from which programs read
 - File descriptor: `0`
- ◆ Standard Output (`stdout`)
 - default place to which programs write
 - File descriptor: `1`
- ◆ Standard Error (`stderr`)
 - default place where errors are reported
 - File descriptor: `2`

Standard I/O (2)

- ◆ For terminal,
 - The default standard input is the keyboard,
 - The default standard output is the display
 - The default standard error is the display.
- ◆ To demonstrate -- **cat**
 - Echoes everything you typed in with an <enter>
 - Quits when you press **Ctrl-d** at a new line -- (**EOF**)
- ◆ To redirect (change) the default
 - Use **<** to redirect standard input (the same as **0<**)
 - Use **>** to redirect standard output (the same as **1>**)
 - Use **2>** to redirect standard error

Redirecting Standard Output

◆ `cat file1 file2 > file3`

- concatenates file1 and file2 into file3
- file3 is created if not there

◆ `cat file1 file2 >| file3`

- file3 is clobbered if there

◆ `cat file1 file2 >> file3`

- warning if file3 is not there
- file3 is appended to if it is there

◆ `cat > file3`

- file3 is created from whatever user provides from standard input

Redirecting Standard Error

- ◆ To write standard output and standard error into different files:

`compute[1] > cat myfile > yourfile 2> yourerrorfile`

- ◆ Generally direct standard output and standard error to the same place:

`compute[2] > cat myfile &> yourfile`

- ❖ If `myfile` exists, it is copied into `yourfile`

- ❖ If `myfile` does not exist, an error message
`cat: myfile: No such file or directory`
is copied into `yourfile`

- ◆ A more general way is

- `cat myfile > yourfile 2>&1`

- stdout goes to `yourfile` and stderr goes to where stdout goes

Redirecting Standard Input

- ◆ `compute[1] > cat < oldfile > newfile`
- ◆ A more useful example:
 - `compute[2] > tr string1 string2`
 - ❖ read from standard input.
 - ❖ character *n* of `string1` translated to character *n* of `string2`.
 - ❖ results written to standard output.
 - Example of use:
 - `compute[3] > tr aeiou eoiaa <file1 >file2`
 - `compute[4] > tr eoiaa aeiou <file2 >file3`
 - `compute[5] > tr a-z A-Z < file1 > file2`

/dev/null

◆ /dev/null

- A virtual file that is always empty.
- Copy things to here and they disappear.
 - ❖ `cp myfile /dev/null`
- Copy from here and get an empty file.
 - ❖ `cp /dev/null myfile`
- Redirect error messages to this file
 - ❖ `ls -l > recordfile 2> /dev/null`
 - ❖ Basically, all error messages are discarded.

Filters (1)

- ◆ Filters are programs that:
 - Read from `stdin`.
 - Modify it (may do nothing).
 - Write the results to `stdout`.
- ◆ Filters typically do not need user input.
- ◆ Example:
 - `tr` (translate):
 - ❖ Read `stdin`
 - ❖ Echo to `stdout`, translating some specified characters
- ◆ Many filters can also take file names as operands for input, instead of using `stdin`.

Filters (2)

◆ **grep patternstr:**

- Read **stdin** and write lines containing **patternstr** to **stdout**

compute[1] > grep "unix is easy" < myfile1 > myfile2

- Write all lines of **myfile1** containing phrase ***unix is easy*** to **myfile2**

◆ **wc:**

- Count the number of chars/words/lines on **stdin**
- Write the resulting statistics to **stdout**

◆ **sort:**

- Sort all the input lines in alphabetical order and write to the standard output.

Pipes

◆ The pipe:

- Connects **stdout** of one program with **stdin** of another
- General form:

command1 | command2

- **stdout** of command1 used as **stdin** for command2
- Example:

compute[1] > cat readme.txt | grep unix | wc -l

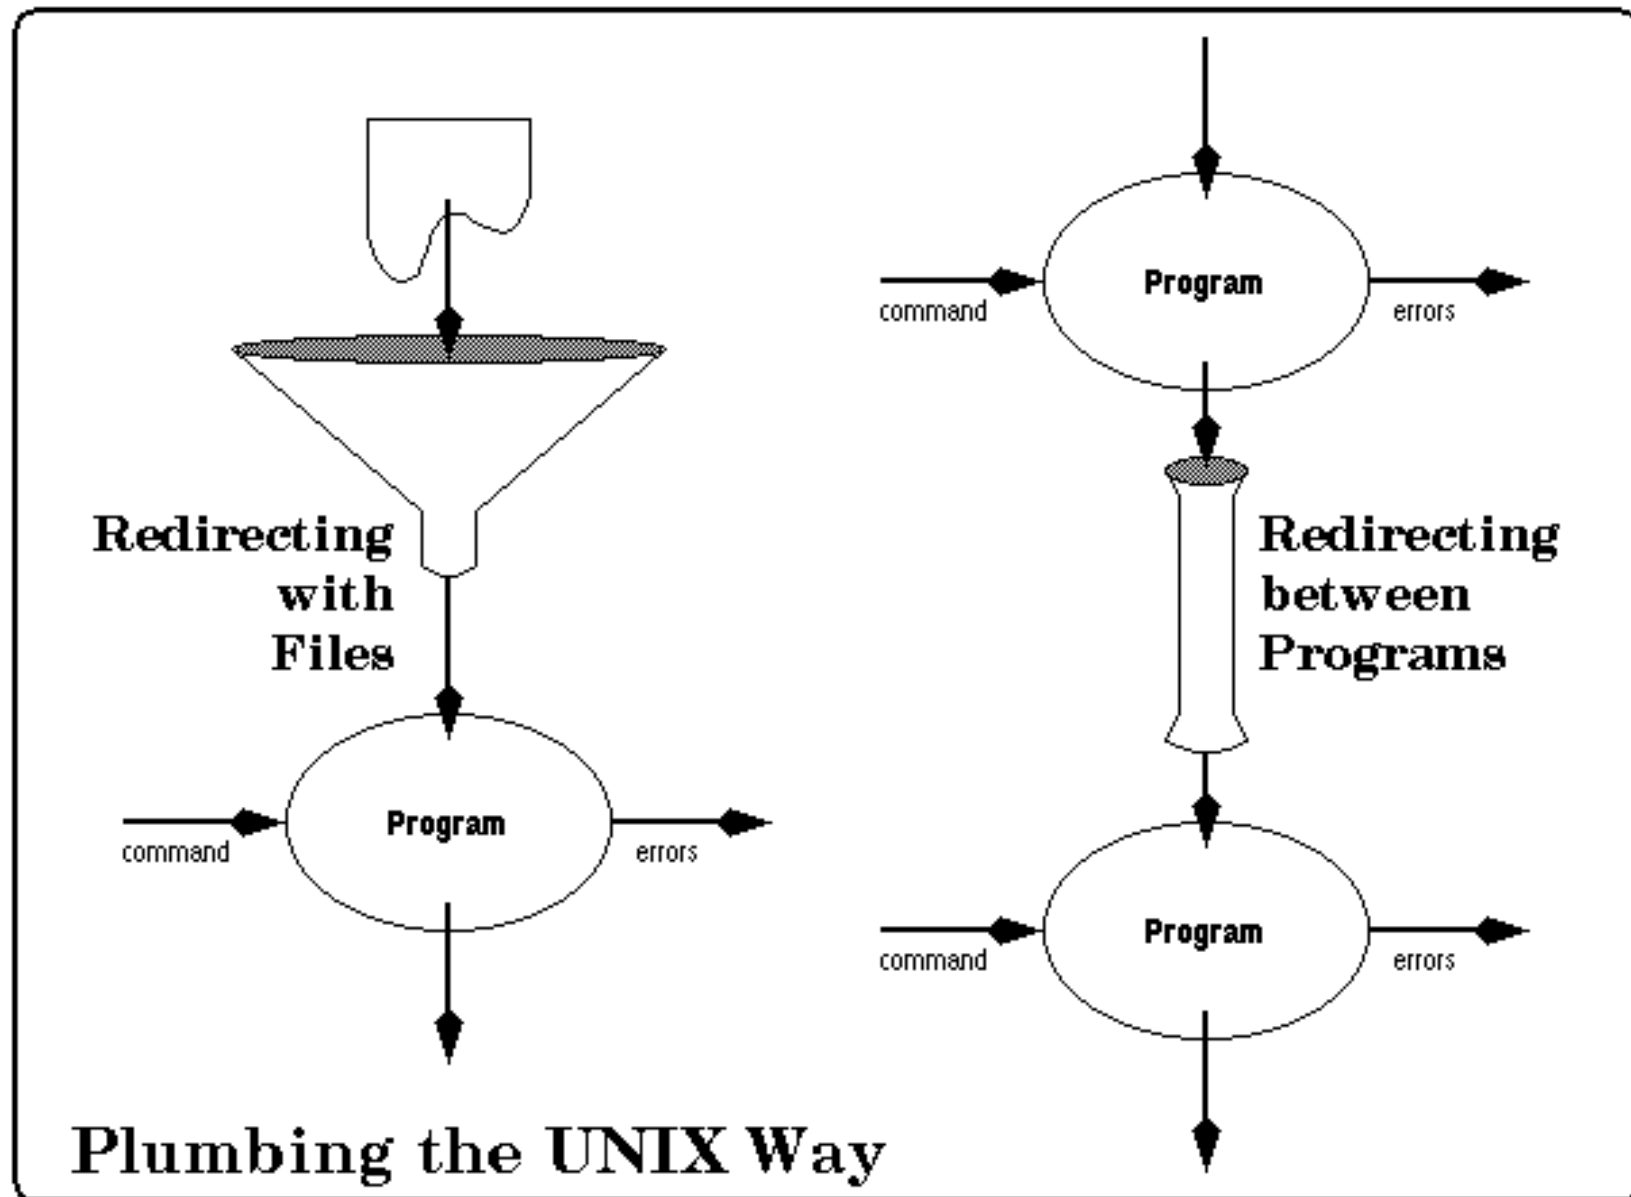
◆ An alternative way (not efficient) is to:

compute[2] > grep unix < readme.txt > tmp

compute[3] > wc -l < tmp

◆ Can also pipe **stderr**: **command1 |& command2**

Redirecting and Pipes (1)



Redirecting and Pipes (2)

- ◆ Note: The name of a command always comes first on the line.
- ◆ There may be a tendency to say:
`compute[1] > readme.txt > grep unix | wc -l`
 - This is WRONG!!!
 - Your shell will go looking for a program named `readme.txt`
- ◆ To do it correctly, many alternatives!
`compute[2] > cat readme.txt | grep unix | wc -l`
`compute[3] > grep unix < readme.txt | wc -l`
`compute[4] > grep unix readme.txt | wc -l`
`compute[5] > grep -c unix readme.txt`

The tee Command

- ◆ **tee** - replicate the standard output
 - `cat readme.txt | tee myfile`
 - `tee -a` (append)

