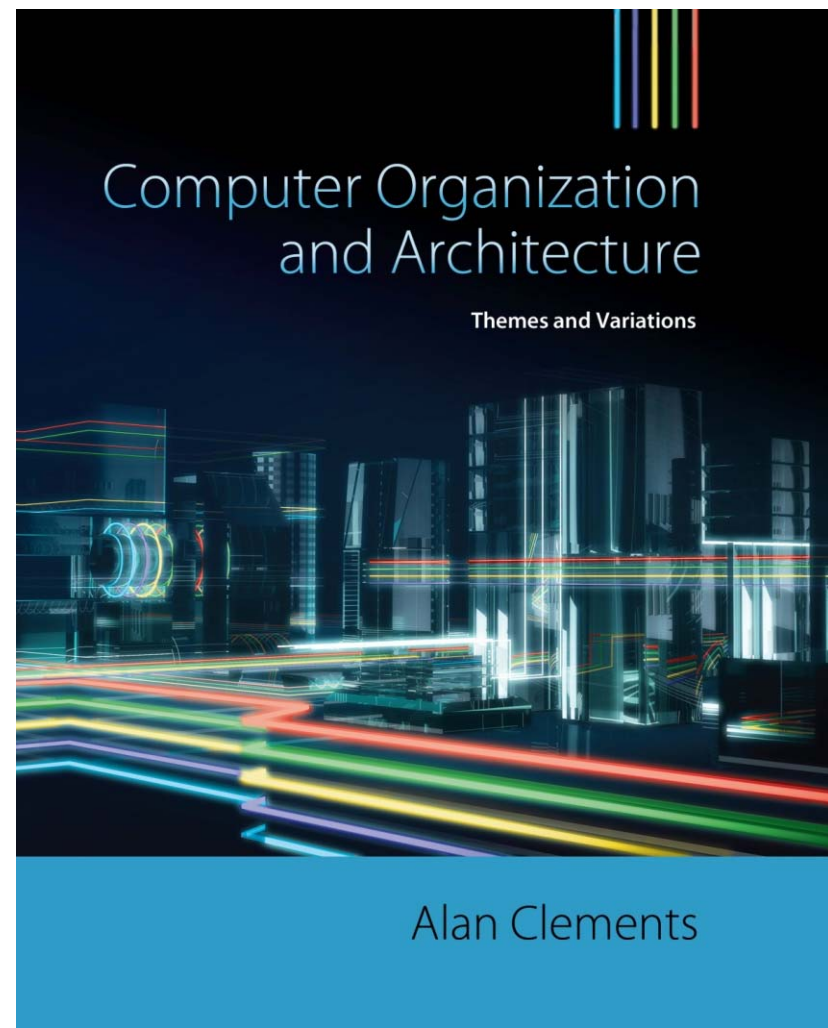


# Part 0x7

## CHAPTER 3

### Architecture and Organization



1

These slides are being provided with permission from the copyright for in-class (CS2208B) use only. The slides must not be reproduced or provided to anyone outside of the class.

All download copies of the slides and/or lecture recordings are for personal use only. Students must destroy these copies within 30 days after receipt of final course evaluations.

# Instruction Encoding

## An Insight into the ARM's Architecture

- ❑ The branch instruction (Figure 3.41) has
  - an 8-bit op-code
  - a 24-bit **signed** program-counter relative **offset** (*word address offset*).
- ❑ Converting the 24-bit **word** offset to real **byte** address:
  - shift left twice the 24-bit **word** offset to convert the **word-offset** address to a **byte-offset** address,
  - **sign-extended** to 32 bits,
  - added it to the current value of the program counter • • •  
(*the result is in the range  $PC \pm 32\text{ MBytes}$* ).

Do not forget the pipelining effect

Basically, it is the number of instructions away from the current location (after considering the pipelining effect.

FIGURE 3.41

Encoding ARM's branch and branch-with-link instructions



The L-bit is 0 for a branch instruction and 1 for a branch with link instruction.

The 24-bit word offset is shifted left twice to create a 26-bit byte offset.





# Instruction Encoding

## An Insight into the ARM's Architecture

- ❑ Figure 3.26 illustrates the structure of the ARM's **data processing** instructions and demonstrates how bit 25 is used to control the nature of the second source operand.

### Shift type

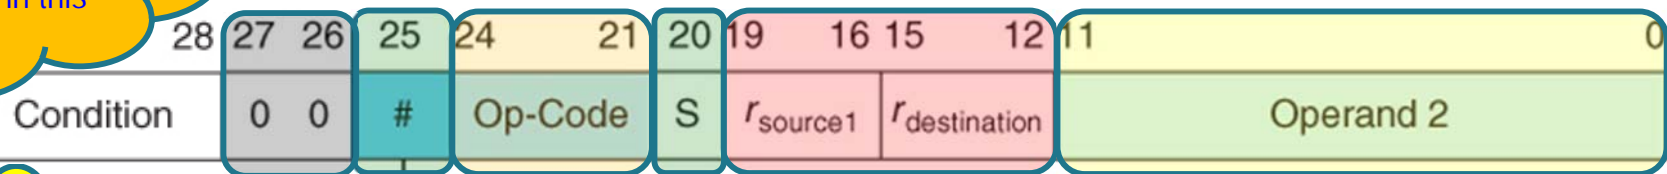
00 = logical left  
01 = logical right  
10 = Arithmetic right  
11 = rotate right

To reduce the course workload, **Multiplication** encoding/decoding will not be included in this course.

1 0 1 B/BL instructions

26

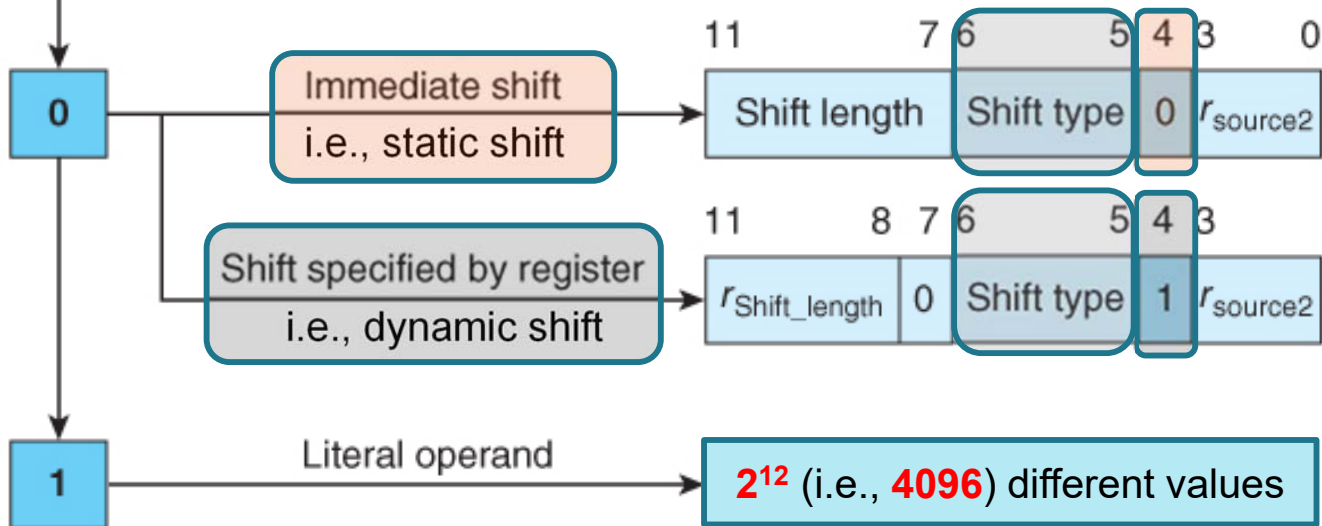
Encoding the ARM's data processing instructions



**Multiplication** instructions are belong to the **data processing** instructions

0000 = AND  
0001 = EOR  
0010 = SUB  
0011 = RSB  
0100 = ADD  
0101 = ADC  
0110 = SBC  
0111 = RSC  
1000 = TST  
1001 = TEQ  
1010 = CMP  
1011 = CMN  
1100 = ORR  
1101 = MOV  
1110 = BIC  
1111 = MNV

Multiplication instructions are encoded by setting bits 25, 26, 27 to zero and bits 4 and 7 to one.



The rotate right through carry (RRX) is encoded as rotate right with zero shift.



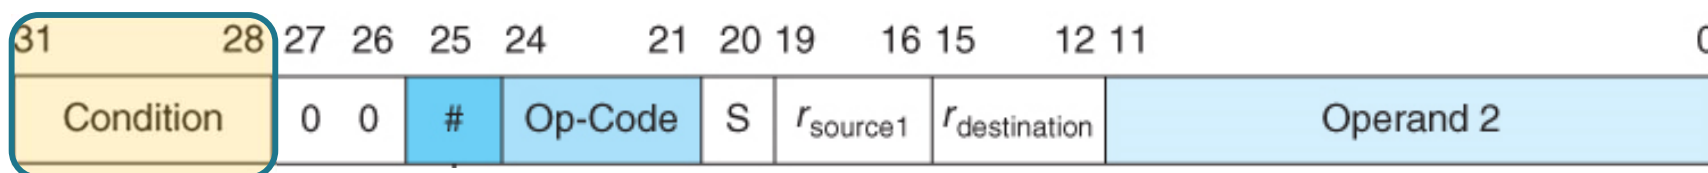
## ARM's Flow Control Instructions (Conditional Execution)

- ❑ One of **ARM**'s most unusual features is that each instruction can be *conditionally executed*
  - associating an instruction with a logical condition.
    - If the stated condition is true, the instruction is executed.
    - Otherwise it is bypassed (*squashed*).
- ❑ Assembly language programmers indicate the conditional execution mode by appending the appropriate condition to a mnemonic (*mnemonic is a text abbreviation for an operation code*).
- ❑ for example,

**ADDEQ**   **r1**, r2, r3   ; IF Z = 1 THEN [r1] <- [r2] + [r3]

specifies that the addition is performed only if the Z-bit is set because a previous result was zero.

**FIGURE 3.26**   Encoding the ARM's data processing instructions



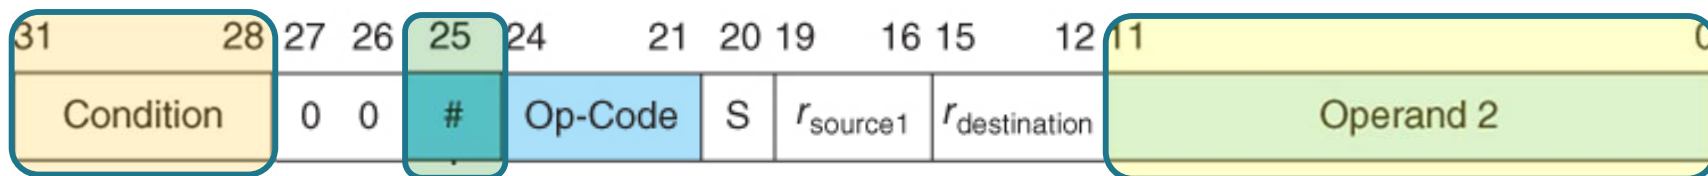
## ARM's Flow Control Instructions (Conditional Execution)

- ❑ There is nothing to stop you **combining** *conditional execution* and *shifting* because the branch and shift fields of an instruction are independent.
- ❑ You can write

**ADDCC** **r1**, r2, r3, LSL r4 ; IF C=0 THEN [r1] ← [r2] + [r3] × 2<sup>[r4]</sup>

FIGURE 3.26

Encoding the ARM's data processing instructions



## ARM's Flow Control Instructions (Conditional Execution)

- ❑ **ARM**'s conditional execution mode makes it easy to implement conditional operations in a high-level language.
- ❑ Consider the following fragment of C code.  
if (P == Q) X = P - Y ;
- ❑ If we assume that r1 contains P,  
r2 contains Q,  
r3 contains X, and  
r4 contains Y, then we can write

```
CMP      r1,r2           ;compare P == Q
SUBEQ    r3,r1,r4        ;if (P == Q) then r3 = r1 - r4
```

- ❑ Notice how simple this operation is implemented without using a branch instruction
  - In this case the subtraction is squashed if the comparison is false



## ARM's Flow Control Instructions (Conditional Execution)

- Now consider a more complicated example of a C construct with a compound predicate:

```
if ((a == b) && (c == d)) e++;
```

- We can write

```
CMP      r0,r1      ;compare a == b
CMPEQ    r2,r3      ;if a == b then test c == d
ADDEQ    r4,r4,#1   ;if a == b AND c == d THEN increment e
```

- The first line, `CMP r0,r1`, compares a and b.
- The next line, `CMPEQ r2,r3`, executes a conditional comparison only if the result of the first line was true (i.e., `a == b`). (**short circuit**)
- The third line, `ADDEQ r4,r4,#1`, is executed only if the previous line was true (i.e., `c == d`) to implement the `e++`.

## ARM's Flow Control Instructions (Conditional Execution)

❑ You can also handle some testing with multiple conditions.

❑ Consider:

```
if (a == b) e = e + 4;
if (a < b)  e = e + 7;
if (a > b)  e = e + 12;
```

We can use conditional execution to implement this as

```
CMP      r0,r1                ;compare a == b
ADDEQ    r4,r4,#4             ;if a == b then e = e + 4
ADDLT    r4,r4,#7             ;if a < b  then e = e + 7
ADDGT    r4,r4,#12            ;if a > b  then e = e + 12
```

❑ Without the conditional execution, we can implement it as follow:

```
      CMP r0,r1                ;compare a == b
      BNE NotEqual
Equal  ADD r4,r4,#4             ;if a == b then e = e + 4
      B   AfterAll
NotEqual BLT LessThan
      ADD r4,r4,#12            ;if a > b  then e = e + 12
      B   AfterAll
LessThan ADD r4,r4,#7          ;if a < b  then e = e + 7
AfterAll ...
```













# Instruction Encoding

ARM Instruction: CMPGT r3, r5

Condition = 1100 (GT)

Op-Code = 1010 (i.e., CMP)

**S = 1 (update flags)**

$r_{destination} = 0000$  (must be zeros)

$r_{source1} = 0011$  (first operand)

# = 0 (second operand not a constant)

Operand 2 (bit number 4 = 0)

$r_{source2} = 0101$

shift type = 00 (logical left)

shift length = 00000

3 3 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0  
1 0 9 8 7 6 5 4 3 2 1 0 9 8 7 6 5 4 3 2 1 0 9 8 7 6 5 4 3 2 1 0  
1 1 0 0 0 0 1 0 1 0 1 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1

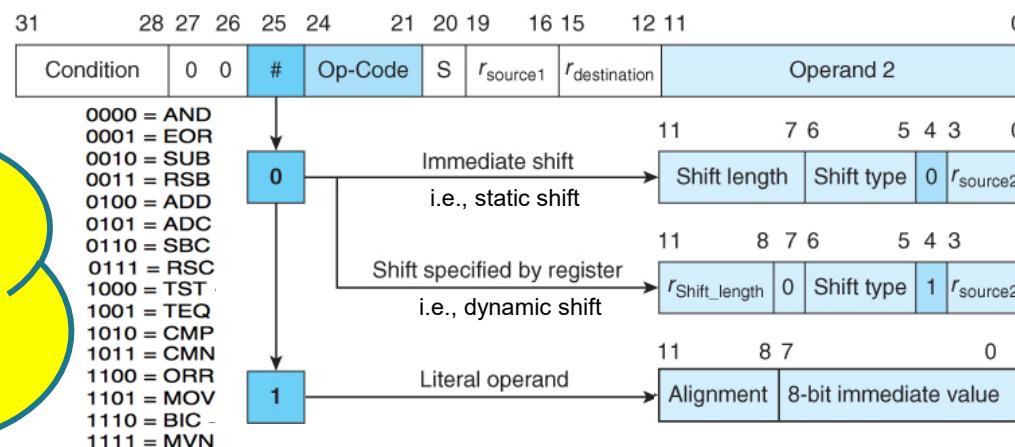
TABLE 3.2

ARM's Conditional Execution and Branch Control Mnemonics

Encoding	Mnemonic	Branch on Flag Status	Execute on condition
0000	EQ	Z set	Equal (i.e., zero)
0001	NE	Z clear	Not equal (i.e., not zero)
0010	CS	C set	Unsigned higher or same
0011	CC	C clear	Unsigned lower
0100	MI	N set	Negative
0101	PL	N clear	Positive or zero
0110	VS	V set	Overflow
0111	VC	V clear	No overflow
1000	HI	C set and Z clear	Unsigned higher
1001	LS	C clear or Z set	Unsigned lower or same
1010	GE	N set and V set, or N clear and V clear	Greater or equal
1011	LT	N set and V clear, or N clear and V set	Less than
1100	GT	Z clear, and either N set and V set, or N clear and V clear	Greater than
1101	LE	Z set, or N set and V clear, or N clear and V set	Less than or equal
1110	AL		Always (default)
1111	NV		Never (reserved)

FIGURE 3.26

Encoding the ARM's data processing instructions



**0xC1530005**

In all test-and-compare instructions, i.e., **TST**, **TEQ**, **CMP**, and **CMN**, the destination register field **MUST BE encoded as 0000** and **S = 1**

# Instruction Encoding

ARM Instruction:

MOV PC, LR

Condition = 1110 (always – unconditional)

Op-Code = 1101 (i.e., MOV)

S = 0 (not MOVS)

$r_{\text{destination}} = 1111$  (PC)

$r_{\text{source1}} = 0000$  (must be zeros)

# = 0 (second operand not a constant)

Operand 2 (bit number 4 = 0)

$r_{\text{source2}} = 1110$

shift type = 00 (logical left)

shift length = 00000

3 3 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0  
1 0 9 8 7 6 5 4 3 2 1 0 9 8 7 6 5 4 3 2 1 0  
1 1 1 0 0 0 1 1 0 1 0 0 0 0 1 1 1 1 0 0 0 0 0 0 1 1 1 0

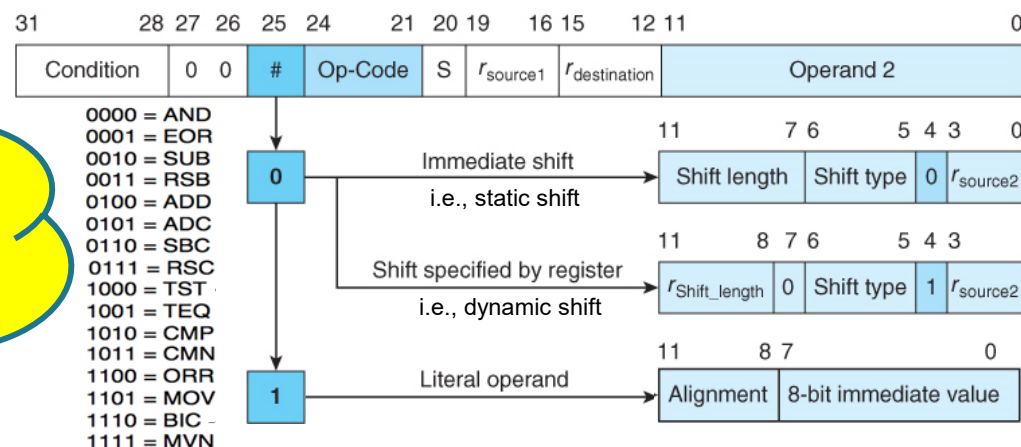
TABLE 3.2

ARM's Conditional Execution and Branch Control Mnemonics

Encoding	Mnemonic	Branch on Flag Status	Execute on condition
0000	EQ	Z set	Equal (i.e., zero)
0001	NE	Z clear	Not equal (i.e., not zero)
0010	CS	C set	Unsigned higher or same
0011	CC	C clear	Unsigned lower
0100	MI	N set	Negative
0101	PL	N clear	Positive or zero
0110	VS	V set	Overflow
0111	VC	V clear	No overflow
1000	HI	C set and Z clear	Unsigned higher
1001	LS	C clear or Z set	Unsigned lower or same
1010	GE	N set and V set, or N clear and V clear	Greater or equal
1011	LT	N set and V clear, or N clear and V set	Less than
1100	GT	Z clear, and either N set and V set, or N clear and V clear	Greater than
1101	LE	Z set, or N set and V clear, or N clear and V set	Less than or equal
1110	AL		Always (default)
1111	NV		Never (reserved)

FIGURE 3.26

Encoding the ARM's data processing instructions



**0xE1A0F00E**

In all moving instructions, i.e., MOV, and MVN, the source<sub>1</sub> register field **MUST BE encoded as 0000**

**Shift type**

00 = logical left  
01 = logical right  
10 = arithmetic right  
11 = rotate right



## Handling Literals

- ❑ In **ARM**, *operand 2* can be a literal.

```
ADD r0,r1,#7      ;adds 7 to r1 and puts the result in r0.
MOV r3,#25        ;moves 25 into r3.
```

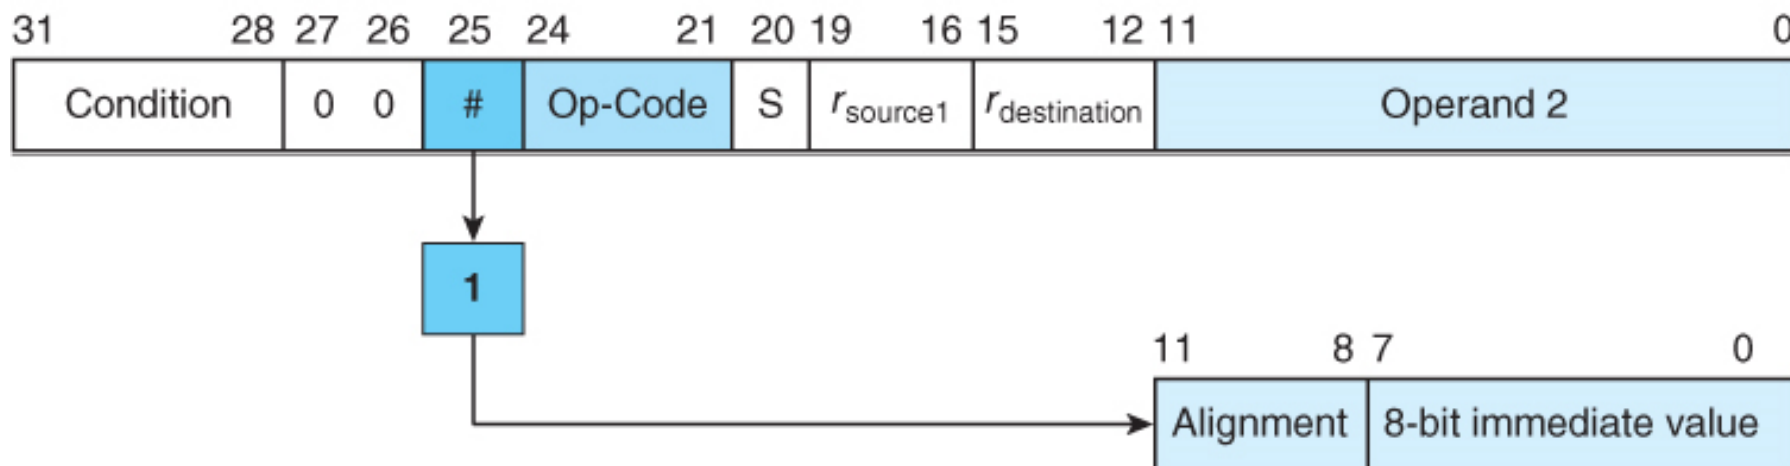
- ❑ *What is the range of such literals?*

- *Operand 2* is a 12-bits field, i.e., it can encode **4096** different values
  - **ARM** encodes these 12-bits as a value from 0 to 255 (i.e., 8-bits) to be rotated (aligned) according to the value of the other bits (i.e., 4-bits)

- ❑ Figure 3.28 illustrate the format of **ARM**'s instructions with a literal operand.

**FIGURE 3.28**

Diagram of ARM's literal operand encoding

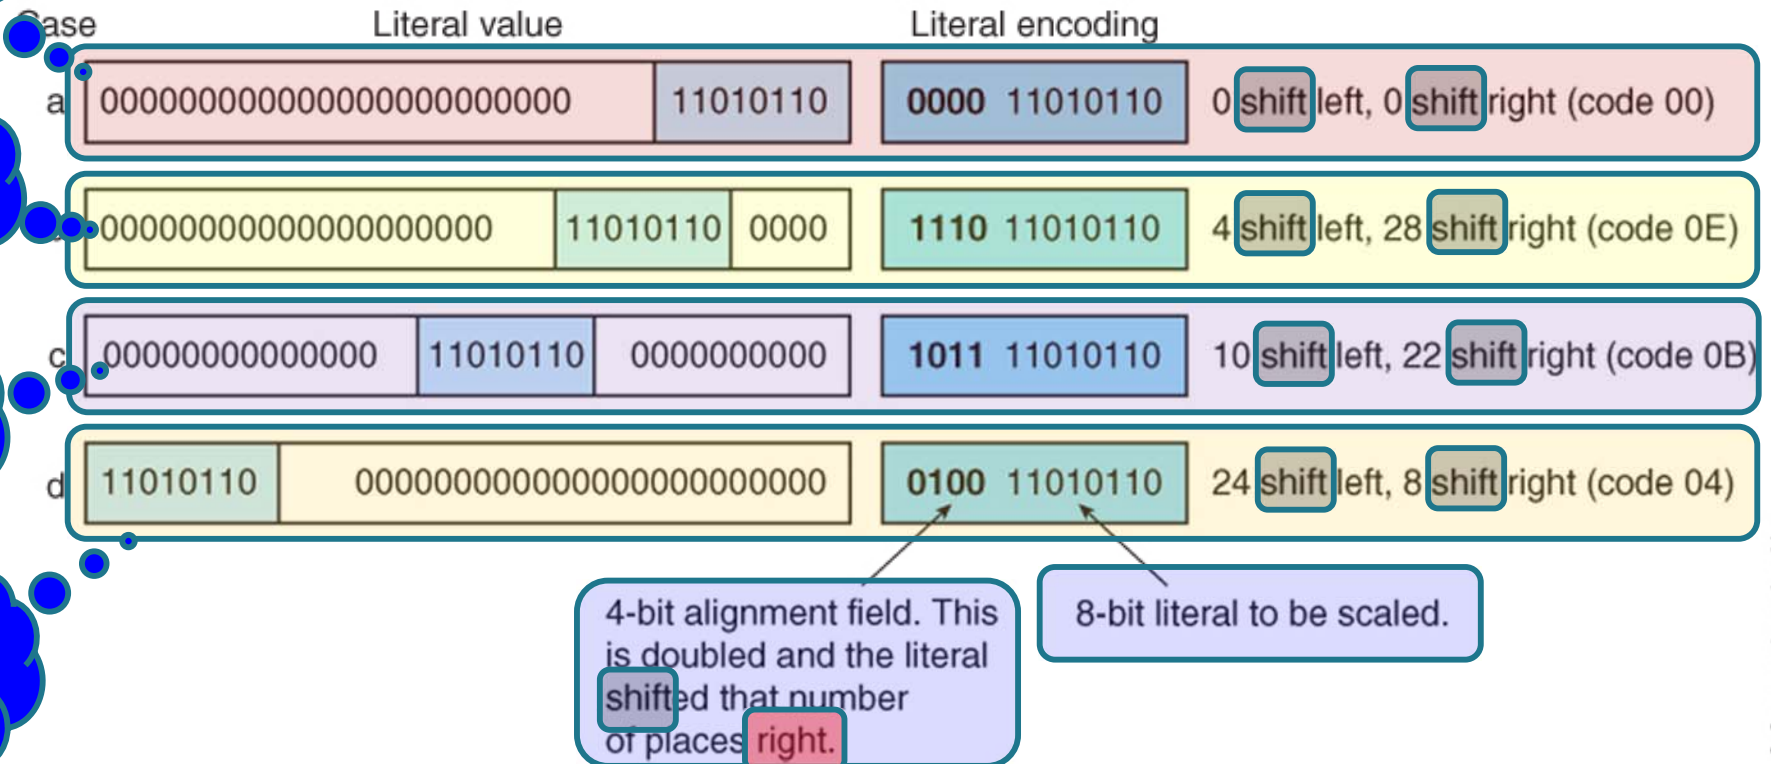




# Handling Literals

Replace the word "shift" in this slide by "rotate".  
The book is wrong!!

**FIGURE 3.29** ARM's literal operand encoding



The number of actual rotations is ALWAYS even (i.e., can not be odd at all).

Also called the  
4 bits *alignment*

## Handling Literals

You need to know how to  
*decode* and *encode* literals

Encoded literal	Scale value	#of rotations <b>right</b> = 2 × Scale value	# of rotations <b>left</b> = 32 - 2 × Scale value	Decoded literal
<b>0000</b> mnop wxyz	<b>0</b>	0	(32) <sub>10</sub>	0000 0000 0000 0000 0000 0000 mnop wxyz
<b>1111</b> mnop wxyz	<b>(15)<sub>10</sub></b>	(30) <sub>10</sub>	2	0000 0000 0000 0000 0000 00mn opwx yz00
<b>1110</b> mnop wxyz	<b>(14)<sub>10</sub></b>	(28) <sub>10</sub>	4	0000 0000 0000 0000 0000 mnop wxyz 0000
<b>1101</b> mnop wxyz	<b>(13)<sub>10</sub></b>	(26) <sub>10</sub>	6	0000 0000 0000 0000 00mn opwx yz00 0000
<b>1100</b> mnop wxyz	<b>(12)<sub>10</sub></b>	(24) <sub>10</sub>	8	0000 0000 0000 0000 mnop wxyz 0000 0000
<b>1011</b> mnop wxyz	<b>(11)<sub>10</sub></b>	(22) <sub>10</sub>	(10) <sub>10</sub>	0000 0000 0000 00mn opwx yz00 0000 0000
<b>1010</b> mnop wxyz	<b>(10)<sub>10</sub></b>	(20) <sub>10</sub>	(12) <sub>10</sub>	0000 0000 0000 mnop wxyz 0000 0000 0000
<b>1001</b> mnop wxyz	<b>9</b>	(18) <sub>10</sub>	(14) <sub>10</sub>	0000 0000 00mn opwx yz00 0000 0000 0000
<b>1000</b> mnop wxyz	<b>8</b>	(16) <sub>10</sub>	(16) <sub>10</sub>	0000 0000 mnop wxyz 0000 0000 0000 0000
<b>0111</b> mnop wxyz	<b>7</b>	(14) <sub>10</sub>	(18) <sub>10</sub>	0000 00mn opwx yz00 0000 0000 0000 0000
<b>0110</b> mnop wxyz	<b>6</b>	(12) <sub>10</sub>	(20) <sub>10</sub>	0000 mnop wxyz 0000 0000 0000 0000 0000
<b>0101</b> mnop wxyz	<b>5</b>	(10) <sub>10</sub>	(22) <sub>10</sub>	00mn opwx yz00 0000 0000 0000 0000 0000
<b>0100</b> mnop wxyz	<b>4</b>	8	(24) <sub>10</sub>	mnop wxyz 0000 0000 0000 0000 0000 0000
<b>0011</b> mnop wxyz	<b>3</b>	6	(26) <sub>10</sub>	opwx yz00 0000 0000 0000 0000 0000 00mn
<b>0010</b> mnop wxyz	<b>2</b>	4	(28) <sub>10</sub>	wxyz 0000 0000 0000 0000 0000 0000 mnop
<b>0001</b> mnop wxyz	<b>1</b>	2	(30) <sub>10</sub>	yz00 0000 0000 0000 0000 0000 00mn opwx



# Handling Literals

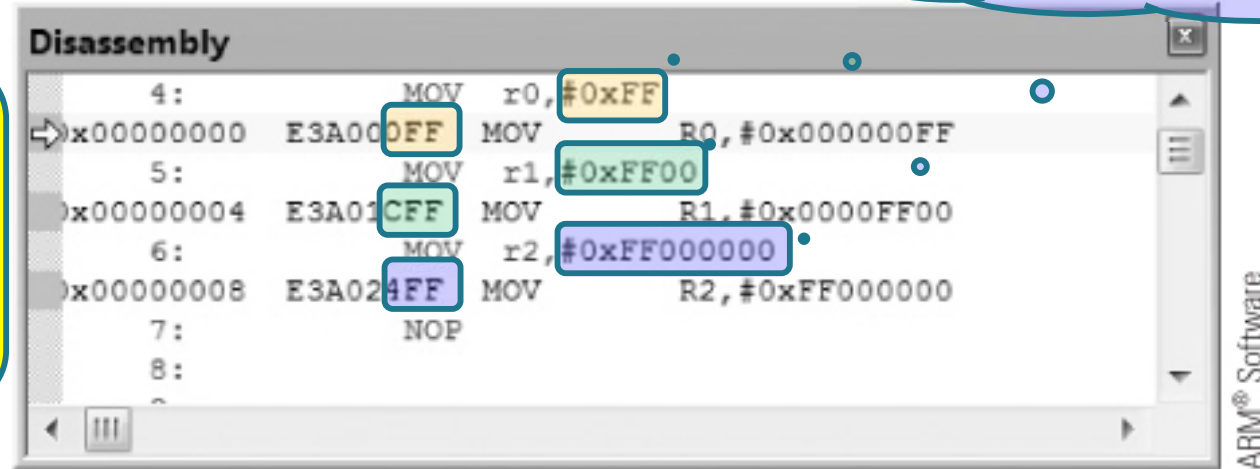
8-bit immediate value = 0xFF  
and the ZERO rotation left

8-bit immediate value = 0xFF  
and the 8 rotations left

8-bit immediate value = 0xFF  
and the 24 rotations left

FIGURE 3.30

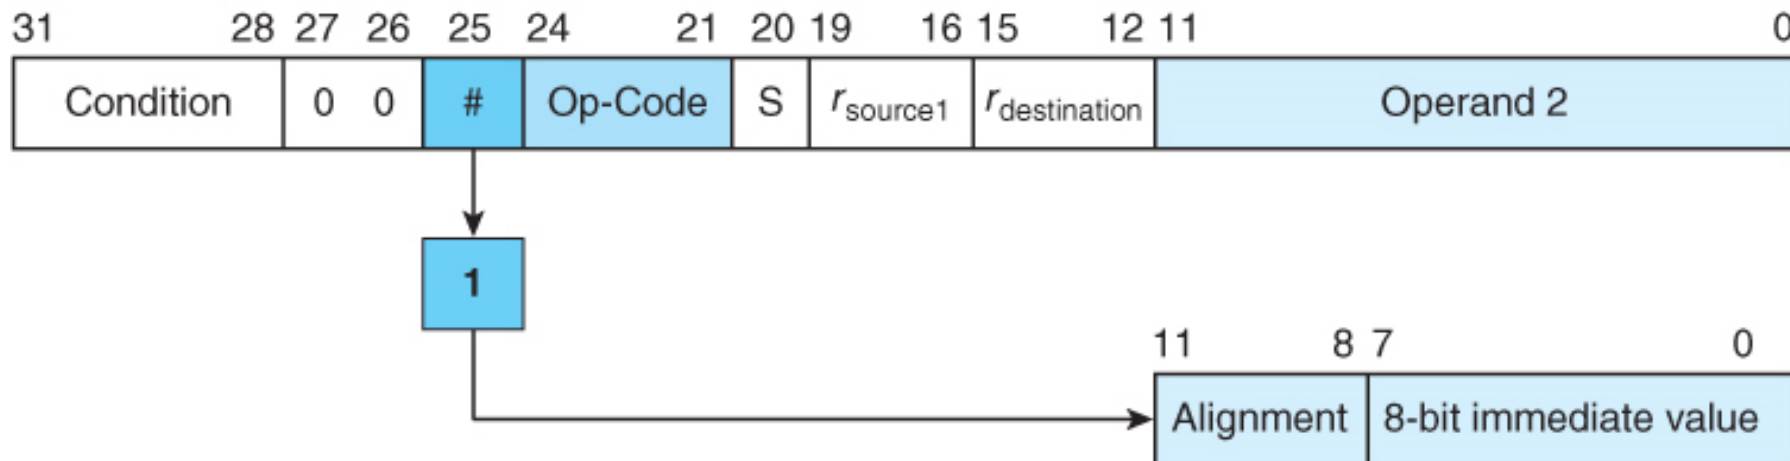
Example of ARM literal encoding



If the value can not be represented that way, you will get an error message saying: "cannot be represented by 0-255 and a rotation."

FIGURE 3.28

Diagram of ARM's literal operand encoding



© Cengage Learning 2014

# Handling Literals

If the literal value is 0x128,  
what is the 0-to-255 value in decimal and the align code in decimal?

Convert the literal into 32-bit binary value.

0x128 → 0x0000 0000 0000 0000 0000 0001 0010 1000

Identify the 1<sup>st</sup> one to the left and the 1<sup>st</sup> one to the right in the 32-bit binary value

0x128 → 0x0000 0000 0000 0000 0000 0001 0010 1000

If the distance between these two ones is *less than or equal 8, inclusive*, or *more than or equal 24, inclusive*, it means that you will be able to encode the number as a value from 0 to 255 and an align code.

0x128 → 0x0000 0000 0000 0000 0000 0001 0010 1000

If the distance is less than 8, use some of the left and right zeros to get the 0-to-255 value (8 bits in total). *Make sure that the number of the other zeros to the left and to the right are even.*

0x128 → 0x0000 0000 0000 0000 0000 0001 0010 1000

The 0-to-255 value is 01 0010 10<sub>2</sub> → 74<sub>10</sub>

The value 01 0010 10<sub>2</sub> needs to be rotated to the left 2 times to become 0x128, which is equivalent to 30 times rotation to the right; hence the align code is 15 (*see slide 107*)

# Instruction Encoding

ARM Instruction: ORRGTS **r1**, r2, #0xAA00

Condition = 1100 (Greater than)

Op-Code = 1100 (i.e., ORR)

S = 1 (ORRGTS)

r<sub>destination</sub> = 0001 (destination *operand*)

r<sub>source1</sub> = 0010 (first *operand*)

# = 1 (second operand is a constant)

Operand 2 (*to be 0-255 and a rotation*)

8-bit immediate value = 0xAA

rotations left = 8

equivalent to 24 rotations right

Half of the rotations right = 12

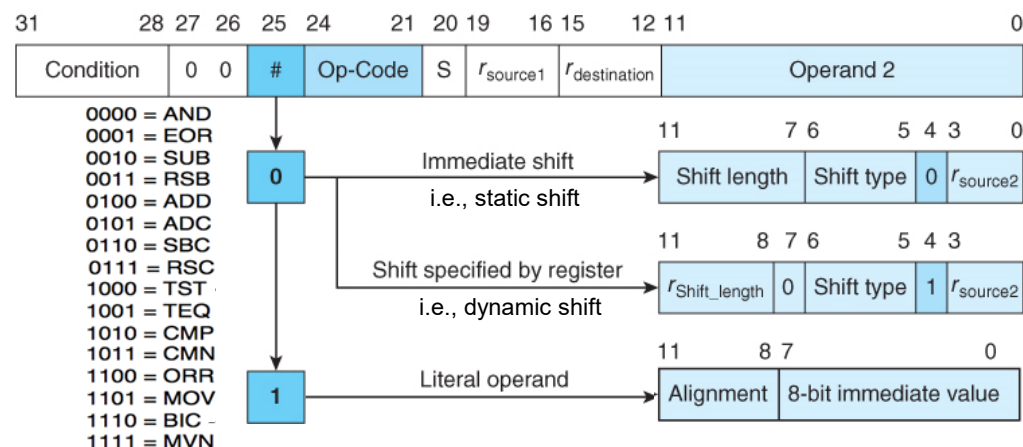
i.e., 1100 in binary

3 3 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 1 1 0 0 0 0 0 0 0 0  
1 0 9 8 7 6 5 4 3 2 1 0 9 8 7 6 5 4 3 2 1 0  
1 1 0 0 0 0 1 1 1 0 0 0 1 0 0 0 1 1 0 0 1 0 1 0 1 0 1 0

TABLE 3.2 ARM's Conditional Execution and Branch Control Mnemonics			
Encoding	Mnemonic	Branch on Flag Status	Execute on condition
0000	EQ	Z set	Equal (i.e., zero)
0001	NE	Z clear	Not equal (i.e., not zero)
0010	CS	C set	Unsigned higher or same
0011	CC	C clear	Unsigned lower
0100	MI	N set	Negative
0101	PL	N clear	Positive or zero
0110	VS	V set	Overflow
0111	VC	V clear	No overflow
1000	HI	C set and Z clear	Unsigned higher
1001	LS	C clear or Z set	Unsigned lower or same
1010	GE	N set and V set, or N clear and V clear	Greater or equal
1011	LT	N set and V clear, or N clear and V set	Less than
1100	GT	Z clear, and either N set and V set, or N clear and V clear	Greater than
1101	LE	Z set, or N set and V clear, or N clear and V set	Less than or equal
1110	AL		Always (default)
1111	NV		Never (reserved)

FIGURE 3.26 Encoding the ARM's data processing instructions

**0xC3921CAA**



Shift type  
00 = logical left  
01 = logical right  
10 = arithmetic right  
11 = rotate right

