

Chapter 3

Formatted Input/Output

The `printf` Function

- The `printf` function must be supplied with a *format string*, followed by any values that are to be inserted into the string during printing:

```
printf(string, expr1, expr2, ...);
```

- The format string may contain both ordinary characters and conversion specifications, which begin with the `%` character.
- A conversion specification is a placeholder representing a value to be filled in during printing.
 - `%d` is used for `int` values
 - `%f` is used for `float` values

The `printf` Function

- Ordinary characters in a format string are printed as they appear in the string; conversion specifications are replaced.

- Example:

```
int i, j;  
float x, y;
```

```
i = 10;
```

```
j = 20;
```

```
x = 43.2892f;
```

```
y = 5527.0f;
```

```
printf("i = %d, j = %d, x = %f, y = %f\n", i, j, x, y);
```

- Output:

```
i = 10, j = 20, x = 43.289200, y = 5527.000000
```

The `printf` Function

- Compilers aren't required to check that the number of conversion specifications in a format string matches the number of output items.
- Too many conversion specifications: *but possibly it would work*

```
printf("%d %d\n", i);    /*** WRONG ***/
```

- Too few conversion specifications:

```
printf("%d\n", i, j);    /*** WRONG ***/
```

*↑
this would run, but j will not be printed.*

The `printf` Function

- Compilers aren't required to check that a conversion specification is appropriate.
- If the programmer uses an incorrect specification, the program will produce meaningless output:

```
printf("%f %d\n", i, x);    /*** WRONG ***/
```

↑
this will run, might get some awkward result anyway

Conversion Specifications

- A conversion specification can have the form $\%m.pX$ or $\%-m.pX$, where m and p are integer constants and X is a letter.
- Both m and p are optional; if p is omitted, the period that separates m and p is also dropped.
- In the conversion specification $\%10.2f$, m is 10, p is 2, and X is f .
- In the specification $\%10f$, m is 10 and p (along with the period) is missing, but in the specification $\%.2f$, p is 2 and m is missing.

Conversion Specifications

- The *minimum field width*, m , specifies the minimum number of characters to print.
- If the value to be printed requires fewer than m characters, it is right-justified within the field.
 - `%4d` displays the number 123 as `•123`. (• represents the space character.) *it would fill something to fill the space*
- If the value to be printed requires more than m characters, the field width automatically expands to the necessary size.
- Putting a minus sign in front of m causes left justification.
 - The specification `%-4d` would display 123 as `123•`.

Conversion Specifications

- The meaning of the *precision*, p , depends on the choice of X , the *conversion specifier*.
- The d specifier is used to display an integer in decimal form.
 - p indicates the minimum number of digits to display (extra zeros are added to the beginning of the number if necessary).
 - If p is omitted, it is assumed to be 1.

Conversion Specifications

- Conversion specifiers for floating-point numbers:
 - e — Exponential format. p indicates how many digits should appear after the decimal point (the default is 6). If p is 0, no decimal point is displayed.
 - f — “Fixed decimal” format. p has the same meaning as for the e specifier.
 - g — Either exponential format or fixed decimal format, depending on the number’s size. p indicates the maximum number of significant digits to be displayed. The g conversion won’t show trailing zeros. If the number has no digits after the decimal point, g doesn’t display the decimal point.

Program: Using `printf` to Format Numbers

- The `tprintf.c` program uses `printf` to display integers and floating-point numbers in various formats.

Chapter 3: Formatted Input/Output

tprintf.c

```
/* Prints int and float values in various formats */  
  
#include <stdio.h>  
  
int main(void)  
{  
    int i;  
    float x;  
  
    i = 40;  
    x = 839.21f;  
  
    printf("|%d|%5d|%-5d|%5.3d|\n", i, i, i, i);  
    printf("|%10.3f|%10.3e|%-10g|\n", x, x, x);  
  
    return 0;  
}
```

- Output:

Five spaces

```
|40|_ _ _40|40_ _ _|_ _ _040|  
|_ _ _839.210|_ _ _8.392e+02|839.21_ _ _|
```

3 7 spaces 1 space

Escape Sequences

- The `\n` code that used in format strings is called an *escape sequence*.
- Escape sequences enable strings to contain nonprinting (control) characters and characters that have a special meaning (such as ").
- A partial list of escape sequences:

Alert (bell)	<code>\a</code>
Backspace	<code>\b</code>
New line	<code>\n</code>
Horizontal tab	<code>\t</code>

Escape Sequences

- A string may contain any number of escape sequences:

```
printf("Item\tUnit\tPurchase\n\tPrice\tDate\n");
```

- Executing this statement prints a two-line heading:

```
Item      Unit      Purchase
          Price    Date
```

Escape Sequences

- Another common escape sequence is `\"`, which represents the `"` character:

```
printf("\\"Hello!\");  
/* prints "Hello!" */
```

- To print a single `\` character, put two `\` characters in the string:

```
printf("\\");  
/* prints one \ character */
```

The `scanf` Function

- `scanf` reads input according to a particular format.
- A `scanf` format string may contain both ordinary characters and conversion specifications.
- The conversions allowed with `scanf` are essentially the same as those used with `printf`.

The `scanf` Function

- In many cases, a `scanf` format string will contain only conversion specifications:

```
int i, j;  
float x, y;  
  
scanf("%d%d%f%f", &i, &j, &x, &y);
```

initialization ✓

- Sample input:

```
1 -20 .3 -4.0e3
```

`scanf` will assign 1, -20, 0.3, and -4000.0 to `i`, `j`, `x`, and `y`, respectively.

The `scanf` Function

- When using `scanf`, the programmer must check that the number of conversion specifications matches the number of input variables and that each conversion is appropriate for the corresponding variable.
- Another trap involves the `&` symbol, which normally precedes each variable in a `scanf` call.
- The `&` is usually (but not always) required, and it's the programmer's responsibility to remember to use it.

How `scanf` Works

- `scanf` tries to match groups of input characters with conversion specifications in the format string.
- For each conversion specification, `scanf` tries to locate an item of the appropriate type in the input data, skipping blank space if necessary.
- `scanf` then reads the item, stopping when it reaches a character that can't belong to the item.
 - If the item was read successfully, `scanf` continues processing the rest of the format string.
 - If not, `scanf` returns immediately.

How scanf Works

- As it searches for a number, `scanf` ignores *white-space characters* (space, horizontal and vertical tab, form-feed, and new-line).

- A call of `scanf` that reads four numbers:

```
scanf ("%d%d%f%f", &i, &j, &x, &y);
```

- The numbers can be on one line or spread over several lines:

```
1
-20    .3
      -4.0e3
```

- `scanf` sees a stream of characters (␣ represents new-line):

```
••1␣-20•••.3␣•••-4.0e3␣
ssrsrrrrsssrssssrrrrrrr (s = skipped; r = read)
```

- `scanf` “peeks” at the final new-line without reading it.

How `scanf` Works

- When asked to read an integer, `scanf` first searches for a digit, a plus sign, or a minus sign; it then reads digits until it reaches a nondigit.
- When asked to read a floating-point number, `scanf` looks for
 - a plus or minus sign (optional), followed by
 - digits (possibly containing a decimal point), followed by
 - an exponent (optional). An exponent consists of the letter `e` (or `E`), an optional sign, and one or more digits.
- `%e`, `%f`, and `%g` are interchangeable when used with `scanf`.

How `scanf` Works

- When `scanf` encounters a character that can't be part of the current item, the character is “put back” to be read again during the scanning of the next input item or during the next call of `scanf`.

How `scanf` Works

- Sample input:
1-20.3-4.0e3x
- The call of `scanf` is the same as before:

```
scanf("%d%d%f%f", &i, &j, &x, &y);
```
- Here's how `scanf` would process the new input:
 - `%d`. Stores 1 into `i` and puts the `-` character back.
 - `%d`. Stores -20 into `j` and puts the `.` character back.
 - `%f`. Stores 0.3 into `x` and puts the `-` character back.
 - `%f`. Stores -4.0×10^3 into `y` and puts the new-line character back.

Ordinary Characters in Format Strings

- When it encounters one or more white-space characters in a format string, `scanf` reads white-space characters from the input until it reaches a non-white-space character (which is “put back”).
- When it encounters a non-white-space character in a format string, `scanf` compares it with the next input character.
 - If they match, `scanf` discards the input character and continues processing the format string.
 - If they don’t match, `scanf` puts the offending character back into the input, then aborts.

Ordinary Characters in Format Strings

- Examples:
 - If the format string is `"%d/%d"` and the input is `•5/•96`, `scanf` succeeds.
 - If the input is `•5•/•96`, `scanf` fails, because the `/` in the format string doesn't match the space in the input.
- To allow spaces after the first number, use the format string `"%d /%d"` instead.

Confusing `printf` with `scanf`

- Although calls of `scanf` and `printf` may appear similar, there are significant differences between the two.
- One common mistake is to put `&` in front of variables in a call of `printf`:

```
printf("%d %d\n", &i, &j);    /*** WRONG ***/
```

↑
this will print
the address of these
values. aka. garbage.

Confusing `printf` with `scanf`

- Incorrectly assuming that `scanf` format strings should resemble `printf` format strings is another common error.
- Consider the following call of `scanf`:

```
scanf("%d, %d", &i, &j);
```

- `scanf` will first look for an integer in the input, which it stores in the variable `i`.
- `scanf` will then try to match a comma with the next input character.
- If the next input character is a space, not a comma, `scanf` will terminate without reading a value for `j`.

Confusing `printf` with `scanf`

- Putting a new-line character at the end of a `scanf` format string is usually a bad idea.
- To `scanf`, a new-line character in a format string is equivalent to a space; both cause `scanf` to advance to the next non-white-space character.
- If the format string is "`%d\n`", `scanf` will skip white space, read an integer, then skip to the next non-white-space character.
- A format string like this can cause an interactive program to “hang.”

Program: Adding Fractions

- The `addfrac.c` program prompts the user to enter two fractions and then displays their sum.
- Sample program output:

Enter first fraction: 5/6

Enter second fraction: 3/4

The sum is 38/24

Chapter 3: Formatted Input/Output

addfrac.c

```
/* Adds two fractions */

#include <stdio.h>

int main(void)
{
    int num1, denom1, num2, denom2, result_num, result_denom;

    printf("Enter first fraction: ");
    scanf("%d/%d", &num1, &denom1);

    printf("Enter second fraction: ");
    scanf("%d/%d", &num2, &denom2);

    result_num = num1 * denom2 + num2 * denom1;
    result_denom = denom1 * denom2;
    printf("The sum is %d/%d\n", result_num, result_denom)

    return 0;
}
```