




Finite State Machines

State Minimization

Chapter 5

State Minimization


Consider:



Is this a minimal machine?


State Minimization

Step (1): Get rid of **unreachable** states.



State 3 is unreachable.

Step (2): Get rid of **redundant** states.




States 2 and 3 are redundant.

Getting Rid of Unreachable States

We can't easily find the unreachable states directly.
But we can find the **reachable** ones and determine the unreachable ones from there.


An algorithm for finding the reachable states:



Getting Rid of Redundant States

Intuitively, two states are **equivalent** to each other (and thus one is redundant) if all strings in Σ^* have the same fate, regardless of which of the two states the machine is in. But how can we tell this?


The simple case:



Two states have identical sets of transitions out.

Getting Rid of Redundant States

The harder case:



The outcomes in states 2 and 3 are the same, even though the states aren't.



Finding an Algorithm for Minimization

Capture the notion of equivalence classes of strings with respect to a language.

Prove that we can always find a (unique up to state naming) deterministic FSM with a number of states equal to the number of equivalence classes of strings.

Describe an algorithm for finding that deterministic FSM.

Defining Equivalence for Strings

We want to capture the notion that two strings are equivalent or **indistinguishable** with respect to a language L if, no matter what is tacked on to them on the right, either they will both be in L or neither will. Why is this the right notion? Because it corresponds naturally to what the states of a recognizing FSM have to remember.

Example:

(1) a b a b a b

(2) b a a b a b

Suppose $L = \{w \in \{a, b\}^* : |w| \text{ is even}\}$. Are (1) and (2) equivalent?

Suppose $L = \{w \in \{a, b\}^* : \text{every } a \text{ is immediately followed by } b\}$.
Are (1) and (2) equivalent?



Defining Equivalence for Strings

Equivalent, or **indistinguishable**, string with respect to L :

$$\forall x, y \in \Sigma^*: \quad x \approx_L y \quad \text{iff} \quad \forall z \in \Sigma^* (xz \in L \text{ iff } yz \in L).$$

\approx_L is an Equivalence Relation

\approx_L is an **equivalence** relation because it is:

- Reflexive: $\forall x \in \Sigma^* (x \approx_L x)$, because:
$$\forall x, z \in \Sigma^* (xz \in L \leftrightarrow xz \in L).$$
- Symmetric: $\forall x, y \in \Sigma^* (x \approx_L y \rightarrow y \approx_L x)$, because:
$$\forall x, y, z \in \Sigma^* ((xz \in L \leftrightarrow yz \in L) \leftrightarrow (yz \in L \leftrightarrow xz \in L)).$$
- Transitive: $\forall x, y, z \in \Sigma^* (((x \approx_L y) \wedge (y \approx_L w)) \rightarrow (x \approx_L w))$,
because:
$$\forall x, y, z \in \Sigma^* (((xz \in L \leftrightarrow yz \in L) \wedge (yz \in L \leftrightarrow wz \in L)) \rightarrow (xz \in L \leftrightarrow wz \in L)).$$

\approx_L is an Equivalence Relation

Because \approx_L is an equivalence relation:

- No equivalence class of \approx_L is empty.
- Each string in Σ^* is in exactly one equivalence class of \approx_L .



An example of \approx_L

$$\Sigma = \{a, b\}$$

$$L = \{w \in \Sigma^* : |w| \text{ is even}\}$$

The equivalence classes of \approx_L :

$[\varepsilon, aa, ab, ba, bb, aaaa, aaab, \dots]$ – even length

$[a, b, aaa, aab, aba, abb, \dots]$ – odd length

Another example

$$\Sigma = \{a, b\}$$

$$L = \{w \in \Sigma^*: \text{every } a \text{ is immediately followed by } b\}$$

The equivalence classes of \approx_L :

$[\epsilon, b, abb, \dots]$ [all strings in L].

$[a, abbba, \dots]$ [all strings that end in a and have no prior a that is not followed by a b].

$[aa, abaa, \dots]$ [all strings that contain at least one instance of aa].



Yet Another Example of \approx_L

$$\Sigma = \{a, b\}$$

$$L = aab^*a$$

The equivalence classes of \approx_L :

When More Than One Class Contains Strings in L


$$\Sigma = \{a, b\}$$


$$L = \{w \in \Sigma^* : \text{no two adjacent characters are the same}\}$$

The equivalence classes of \approx_L :

[0] $[\epsilon]$ 

[1] $[a, aba, ababa, \dots]$ 


[2] $[b, ab, bab, abab, \dots]$ 

[3] $[aa, abaa, ababb\dots]$ 

Does \approx_L Always Have a Finite Number of Equivalence Classes?

$$\Sigma = \{a, b\}$$

$$L = \{a^n b^n, n \geq 0\}$$

ε	aa	aaaa
a	aba	aaaaa
b	aaa	

The equivalence classes of \approx_L :



Containment of Strings

Definition:

A state q of a DFSA M **contains** the set of strings s such that M lands in state q after reading s , where M is assumed to have started in its start state.

The Best We Can Do

Theorem 5.4: Let L be a regular language and let M be a DFSA that accepts L . The number of states in M is greater than or equal to the number of equivalence classes of \approx_L .

Proof: Suppose that the number of states in M were less than the number of equivalence classes of \approx_L . Then, by the pigeonhole principle, there must be at least one state q that contains strings from at least two equivalence classes of \approx_L . But then M 's future behavior on those strings will be identical, which is not consistent with the fact that they are in different equivalence classes of \approx_L .

The Best Is Unique

Theorem 5.5: Let L be a regular language over some alphabet Σ . Then there is a DFSM M that accepts L and that has precisely n states where n is the number of equivalence classes of \approx_L . Any other FSM that accepts L must either have more states than M or it must be equivalent to M except for state names.

Proof: (by construction)

$M = (K, \Sigma, \delta, s, A)$, where:

- K contains n states, one for each equivalence class of \approx_L .
- $s = [\varepsilon]$, the equivalence class of ε under \approx_L .
- $A = \{[x] : x \in L\}$.
- $\delta([x], a) = [xa]$. In other words, if M is in the state that contains some string x , then, after reading the next symbol, a , it will be in the state that contains xa .

Proof, Continued

We must show that:

- **K is finite.** Since L is regular, it is accepted by some DFSA M' . M' has some finite number of states m . By Theorem 5.4 (see above), $n \leq m$. So K is finite.
- **δ is a function.** In other words, it is defined for all (state, input) pairs and it produces, for each of them, a unique value. The construction defines a value of δ for all (state, input) pairs. The fact that the construction guarantees a unique such value follows from the definition of \approx_L .

Proof, Continued

- $L = L(M)$. To prove this, we show first that $\forall s, t \ (([\epsilon], st) \vdash_M^* ([s], t))$.

We do this by induction on $|s|$.

If $|s| = 0$ then we have $([\epsilon], \epsilon t) \vdash_M^* ([\epsilon], t)$, which is true since M simply makes zero moves.

Proof, Continued

Assume that the claim is true if $|s| = k$. Then we consider what happens when $|s| = k+1$. $|s| \geq 1$, so we can let $s = yc$ where $y \in \Sigma^*$ and $c \in \Sigma$. We have:

- M reads the first k characters:
 $([\varepsilon], yct) \vdash_M^* ([y], ct)$ (induction hypothesis, since $|y| = k$).
- M reads one more character:
 $([y], ct) \vdash_M^* ([yc], t)$ (definition of δ_M).
- Combining those two, after M has read $k+1$ characters:
 $([\varepsilon], st) = ([\varepsilon], yct) \vdash_M^* ([y], ct) \vdash_M^* ([yc], t) = ([s], t)$

Proof, Continued

So we proved :

$$\forall s, t \ (([\epsilon], st) \vdash_M^* ([s], t)).$$

Let t be ϵ . Let s be any string in Σ^* . By the above:

$$([\epsilon], s) \vdash_M^* ([s], \epsilon).$$

M accepts s iff $[s] \in A = \{[x] : x \in L\}$ iff $s \in L$


So M accepts precisely those strings that are in L , that is,
 $L(M) = L$

Proof, Continued

- There exists no smaller machine $M^\#$ that also accepts L . This follows directly from Theorem 5.4 (see above), which says that the number of equivalence classes of \approx_L imposes a lower bound on the number of states in any DFSA that accepts L .
- There is no different machine $M^\#$ that also has n states and that accepts L .

Example

- $L = \{w \in \{a,b\}^* : \text{no adjacent characters are the same}\}$
- equivalence classes of \approx_L :
- [0] $[\epsilon]$ = start state
- [1] $[a, ba, aba, baba, \dots]$
- [2] $[b, ab, bab, abab, \dots]$
- [3] $[aa, bb, abaa, ababb\dots]$



$$\delta([x], a) = [xa]$$

The Myhill-Nerode Theorem

Theorem 5.6 (Myhill-Nerode): A language is regular iff the number of equivalence classes of \approx_L is finite.

Proof: Show the two directions of the implication:

L regular \rightarrow the number of equivalence classes of \approx_L is finite: If L is regular, then there exists some FSM M that accepts L . M has some finite number of states m . The cardinality of $\approx_L \leq m$. So the cardinality of \approx_L is finite.


The number of equivalence classes of \approx_L is finite $\rightarrow L$ regular: If the cardinality of \approx_L is finite, then the construction that was described in the proof of the previous theorem will build an FSM that accepts L . So L must be regular.

So Where Do We Stand?

1. We know that for any regular language L there exists a minimal accepting machine M_L .
2. We know that $|K|$ of M_L equals the number of equivalence classes of \approx_L .
3. We know how to construct M_L from \approx_L .
4. We know that M_L is unique up to the naming of its states.

But is this good enough?

Consider:





Minimizing an Existing DFSM (Without Knowing \approx_L)

Two approaches:

- Begin with M and collapse redundant states, getting rid of one at a time until the resulting machine is minimal.
- Begin by overclustering the states of L into just two groups, accepting and nonaccepting. Then iteratively split those groups apart until all the distinctions that L requires have been made.


The Overclustering Approach

We need a definition for “equivalent”, i.e., **mergeable** states.

Define $q \equiv p$ iff for all strings $w \in \Sigma^*$, either w drives M to an accepting state from both q and p or it drives M to a rejecting state from both q and p .

An Example

$$\Sigma = \{a, b\} \quad L = \{w \in \Sigma^* : |w| \text{ is even}\}$$



$$q_2 \equiv q_3$$



Constructing \equiv as the Limit of a Sequence of Approximating Equivalence Relations \equiv^n

(Where n is the length of the input strings that have been considered so far)

Consider input strings, starting with ε , and increasing in length by 1 at each iteration. Start by way overgrouping states. Then split them apart as it becomes apparent (with longer and longer strings) that their behavior is not identical.

Constructing \equiv_n

- $p \equiv^0 q$ iff they behave equivalently when they read ϵ . In other words, if they are **both accepting or both rejecting** states.
- $p \equiv^1 q$ iff they behave equivalently when they read any string of length 1, i.e., if any single character sends both of them to an accepting state or both of them to a rejecting state. Note that this is equivalent to saying that any single character sends them to states that are \equiv^0 to each other.
- $p \equiv^2 q$ iff they behave equivalently when they read any string of length 2, which they will do if, when they read the first character they land in states that are \equiv^1 to each other. By the definition of \equiv^1 , they will then yield the same outcome when they read the single remaining character.
- And so forth.

Constructing \equiv (cont'd)

More precisely, $\forall p, q \in K$ and any $n \geq 1$, $q \equiv^n p$
iff:

1. $q \equiv^{n-1} p$, and
2. $\forall a \in \Sigma (\delta(p, a) \equiv^{n-1} \delta(q, a))$


MinDFSM

MinDFSM(*M*: DFSM) =

1. *classes* := {*A*, *K-A*}; 🗨️
2. While $\exists E \in \text{classes}, p, q \in E, c \in \Sigma$ with $[\delta(p, c)] \neq [\delta(q, c)]$ do
3. split *E* such that $[p] \neq [q]$ 🗨️
4. remove *E* from *classes*
5. add the subclasses of *E* to *classes*
6. return $M^* = (\text{classes}, \Sigma, \delta, [s_M], \{[q]: q \in A_M\})$, where δ_{M^*} is constructed as follows:
 if $\delta_M(q, c) = p$, then $\delta_{M^*}([q], c) = [p]$

An Example

$$\Sigma = \{a, b\}$$




$$\equiv^0 = \{ [1,3,5,6], [2,4] \}$$

$$\equiv^1 = \{ [1,3,5], [6], [2,4] \}$$

$$\equiv^2 = \{ [1,3,5], [2], [4], [6] \}$$

The Result



Summary

- Given any regular language L , there exists a minimal DFSA M that accepts L .
- M is unique up to the naming of its states.
- Given any DFSA M , there exists an algorithm *minDFSA* that constructs a minimal DFSA that also accepts $L(M)$.