



Western
UNIVERSITY • CANADA

Chapter 5b – Algorithms

Spring 2023

Algorithms

- [https://en.wikipedia.org/wiki/Scheduling_\(computing\)](https://en.wikipedia.org/wiki/Scheduling_(computing))
- How does the scheduler choose which process in the ready queue to run next?
 - CPU utilization – Can we keep the CPU(s) as busy as possible?
 - Throughput – How many processes can we run in a time frame? The more the merrier
 - Turnaround time – How long are processes taking? From **New** to **Terminated**
 - Waiting time – How long are processes waiting in the ready queue
 - Response time – How long are processes taking to just start giving responses?

Algorithms

- Fairness – The general goal is to:
 - Maximize CPU utilization
 - Maximize throughput
 - Minimize turnaround time
 - Minimize waiting time
 - Minimize response time
 - Some systems may value some criteria more than others
 - No matter what, we don't want any processes to be starved out (never gets run!)
- o CPU time*
↓

Algorithms

- First-Come, First-Served Scheduling (FIFO) (Queue)
- Last-Come, First Served Scheduling (LIFO) (Stack)
- Shortest-Job-First Scheduling
- Round-Robin Scheduling
- Priority Scheduling
- Lottery Scheduling
- Multilevel Queue Scheduling
- Multilevel Feedback Queue Scheduling

First-Come, First-Served

- The first process to enter the ready queue gets to run first
- A process continues to run until it needs I/O or until it terminates
- Is this preemptive or nonpreemptive? ✓
- Pro – The simplest algorithm to understand and implement
- Con – The waiting time can be long

First-Come, First-Served

- Consider the following processes in the following order:

Process	Burst time (milliseconds)
P ₁	24
P ₂	3
P ₃	3

- The Gantt chart for the schedule is:



- Waiting time for P₁ = 0, P₂ = 24, P₃ = 27. Average is $(0 + 24 + 27)/3 = 17$
- Turnaround time for P₁ = 24, P₂ = 27, P₃ = 30. Average is $(24 + 27 + 30)/3 = 27$

First-Come, First-Served

- Suppose we had a different order:

Process	Burst time (milliseconds)
P ₂	3
P ₃	3
P ₁	24

- The Gantt chart for the schedule is:



- Waiting time for **P₁ = 6**, **P₂ = 0**, **P₃ = 3**. Average is **(6 + 0 + 3)/3 = 3**
- Turnaround time for **P₁ = 30**, **P₂ = 3**, **P₃ = 6**. Average is **(30 + 3 + 6)/3 = 14**

First-Come, First-Served

- Suppose we had a different order:

Process	Burst time (milliseconds)
P ₂	3
P ₁	24
P ₃	3

- What does the Gantt chart look like?



- What is the waiting time for P₁, P₂, P₃? Average waiting time? $(0 + 3 + 27) / 3 = 10$
- What is the turnaround time for P₁, P₂, P₃? Average turnaround time? $(3 + 27 + 30) / 3 = 20$

First-Come, First-Served

- Summary:

	P ₁ , P ₂ , P ₃	P ₂ , P ₃ , P ₁	P ₂ , P ₁ , P ₃
Average Waiting Time	17	3	$(3 + 0 + 27)/3 = 10$
Average Turnaround Time	27	14	$(27 + 3 + 30)/3 = 20$

- Convoy effect – Short processes stuck behind long processes
 - We want to avoid this!

First-Come, First-Served

- We should be considering arrival time

Process	Arrival time	Burst time (milliseconds)
P ₁	0	24
P ₂	1	3
P ₃	2	3

- The Gantt chart for the schedule is:



- Waiting time for P₁ = 0, P₂ = 24-1 = 23, P₃ = 27-2 = 25. Average is $(0 + 23 + 25)/3 = 16$
- Turnaround time for P₁ = 24, P₂ = 27-1 = 26, P₃ = 30-2 = 28. Average is $(24 + 26 + 28)/3 = 26$

First-Come, First-Served

- Suppose we had a different order

Process	Arrival time	Burst time (milliseconds)
P ₂	0	3
P ₁	1	24
P ₃	2	3

- What does the Gantt chart look like?



- What is the waiting time for P₁, P₂, P₃? Average waiting time? $[0 + (3-2) + (24-1)] / 3 = 9$
- What is the turnaround time for P₁, P₂, P₃? Average turnaround time? $[3 + (24-1) + (3-2)] / 3 = 19$

Last-Come, First Served

- The latest process to enter the ready queue gets to run first
- Pro – Improves response time for newly created processes
- Con – Risk of starvation. Some processes may never be picked up

Shortest-Job-First

- (More accurate to call it "Shortest-next-CPU-burst", since we are based off CPU bursts, not the entire job)
- Associate with each process the length of its next CPU burst
 - Use these lengths to schedule the process with the shortest time
 - Generally, nonpreemptive but preemptive is, of course, possible
- Pro – If the CPU burst time is accurate, SJF is optimal. Guaranteed to provide the minimum average waiting time for a given set of processes
- Con – How do we determine the length of the next CPU burst? We could ask the user/developer to set it, or we could make an estimate based on previous history


2333}

Shortest-Job-First

- Consider the following processes

Process	Burst time (milliseconds)
P ₁	6
P ₂	8
P ₃	7
P ₄	3

- What does the Gantt chart look like?



- Waiting time for P₁ = 3, P₂ = 16, P₃ = 9, P₄ = 0. Average is $(3 + 16 + 9 + 0)/4 = 7$
- Turnaround time for P₁ = 9, P₂ = 24, P₃ = 16, P₄ = 3. Average is $(9 + 24 + 16 + 3)/4 = 13$

Shortest-Job-First

- With arrival time

Process	Arrival time	Burst time (milliseconds)
P ₁	0	6
P ₂	1	8
P ₃	2	7
P ₄	3	3

- What does the Gantt chart look like?
- What is the waiting time for P₁, P₂, P₃, P₄? Average waiting time?
- What is the turnaround time for P₁, P₂, P₃, P₄? Average turnaround time?

Shortest-Job-First

- How do we estimate the next CPU burst? **Exponential smoothing** (or Exponential averaging)
 - It probably won't be the same as the previous CPU burst, but it will probably be similar
 1. t_n = actual length of n^{th} CPU burst
 2. τ_{n+1} = predicted value for the next CPU burst
 3. α , $0 \leq \alpha \leq 1$
 4. Define:

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$$

Shortest-Job-First

- How do we estimate the next CPU burst?
 - Set α to 0 – The most recent history has no bearing. Pick an estimate and never update it

$$\tau_{n+1} = \tau_n$$


- Set α to 1 – Only the most recent history has any bearing

$$\tau_{n+1} = t_n$$

- More commonly, set α to $\frac{1}{2}$ so we get the most recent history and all the past history
 - Set α to 0.5 – 1.0 to put higher weight on recent history, 0 – 0.5 to put higher weight on past history

Shortest-Job-First

- How do we estimate the next CPU burst?
- Pick any value for τ_0 to get us started (say, 10) set α to $\frac{1}{2}$



$$T_{n+1} = \frac{1}{2} \times 6 + \frac{1}{2} \times 10$$

$$= 8$$

$$T_{n+1} = \frac{1}{2} \times 4 + \frac{1}{2} \times 8 = 6$$

CPU burst (t_i)	6	4	6	4	13	13	...		
"guess" (τ_i)	10	8	6	6	5	9	11	12	...



Shortest-Job-First

- How do we estimate the next CPU burst?
- What if we give higher weight to the most recent burst and set α to 0.75

t_i		6	4	6	4	13	13	13
τ_0	10	4.5+2.5 7	3+1.8 4.8	4.5+1.2 5.7	3+1.4 4.4	9.8+1.1 10.9	9.8+2.8 12.6	9.8+3.2 13

- What if we give higher weight to past history and set α to 0.25


t_i		6	4	6	4	13	13	13
τ_0	10							

Shortest-Job-First

- With arrival time and preemption ("shortest-remaining-time-first")

Process	Arrival time	Burst time (milliseconds)
P ₁	0	8
P ₂	1	4
P ₃	2	9
P ₄	3	5

- What does the Gantt chart look like? P₁ has 7ms left but P₂ has 4:



- P₁ = 10-1, P₂ = 1-1, P₃ = 17-2, P₄ = 5-3 Average waiting time is 6.5
- P₁ = 17-0, P₂ = 5-1, P₃ = 26-2, P₄ = 10-3? Average turnaround is 13

Round-Robin


- Similar to FCFS scheduling except we add a preemption using a timer
- Each process gets a small unit of CPU time (time quantum), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- Pro – If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- Con – If q is very high, this is just FCFS with extra steps! If q is very low, you may spend more time context switching than actually executing tasks! Typically, $q \sim 10-100$ milliseconds. Context switches < 10 microseconds

Round-Robin

- Consider the following processes with time quantum 4

Process	Burst time (milliseconds)
P ₁	24
P ₂	3
P ₃	3


- What does the Gantt chart look like?




- Waiting time for P₁ = 10-4, P₂ = 4, P₃ = 7. Average is $(6 + 4 + 7)/3 = 5.66$
- Turnaround time for P₁ = 30, P₂ = 7, P₃ = 10. Average is $(30+7+10)/3 = 15.66$

Round-Robin

- A smaller quantum generally means more context switches which generally means longer turnaround time
- Ideally, ~80% of CPU bursts should be shorter than the quantum



quantum	context switches
12	0
6	1
1	9



process	time
P_1	6
P_2	3
P_3	1
P_4	7

Priority

- Assign a priority number (integer) to each process
- The CPU is allocated to the process with the highest priority (We will treat the smallest integer as highest priority, but this could be implemented the other way too)
- Generally, a preemptive algorithm but nonpreemptive could be possible
- FCFS is priority scheduling where all processes have the same priority
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time (The larger the CPU burst, the lower the priority, and vice versa)
- Pro – More control over CPU scheduling
- Con – Starvation is possible and needs to be accounted for *most resources are put to higher priority.*

Priority

- Consider the following processes with the following priorities

Process	Burst time (milliseconds)	Priority
P ₁	10	3
P ₂	1	1
P ₃	2	4
P ₄	1	5
P ₅	5	2

- What does the Gantt chart look like?




- Average waiting time is 8.2. Average turnaround time is 12

Priority

- Consider the following processes with the following priorities with round-robin with time quantum 2

Process	Burst time (milliseconds)	Priority
P ₁	4	3
P ₂	5	2
P ₃	8	2
P ₄	7	1
P ₅	3	3

- What does the Gantt chart look like?



if two processes have same priority, then it would run FCFS, then round-robin.

Priority

- A process that is ready to run but may be in the ready queue indefinitely
 - It may EVENTUALLY run (e.g. 2am on a Sunday)
 - The computer may reboot or crash and the process may never get run
- Implement aging – As time progresses, increase the priority of the process
avoid of starvation.



Lottery

- The scheduler distributes "lottery tickets" to all the processes
- The scheduler then randomly selects a lottery number
- The process assigned the lottery number is the next process to run
- E.g. 100 tickets for 3 processes: Process A gets 50%, Process B gets 15%, Process C gets 35%
- Pro – If each process gets tickets, no process will ever starve *get a chance to run*
- Con – If there are many tickets for many processes, this may be inefficient


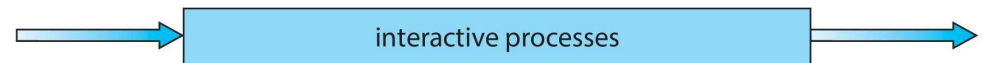

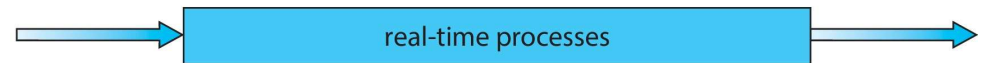
Multilevel Queue

- The ready queue consists of multiple queues. This is how ready queues are usually implemented
- Multilevel queue scheduler requires: A set of queues, a scheduling algorithms for each queue, a method used to determine which queue a process will enter when that process needs service, and scheduling among the queues
- Pro – The best of all worlds. Pick the most appropriate scheduling algorithm for the queue
- Con – Starvation is still possible. High priority processes may use up all the CPU time

Multilevel Queue



highest priority




lowest priority

Multilevel Feedback Queue

- To address the issue of starvation in a multilevel queue, provide an aging mechanism so processes can be promoted to higher priority queues
- To address the issue of processes using too much CPU time, provide a mechanism to demote them to a lower priority queue

Multilevel Feedback Queue

- Consider an example with three queues
 - Q_0 – RR with time quantum 8 milliseconds
 - Q_1 – RR time quantum 16 milliseconds
 - Q_2 – FCFS
- A new process enters queue Q_0 and receives 8 milliseconds
- If it does not finish in 8 milliseconds, the process is moved to queue Q_1
- At Q_1 job is again served in RR and receives 16 additional milliseconds
- If it still does not complete, it is preempted and moved to queue Q_2





Western
UNIVERSITY • CANADA