

A decorative graphic on the left side of the slide consisting of white and light blue lines and circles, resembling a circuit board or a network diagram.

# WEEK 10

TRANSACTION – CONCURRENCY CONTROL

# STUDENT OBJECTIVES

- Upon completion of this video, you should be able to:
  - Identify the 3 main problems with interleaving schedules
  - Distinguish between a serial and a non serial schedule
  - Determine if a schedule is serializable

# CONCURRENCY CONTROL

- We want to allow multiple transactions to get at all the data and resources at the same time.
- Objective is to ensure **serializability** of transactions in multi-user database environment (interleaving of transactions where the resulting state is the same as one the states that would occur if the transactions were done serially in some order).
- **Schedule:** A schedule  $S$  of  $n$  transactions  $T_1, T_2, \dots, T_N$  is an ordering of the operations of the transactions. Each operation of an individual transaction  $T_i$  in  $S$  must appear in the same order that it appears in  $T_i$ . The scheduler interleaves the execution of database operations to ensure serializability

# SAMPLE SCHEDULES:

Schedule 1:

| Time | Transaction | Step             | Stored Value |
|------|-------------|------------------|--------------|
| 1    | T1          | Read Bal         | 35           |
| 2    | T1          | Bal = $35 + 100$ |              |
| 3    | T1          | Write Bal        | 135          |
| 4    | T1          | Commit           |              |
| 5    | T2          | Read Bal         | 135          |
| 6    | T2          | Bal = $135 - 30$ |              |
| 7    | T2          | Write Bal        | 105          |
| 8    | T2          | Commit           |              |

Schedule 2:

| Time | Transaction | Step            | Stored Value |
|------|-------------|-----------------|--------------|
| 1    | T2          | Read Bal        | 35           |
| 2    | T2          | Bal = $35 - 30$ |              |
| 3    | T2          | Write Bal       | 5            |
| 4    | T2          | Commit          |              |
| 5    | T1          | Read Bal        | 5            |
| 6    | T1          | Bal = $5 + 100$ |              |
| 7    | T1          | Write Bal       | 105          |
| 8    | T1          | Commit          |              |

### Schedule 3:

| Time | Transaction | Step         | Stored Value |
|------|-------------|--------------|--------------|
| 1    | T1          | Read Bal     | 35           |
| 2    | T2          | Read Bal     | 35           |
| 3    | T1          | Bal = 35+100 |              |
| 4    | T2          | Bal = 35-30  |              |
| 5    | T1          | Write Bal    | 135          |
| 6    | T2          | Write Bal    | 5            |
| 7    | T1          | Commit       |              |
| 8    | T2          | Commit       |              |

**NOTICE THE PROBLEM  
IN SCHEDULE 3!**

# CONFLICT

- 2 operations in a schedule *conflict* if
  - 1) they belong to different transactions,
  - 2) they access the same data item X
  - 3) at least one of the transactions issues a WRITE(X).
- 3 Main Problems:
  - Lost Updates
  - Uncommitted Data
  - Inconsistent Retrievals

# EXAMPLE: LOST UPDATE

- T1: Homer puts in 100 dollars to the savings account ( $Bal = Bal + 100$ )
- T2: Marge takes out 30 dollars from the savings account ( $Bal = Bal - 30$ )

## Correct Schedule

| Time | Transaction | Step             | Stored Value |
|------|-------------|------------------|--------------|
| 1    | T1          | Read Bal         | 35           |
| 2    | T1          | $Bal = 35 + 100$ |              |
| 3    | T1          | Write Bal        | 135          |
| 4    | T1          | Commit           |              |
| 5    | T2          | Read Bal         | 135          |
| 6    | T2          | $Bal = 135 - 30$ |              |
| 7    | T2          | Write Bal        | 105          |
| 8    | T2          | Commit           |              |

## Incorrect Schedule:

| Time | Transaction | Step             | Stored Value      |
|------|-------------|------------------|-------------------|
| 1    | T1          | Read Bal         | 35                |
| 2    | T2          | Read Bal         | 35                |
| 3    | T1          | $Bal = 35 + 100$ |                   |
| 4    | T2          | $Bal = 35 - 30$  |                   |
| 5    | T1          | Write Bal        | 135 ← Lost Update |
| 6    | T2          | Write Bal        | 5                 |
| 7    | T1          | Commit           |                   |
| 8    | T2          | Commit           |                   |

# UNCOMMITTED DATA

Something has happened here, maybe the ATM rollers broke before the money was sucked in

When 2 transactions T1 and T2 are executed, T1 is rolled back. T2 reads uncommitted data (violates the Isolation property)  
Transaction starts reading before the rollback is completed

## Correct Schedule

| Time | Transaction | Step         | Stored Value |
|------|-------------|--------------|--------------|
| 1    | T1          | Read Bal     | 35           |
| 2    | T1          | Bal = 35+100 |              |
| 3    | T1          | Write Bal    | 135          |
| 4    | T1          | Rollback     | 35           |
| 5    | T2          | Read Bal     | 35           |
| 6    | T2          | Bal = 35-30  |              |
| 7    | T2          | Write Bal    | 5            |
| 8    | T2          | Commit       |              |

## Incorrect Schedule:

| Time | Transaction | Step         | Stored Value              |
|------|-------------|--------------|---------------------------|
| 1    | T1          | Read Bal     | 35                        |
| 2    | T1          | Bal = 35+100 |                           |
| 3    | T1          | Write Bal    | 135                       |
| 4    | T2          | Read Bal     | 135 ←<br>Uncommitted Data |
| 5    | T2          | Bal = 135-30 |                           |
| 6    | T1          | Rollback     | 35                        |
| 7    | T2          | Write Bal    | 105                       |
| 8    | T2          | Commit       |                           |



# INCONSISTENT RETRIEVALS

- Occurs when a transaction reads some data before they are changed and other data after they are changed.

- For example if T1 calculates the sum of all inventory:

***SELECT SUM(Quantity\_on\_hand) FROM inventory***

at the same time as T2 is update inventory for some items, the total at the end will be wrong.

# SCHEDULER

- Want to ensure that once a transaction  $T$  is committed, it should never be necessary to roll back  $T$ . (A recoverable schedule)
- Recoverable schedules may have cascading rollbacks that allows uncommitted transactions to be rolled back because it read an item from a transaction that aborted (try to avoid this because it is time consuming!)

- **Scheduler:** A scheduler sets the order that concurrent transactions are executed. It interweaves the operations to ensure serialization.
- A schedule can use a number of methods, we will look at 3: **locking**, **time stamping** and **optimistic**.
  - It **MUST** preserve the order of operations within the original transactions
  - It makes some system component schedules
  - It can't really make a planned schedule because it doesn't know what transactions the users are going to submit and in what order.

**QUESTION: If the scheduler just had to ensure serialization, what would it always do to ensure this?**

**ANSWER: Never ever let them mix, always finish one transaction before doing the next transaction**

# SERIALIZABILITY

- Suppose 2 users (e.g. airline clerks) submit the 2 transactions: T1 and T2 at approximately the same time: 3 resolutions:
  - Execute all the transaction of T1 in sequence followed by T2
  - Execute all the transactions of T2 in sequence followed by T1
  - Some interleaving of operations may be allowed

# EXAMPLE

- Say Account A has 1000 and Account B has 2000
  - T1: Transfer \$50 from account A to account B
  - T2: Transfer 10% of account A to account B
- NOTE: Even if we do the transactions as atomic units we get 2 different answers depending on the order we execute them!!
- **THAT IS OKAY**, either order: T1 then T2 OR T2 then T1 produces a serial schedule

- **Serializability Theory:** determines which schedules are "correct" and those that are not and tries to allow only correct schedules
- **Serial:** A schedule is serial if transactions are executed consecutively with out interleaving. Every serial schedule is considered correct
- **Non-serial:** A schedule is non-serial if it allows interleaving of transactions

**QUESTION: Why not make all schedules Serial?**

**ANSWER: Too slow, we MUST allow some interleaving**

- A schedule of transactions is serializable if it is equivalent to some serial schedule.

**QUESTION:** Is *result equivalence* (a non-serial schedule that produces the same final database state as a serial schedule) enough?

**ANSWER:** No! It might have just been by accident!

**Example:**

If X starts at 100, everything is fine with both schedules but not with other numbers!

Example 1

| Schedule 1   | Schedule 2   |
|--------------|--------------|
| Read X       | Read X       |
| $X = X + 10$ | $X = X - 6$  |
| Write x      | Write X      |
| Read X       | Read X       |
| $X = X - 5$  | $X = X + 11$ |
| Write X      | Write X      |

Example 2

| Schedule 1   | Schedule 2    |
|--------------|---------------|
| Read X       | Read X        |
| $X = X + 10$ | $X = X * 1.1$ |
| Write X      | Write X       |



# CONFLICT EQUIVALENCE

- **Conflict equivalence:** if the order of any 2 conflicting operations (operations from different transactions, using the same data and one of them is a write operation) is the same in both schedules
- **Testing for Conflict Serializability** (using a serialization graph)
  - for each transaction  $T_i$  participating in schedule  $S$  create a node labeled  $T_i$  in the graph
  - for each case in  $S$  where  $T_j$  executes a read ( $X$ ) after a write( $X$ ) executed by  $T_i$  create an edge ( $T_i \rightarrow T_j$ )
  - for each case in  $S$  where  $T_j$  executes write( $X$ ) after  $T_i$  executes a read( $X$ ) create an edge ( $T_i \rightarrow T_j$ )
  - for each case in  $S$  where  $T_j$  executes a write ( $X$ ) after  $T_i$  executes a write( $X$ ) create an edge ( $T_i \rightarrow T_j$ )
  - The schedule is serializable if and only if the precedence graph has no cycles



# EXAMPLE 1:

| Time | Transaction | Graph   |
|------|-------------|---|
| T1   | Read x      | <p>Question: Why don't the red Read and Write of y matter in this case?</p> |
| T1   | $x = x - n$ |   |
| T1   | Write x     |   |
| T1   | Read y      |   |
| T1   | $y = y + n$ |   |
| T1   | Write y     |   |
| T2   | Read x      |   |
| T2   | $x = x + M$ |   |
| T2   | Write x     |   |

## EXAMPLE 2:



| Time | Transaction | Graph   |
|------|-------------|---|
| T1   | Read x      | <pre>graph LR; T1((T1)) -- green --&gt; T2((T2)); T2 -- purple --&gt; T1;</pre> |
| T1   | $x = x - n$ |   |
| T2   | Read x      |   |
| T2   | $x = x + M$ |   |
| T1   | Write x     |   |
| T1   | Read y      |   |
| T2   | Write x     |   |
| T1   | $y = y + n$ |   |
| T1   | Write y     |   |

### EXAMPLE 4:

| Time | Transaction | Graph |
|------|-------------|-------|
| T2   | Read z      |       |
| T2   | Read y      |       |
| T2   | Write y     |       |
| T3   | Read y      |       |
| T3   | Read z      |       |
| T1   | Read x      |       |
| T1   | Write x     |       |
| T3   | Write y     |       |
| T3   | Write z     |       |
| T2   | Read x      |       |
| T1   | Read y      |       |
| T1   | Write y     |       |
| T2   | Write y     |       |

CS319

19

# EXAMPLE 4:

| Time | Transaction | Graph |
|------|-------------|-------|
| T2   | Read z      |       |
| T2   | Read y      |       |
| T2   | Write y     |       |
| T3   | Read y      |       |
| T3   | Read z      |       |
| T1   | Read x      |       |
| T1   | Write x     |       |
| T3   | Write y     |       |
| T3   | Write z     |       |
| T2   | Read x      |       |
| T1   | Read y      |       |
| T1   | Write y     |       |
| T2   | Write y     |       |

CS319

19

