

CS3350B Computer Organization

Chapter 3: CPU Control & Datapath

Part 1: Introduction to MIPS

Iqra Batool

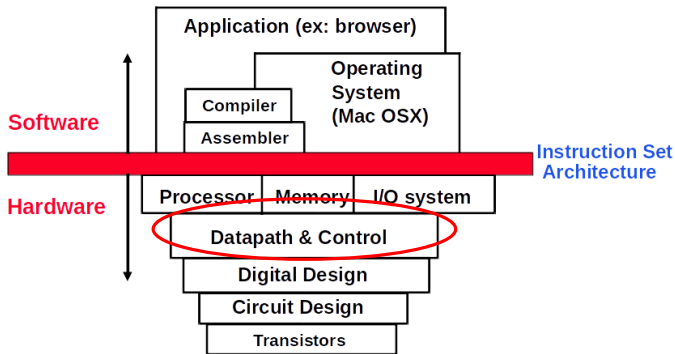
Department of Computer Science
University of Western Ontario, Canada

Monday February 12, 2024

Outline

- 1 Overview
- 2 MIPS Assembly
- 3 Instruction Fetch & Instruction Decode
- 4 MIPS Instruction Formats
- 5 Aside: Program Memory Space

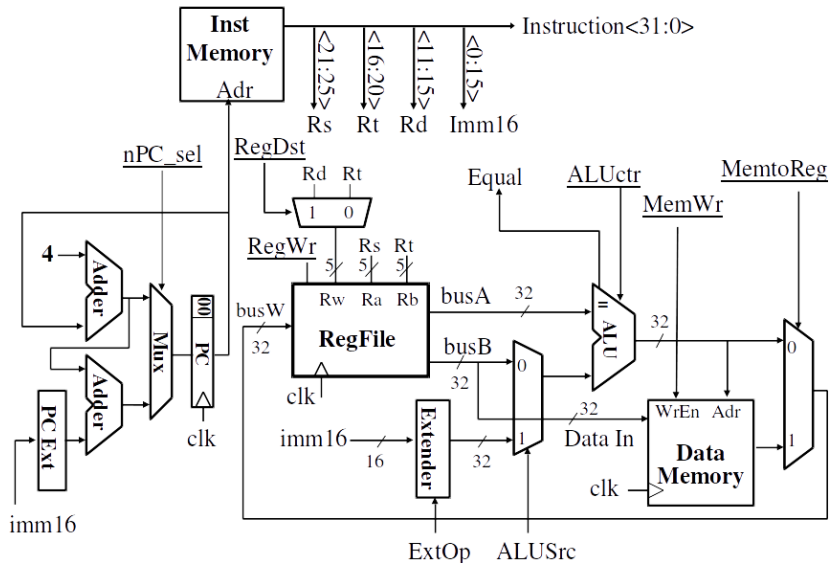
Layers of Abstraction



We will put together the combinational circuits and state circuits to see how a CPU actually works.

- How does data flow through the CPU?
- How are the path and circuits (MUX, ALU, registers) controlled?

Preview: MIPS Datapath



MIPS ISA

MIPS ISA: Microprocessor without Interlocked Pipelined Stages.

- ↳ ISA: The language of the computer.
- ↳ Microprocessor: a CPU.
- ↳ **Interlocking:** To come in chapter 4.
- ↳ **Pipelined Stages:** The data path is broken into a ubiquitous 5-stage pipeline (also chapter 4).

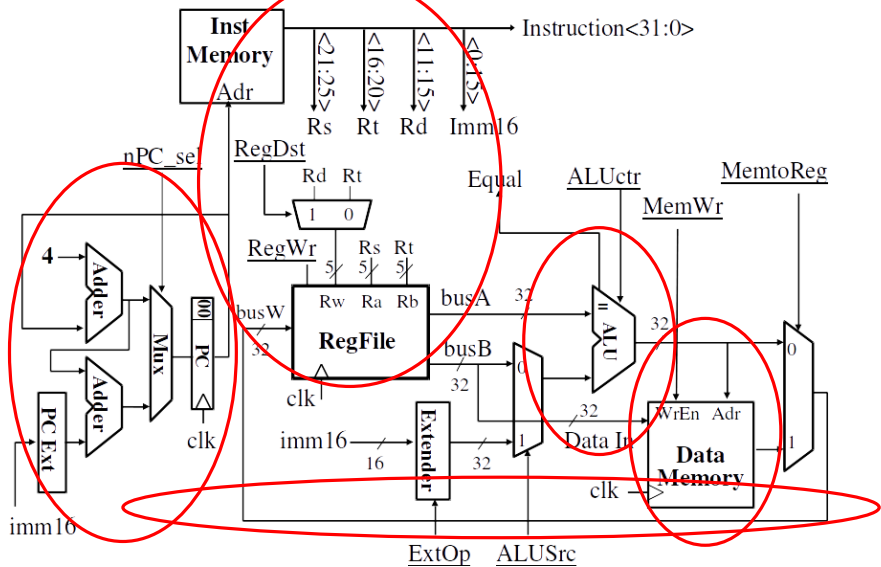
MIPS is a RISC ISA.

- **RISC:** Reduced Instruction Set Computer.
 - ↳ Provided instructions are simple, datapath is simple.
- Contrast with **CISC:** Complex Instruction Set Computer.
 - ↳ Instructions can be very “meta”, each performing many lower-level instructions.

The 5-Stages of the Datapath

- 1 **IF**: Instruction Fetch
- 2 **ID**: Instruction Decode
- 3 **EX/ALU**: Execute/Arithmetic
- 4 **MEM**: Access Memory
- 5 **WB**: Write-back result

MIPS Datapath, Spot The Stages



Coupling of ISA and Datapath

A datapath must be built to satisfy the requirements of the ISA.

- Instructions in the ISA determine what is needed internally.
- Circuitry limit possible instructions.
- Built from combinational blocks composed together.
- Very hard to decouple the two.

We begin by looking at the MIPS ISA before looking at the datapath components.

- ↳ We need a common language to discuss the datapath and give concrete examples.

Layers of an ISA

Start at high-level and work down.

- MIPS assembly
- MIPS instruction formats
- MIPS instruction binary
 - ↳ *Everything* on a computer is a number
 - ↳ Recall instruction memory cache (banked L1 cache)

MIPS assembly is a type of **RTL**: Register transfer language.

- Everything in MIPS is specified by registers and movement between them or between registers and memory.
- Most often, we abstract away the concept of caches here and assume CPU talks directly with memory.
 - ↳ In reality, the circuitry of the cache automatically abstracts away the memory hierarchy and handles cache hits/misses.

Outline

- 1 Overview
- 2 MIPS Assembly
- 3 Instruction Fetch & Instruction Decode
- 4 MIPS Instruction Formats
- 5 Aside: Program Memory Space

MIPS Assembly: The Basics

Registers:

- 32 general purpose 32-bit integer registers, denoted \$0–\$31.
- \$0 always holds the value 0.
- \$31 is reserved as the link register: stores the the point in instruction memory to return to after a function call.
- \$PC holds the **program counter** – address of current instruction
- \$HI & \$LO store results of multiplication/division

Memory:

- 32-bit words and 32-bit memory addresses.
- Byte-addressable memory.
- Indexed like a big array of bytes: Mem[0], Mem[1024], Mem[32768].

MIPS Assembly: RTL Examples

3-Operand Arithmetic:

- `add $8, $9, $10` \equiv $R[8] \leftarrow R[9] + R[10];$
- `sub $8, $9, $10` \equiv $R[8] \leftarrow R[9] - R[10];$

2-Operand Arithmetic (Immediate Arithmetic):

- `addi $8, $9, 127` \equiv $R[8] \leftarrow R[9] + 127;$
- `addi $8, $9, 913` \equiv $R[8] \leftarrow R[9] + 913;$
- `addi $8, $9, -6` \equiv $R[8] \leftarrow R[9] - 6;$

Data Transfer (Memory Accesses):

- `lw $13, 32($10)` \equiv $R[13] \leftarrow \text{Mem}[R[10] + 32];$
- `sw $13, 8($10)` \equiv $\text{Mem}[R[10] + 8] \leftarrow R[13];$

MIPS Assembly: 3 Operand Arithmetic

$\text{op } \$rd, \$rs, \$rt \quad \equiv \quad \$rd = \$rs \text{ op } \rt

- $\$rd$ is the destination register.
- $\$rs$ is the (first) source register.
- $\$rt$ is the second source register.
- op is some arithmetic operation:
 - ↳ add, addu, sub, subu, and, or, xor, ...

These instructions *assume* an interpretation of the bits stored in the register.

- Programmer/compiler must choose proper instruction for data.
- add vs addu

MIPS Assembly: 2 Operand Arithmetic

$\text{op } \$rt, \$rs, \text{imm} \equiv \$rt = \$rs \text{ op } \text{imm}$

- $\$rt$ is the destination register.
- $\$rs$ is the source register.
- imm is an **immediate**—a number whose value is hard-coded into the instruction.
 - ↳ C Example: `int i = j + 12;`
- op is some arithmetic operations:
 - ↳ `addi, addiu, subiu, andi, ori, xori, ...`
 - ↳ `sll, srl` (logical); `sla, sra` (arithmetic) (shifts are really a special case of R-type instr.)

These instructions *assume* an interpretation of the bits stored in the register.

- Programmer/compiler must choose proper instruction for data.
- Signed vs unsigned arithmetic. Logical vs arithmetic shifts.


MIPS Assembly: Data Transfer

`lw $rt, offset($rs)` \equiv $\$rt = \text{Mem}[\$rs + \text{offset}]$
`sw $rt, offset($rs)` \equiv $\text{Mem}[\$rs + \text{offset}] = \rt

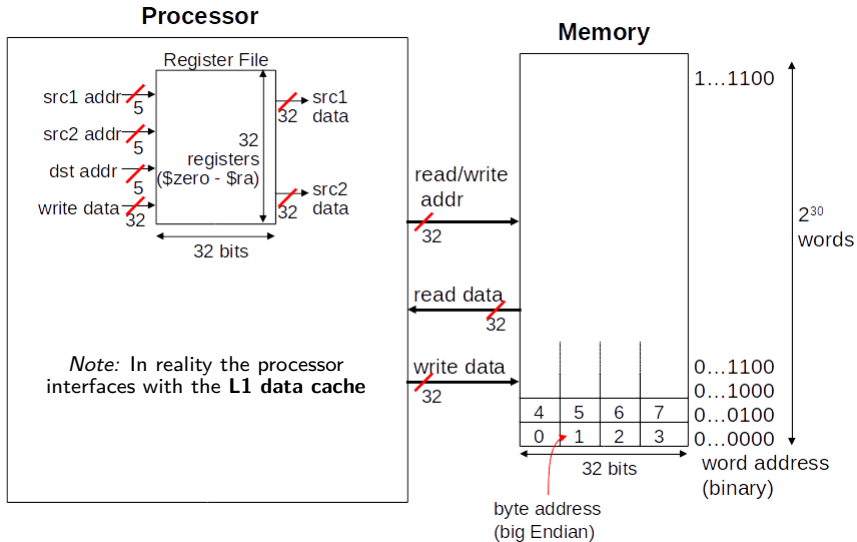
- `$rt` is the “value” register.
- `$rs` is the “address” register.
- `offset` is an **immediate**.

It is also possible to load and store bytes, **halfwords**:

- `lb`, `sb` (byte);
- `lwr`, `swr` (least-significant halfword);
- `lwl`, `swl` (most-significant halfword).



MIPS: A View of Memory



Aside: Endianness Defined

Endianness: The ordering of multiple bytes which are intended to be interpreted together as a single number.

- Important in memory layout, digital signals, networks, etc.

Consider the number: 0xAABBCCDD

Little-Endian: The least-significant byte is stored/sent first.

- Ordering: 0xDD, 0xCC, 0xBB, 0xAA

Big-Endian: The most-significant byte is stored/sent first.

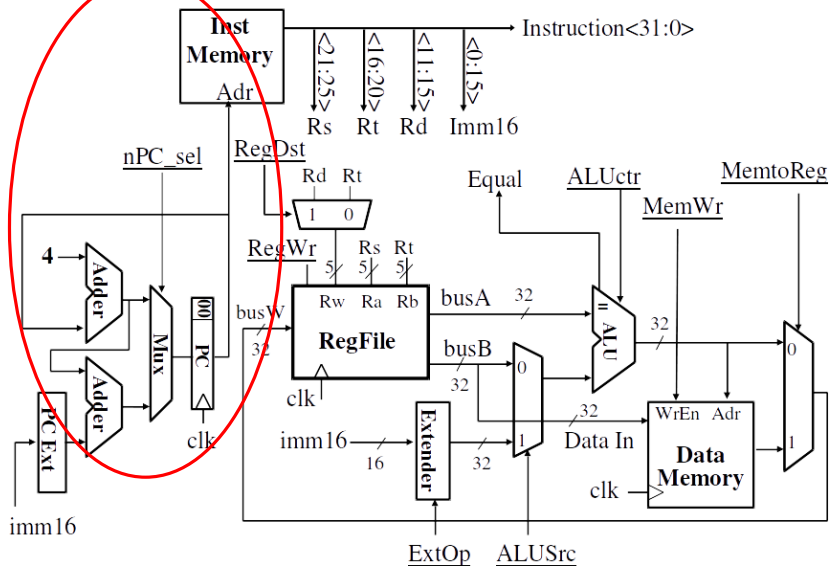
- Ordering: 0xAA, 0xBB, 0xCC, 0xDD
- MIPS is big-endian.

Big-endian conceptually easier but little-endian has performance benefits. In reality, hardware handles all conversions to and from, so we *rarely* care.

Outline

- 1 Overview
- 2 MIPS Assembly
- 3 Instruction Fetch & Instruction Decode
- 4 MIPS Instruction Formats
- 5 Aside: Program Memory Space

MIPS Datapath, Instruction Fetch



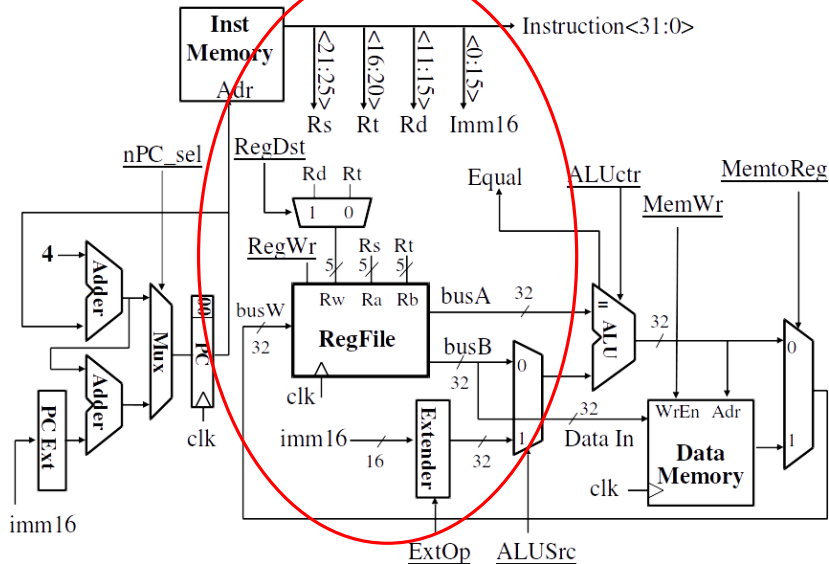
Instruction Fetch

IF — Instruction Fetch

- The simplest part of the datapath.
- One job: fetch the next instruction to execute from memory.
- *Banked L1 cache* \implies separate cache just for instructions

1. On the clock's rising edge, update the value of **PC**—program counter.
 - ↳ Essentially the index into instruction memory.
2. *Fetch* the instruction from memory and pass it to next stage: instruction decode.
3. Prepare for next instruction: calculate $PC + 4$ (4 bytes since instructions are word-aligned).

MIPS Datapath, Instruction Decode



Instruction Decode

ID — Instruction Decode

- Break the *binary value* of an instruction into its parts and decide what to do.
 - ↳ Recall: instructions eventually get compiled down to *bytecode* (i.e. binary).
 - Get the values ready for arithmetic: registers, immediates.
1. Break the instruction into individual bit segments.
 2. Access operand values from registers.
 3. Extend immediate to 32 bits (if using).

But how to break the instruction?

Aside: MIPS Special Register Names

- \$zero: the zero-valued register (\$0)
- \$at: reserved for compiler (\$1)
- \$v0, \$v1: result values (\$2, \$3)
- \$a0 - \$a3: arguments (\$4-\$7)
- \$t0 - \$t9: temporaries (\$8-\$15, \$24, \$25)
 - ↳ Can be overwritten by callee
- \$s0 - \$s7: saved (\$16-\$23)
 - ↳ Must be saved/restored by callee
- \$gp: global pointer for static data (\$28)
- \$sp: stack pointer (\$29)
- \$fp: frame pointer (\$30)
- \$ra: return address (\$31)

Outline

- 1 Overview
- 2 MIPS Assembly
- 3 Instruction Fetch & Instruction Decode
- 4 MIPS Instruction Formats**
- 5 Aside: Program Memory Space

MIPS Instruction Formats

Every instruction in MIPS is 32-bits.

- A memory word is 32 bits, after all.

All instructions belong to 3 pre-defined formats:

- **R-Type**: “Register”
- **I-Type**: “Immediate”
- **J-Type**: “Jump”

Each format defines how those 32 bits of instruction data are broken up into individual “bit-fields” and how they are interpreted during ID stage.

- The first 6 bits always encode the **opcode**.
- The opcode determines the type of instruction and format of the remaining bits.

R-Type Instructions

R-Type instructions usually have 3 registers as its operands.

- ↳ “Register type”.
- ↳ General arithmetic operations.

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

- op — the opcode.
- rs — first source register.
- rt — second source register.
- rd — destination register.
- shamt — shift amount; used for shift instructions, 0 otherwise.
- funct — the arithmetic *function* the ALU should perform.

R-Type Examples 1

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

add \$t0, \$s1, \$s2

op	\$s1	\$s2	\$t0	shamt	add
0	17	18	8	0	32
000000	10001	10010	01000	00000	100000

sub \$t0, \$s1, \$s2

op	\$s1	\$s2	\$t0	shamt	sub
0	17	18	8	0	34
000000	10001	10010	01000	00000	100010

- For R-Type instructions the opcode and funct *together* determine the operations to perform.

R-Type Examples 2

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

`sll $s0, $t0, 4`

op	rs	\$t0	\$s0	4	shift left
0	0	8	16	4	0
000000	00000	01000	10000	00100	000000

Note: shift instruction have two registers and an immediate, but are *not* immediate instructions.

- Here, the allowed value of the shift amount is only 5 bits, not 16 bits as in an immediate-type instruction.

R-Type Examples 3: Bit-Wise Logical Operations

- Useful to mask (remove) bits in a word.

↳ Select some bits, clear others to 0.

and \$t0, \$t1, \$t2

\$t2	0000 0000 0000 0000 0000 1101 1100 0000
\$t1	0000 0000 0000 0000 0011 1100 0000 0000
\$t0	0000 0000 0000 0000 0000 1100 0000 0000

- Useful to include bits in a word.

↳ Set some bits to 1, leave others unchanged.

or \$t0, \$t1, \$t2


\$t2	0000 0000 0000 0000 0000 1101 1100 0000
\$t1	0000 0000 0000 0000 0011 1100 0000 0000
\$t0	0000 0000 0000 0000 0011 1101 1100 0000

I-Type Instructions

I-Type instructions always have 2 registers and an **immediate**.

↳ “Immediate type”.

↳ Immediate arithmetic, data transfer, branch.



- op — the opcode.
- rs — first source register.
- rt — second source (or destination) register.
- imm — the immediate/constant.

I-Type Examples 1

op	rs	rt	immediate
6 bits	5 bits	5 bits	16 bits

`addi $t1, $t0, 10`

op	rs	rt	immediate
8	\$t0	\$t1	10
001000	01000	01001	00000000000001010

`addiu $t1, $t0, 10`

op	rs	rt	immediate
9	\$t0	\$t1	10
001001	01000	01001	00000000000001010

Note: unsigned instructions will *not* signal exception on overflow.

I-Type Examples 2

op	rs	rt	immediate
6 bits	5 bits	5 bits	16 bits

lw \$t1, 12(\$t0)

op	rs	rt	immediate
35	\$t0	\$t1	12
100011	01000	01001	00000000000001100

sw \$t1, 32(\$t0)

op	rs	rt	immediate
43	\$t0	\$t1	32
101011	01000	01001	0000000000100000

I-Type Examples 3

op	rs	rt	immediate
6 bits	5 bits	5 bits	16 bits


```
bne $t0, $t1, 24          if ($t0 != $t1)
                           PC = PC + 4 + (24 << 2);
                           else
                           PC = PC + 4;
```

op	rs	rt	immediate
5	\$t0	\$t1	24
000101	01000	01001	00000000000011000

J-Type Instructions

J-Type instructions have just one big immediate, called a **target**.

- ↳ “Jump type”.
- ↳ Only two instructions: j (jump) and jal (jump and link).



- op — the opcode.
- target — the target memory address to jump to.

Note: target is always multiplied by 4 before being applied to program counter...


J-Type Instructions and Pseudo-Direct Addressing

Pseudo-Direct Addressing: Almost a direct addressing of instruction memory.

- Compiler usually handles the calculation of the exact jump target.


Next value of PC is $\text{target} \times 4$ combined with upper 4 bits of current PC.

$$\text{nPC} = (\text{PC} \& 0xf0000000) \mid (\text{target} \ll 2);$$




Addressing Instruction Memory in MIPS

■ **Pseudo-Direct:** J-Type instructions



■ **PC-Relative:** Branch instructions



Addressing Operands in MIPS


- **Immediate Addressing**, I-Type instruction.



- **Register Addressing**, Almost all instructions.



- **Base Addressing**, Data transfer instructions.



MIPS ISA: Some Important Instructions

Category	Instruction		OP/ func	Example	Meaning
Logic & Arith.	add	R	0/32	add \$s1, \$s2, \$s3	$\$s1 = \$s2 + \$s3$
	subtract	R	0/34	sub \$s1, \$s2, \$s3	$\$s1 = \$s2 - \$s3$
	add immediate	I	8	addi \$s1, \$s2, 6	$\$s1 = \$s2 + 6$
	and/or	R	0/(36/37)	(and/or) \$s1, \$s2, \$s3	$\$s1 = \$s2 (\wedge/\vee) \$s3$
	(and/or) immediate	I	12/13	(andi/ori) \$s1, \$s2, 6	$\$s1 = \$s2 (\wedge/\vee) 6$
	shift right logical	R	0/2	srl \$rd, \$rt, 4	$\$rd = \$rt \gg 4$
	shift right arithmetic	R	0/3	sra \$rd, \$rt, 4	$\$rd = \$rt \ggg 4$
Data Transfer	load word	I	35	lw \$s1, 24(\$s2)	$\$s1 = \text{Memory}(\$s2+24)$
	store word	I	43	sw \$s1, 24(\$s2)	$\text{Memory}(\$s2+24) = \$s1$
	load byte	I	32	lb \$s1, 25(\$s2)	$\$s1 = \text{Memory}(\$s2+25)$
	store byte	I	40	sb \$s1, 25(\$s2)	$\text{Memory}(\$s2+25) = \$s1$
Cond. Branch	br on equal	I	4	beq \$s1, \$s2, L	if $(\$s1 == \$s2)$ go to L
	br on not equal	I	5	bne \$s1, \$s2, L	if $(\$s1 \neq \$s2)$ go to L
	set less than	R	0/42	slt \$s1, \$s2, \$s3	if $(\$s2 < \$s3)$ $\$s1=1$ else $\$s1=0$
	set less than immediate	I	10	slti \$s1, \$s2, 6	if $(\$s2 < 6)$ $\$s1=1$ else $\$s1=0$
Uncond. Jump	jump	J	2	j 250	go to 1000
	jump register	R	0/8	jr \$t1	go to \$t1
	jump and link	J	3	jal 250	go to 1000; $\$ra=PC+4$

Note: knowing the binary values of each bit-field is not necessary, but understanding the semantic meaning of each instruction *is* important.

Full Method Example: C to MIPS

```
void swap(int v[], int k) {  
    int temp;  
    temp = v[k];  
    v[k] = v[k+1];  
    v[k+1] = temp;  
}
```

```
swap:  sll $t1, $a1, 2      # $t1 = k * 4  
      add $t1, $a0, $t1    # $t1 = v+(k*4)  
                                # (address of v[k])  
      lw $t0, 0($t1)       # $t0 (temp) = v[k]  
      lw $t2, 4($t1)       # $t2 = v[k+1]  
      sw $t2, 0($t1)       # v[k] = $t2 (v[k+1])  
      sw $t0, 4($t1)       # v[k+1] = $t0 (temp)  
      jr $ra              # return to calling routine
```

Note: words and int-type are both 32-bits here.

Outline

- 1 Overview
- 2 MIPS Assembly
- 3 Instruction Fetch & Instruction Decode
- 4 MIPS Instruction Formats
- 5 Aside: Program Memory Space

Revisiting Program Basics

- **Frame:** The encapsulation of one method call; arguments, local variables.
 - ↳ “Enclosing subroutine context”.
- (Call) **Stack** (of frames): The stack of method invocations.
 - ↳ Base of stack is the main method, each method call adds a frame to the stack.
- **Heap:** globally allocated data that lives beyond the scope of the frame in which it was allocated.
- **Static Data:** Global data which is stored in a *static* memory address throughout life of program.

MIPS Special Registers

- \$v0, \$v1: result values (\$2, \$3)
- \$a0 - \$a3: arguments (\$4-\$7)
- \$t0 - \$t9: temporaries (\$8-\$15, \$24, \$25)
 - ↳ Can be overwritten by callee
- \$s0 - \$s7: saved (\$16-\$23)
 - ↳ Must be saved/restored by callee
- \$gp: global pointer for static data (\$28)
- \$sp: stack pointer (\$29)
- \$fp: frame pointer (\$30)
 - ↳ The stack pointer *before* the current frame's invocation.
- \$ra: return address (\$31)

Memory Layout in MIPS (and most languages)

Text: program code

Static data: global variables

- ↳ static/global variables, constant arrays, etc.
- ↳ `$gp` initialized to address allowing \pm offsets into this segment

Dynamic data: heap


- ↳ e.g., `malloc` in C, `new` in Java

Stack: “automatic” storage

`$sp` → 7fff fffc_{hex}

`$gp` → 1000 8000_{hex}
1000 0000_{hex}

`pc` → 0040 0000_{hex}
0



Note: In this diagram the higher memory addresses are at top.

Handling the Stack in MIPS

```
sort:  addi $sp, $sp, -20    # make room on stack for 5 registers
      sw $ra, 16($sp)      # save $ra on stack
      sw $s3, 12($sp)      # save $s3 on stack
      sw $s2, 8($sp)       # save $s2 on stack
      sw $s1, 4($sp)       # save $s1 on stack
      sw $s0, 0($sp)       # save $s0 on stack


---


      ...                  # procedure body
      ...                  # call swap a bunch to do bubble sort


---


exit1: lw $s0, 0($sp)       # restore $s0 from stack
      lw $s1, 4($sp)       # restore $s1 from stack
      lw $s2, 8($sp)       # restore $s2 from stack
      lw $s3, 12($sp)      # restore $s3 from stack
      lw $ra, 16($sp)      # restore $ra from stack
      addi $sp, $sp, 20    # restore stack pointer
      jr $ra               # return to calling routine
```