

CS3350B Computer Organization

Chapter 4: Instruction-Level Parallelism

Part 2: Pipeline Hazards

Iqra Batool

Department of Computer Science
University of Western Ontario, Canada

Wednesday March 4, 2024

Outline

1 Overview

2 Structural Hazards

3 Data Hazards

4 Control Hazards

Pros and Cons of Pipelining

- Pipelining overlaps the execution of instructions to keep each stage of the datapath busy at all times.
 - ↳ Improves **throughput** but not **latency**.
 - ↳ Might actually *increase* latency.
- Can increase clock frequency using multi-cycle datapath.
- Ideal speedup can be up to the number of stages.
- Ideal speed up *never* reached.
 - ↳ **Fill time** and **drain time** limits speedup.
 - ↳ Must account for dependencies between results of previous instructions and operands of future instructions.
 - ↳ Sometimes the same hardware is needed simultaneously by different pipeline stages and different instructions (e.g. ID and WB stages).

Categorizing Pipeline Hazards

Structural Hazards

- Conflicts in hardware/circuit use.
- Different stages or different instructions attempt to use same piece of hardware at the same time.

Data Hazards

- Dependencies between the result of an instruction and the input to another instruction.
- Data being used before it is finished being computed or written to memory/registers.

Control Hazards

- Ambiguity in the control flow of the program being executed.
- **Branch instructions**—if/else, loops.
- Take the branch? Don't take the branch? Which instruction follows a branch instruction in the pipeline?

“Resolving” Pipeline Hazards

Not an easy task. Simplest solution: just wait or **stall**.

↳ Any hazard can always be solved by just waiting.

But:

- Ruins potential speedup.
 - ↳ Might end up being slower than a single-cycle datapath.
 - ↳ Since latency can increase in pipelining, with enough stalls becomes slower.
- Increases CPI.
- Works against entire principle of pipelining.
 - ↳ Where's the performance?
- Nonetheless, sometimes it really is the only solution.

Outline

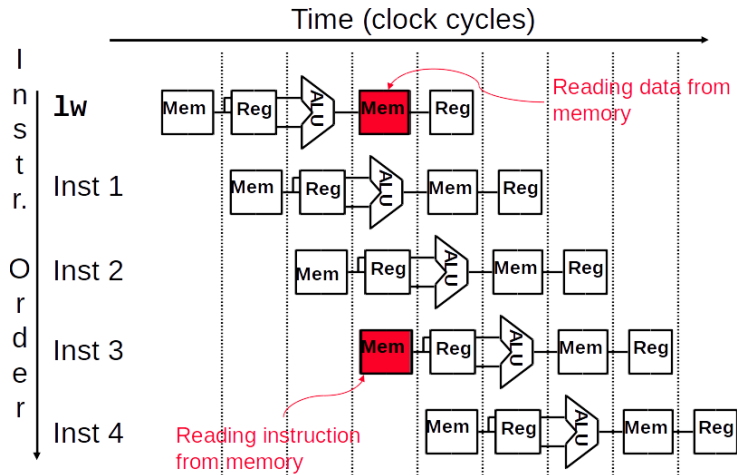
- 1 Overview
- 2 Structural Hazards
- 3 Data Hazards
- 4 Control Hazards

Structural Hazards: Causes and Resolutions

- Structural hazards are caused by two instructions needing to use the same hardware at the same time.
- Easiest to resolve? Just add in redundant hardware.
 - ↳ Works for combinational circuits.
 - ↳ Redundant memory would cause problems in needing to keep both consistent.
- Real structural hazards thus lie in state circuits: registers and memory.
 - ↳ IF stage and MEM stage.
 - ↳ ID stage and WB stage.

Structural Hazards In Memory (1/2)

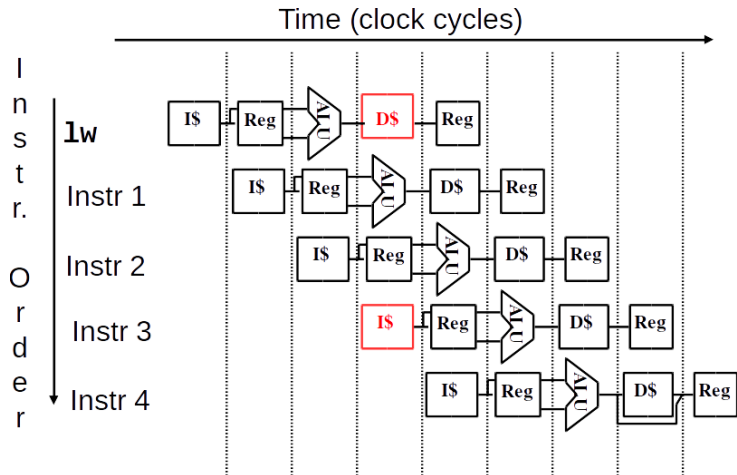
Consider a *unified* L1 cache. Reading instructions and reading/writing data could overlap for pipelined instructions.



Structural Hazards In Memory (2/2)

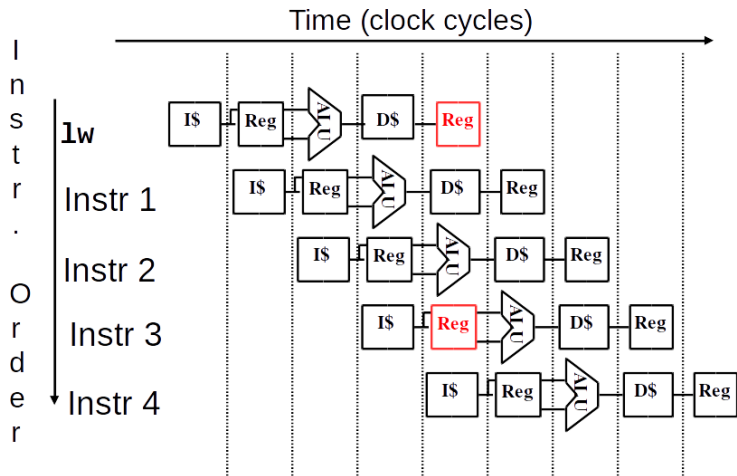
Simple fix: separate instruction memory from data memory.

- Can use a banked cache.



Structural Hazards In Register File (1/2)

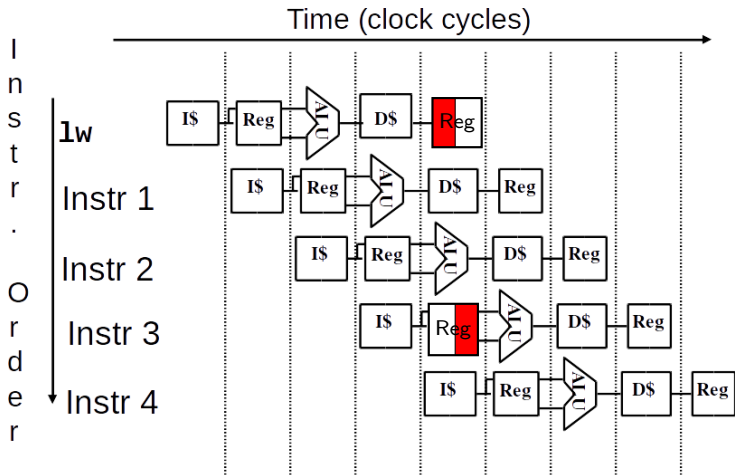
ID stage must read from registers while WB stage must write to registers.



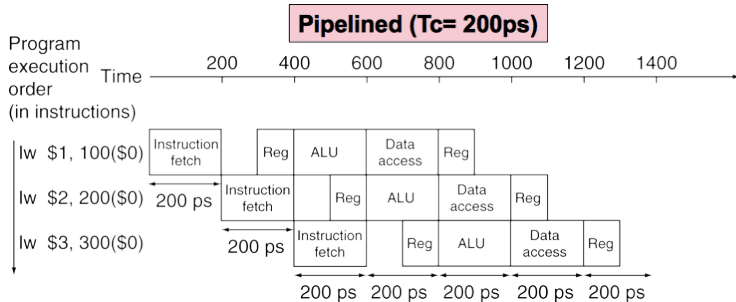
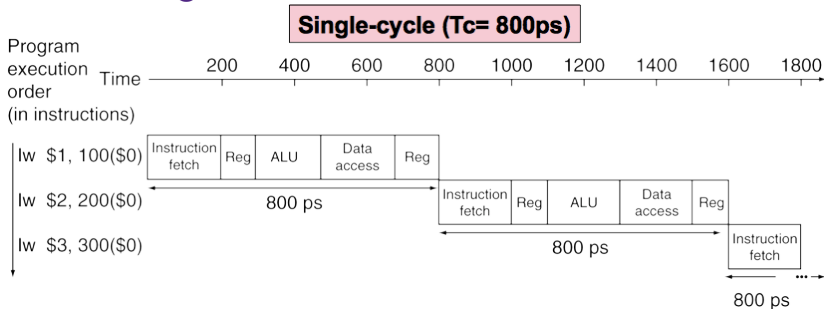
Structural Hazards In Register File (2/2)

In reality, reading from register file is very fast; clock cycle is long enough to allow both ID and WB to occur within a single clock cycle.

- Needs independent read and write ports.



A previous diagram



Outline

- 1 Overview
- 2 Structural Hazards
- 3 Data Hazards
- 4 Control Hazards

Data Hazards: Causes

Data hazards are caused by **dependencies** between instruction.

- **Read After Write** (RAW) dependency. Using the result of one instruction as an operand to a later instruction.
 - **Read After Read** (RAR) dependency. Using the same operand to multiple instructions.
 - **Write After Read** (WAR) dependency. Writing a result to a register after using that register as an operand in some previous instructions.
 - **Write After Write** (WAW) dependency. Writing to the same register as a previous instruction writes to.
-
- RAW is the only *true* dependency.
 - RAR is never a hazard.
 - WAR, WAW only becomes a problem with **out-of-order execution**. (Lecture 14).

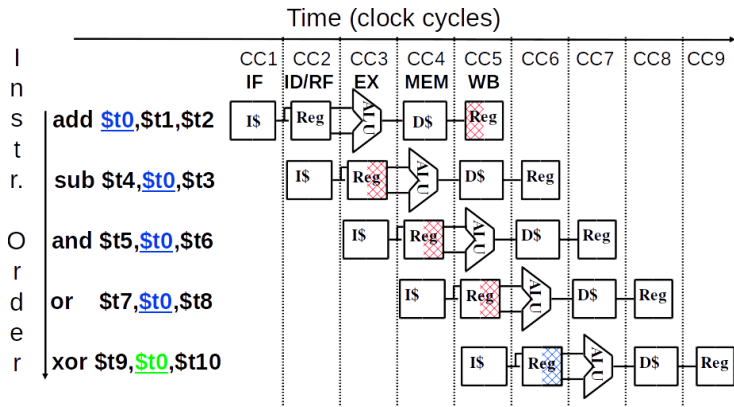
Data Hazards: Resolutions

- **Data hazards** are caused dependencies which *require resolution*.
 - ↳ Not all dependencies create hazards. If the dependent instructions are "far enough away", then there is no hazard.
- Hazards can always be solved by *stalling* the pipeline.
- Hazards can also be solved by special **forwarding** (also called bypass).
- Data hazards are the most common type of hazard.
 - ↳ It's the logical way to write programs: *locality*.
 - ↳ *Make the common case fast*.
 - ↳ Since locality dictates register re-use, datapath should be handle it efficiently. Typically, through forwarding.

Data Hazard Example 1 (1/3)

add produces a result which is then read by sub, and, or, xor.

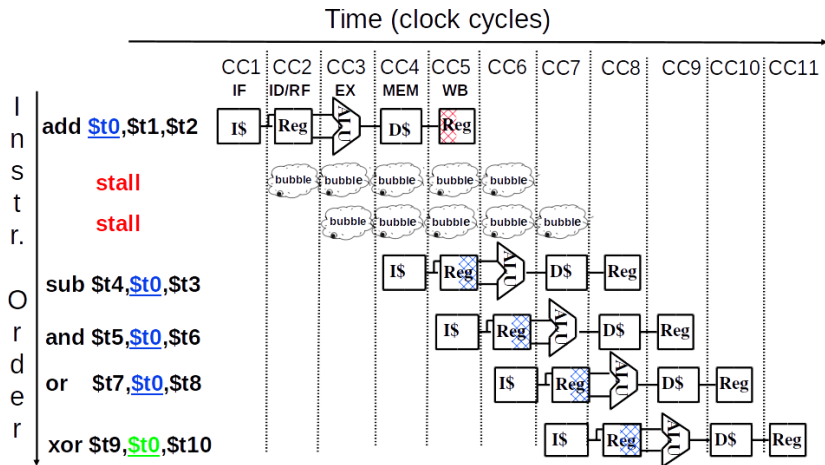
- Read After Write hazard.
- xor is far enough in the future to be okay.
- sub, and need more work.
- or actually already resolved via structural hazard solution



Data Hazard Example 1 (2/3)

Possible (but not great) solution: **stall** the execution.

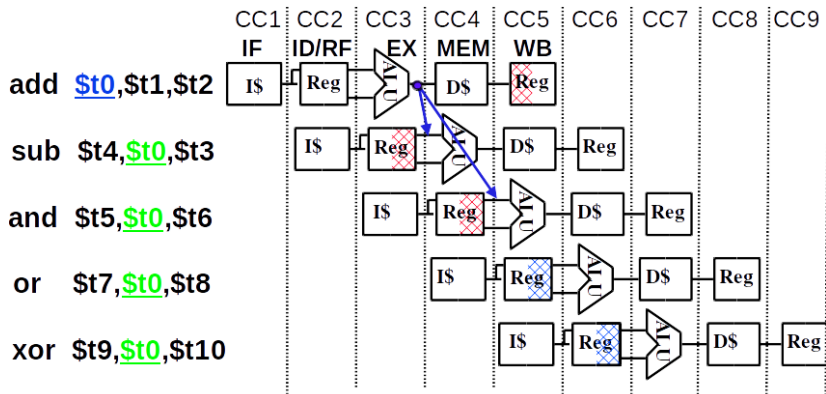
- sub structural hazard already solved.



Data Hazard Example 1 (3/3)

Another possible solution: **forwarding**.



- Implemented via inter-stage registers.
- No more stalls!
- ALU forwarding for add to sub and add to and.
- or structural hazard already solved.



More ALU-ALU Forwarding

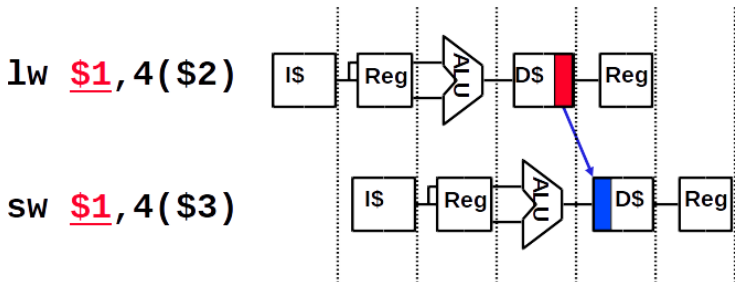
Two kinds of forwarding *into* the ALU:

- From instr. in MEM stage (from ALU/MEM reg): **ALU-ALU** forwd.
- From instr. in WB stage (from MEM/WB ref): **MEM-ALU** forwd.
- Which forwarded value to choose? \implies More control, more MUX.



MEM-MEM Forwarding

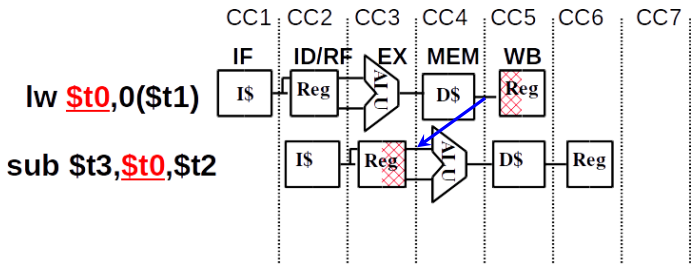
- For efficient memory copies (a common operation) this optimization results in no stalls.
 - ↳ Otherwise, two stalls required.
 - ↳ Eight great ideas in computer arch.: make the common case fast.
- MEM-MEM also used for forwarding an ALU result into a sw
 - ↳ add \$t0, \$t1, \$t4 and sw \$t0 0(\$t2).
 - ↳ \$t0 is MEM-MEM forwarded.



Load-Use Data Hazard

Load-use data hazard, a special kind of RAW hazard.

- Forwarding does not help here, still going backwards in time.
- A **stall** is required.



Implementing a Stall: Pipeline Interlock

Pipeline Interlock—**hardware** detects hazard and stalls the pipeline.

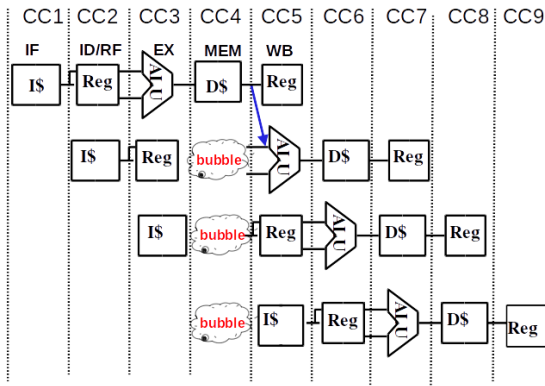
- Quite literally locks the flow of data between stages (locking writes to inter-stage registers).
- Essentially inserts an air bubble into pipeline.

lw \$t0, 0(\$t1)

sub \$t3,\$t0,\$t2

and \$t5,\$t0,\$t4

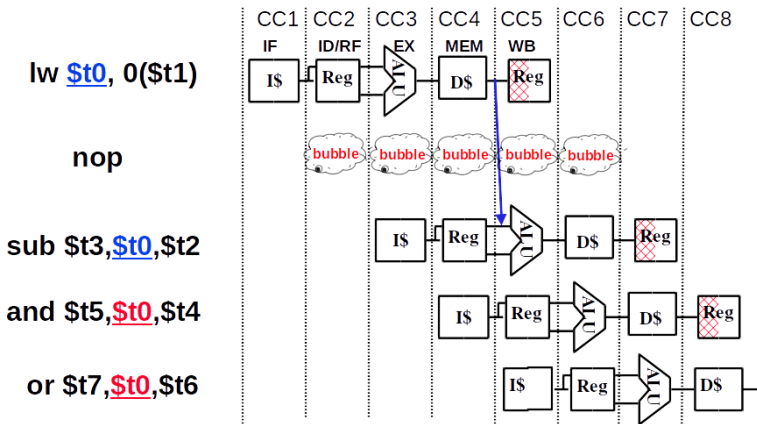
or \$t7,\$t0,\$t6



Implementing a Stall: NOP

NOP—a “no operation” special instruction inserted into instruction flow by **compiler** (or assembler, or assembly programmer).

- Hazards are detected and fixed at compile-time.
- Can be combined with forwarding; MEM-ALU in this case.



Pipeline Interlock vs NOP

- Interlocking requires special circuitry to dynamically detect hazards and stall the datapath.
- `nop` requires extra effort at compile time to detect and resolve hazards.
- Inserted `nop` instructions bloat instruction memory.
- More work at compile time for `nop` insertion but simpler (= faster?) datapath and controller.
- MIPS: Microprocessor *without Interlocked* Pipelined Stages

Data Hazards and Code Structure

Some data hazards are “fake”.

- Only caused by the order of instructions and not a true dependency.
- Re-order code (if possible) so an independent instruction performed instead of a nop.
 - ↳ Where the nop would be inserted is called the **load delay slot**.
 - ↳ Load delay slot can be filled with a nop or an independent instruction.
- Need at least one instruction between `lw` and using the loaded word.

	lw	\$t1, 0(\$t0)		lw	\$t1, 0(\$t0)
	lw	\$t2 , 4(\$t0)		lw	\$t2 , 4(\$t0)
stall	add	\$t3, \$t1, \$t2		lw	\$t4 , 8(\$t0)
	sw	\$t3, 12(\$t0)		add	\$t3, \$t1, \$t2
	lw	\$t4 , 8(\$t0)		sw	\$t3, 12(\$t0)
stall	add	\$t5, \$t1, \$t4		add	\$t5, \$t1, \$t4
	sw	\$t5, 16(\$t0)		sw	\$t5, 16(\$t0)
	13 cycles			11 cycles	

Outline

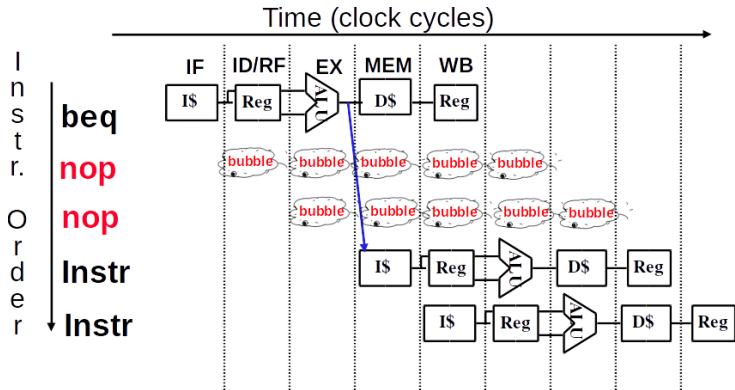
- 1 Overview
- 2 Structural Hazards
- 3 Data Hazards
- 4 Control Hazards**

Control Hazards: Causes and Resolutions

- Control hazards are caused by instructions which change the flow of control.
 - ↳ Branching.
 - ↳ If statements, loops.
- Sometimes called branch hazards.
- Since branch condition (beq, bne) not determined until after EX stage, cannot be *certain* about next instruction to fetch.

Control Hazard Resolution: Wait

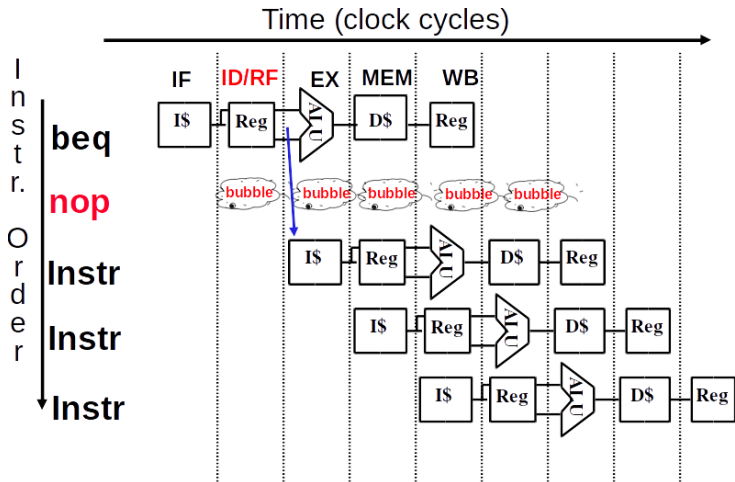
The simplest resolution is to just wait until branch condition is calculated before fetching next instruction.



Control Hazard Resolution: Add a Branch Comparator

Add a special circuit used to calculate branch conditions.

- Now only one stall needed instead of two.
- Similar to load-use hazard we now have a **branch delay slot**.



Delayed Branching


- The **branch delay slot** is the instruction immediately following a branch. Can be a nop or a useful instruction.
- In **delayed branching** the instruction in the branch delay slot is **always** executed whether or not the branch condition holds.
 - ↳ Used in conjunction with a special branch comparator.
 - ↳ Filling the branch delay slot (and other code re-organization) is usually handled by compiler/assembler.
 - ↳ Cannot fill slot with an instruction that influences branch condition.
- Jump instructions also have a delay slot.

<code>addi \$v0, \$0, 1</code>	<code>add \$t0, \$s0, \$s1</code>
<code>add \$t0, \$s0, \$s1</code>	<code>add \$t1, \$s2, \$s3</code>
<code>add \$t1, \$s2, \$s3</code>	<code>beq \$t0, \$t1, L</code>
<code>beq \$t0, \$t1, L</code>	<code>addi \$v0, \$0, 1</code>
<code>⋮</code>	<code>⋮</code>
<code>L: ...</code>	<code>L: ...</code>
	<code># addi executed regardless</code>

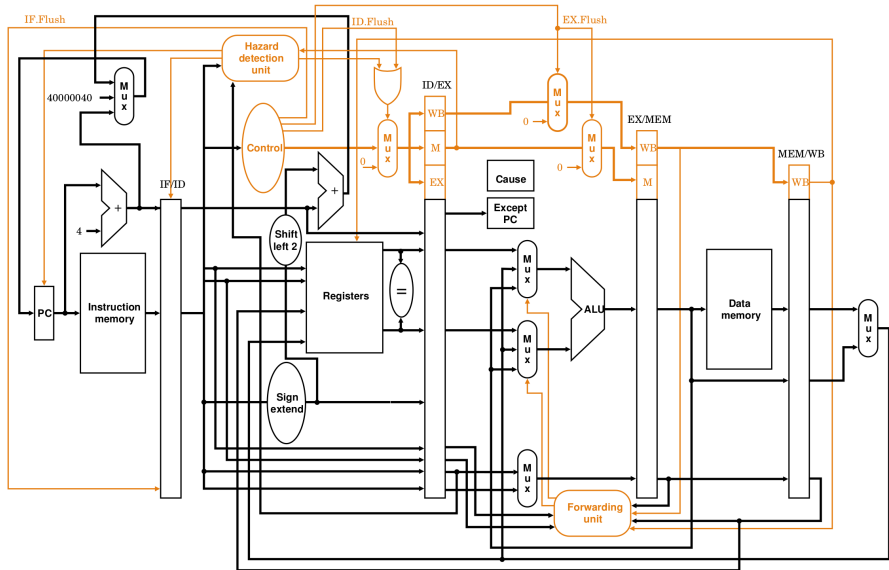
Implementing Branch Prediction

- Branches have exactly two possibilities: taken or not taken.
- In MIPS branches are *statically predicted* to never happen.
- **Dynamic branch prediction** uses run-time information to change prediction between taken or not taken.
 - ↳ Use **branch history** to predict future branches.
 - ↳ Simplest method is to use a *saturated counter*: increment counter if branch actually taken, decrease counter if branch not taken.
 - ↳ Predict based on current count.
 - ↳ More advanced predictors evaluate *patterns* in branch history.
- **Random branch prediction**: statistically 50% correct prediction.

A two-bit saturated counter:



Datapath With Forwarding and Flushing



Hazard Summary

- **Structural hazards** caused by conflicts accessing hardware.
 - ↳ Register access fast enough to happen twice in one clock cycle.
 - ↳ Banked L1 cache for simultaneous instruction and data access.
- **Data hazards** caused by Read After Write (RAW).
 - ↳ ALU-ALU forwarding.
 - ↳ MEM-MEM forwarding (memory copies).
 - ↳ Load-use hazard: stall (load-delay slot) and MEM-ALU forward.
- **Control hazards** caused by branch instructions.
 - ↳ Special branch comparator in ID stage.
 - ↳ Branch delay slot; **delayed branching**.
 - ↳ Branch prediction and **pipeline flush**.
- Compiler handles **nop** insertion to fix hazards.
- Hardware handles fixing hazards with **pipeline interlock**.

Datapath Optimizations Summary

Many possible configurations of datapath.

- Single-cycle to Multi-cycle
- Multi-cycle to Pipelining
- Pipeline with or without banked cache for instr. and data memory
- Pipeline with and without forwarding.
- Pipeline with and without load-use slot.
- Pipeline with and without ID stage branch comparator.
- Pipeline with and without branch delay slot.
- Pipeline with and without branch prediction.

During exercises, read carefully the specification and assumptions of that datapath.