


SVD/LSA as a Neural Network

Singular Value Decomposition

$$M_{m \times n} \approx \dot{M}_{m \times n} = U_{m \times p} D_{p \times p} V^T_{p \times n}$$



Rows of V represent documents

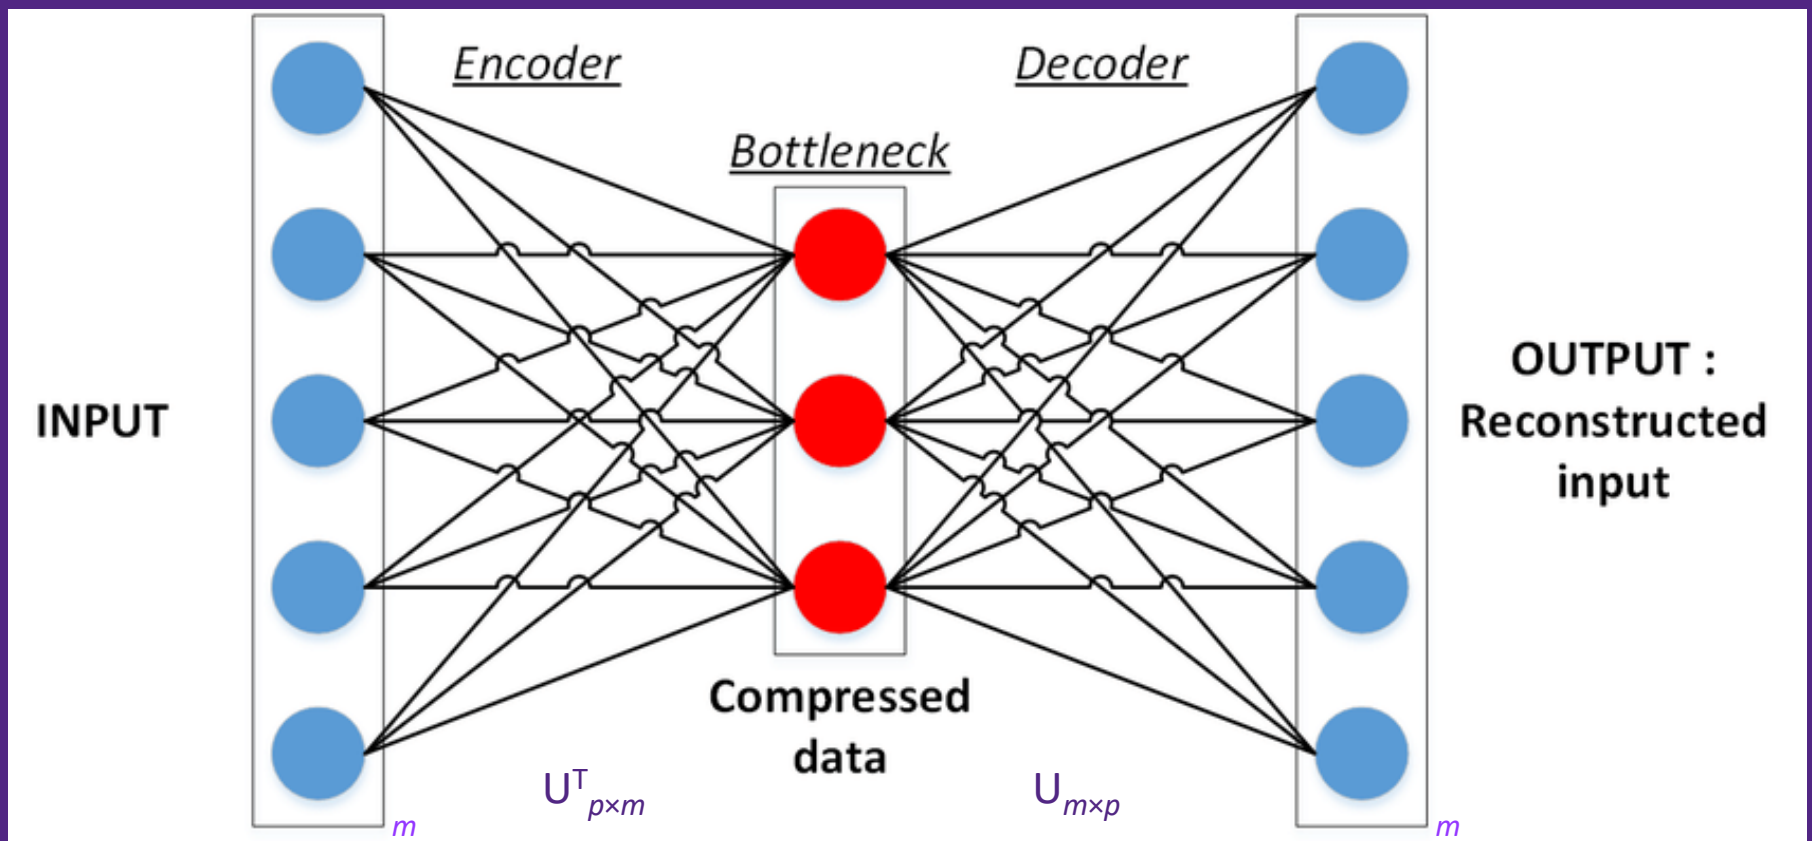
- $(\Sigma^{-1}_{p \times p} U^T_{p \times m}) \dot{M}_{m \times n} = V^T_{p \times n}$
- Each element of column j of V^T is a weighted sum of a column of \dot{M}
- In NN terms:
 - one-layer network
 - m inputs, p outputs
 - fully-connected
 - linear transfer function (no ReLU or anything)
 - weights are $(\Sigma^{-1}_{p \times p} U^T_{p \times m})$

Rows of V represent documents

- $U_{m \times p} \Sigma_{p \times p} V_{p \times n}^T = \dot{M}_{m \times n}$
- Each element of column j of \dot{M} is a weighted sum of a column of V^T
- In NN terms:
 - one-layer network
 - p inputs, m outputs
 - fully-connected
 - linear transfer function (no ReLU or anything)
 - weights are $(U_{m \times p} \Sigma_{p \times p})$ (I'll call them $\tilde{U}_{m \times p}$)

“LSA” Autoencoder

$$\dot{M}_{:,j} = \tilde{U}_{m \times p} \tilde{U}_{p \times m}^T \dot{M}_{:,j}$$



Question: - Are there any issues? (using m neurons in hidden units) A) Yes B) No

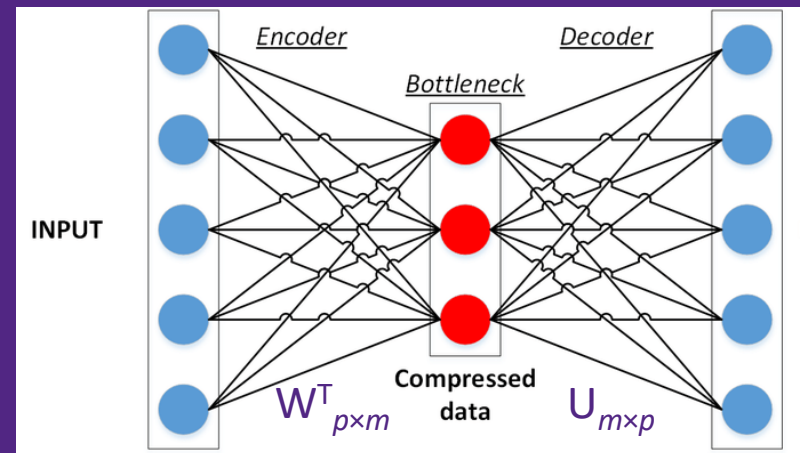
SVD/LSA Versus Autoencoder

- “auto”-encoder meaning “self”-encoder
 - Learns smaller representation for each input (column, document) vector
- Learned weights $\tilde{U}_{m \times p}$ give row (word) representations
- The entries in \tilde{U} will not exactly match the ones you would get from SVD depending on how you train, but they span the same space.
- They’re not “ordered” in terms of importance.
- Full discussion here (optional):
 - <https://arxiv.org/pdf/1804.10253.pdf>

word2vec

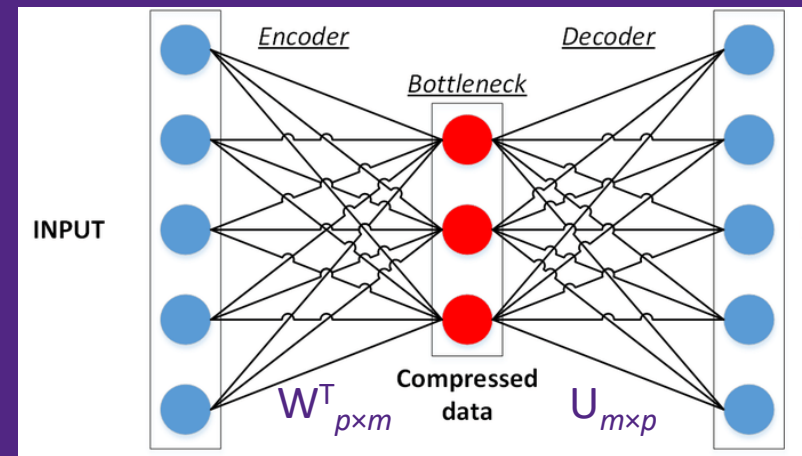
- Simplified from Mikolov, Tomas; et al. “Efficient Estimation of Word Representations in Vector Space”. arXiv:1301.3781 [Optional]
- NN-based word representation learner
 - m input units (one per word)
 - p hidden units (dimension of representation)
 - m output units (one per word)
- **Output not same as input**

$$\zeta(U_{m \times p} W^T_{p \times m} \mathbf{x}_{m \times 1}) = \mathbf{y}_{m \times 1}$$



word2vec

- $\zeta(U_{m \times p} W^T_{p \times m} \mathbf{x}_{m \times 1}) = \mathbf{y}_{m \times 1}$
- \mathbf{x} and \mathbf{y} both represent a word or collection of words
- Matrix $U_{m \times p}$ that we learn will hold our word representations



Choosing x and y :

“Continuous Bag of Words”

- CBoW
 - Vector y (target) is one-hot encoding of a word
 - Vector x (input) is average of one-hot encodings of “context words”
 - Input-output pair created for every position of a “sliding window”

Source Text

The quick brown fox jumps over the lazy dog.

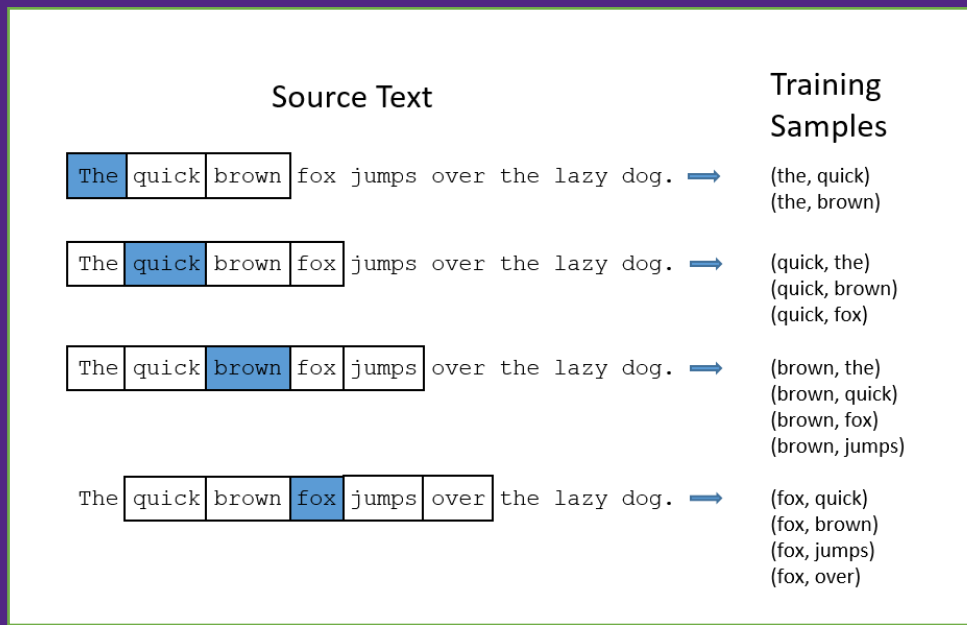
The quick brown fox jumps over the lazy dog.

The quick brown fox jumps over the lazy dog.

The quick brown fox jumps over the lazy dog.


Choosing \mathbf{x} and \mathbf{y} : “Skip-Gram”

- Vector \mathbf{x} (input) is one-hot encoding of a word
- Vector \mathbf{y} (target) is one-hot encoding of a “nearby” word
- Multiple training vectors per sliding window position



“Analogy Task”

- “France is to Paris as Germany is to _____?”
- $\text{Paris} - \text{France} + \text{Germany} = ???$



Structured Representations from Neural Nets – In General

word2vec idea

- Purpose of word2vec was never to “predict an output”
- What matters is the representation created as a “side-effect” of training

Intermediate Representations

- Same intermediate representation \rightarrow same output
- Advantageous for NN to map different inputs that should give same output to similar representation


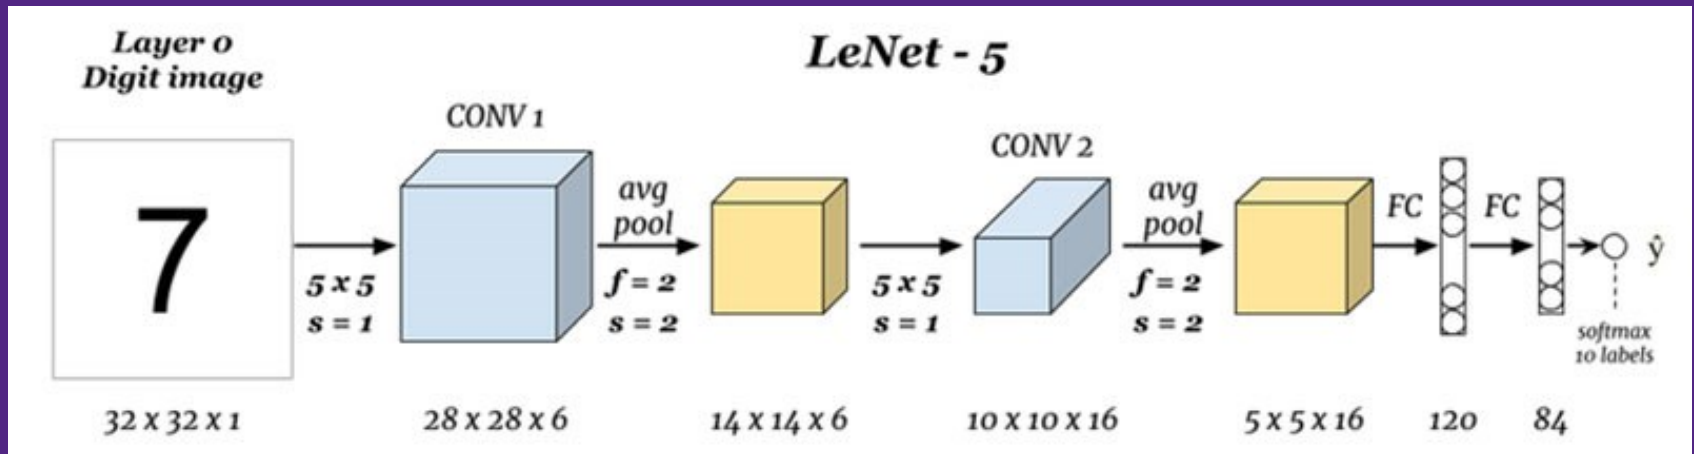


Image Recognition



1024 inputs

Designing Your Own Representations

- Is there a labelling task that can define a useful notion of similarity for you?
- Even if you can't learn that task "100%" you might still learn structure!
 - Recall: Skip Gram model trains on (brown, the), (brown, quick), (brown, fox), (brown, jumps) – impossible to produce a single "right answer" for input "brown"!
- Train a NN that has at least one layer that "compresses" the input.
- Output of that hidden layer, or the hidden layer weights, become your representation

Summary

- We can learn representations that capture relational structure
- LSA/SVD Does this
- LSA/SVD is similar to Autoencoding neural networks
- Other learning tasks and architectures learn different structures
 - CBoW and Skip Grams for word2vec
 - Classification for images
- You can design your own!