


# B-Trees

A B-tree of order  $d$  is a multiway search tree with the following properties:

- The root has at least 2 children and at most  $d$ .
- All internal nodes other than the root have at least  $\left\lceil \frac{d}{2} \right\rceil$  and at most  $d$  children
- All the leaves are at the same level

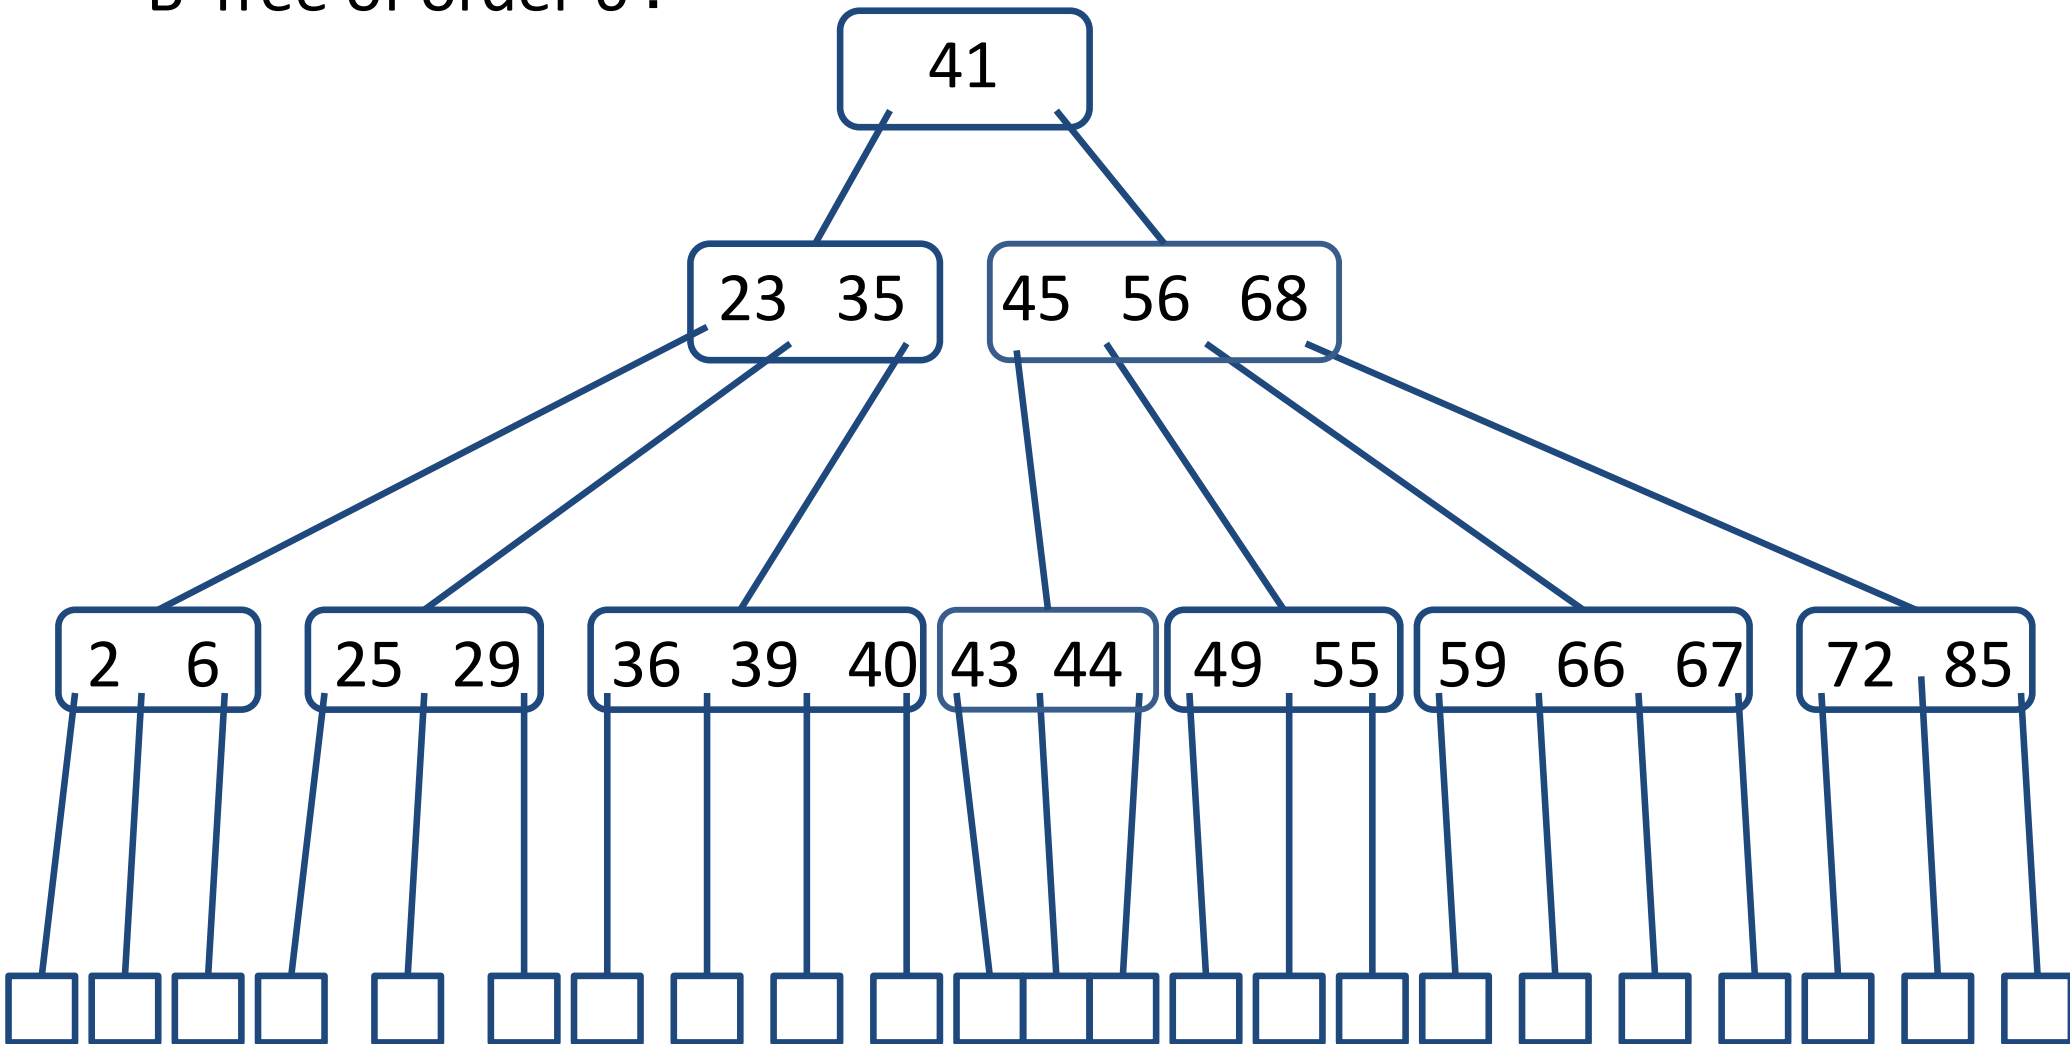
# B-Trees

B-Tree of order 6?



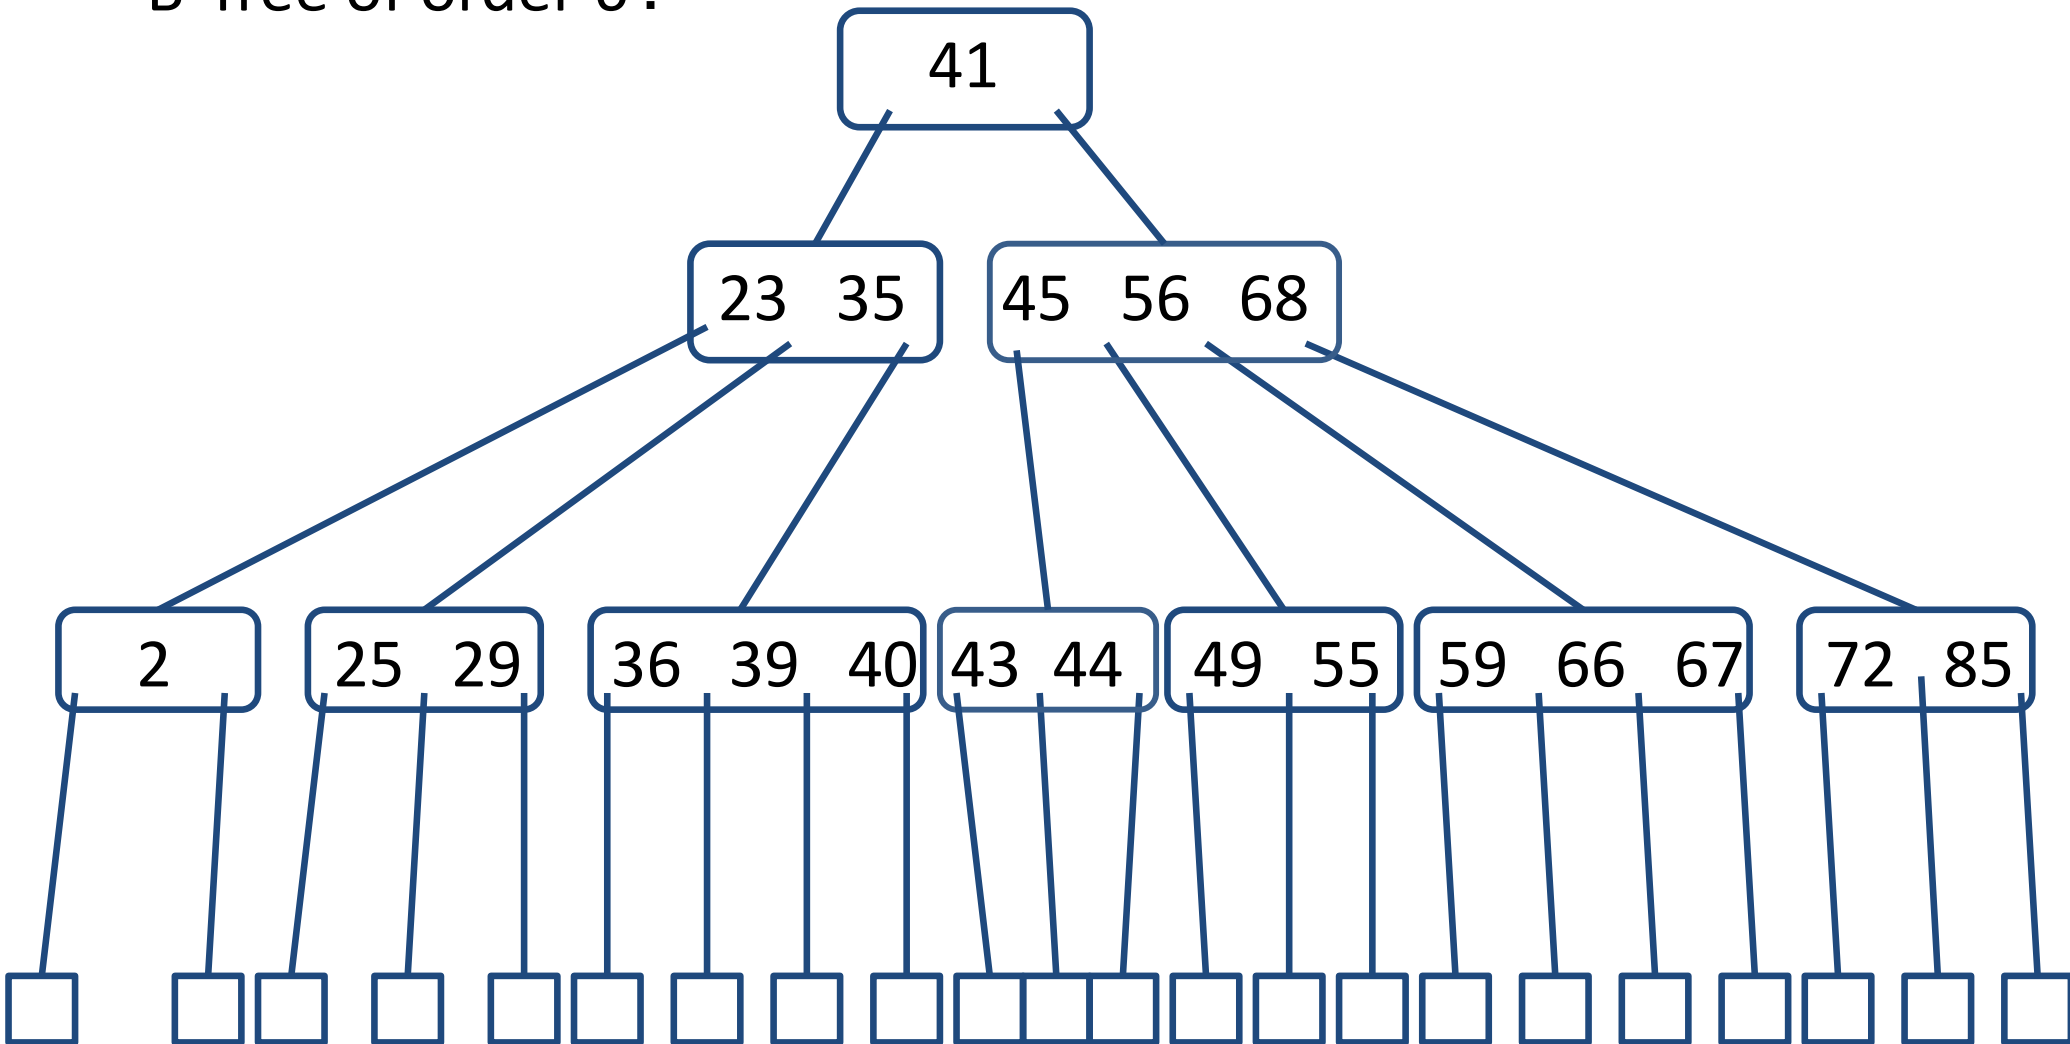
# B-Trees

B-Tree of order 6?



# B-Trees

B-Tree of order 6?




# What is the Maximum Height of a B-Tree?

Height is  $O(\log_d n)$

# B-Trees


B-Tree of order 6



# B-Trees

B-Tree of order 6

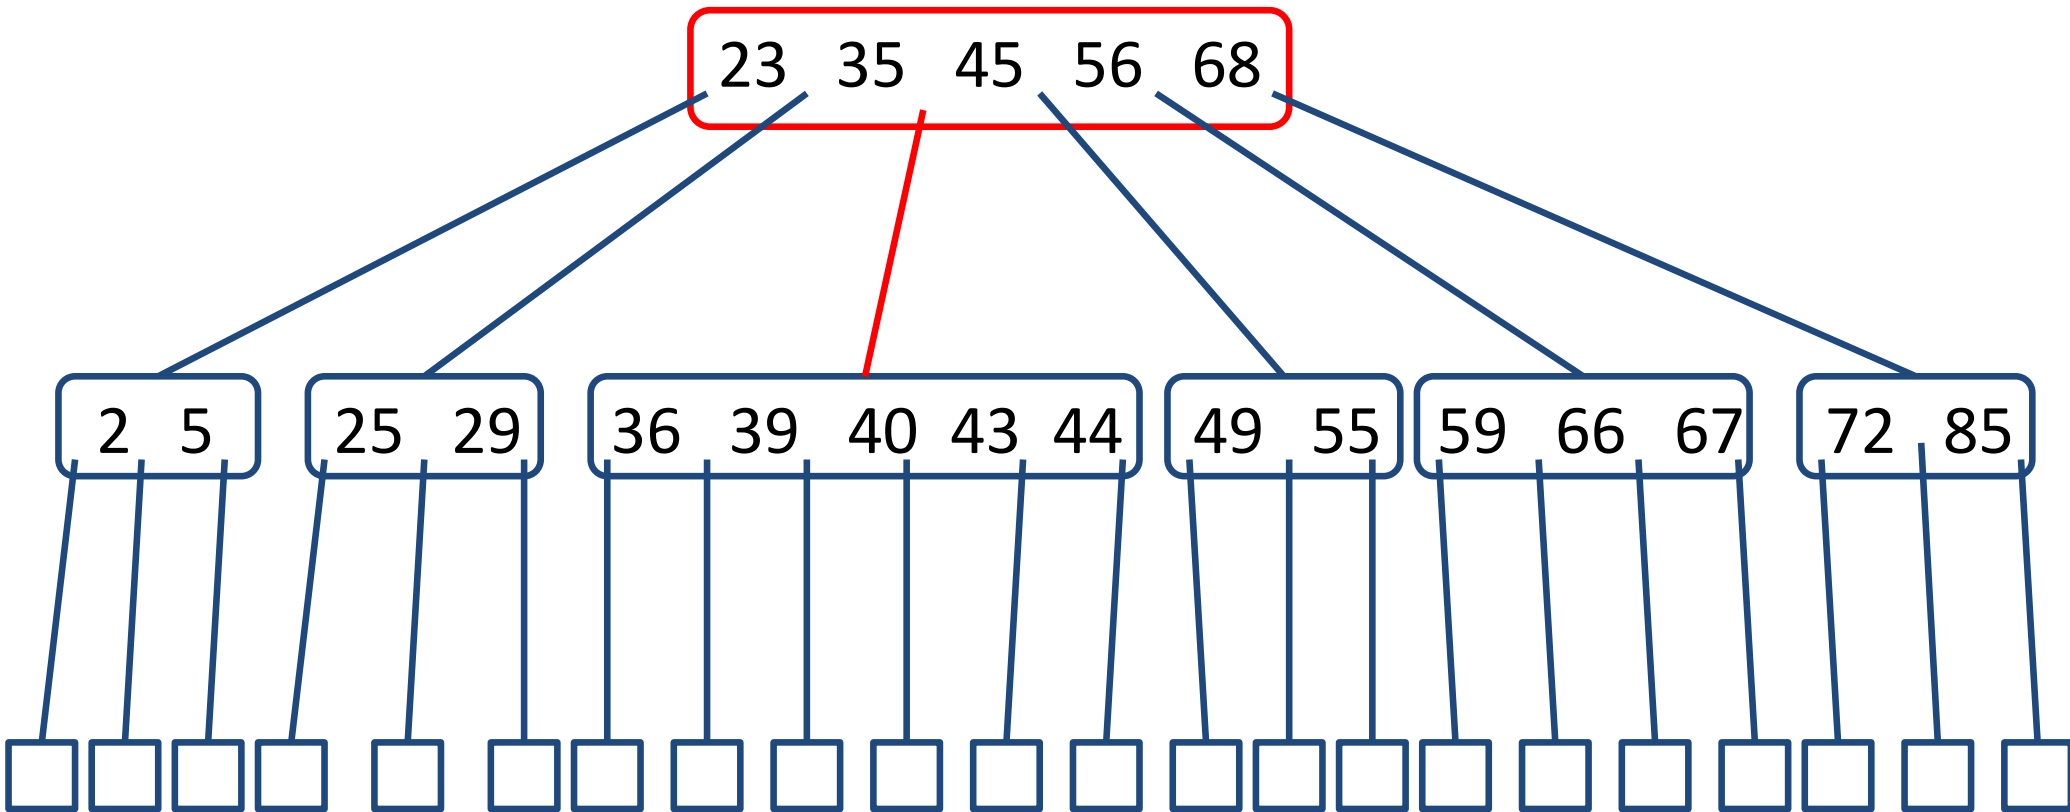
Put 38



# B-Trees

B-Tree of order 6

Put 38

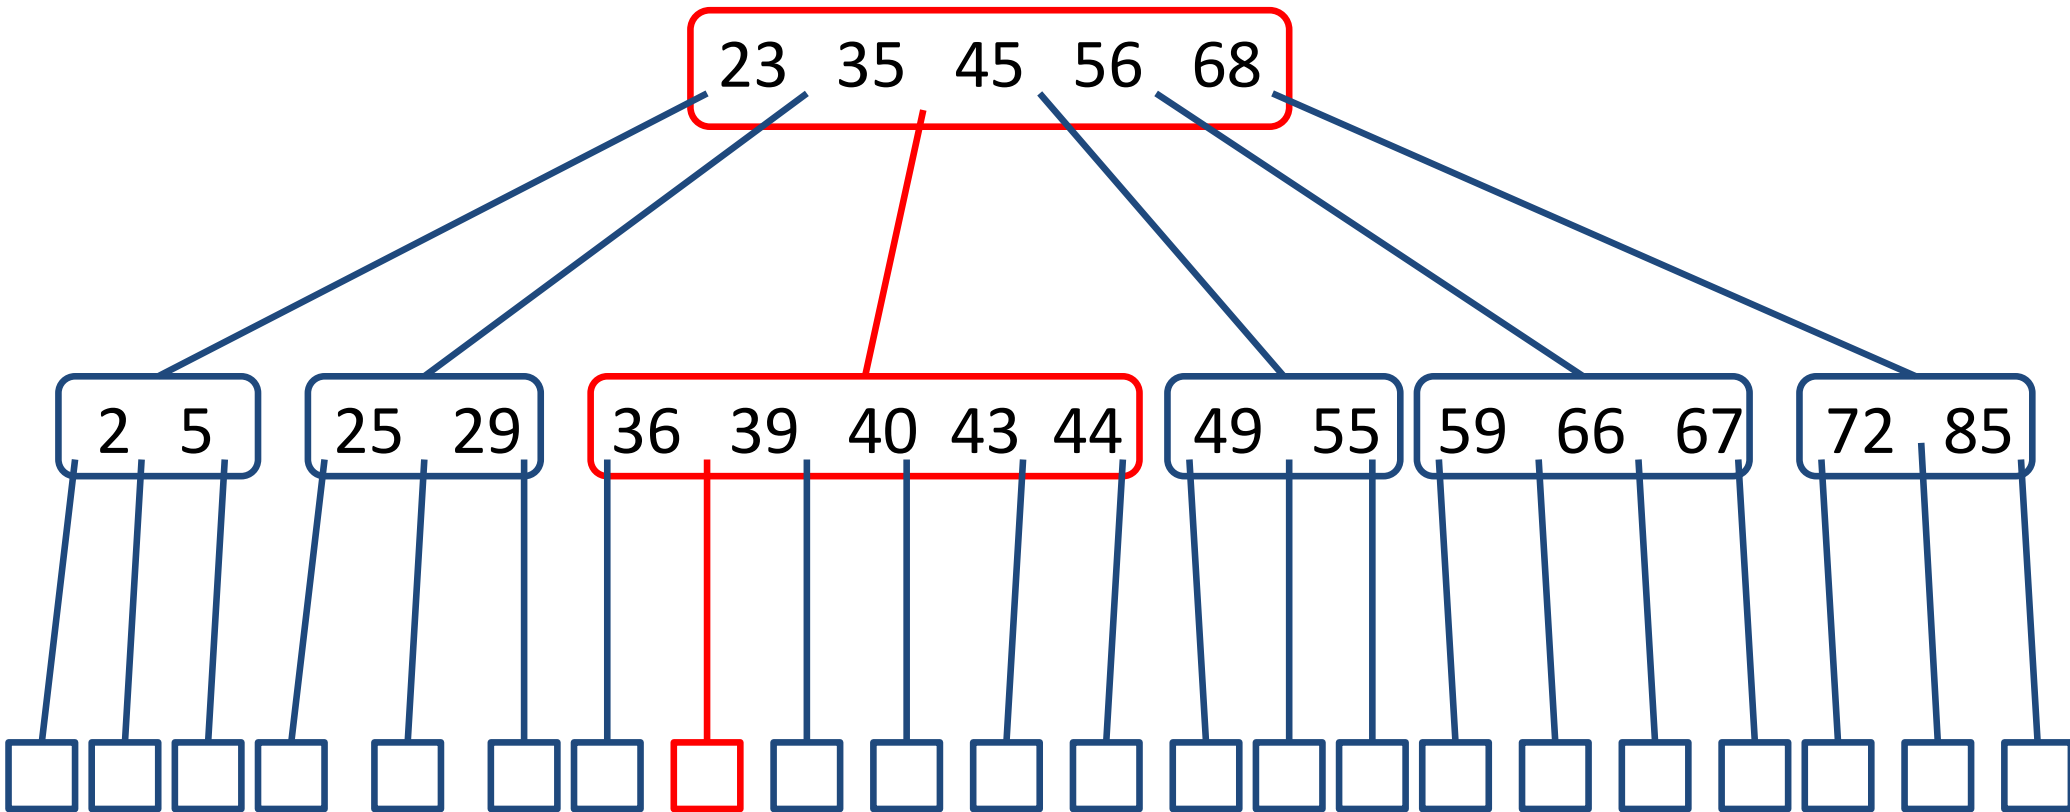




# B-Trees

B-Tree of order 6


Put 38



# B-Trees

B-Tree of order 6


Put 38



# B-Trees

B-Tree of order 6


Put 38



# B-Trees

B-Tree of order 6


Put 38



# B-Trees

B-Tree of order 6


Put 38



# B-Trees

B-Tree of order 6


Put 38



# B-Trees

B-Tree of order 6

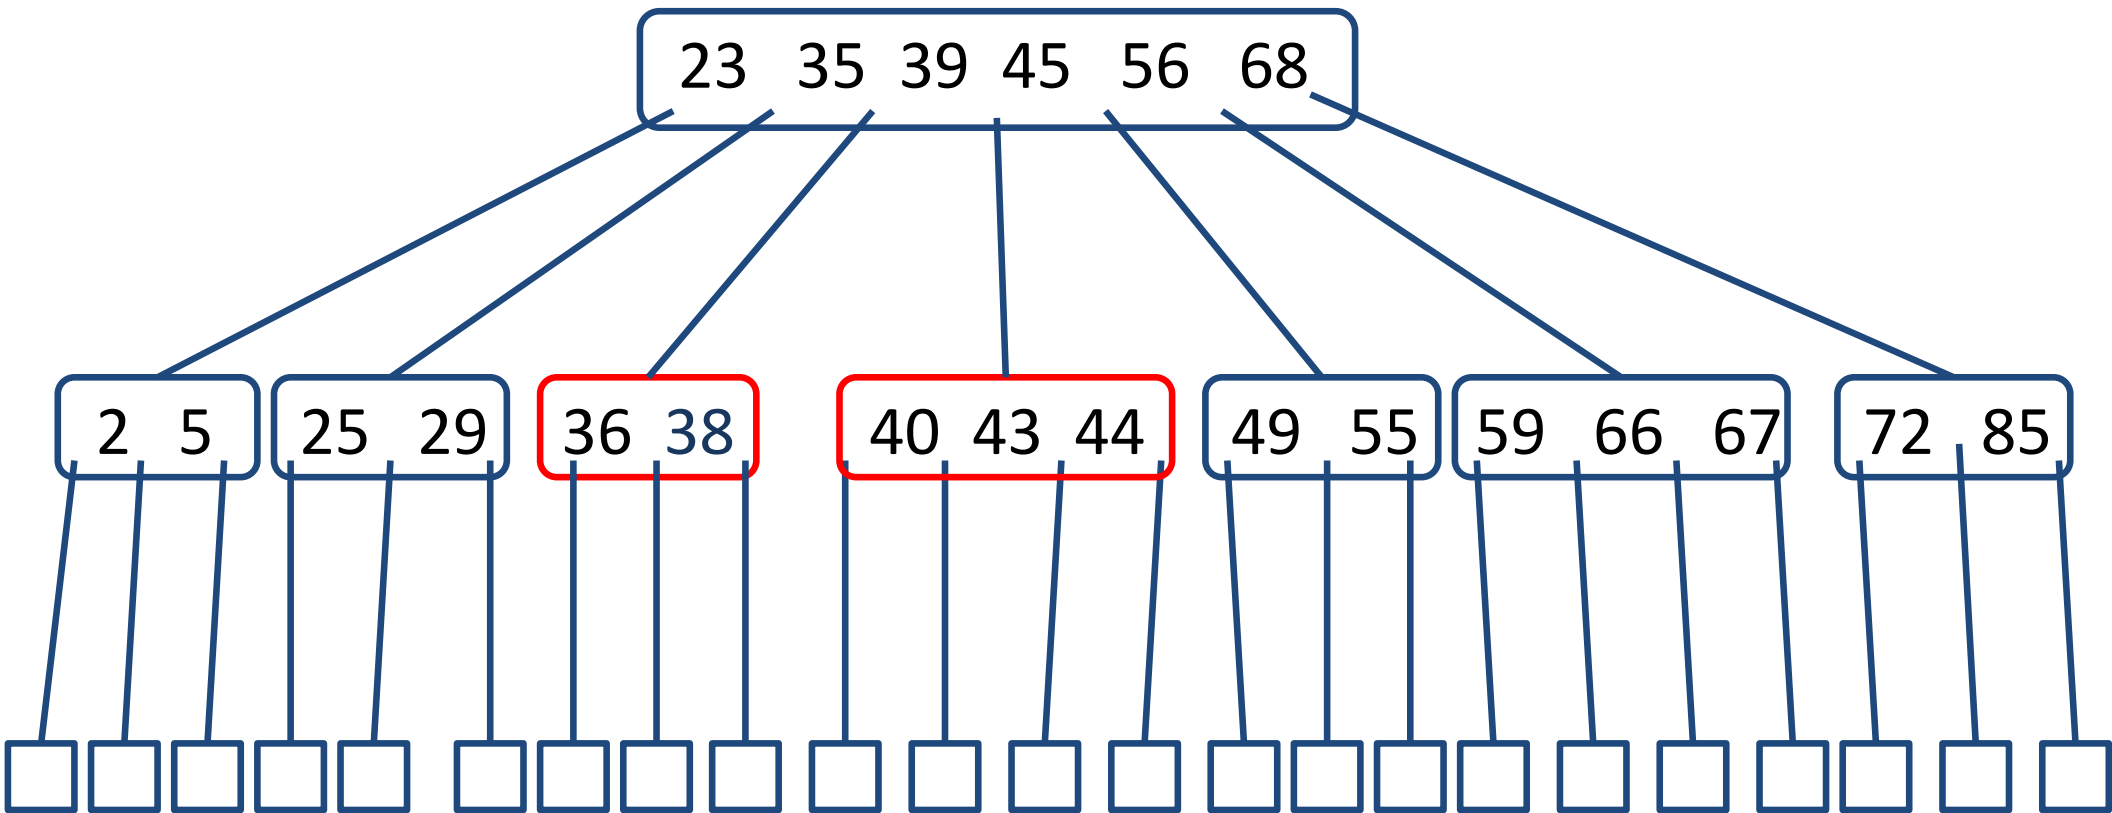
Put 38



# B-Trees

B-Tree of order 6

Put 38

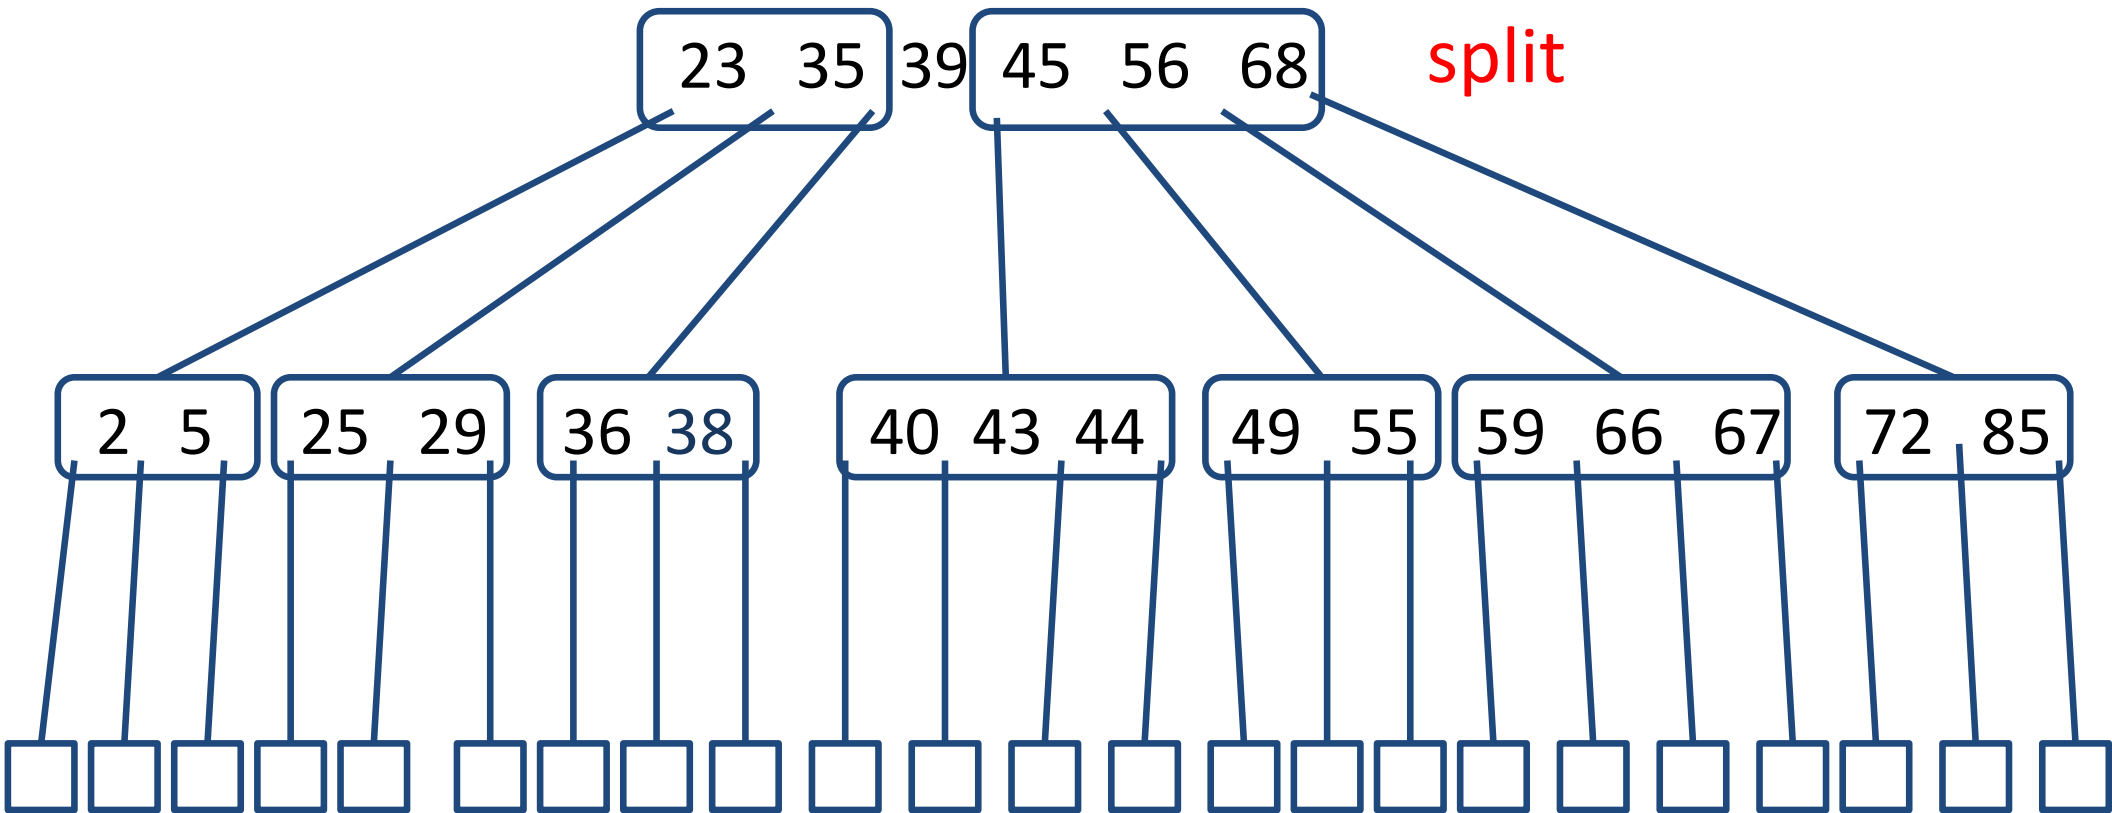




# B-Trees

B-Tree of order 6

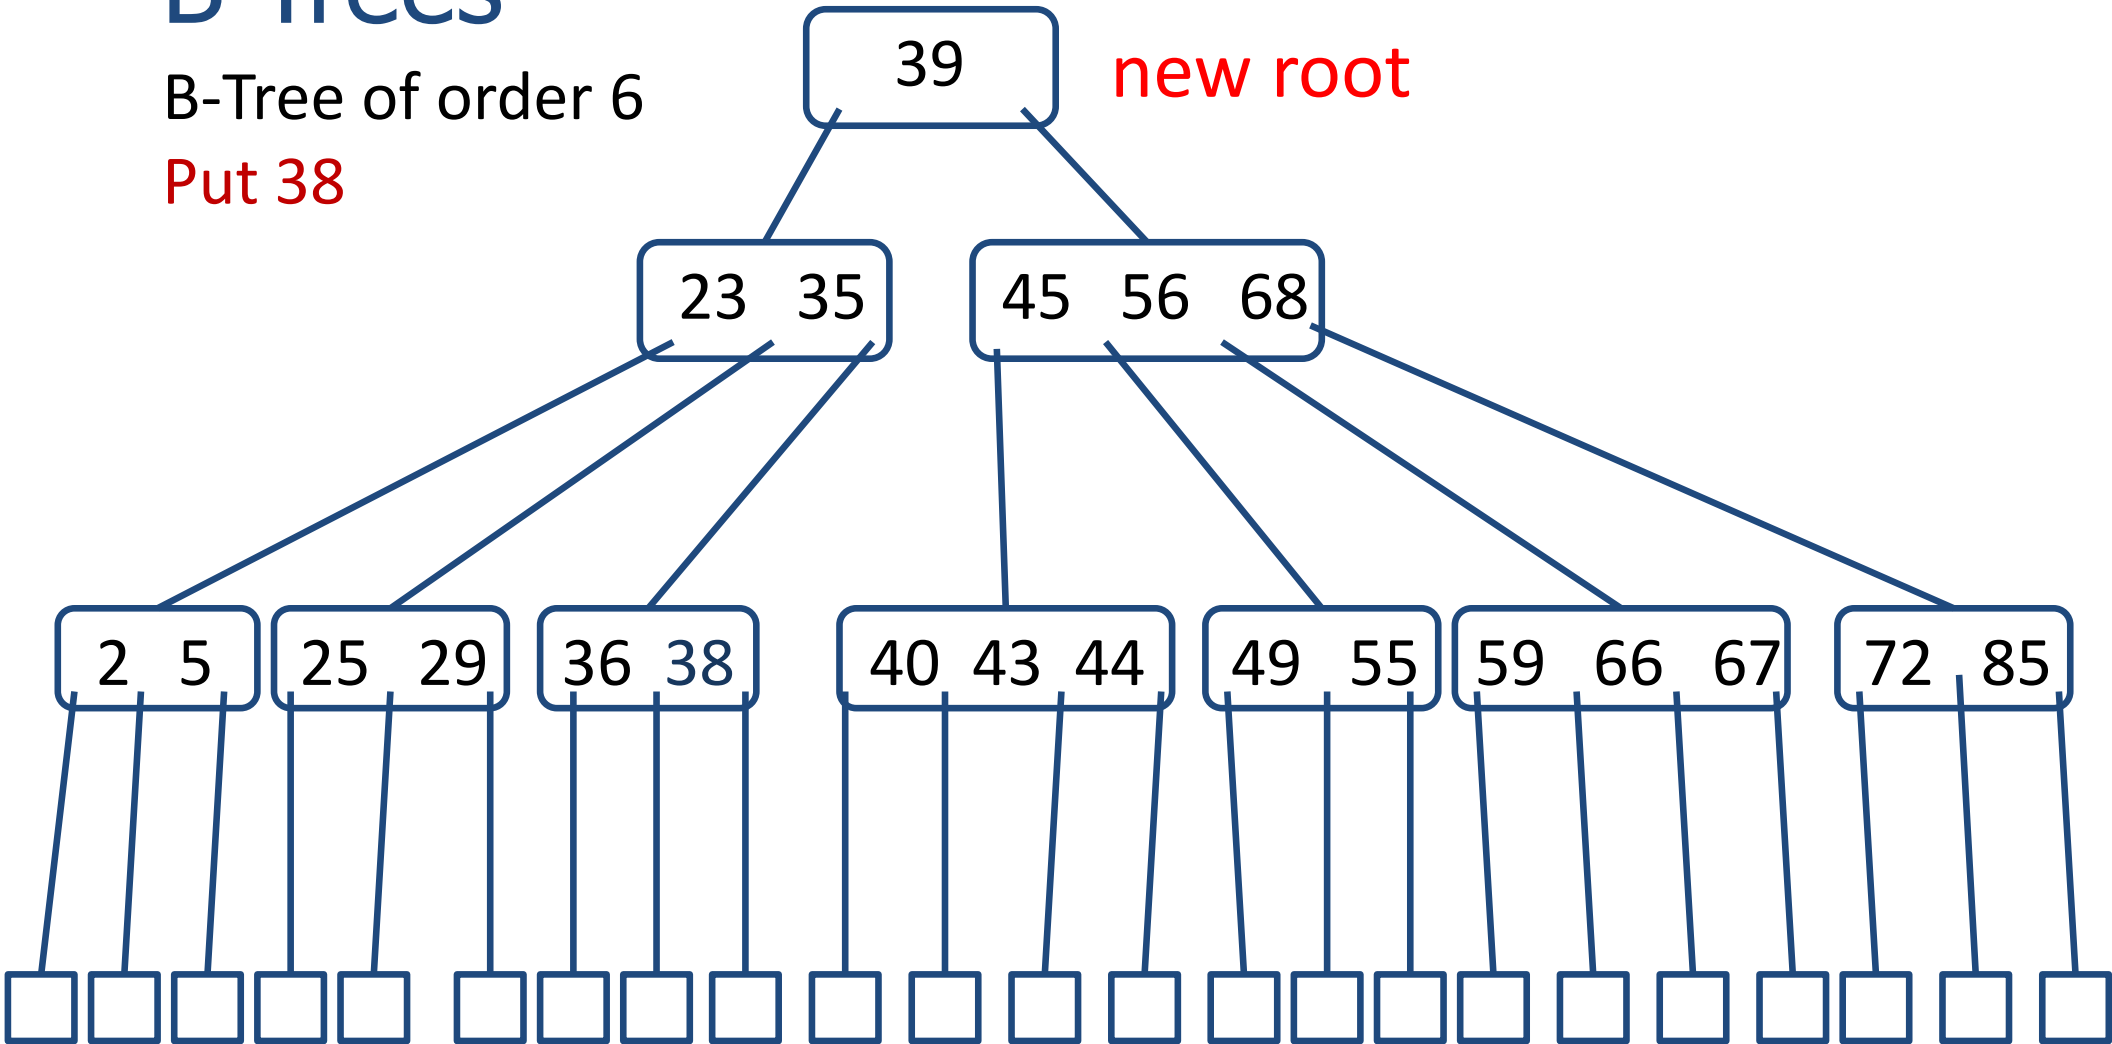
Put 38



# B-Trees

B-Tree of order 6

Put 38



## Algorithm *put* ( $r, k, o$ )

**In:** Root  $r$  of a B-tree, data item  $(k, o)$

**Out:** {Insert data item  $(k, o)$  in the B-tree

Search for  $k$  to find the **lowest** insertion **internal** node  $v$  }

Add the new data item  $(k, o)$  at node  $v$

**while** node  $v$  *overflows* **do** {

**if**  $v$  is the root **then**

        Create a new empty root and set as parent of  $v$


    Split  $v$  around the **middle** key  $k'$ , move  $k'$  to parent, and  
    update parent's children

$v \leftarrow$  parent of  $v$

}

# B-Trees

B-Tree of order 6



## Algorithm *put* ( $r, k, o$ )

**In:** Root  $r$  of a B-tree, data item  $(k, o)$

**Out:** {Insert data item  $(k, o)$  in the B-tree

$O(\log d \times \log_d n)$

Search for  $k$  to find the **lowest** insertion **internal** node  $v$  }

Add the new data item  $(k, o)$  at node  $v$

**while** node  $v$  *overflows* **do** {

**if**  $v$  is the root **then**

        Create a new empty root and set as parent of  $v$

    Split  $v$  around the **middle** key  $k'$ , move  $k'$  to parent, and  
    update parent's children


$v \leftarrow$  parent of  $v$

}

# B-Trees

B-Tree of order 6


Put 38



# B-Trees

B-Tree of order 6


Put 38



# B-Trees

B-Tree of order 6

Put 38





## Algorithm *put* ( $r, k, o$ )

**In:** Root  $r$  of a B-tree, data item  $(k, o)$

**Out:** {Insert data item  $(k, o)$  in the B-tree

Search for  $k$  to find the **lowest** insertion **internal** node  $v$  }  $O(\log d \times \log_d n)$

Add the new data item  $(k, o)$  at node  $v$  }  $O(d)$

**while** node  $v$  *overflows* **do** {

**if**  $v$  is the root **then**

        Create a new empty root and set as parent of  $v$

    Split  $v$  around the **middle** key  $k'$ , move  $k'$  to parent, and  
    update parent's children

$v \leftarrow$  parent of  $v$

}

## Algorithm *put* ( $r, k, o$ )

**In:** Root  $r$  of a B-tree, data item  $(k, o)$

**Out:** {Insert data item  $(k, o)$  in the B-tree

Search for  $k$  to find the **lowest** insertion **internal** node  $v$  }  $O(\log d \times \log_d n)$

Add the new data item  $(k, o)$  at node  $v$  }  $O(d)$

**while** node  $v$  *overflows* **do** {

**if**  $v$  is the root **then**

        Create a new empty root and set as parent of  $v$

    Split  $v$  around the **middle** key  $k'$ , move  $k'$  to parent, and  
    update parent's children

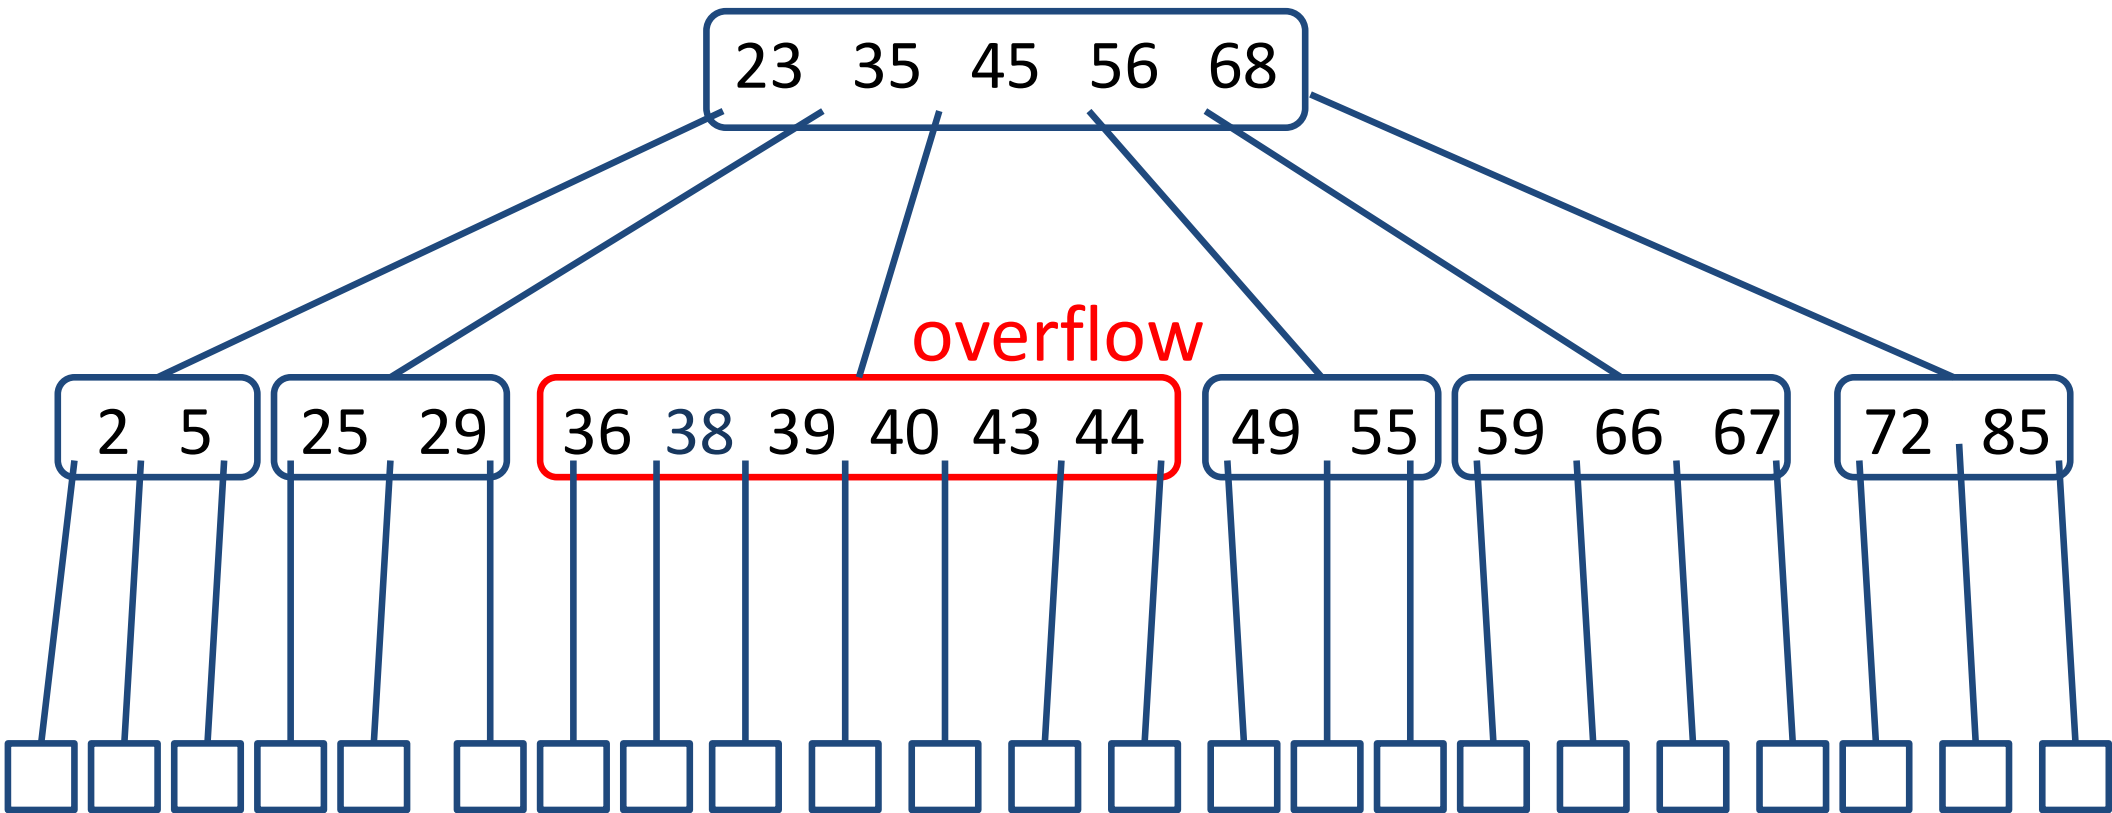
$v \leftarrow$  parent of  $v$

}

# B-Trees

B-Tree of order 6


Put 38



# B-Trees

B-Tree of order 6

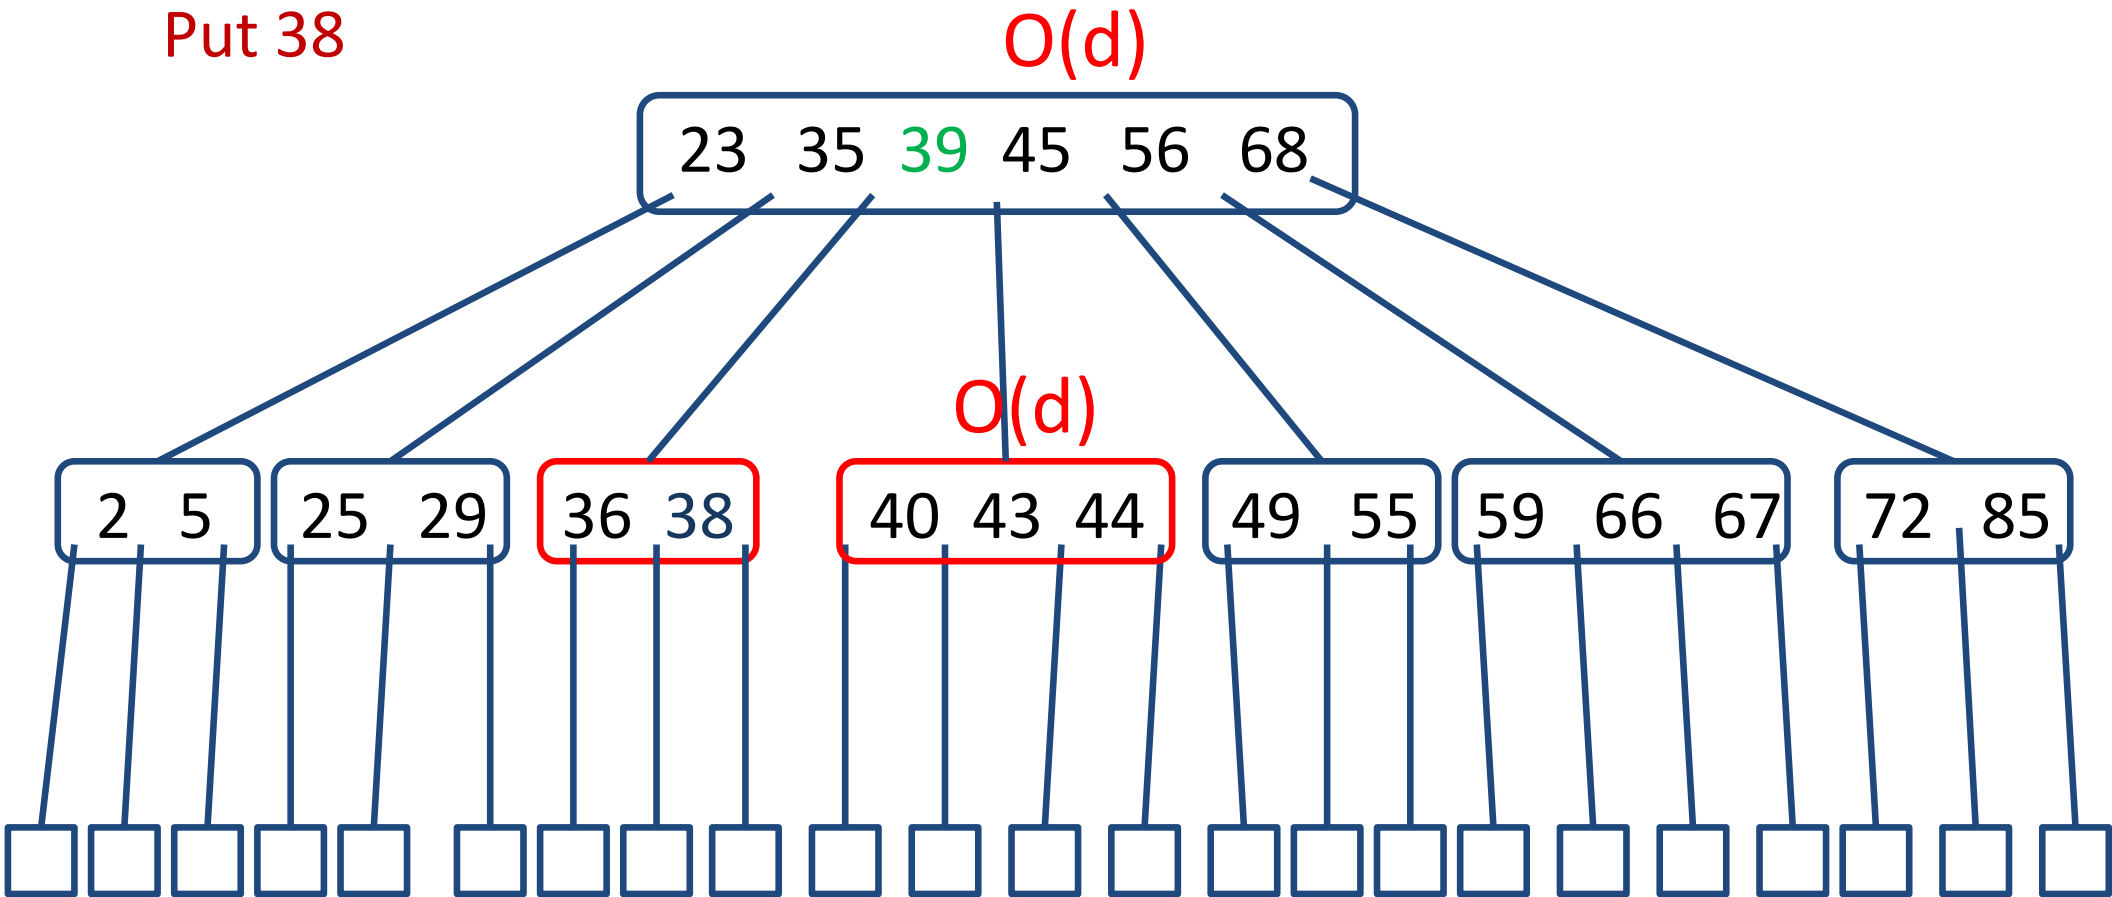
Put 38



# B-Trees

B-Tree of order 6

Put 38



## Algorithm *put* ( $r, k, o$ )

**In:** Root  $r$  of a B-tree, data item  $(k, o)$

**Out:** {Insert data item  $(k, o)$  in the B-tree

Search for  $k$  to find the **lowest** insertion **internal** node  $v$  }  $O(\log d \times \log_d n)$

Add the new data item  $(k, o)$  at node  $v$  }  $O(d)$

**while** node  $v$  *overflows* **do** {

**if**  $v$  is the root **then**

        Create a new empty root and set as parent of  $v$

    Split  $v$  around the **middle** key  $k'$ , move  $k'$  to parent, and  
    update parent's children

$v \leftarrow$  parent of  $v$

}

## Algorithm *put* ( $r, k, o$ )

**In:** Root  $r$  of a B-tree, data item  $(k, o)$

**Out:** {Insert data item  $(k, o)$  in the B-tree

Search for  $k$  to find the **lowest** insertion **internal** node  $v$  }  $O(\log d \times \log_d n)$

Add the new data item  $(k, o)$  at node  $v$  }  $O(d)$

**while** node  $v$  *overflows* **do** {

**if**  $v$  is the root **then**

        Create a new empty root and set as parent of  $v$  }  $O(d)$

    Split  $v$  around the **middle** key  $k'$ , move  $k'$  to parent, and  
    update parent's children

$v \leftarrow$  parent of  $v$

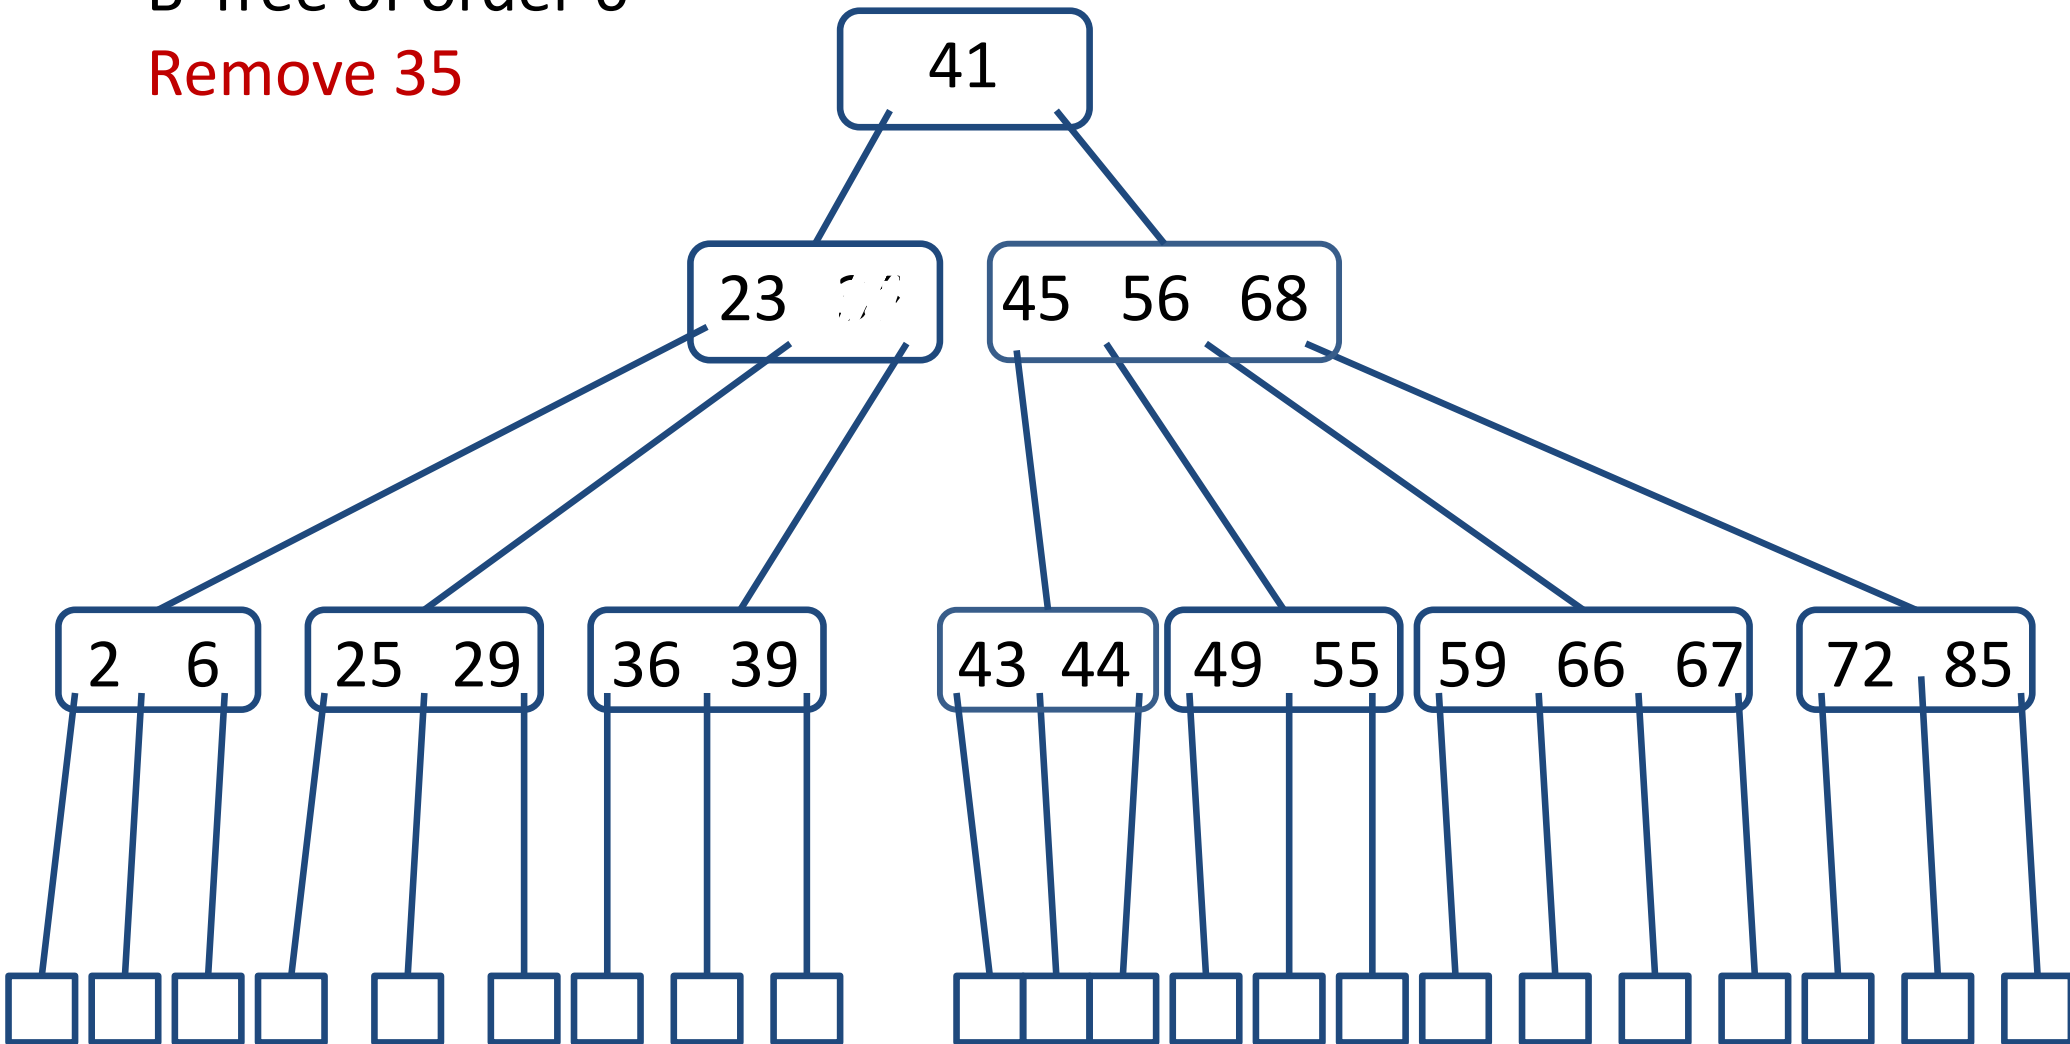
}

Time complexity of *put* is  $O(d \log_d n)$

# B-Trees

B-Tree of order 6

Remove 35

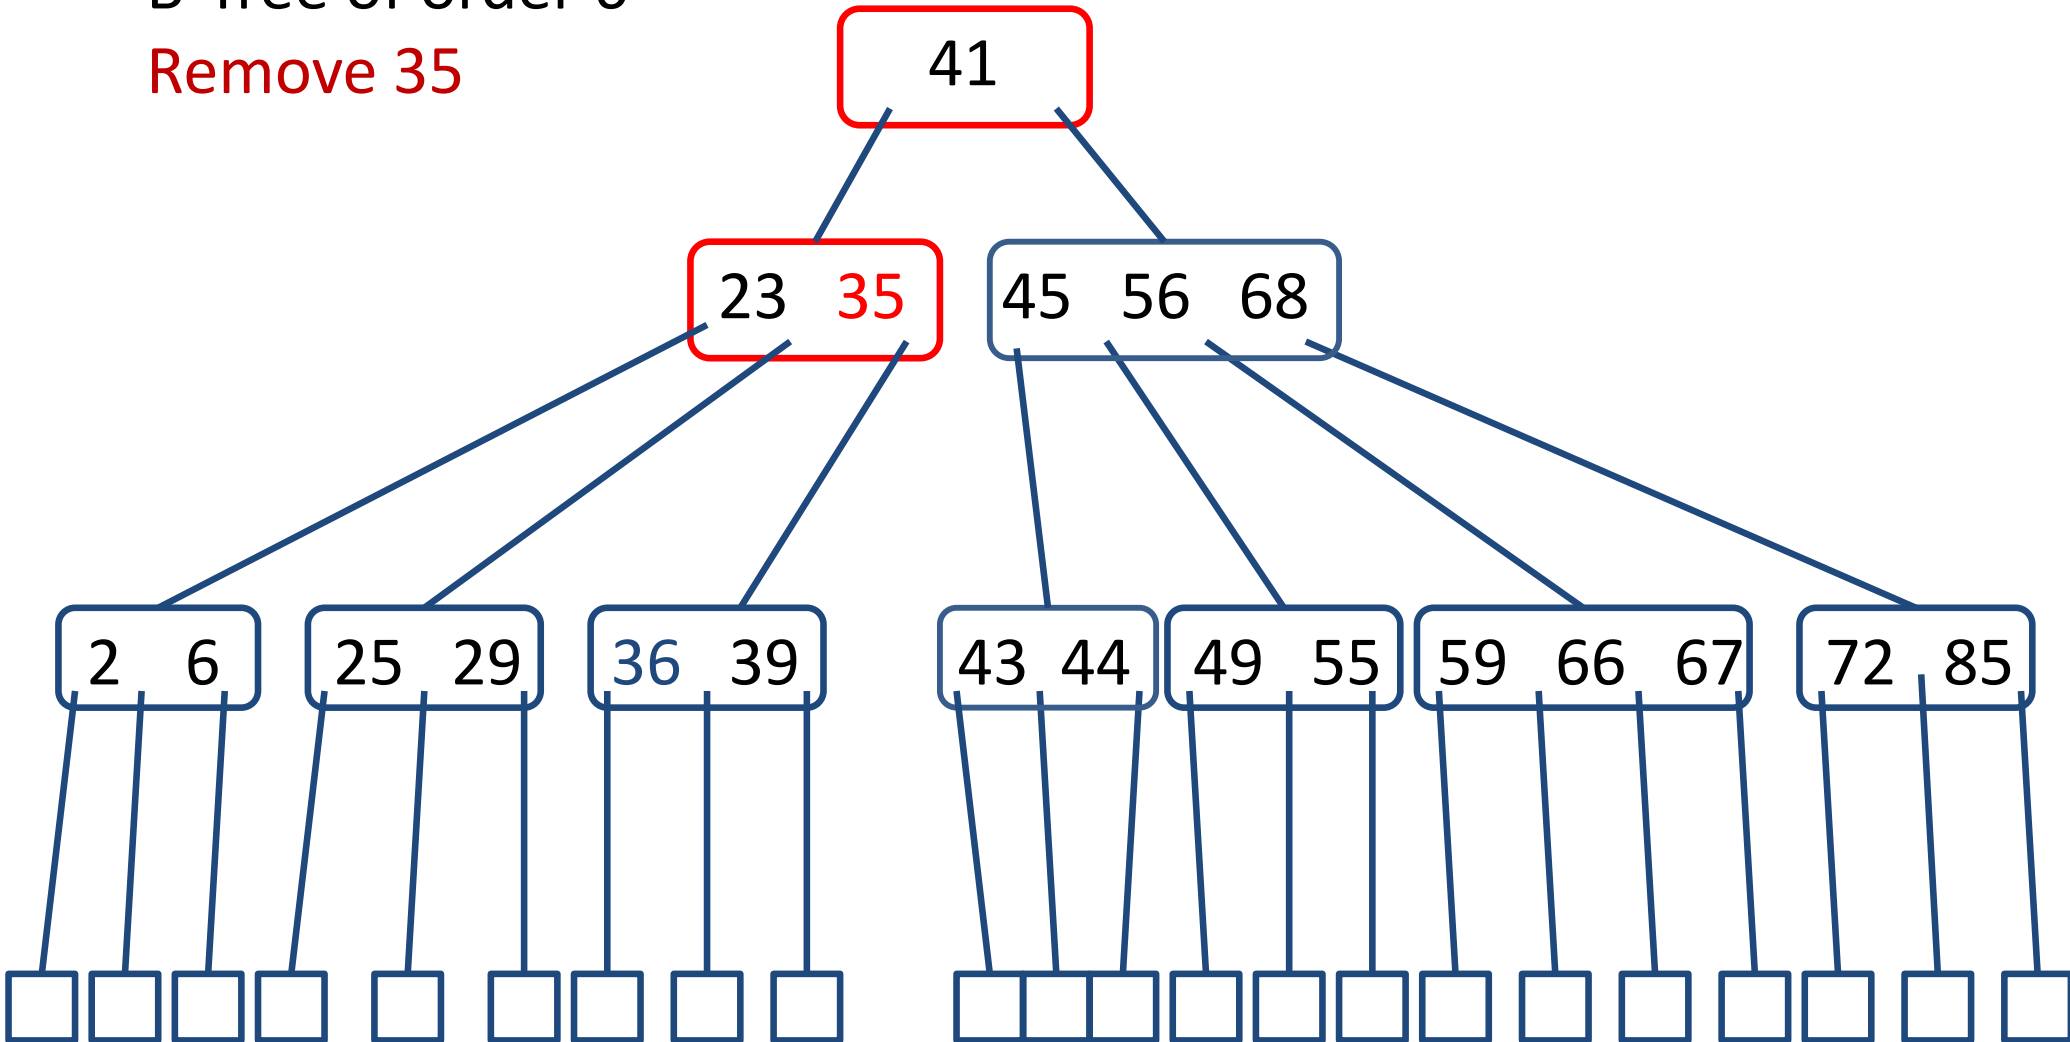




# B-Trees

B-Tree of order 6

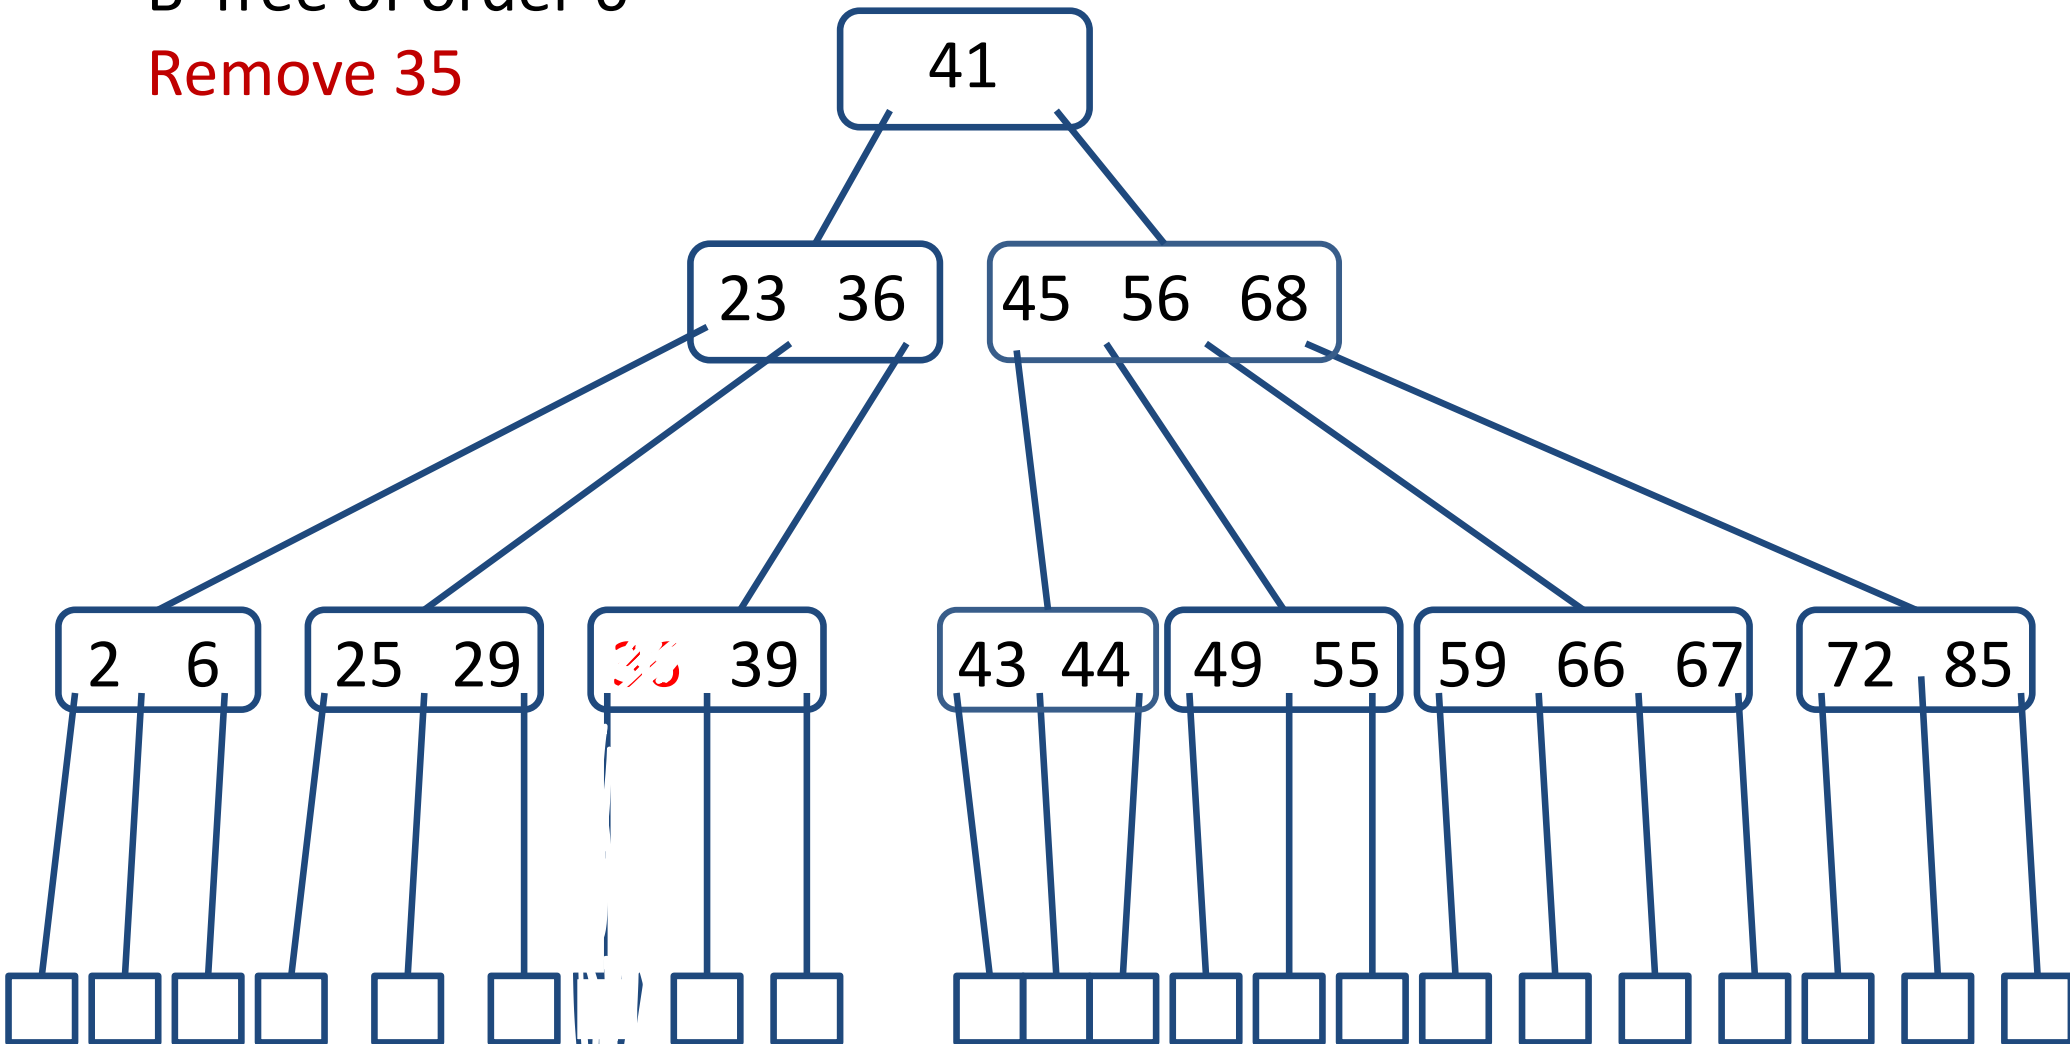
Remove 35



# B-Trees

B-Tree of order 6

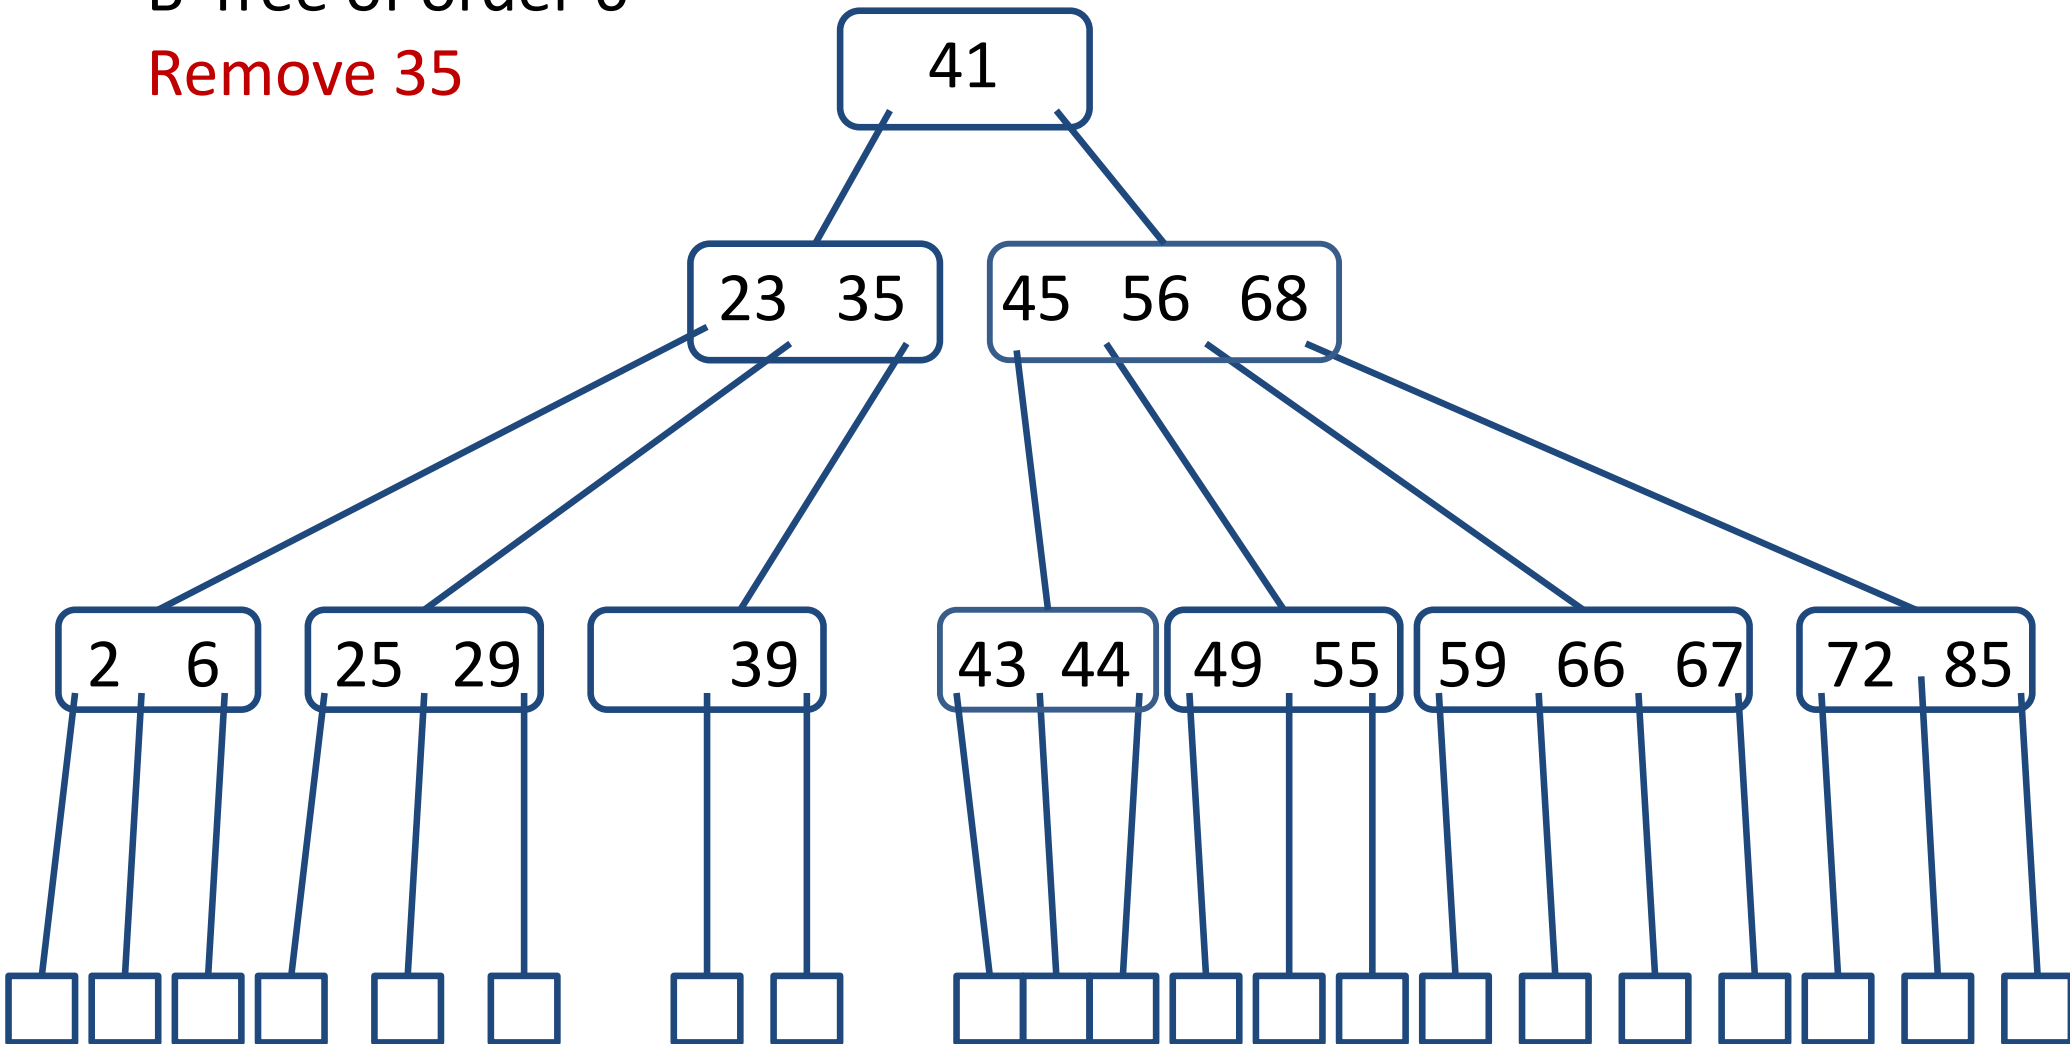
Remove 35



# B-Trees

B-Tree of order 6

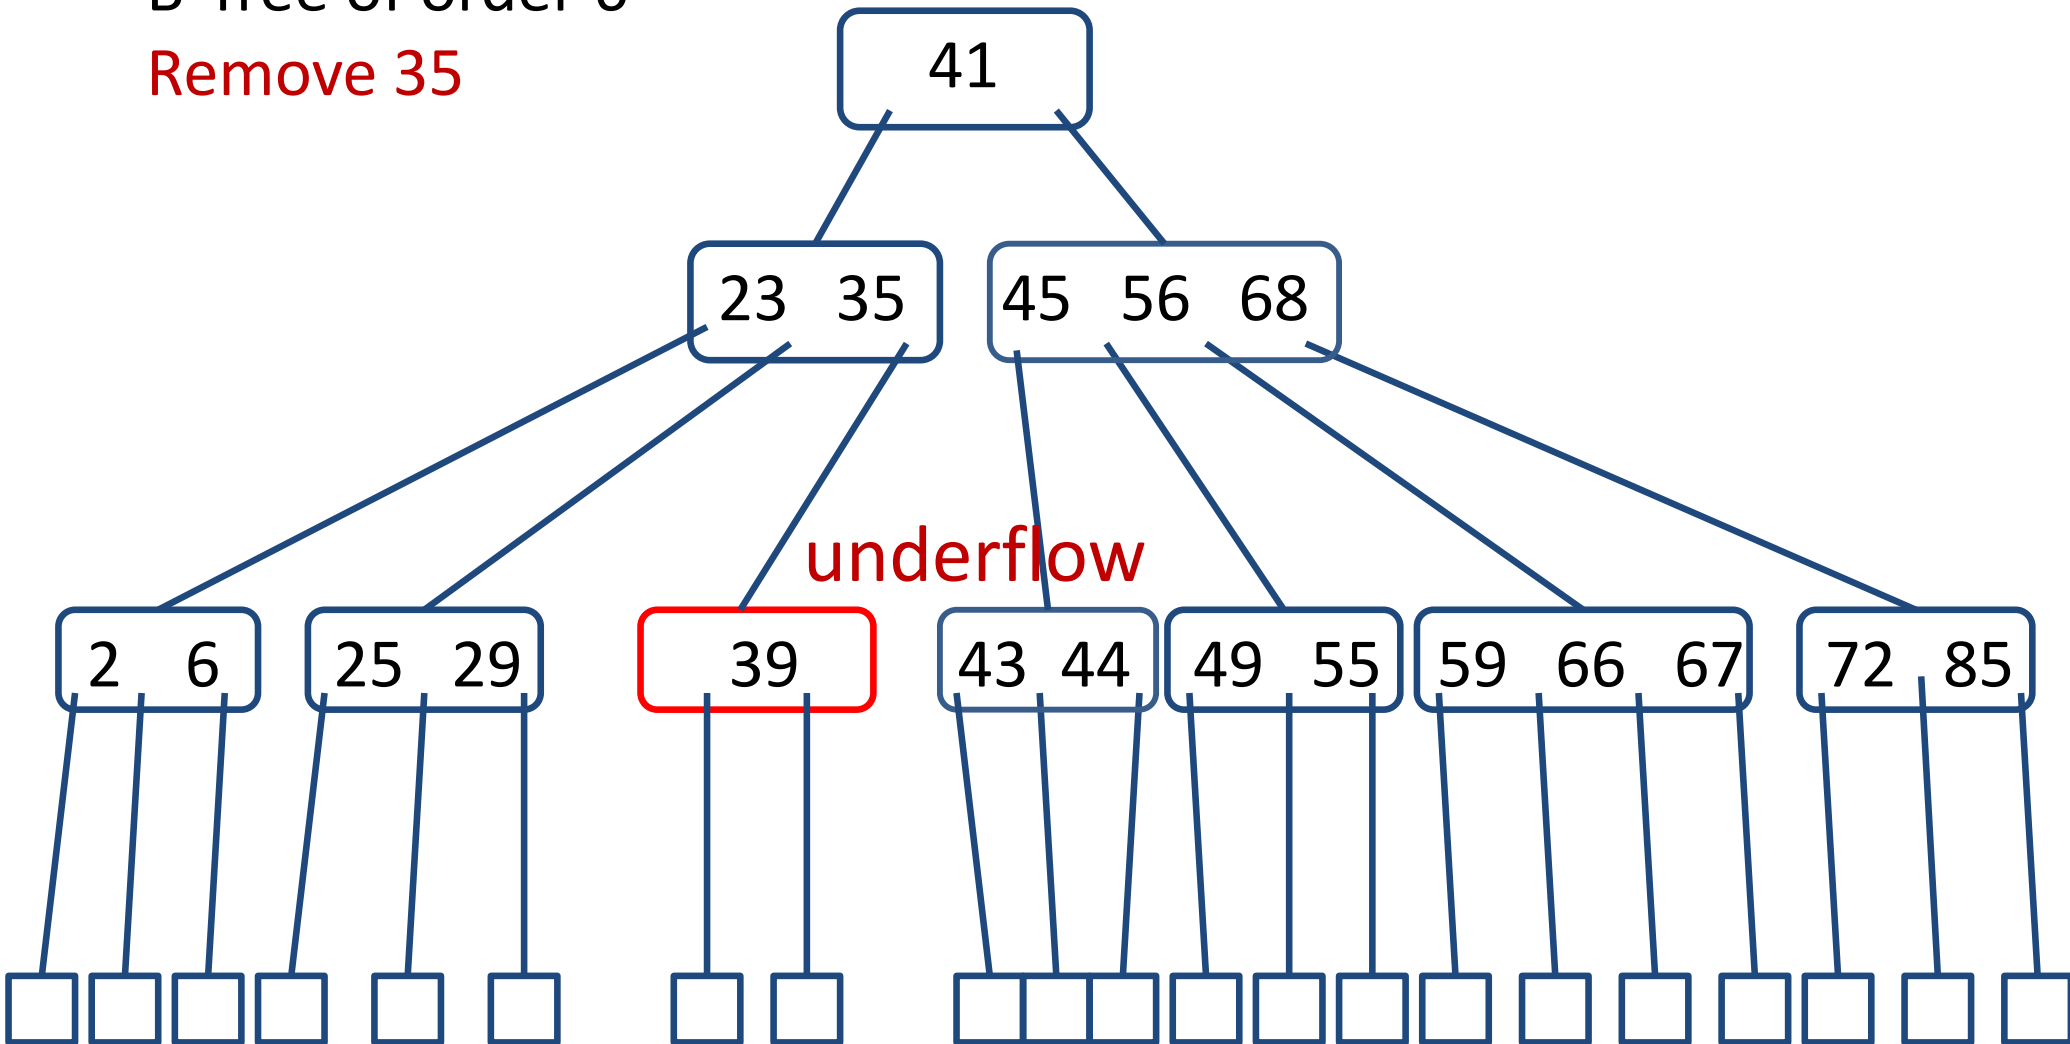
Remove 35



# B-Trees

B-Tree of order 6

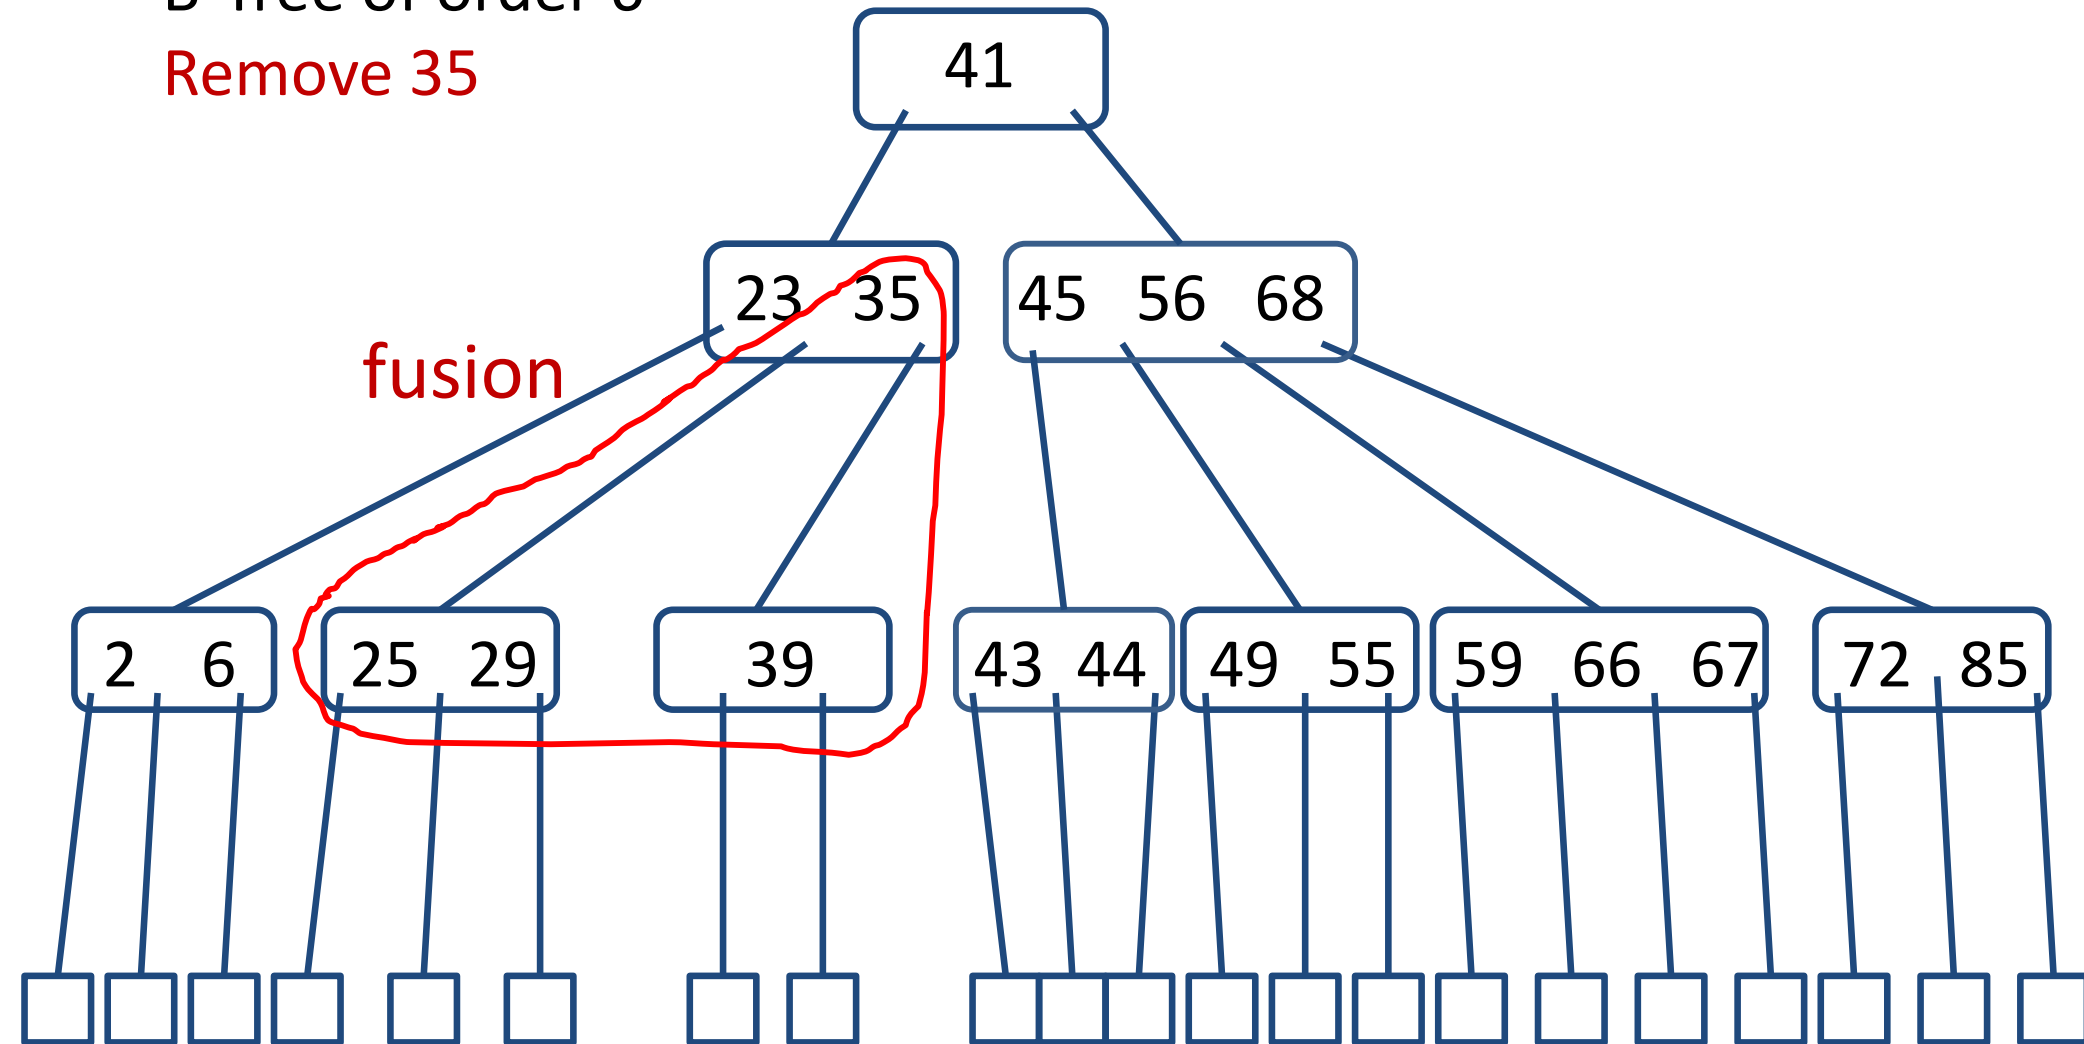
Remove 35



# B-Trees

B-Tree of order 6

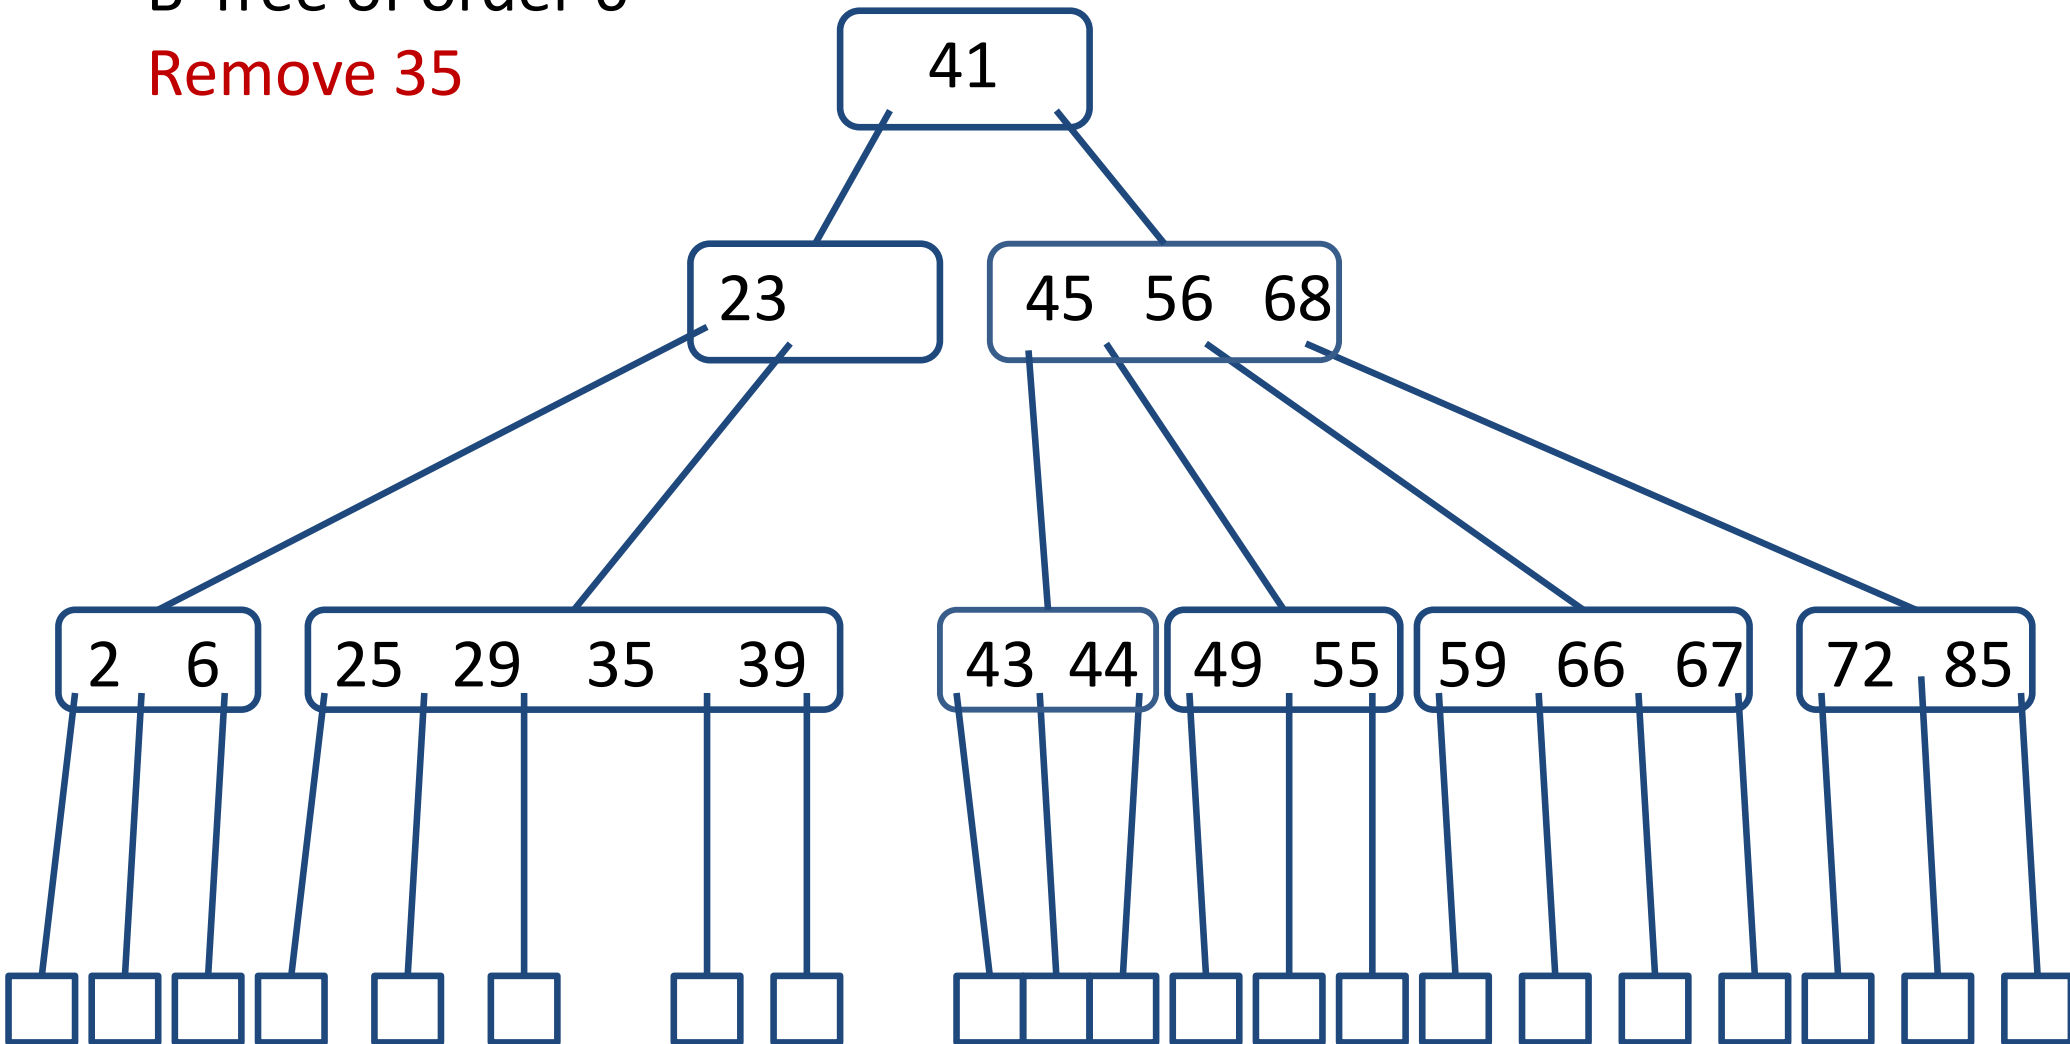
Remove 35



# B-Trees

B-Tree of order 6

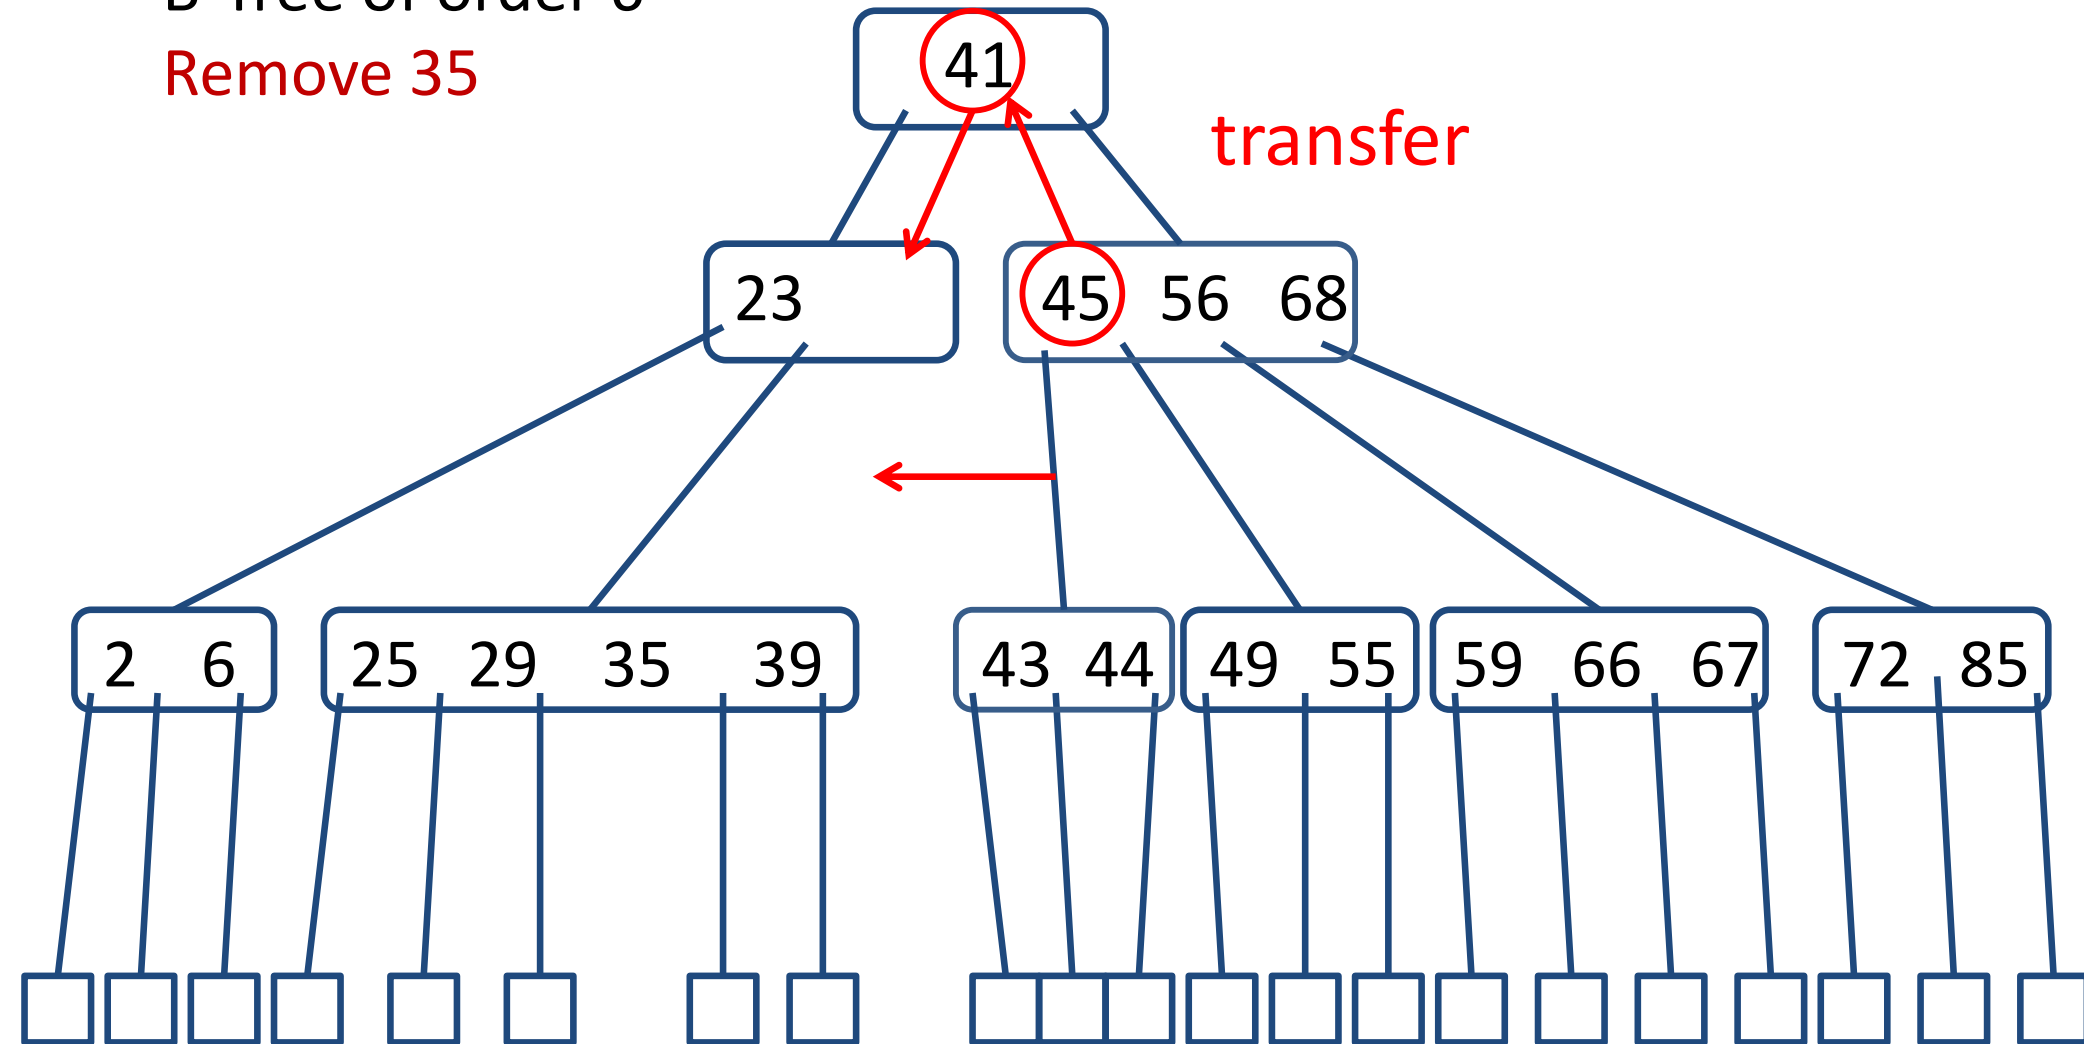
Remove 35



# B-Trees

B-Tree of order 6

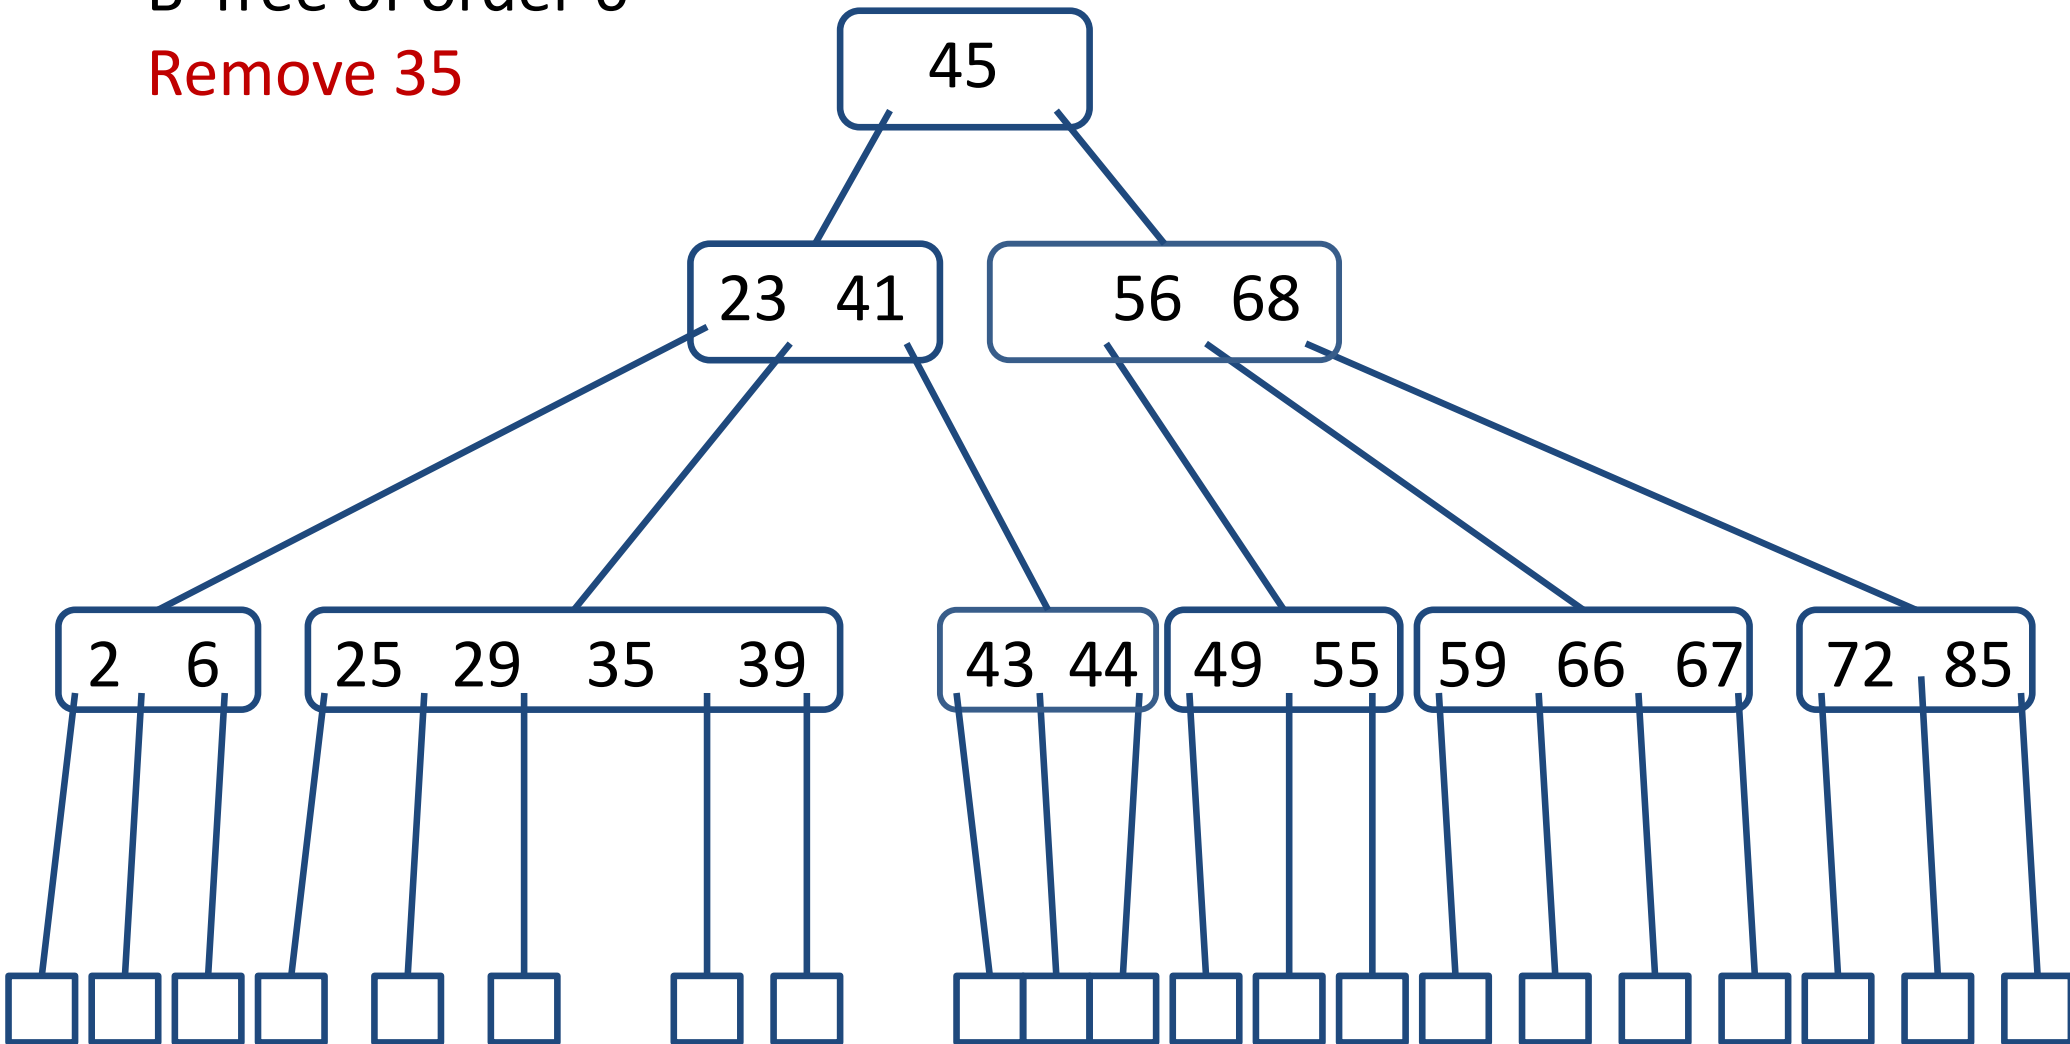
Remove 35



# B-Trees

B-Tree of order 6

Remove 35





**Algorithm** *remove*( $r, k$ )

**In:** Root  $r$  of a B-tree, key  $k$

**Out:** {remove data item with key  $k$  from the tree}

Find the node  $v$  storing key  $k$

Remove  $(k, o)$  from  $v$  replacing it with successor if needed

**while** node  $v$  *underflows* **do** {

**if**  $v$  is the root **then**

        make the first child of  $v$  the new root

**else if** a sibling has more than  $\lceil d/2 \rceil$  keys **then**

        perform a transfer operation

**else** {

        perform a fusion operation

$v \leftarrow$  parent of  $v$

    }

}

## Algorithm *remove*( $r, k$ )

**In:** Root  $r$  of a B-tree, key  $k$

**Out:** {remove data item with key  $k$  from the tree}

Find the node  $v$  storing key  $k$  }  $O(\log d \times \log_d n)$

Remove  $(k, o)$  from  $v$  replacing it with successor if needed

**while** node  $v$  *underflows* **do** {

**if**  $v$  is the root **then**

        make the first child of  $v$  the new root

**else if** a sibling has at least  $\lceil d/2 \rceil$  keys **then**

        perform a transfer operation

**else** {

        perform a fusion operation

$v \leftarrow$  parent of  $v$

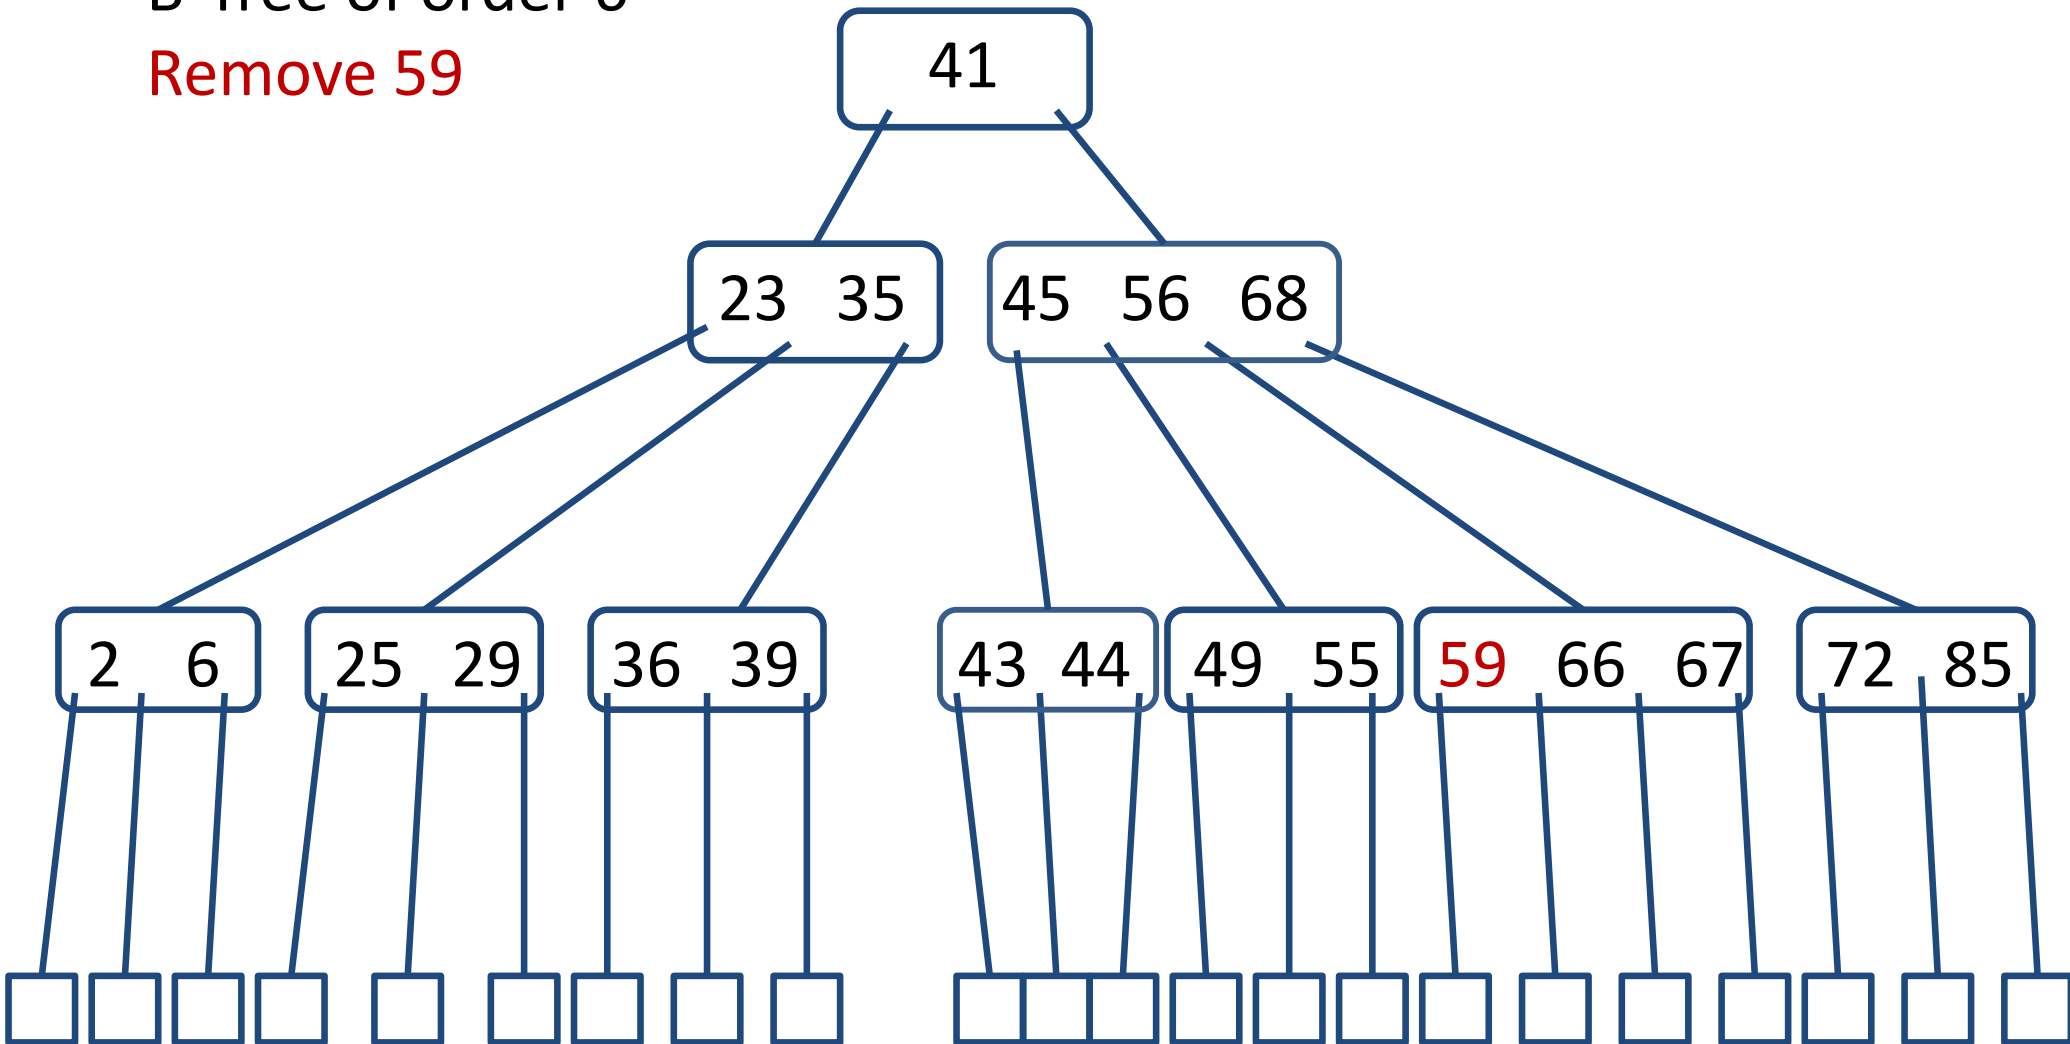
    }

}

# B-Trees

B-Tree of order 6

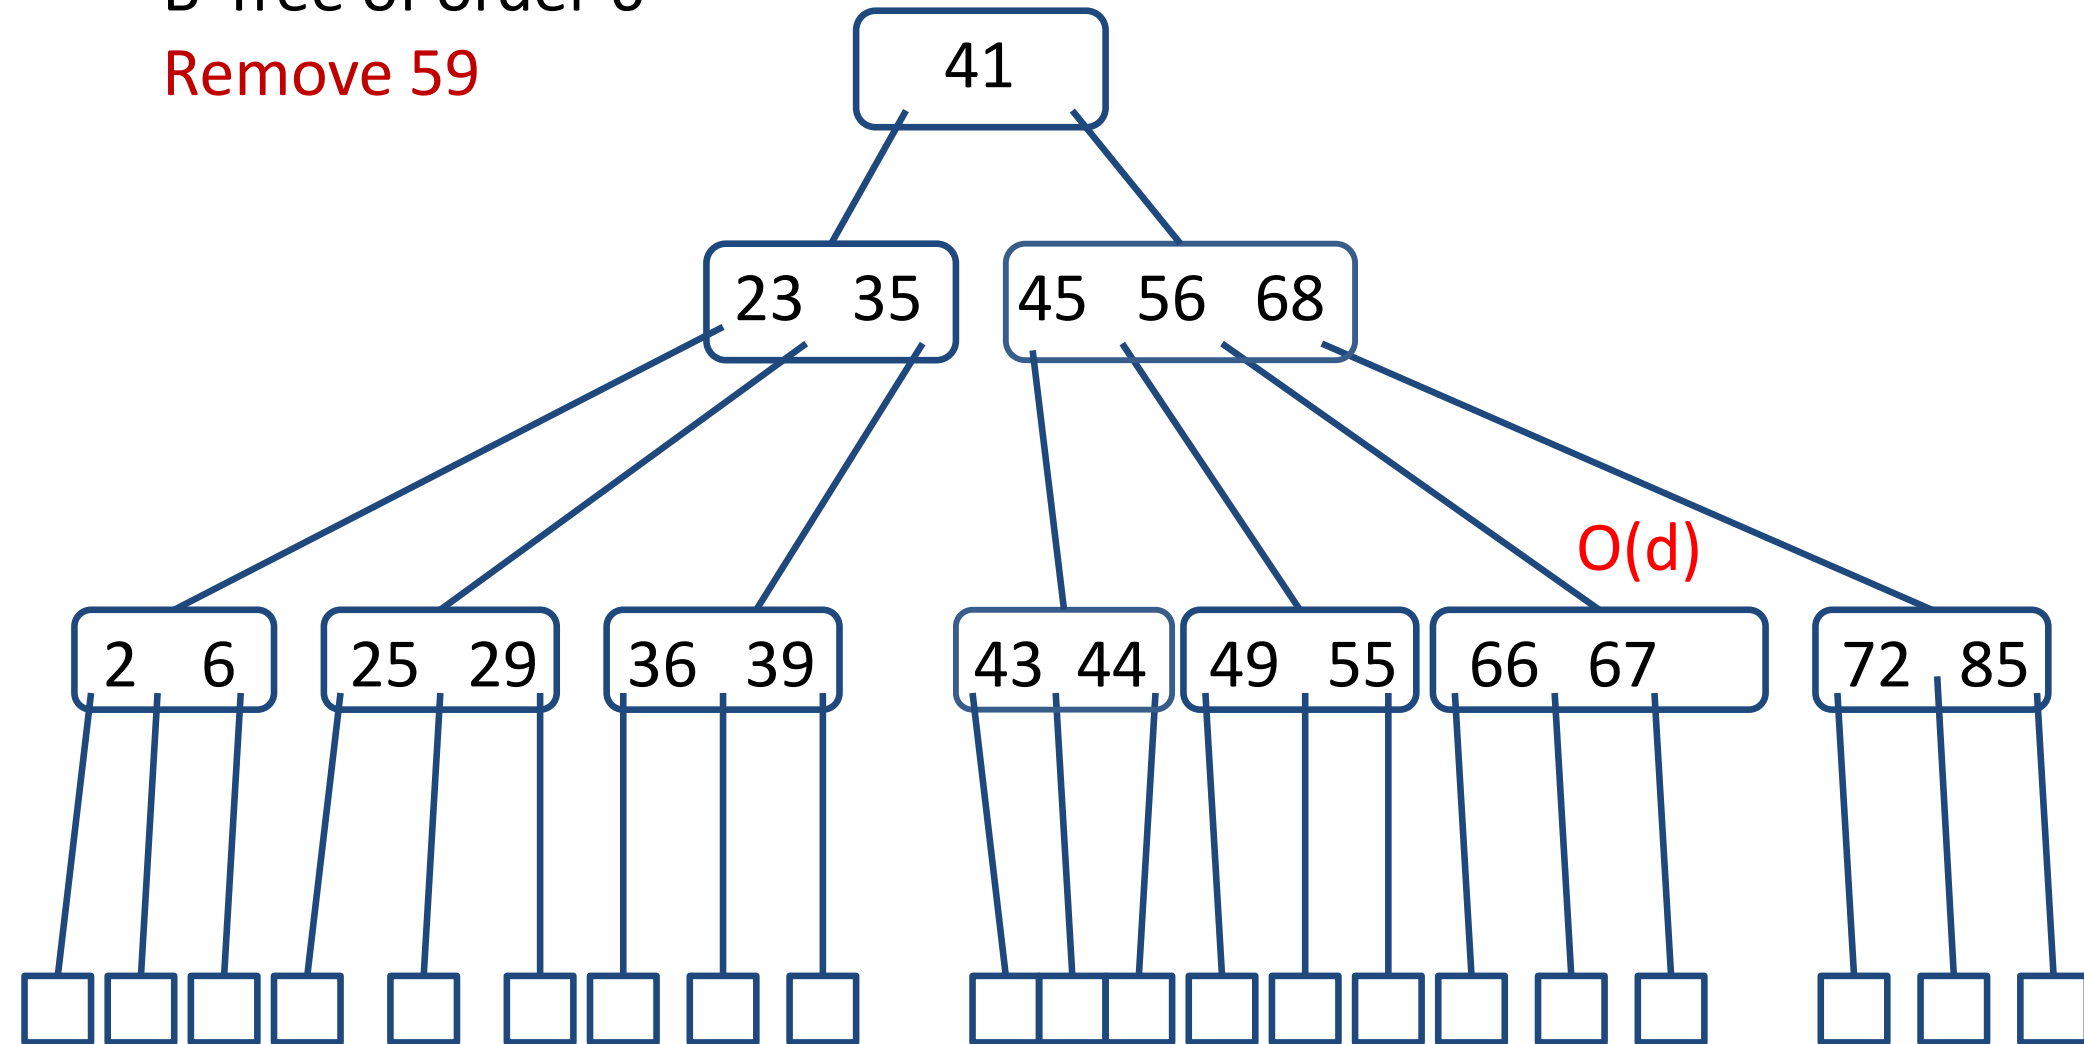
Remove 59



# B-Trees

B-Tree of order 6

Remove 59



## Algorithm *remove(r,k)*

**In:** Root  $r$  of a B-tree, key  $k$

**Out:** {remove data item with key  $k$  from the tree}

Find the node  $v$  storing key  $k$  }  $O(\log d \times \log_d n)$

**Remove** ( $k, o$ ) **from**  $v$  replacing it with successor if needed }

**while** node  $v$  *underflows* **do** {  $O(d)$

**if**  $v$  is the root **then**

        make the first child of  $v$  the new root

**else if** a sibling has at least  $\lceil d/2 \rceil$  keys **then**

        perform a transfer operation

**else** {

        perform a fusion operation

$v \leftarrow$  parent of  $v$

    }

}

**Algorithm** *remove*( $r, k$ )

**In:** Root  $r$  of a B-tree, key  $k$

**Out:** {remove data item with key  $k$  from the tree}

Find the node  $v$  storing key  $k$  }  $O(\log d \times \log_d n)$

Remove  $(k, o)$  from  $v$  **replacing it with successor** if needed }

**while** node  $v$  *underflows* **do** {  $O(d + \log d \times \log_d n)$

**if**  $v$  is the root then

        make the first child of  $v$  the new root

**else if** a sibling has at least  $\lceil d/2 \rceil$  keys **then**

        perform a transfer operation

**else** {

        perform a fusion operation

$v \leftarrow$  parent of  $v$

    }

}

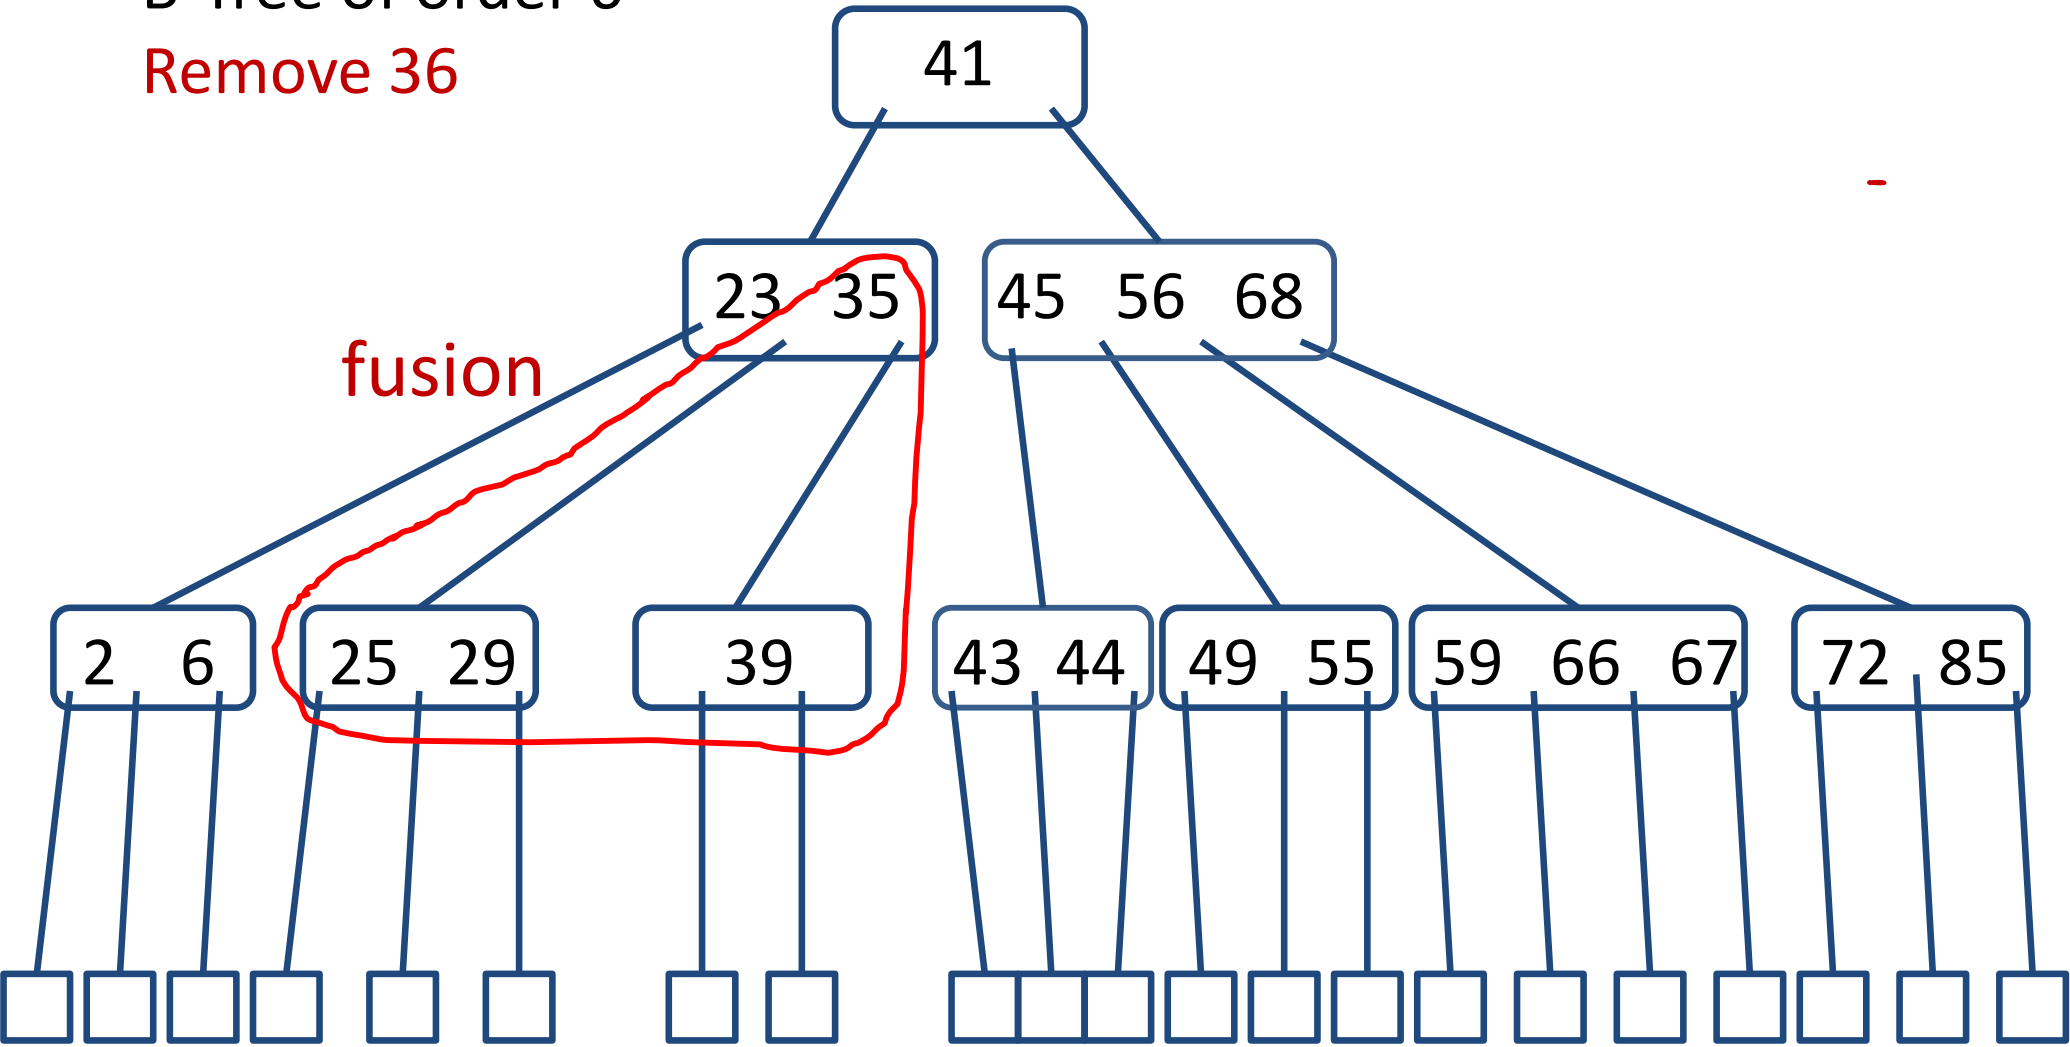
# B-Trees

B-Tree of order 6

Remove 36

fusion

-

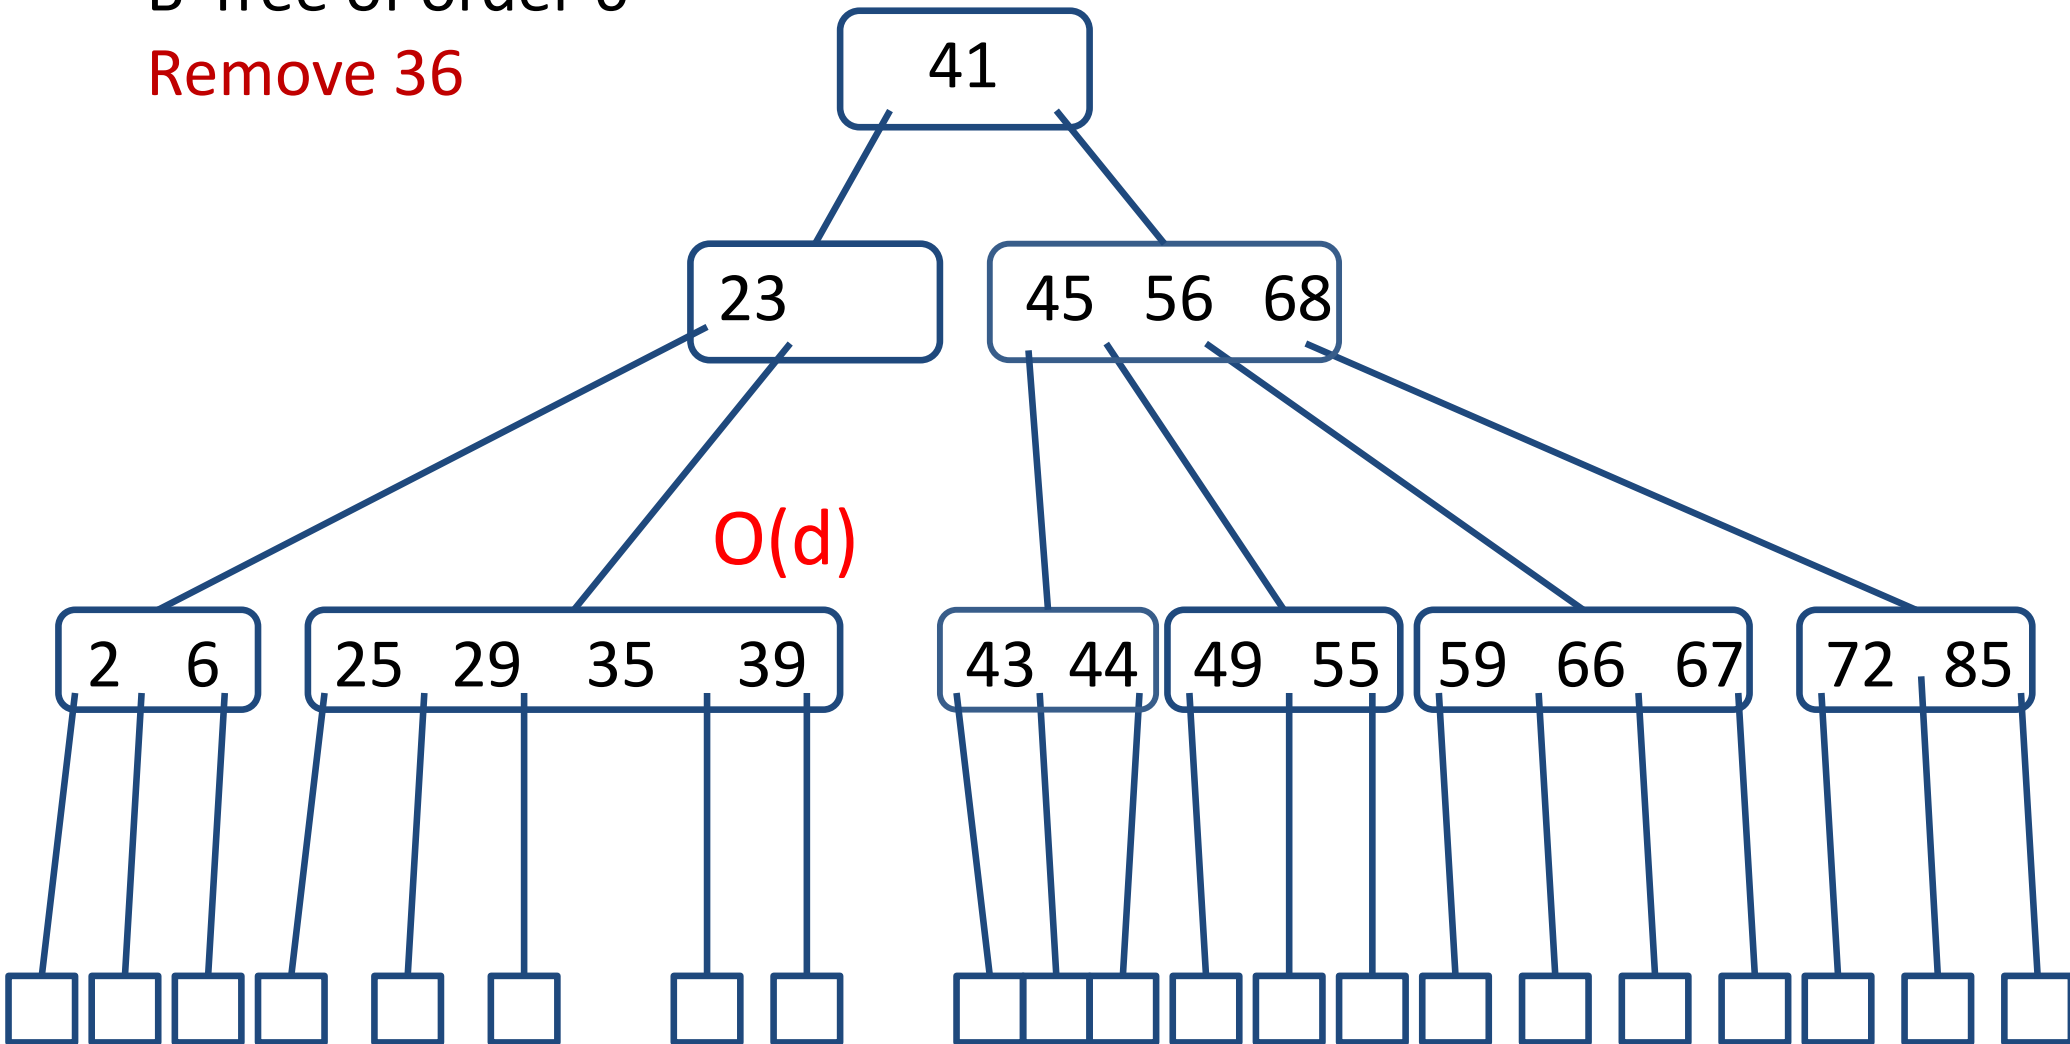


# B-Trees

B-Tree of order 6

Remove 36

$O(d)$

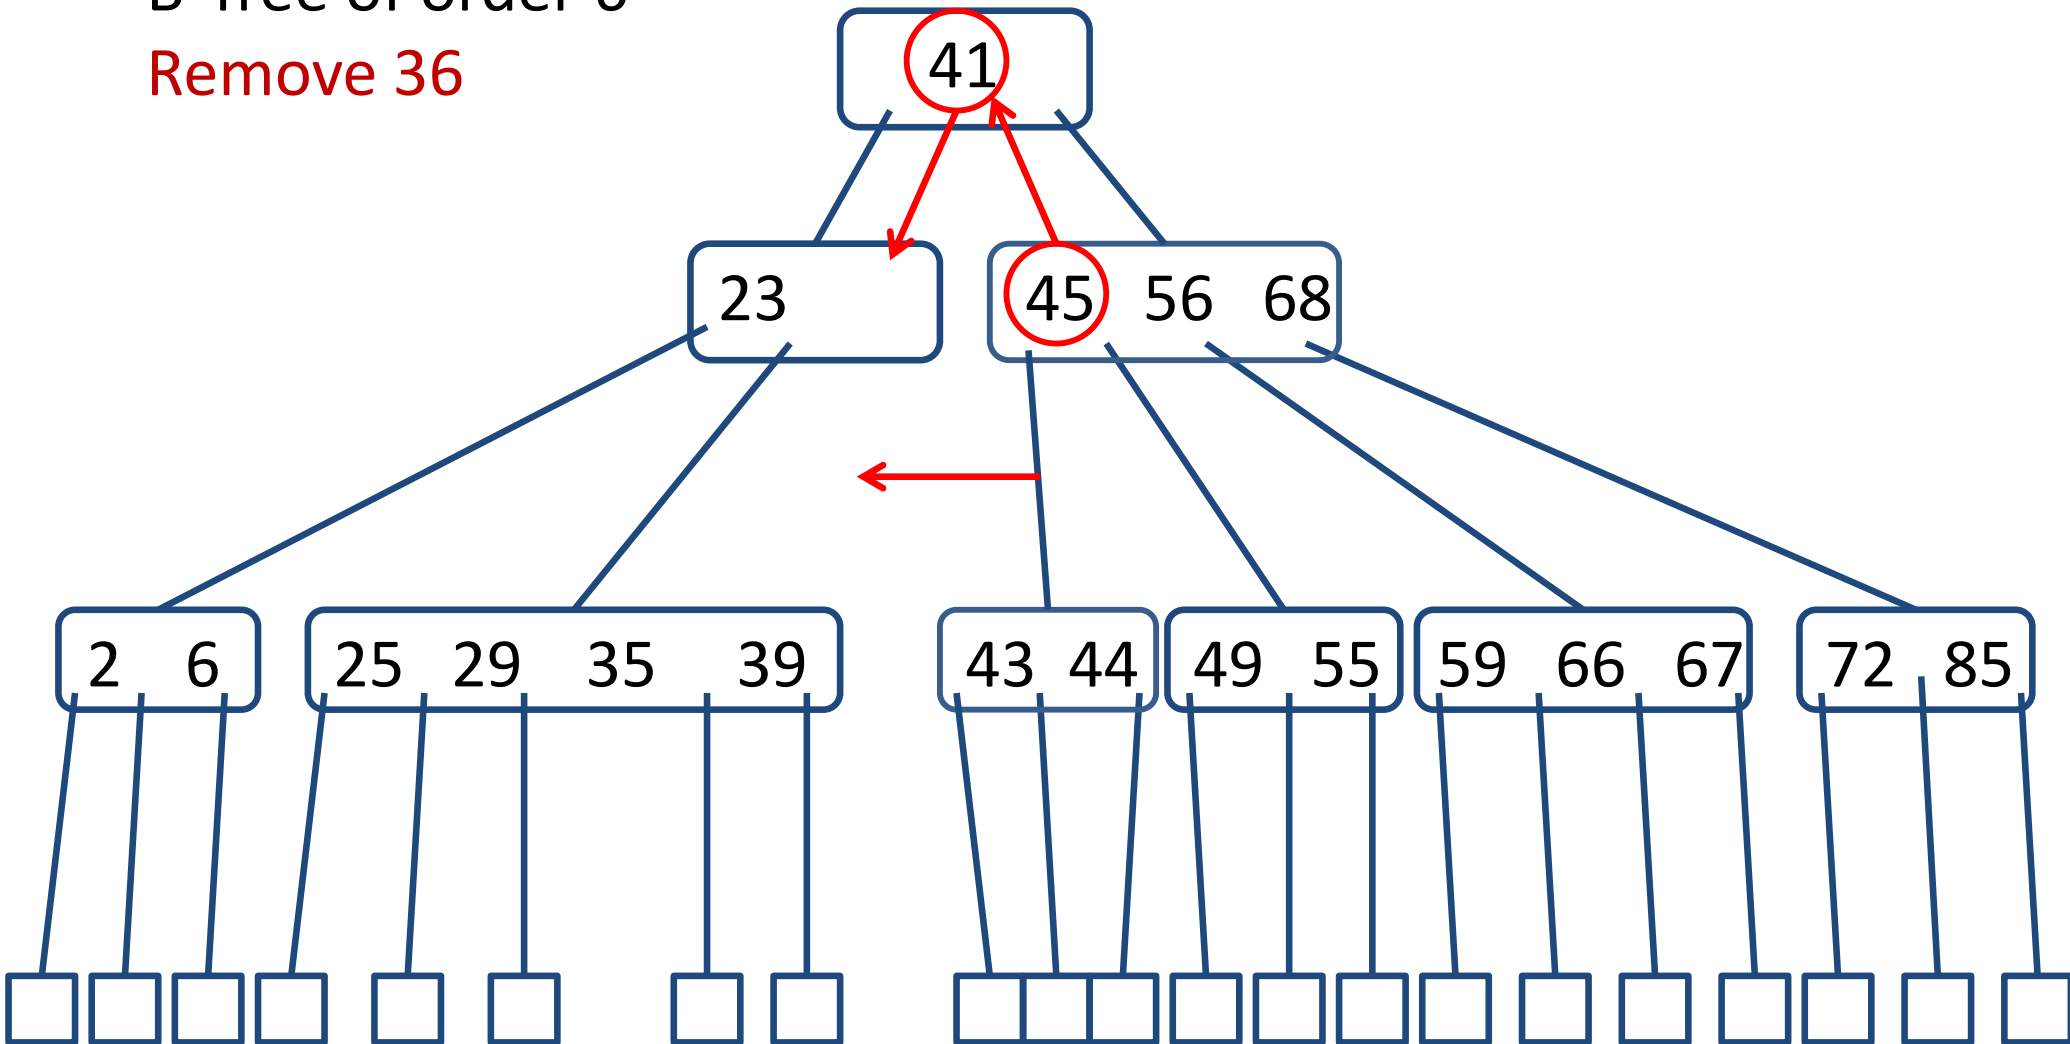




# B-Trees

B-Tree of order 6

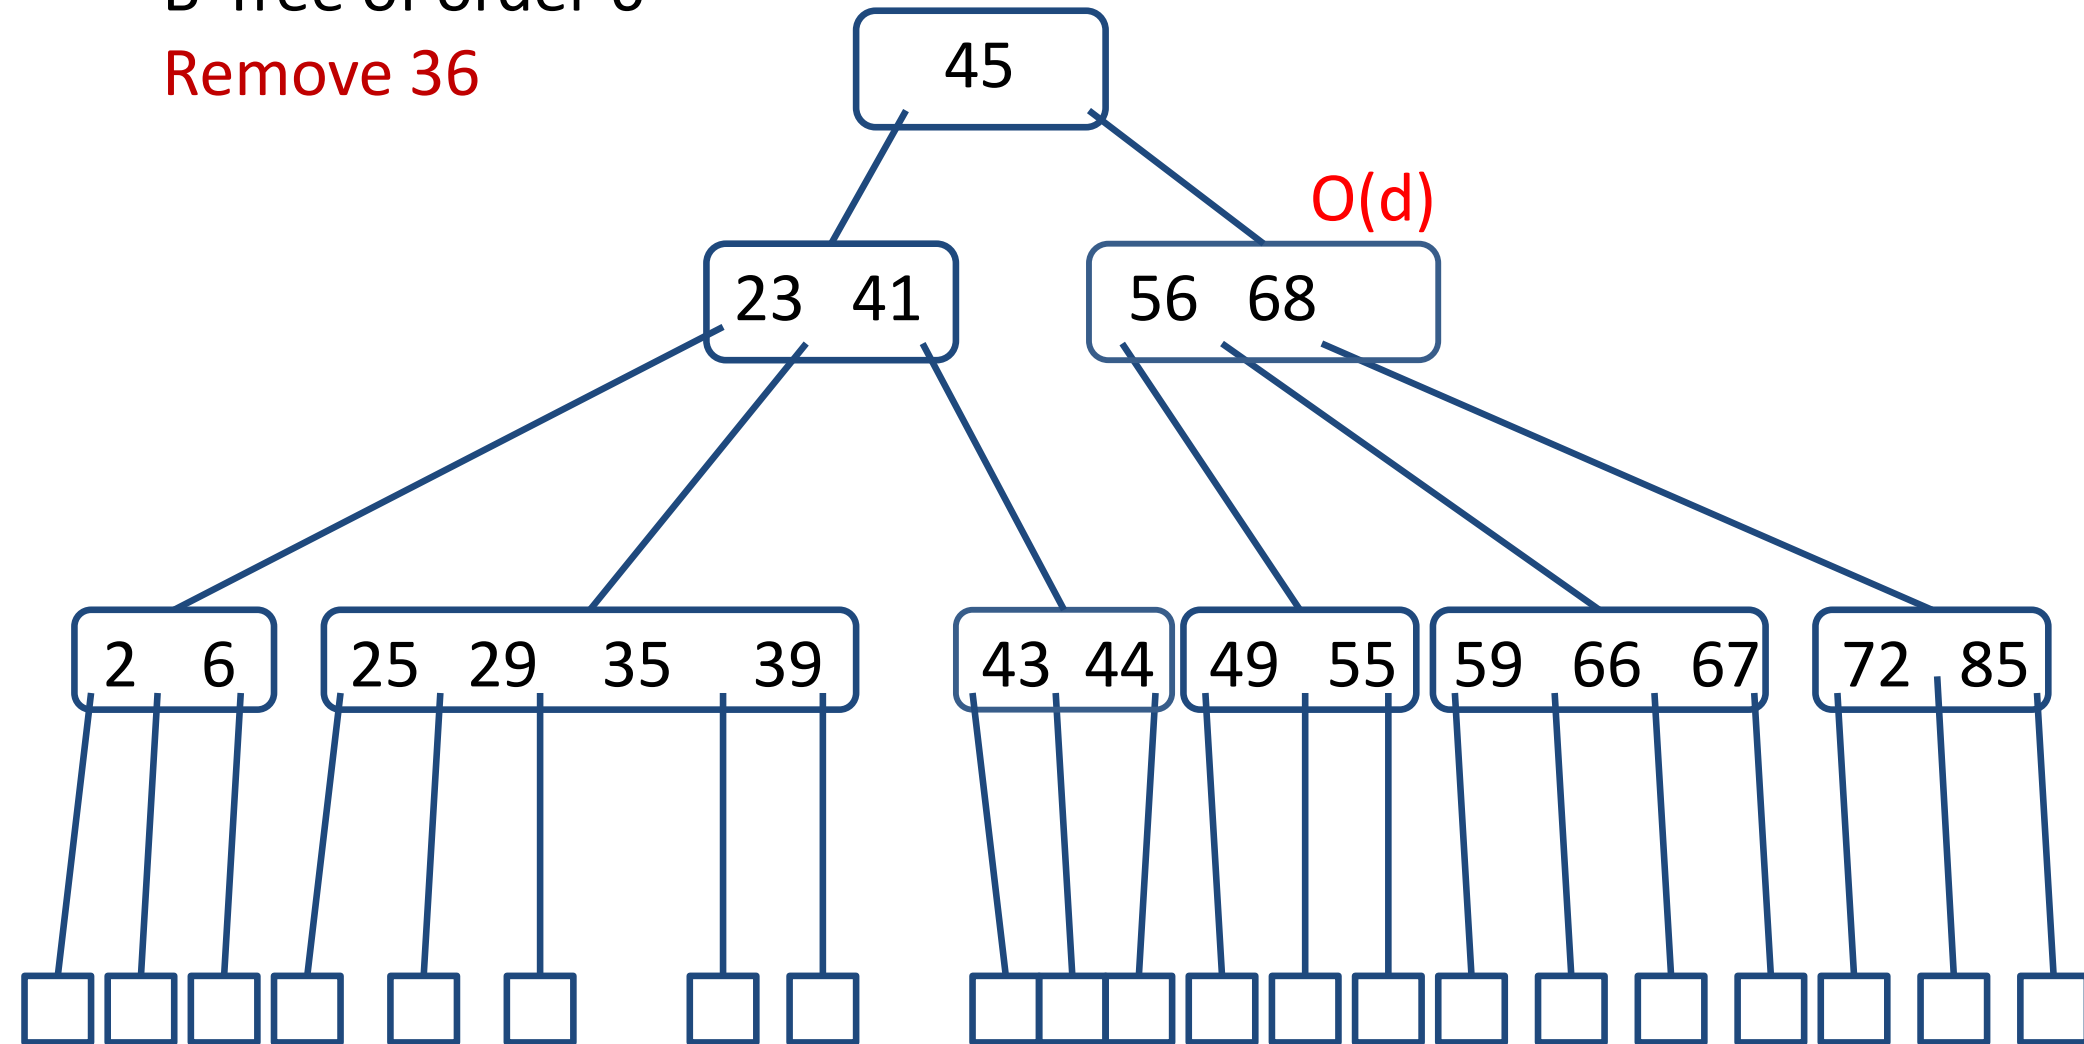
Remove 36



# B-Trees

B-Tree of order 6

Remove 36



$O(d)$

## Algorithm *remove*( $r, k$ )

**In:** Root  $r$  of a B-tree, key  $k$

**Out:** {remove data item with key  $k$  from the tree}

Find the node  $v$  storing key  $k$

}  $O(\log d \times \log_d n)$

Remove  $(k, o)$  from  $v$  replacing it with successor if needed

**while** node  $v$  *underflows* **do** {

$O(d + \log d \times \log_d n)$

**if**  $v$  is the root **then**

        make the first child of  $v$  the new root

**else if** a sibling has at least  $\lceil d/2 \rceil$  keys **then**

        perform a transfer operation

**else** {

        perform a fusion operation

$v \leftarrow$  parent of  $v$

    }

}  $O(d)$

}

**Algorithm** *remove*( $r, k$ )      Time complexity  $O(d \log_d n)$

**In:** Root  $r$  of a B-tree, key  $k$

**Out:** {remove data item with key  $k$  from the tree}

Find the node  $v$  storing key  $k$  }  $O(\log d \times \log_d n)$

Remove  $(k, o)$  from  $v$  replacing it with successor if needed }

**while** node  $v$  *underflows* **do** {  $O(d + \log d \times \log_d n)$

**if**  $v$  is the root **then**

        make the first child of  $v$  the new root

**else if** a sibling has at least  $\lceil d/2 \rceil$  keys **then**

        perform a transfer operation

**else** {

        perform a fusion operation


$v \leftarrow$  parent of  $v$

    }

}


Collection of  $10^9$  records

AVL tree



Height:  $O(\log_2 n)$   
 $\log_2 10^9 \approx 30$


B-tree of degree  $10^3$



Height:  $O(\log_d n)$   
 $\log_{1000} 10^9 = 3$

Collection of  $10^9$  records

AVL tree




Height:  $O(\log_2 n)$   
 $\log_2 10^9 \approx 30$

Number of comparisons  
to find a key

$\approx 30$

B-tree of degree  $10^3$

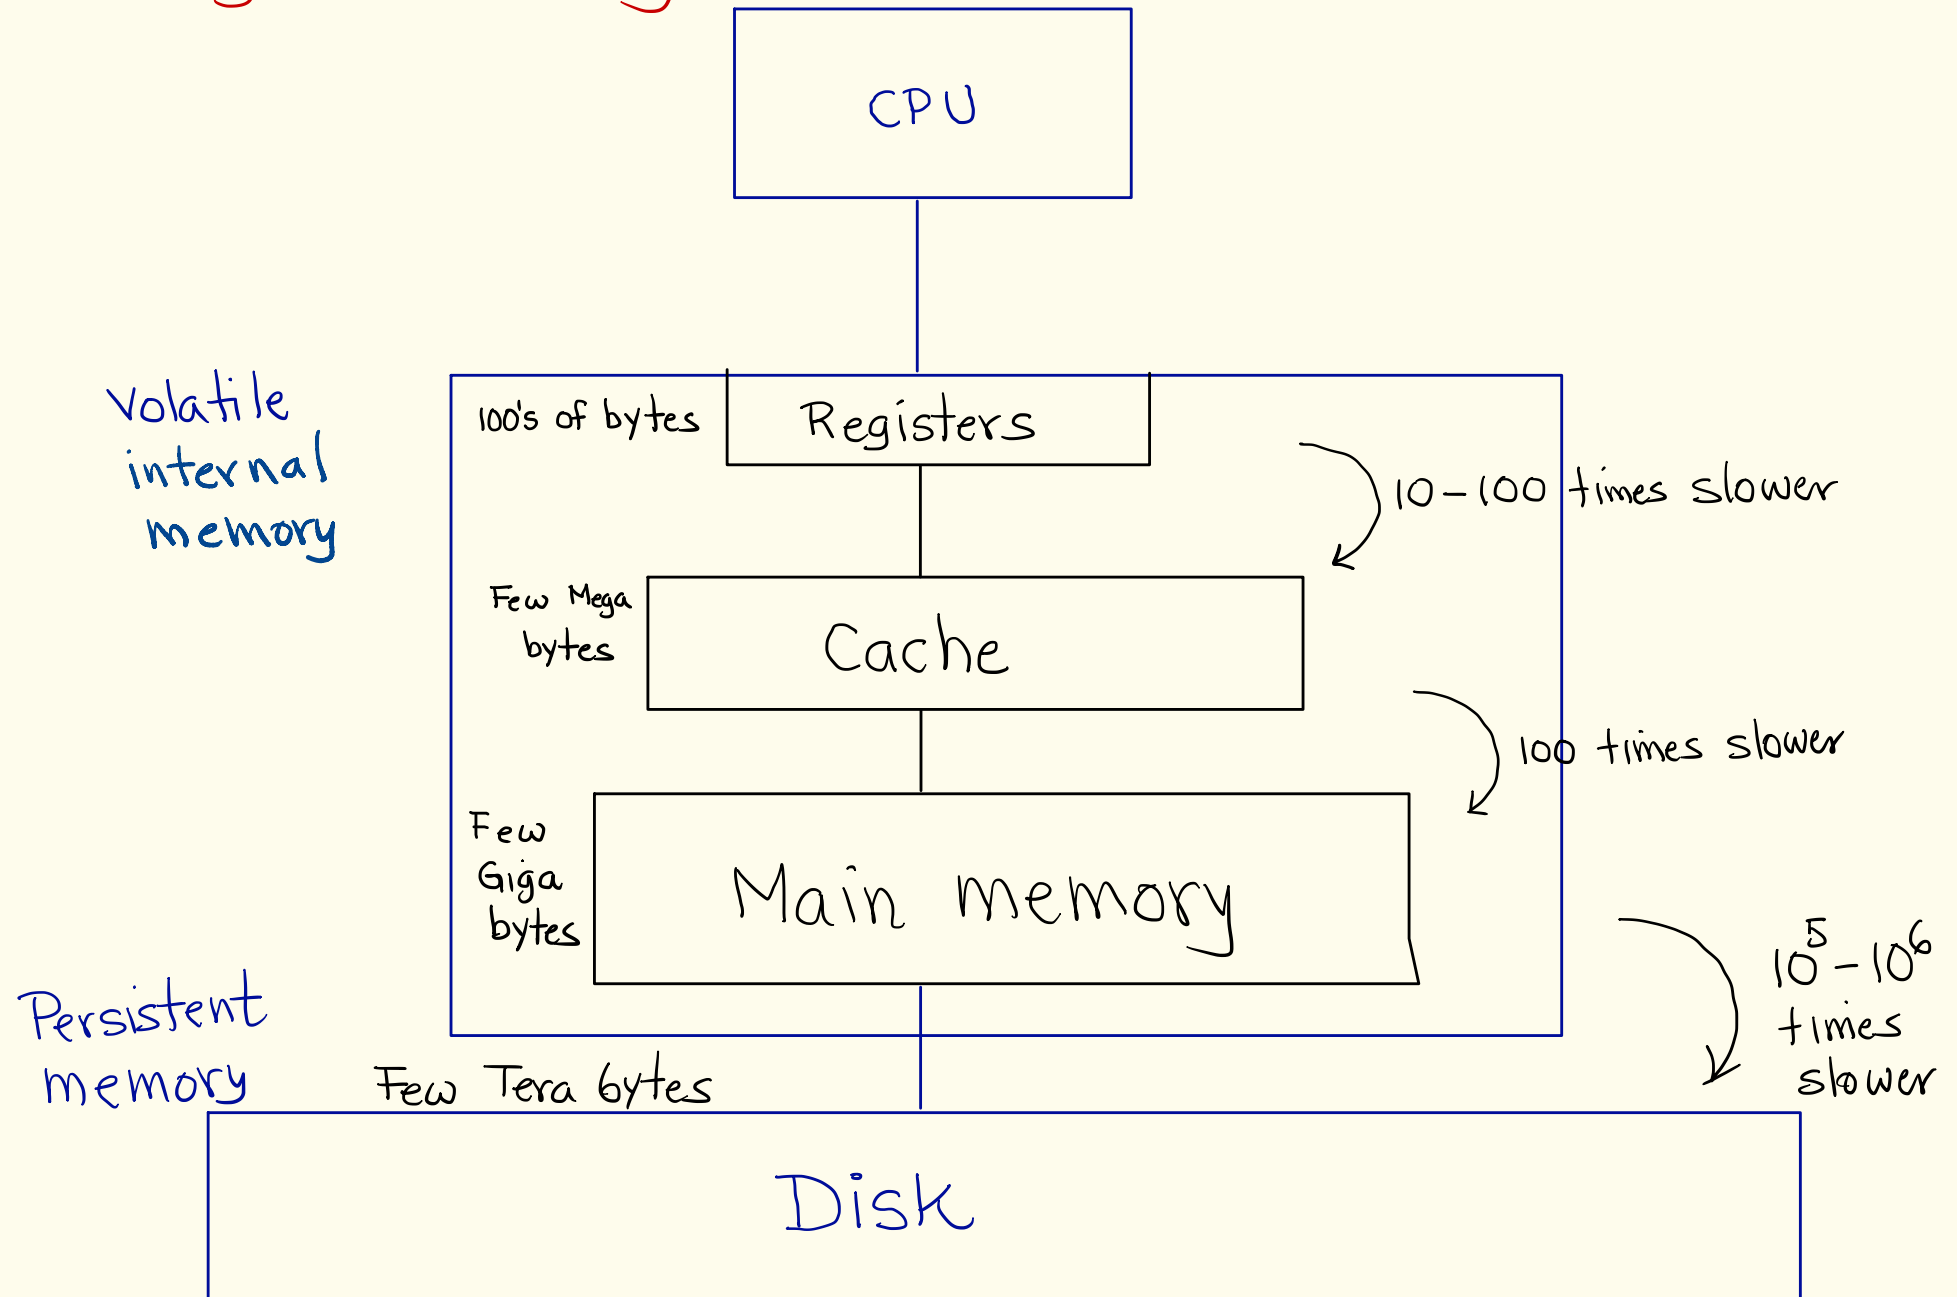
Binary search:  $O(\log_2 d)$   
 $\log_2 10^3 \approx 10$



Height:  $O(\log_d n)$   
 $\log_{1000} 10^9 = 3$


$\approx 30$

# Memory Hierarchy



Collection of  $10^9$  records

AVL tree




$$\text{Height: } O(\log_2 n)$$
$$\log_2 10^9 \approx 30$$

Number of disk accesses: 30

AVL

B-tree of degree  $10^3$



$$\text{Height: } O(\log_a n)$$
$$\log_{1000} 10^9 = 3$$

B-tree

3

So a B-tree is much better.



# Disk Blocks

- Consider the problem of maintaining a large collection of items that does not fit in main memory, such as a typical database.
- In this context, we refer to the external memory is divided into blocks, which we call **disk blocks**.
- The transfer of a block between external memory and primary memory is a **disk transfer** or **I/O**.
- There is a great time difference that exists between main memory accesses and disk accesses
- Thus, we want to minimize the number of disk transfers needed to perform a query or update. We refer to this count as the **I/O complexity** of the algorithm involved.

# Memory Hierarchies

- Computers have a hierarchy of different kinds of memories, which vary in terms of their size and distance from the CPU.
- Closest to the CPU are the internal **registers**. Access to such locations is very fast, but there are relatively few such locations.
- At the second level in the hierarchy are the memory **caches**.
- At the third level in the hierarchy is the **internal memory**, which is also known as main memory or core memory.
- Another level in the hierarchy is the **external memory**, which usually consists of disks.

