

B-Trees


A B-tree of order d is a multiway search tree with the following properties:

- The root has at least 2 children and at most d .
- All internal nodes other than the root have at least $\left\lceil \frac{d}{2} \right\rceil$ and at most d children
- All the leaves are at the same level

B-Trees

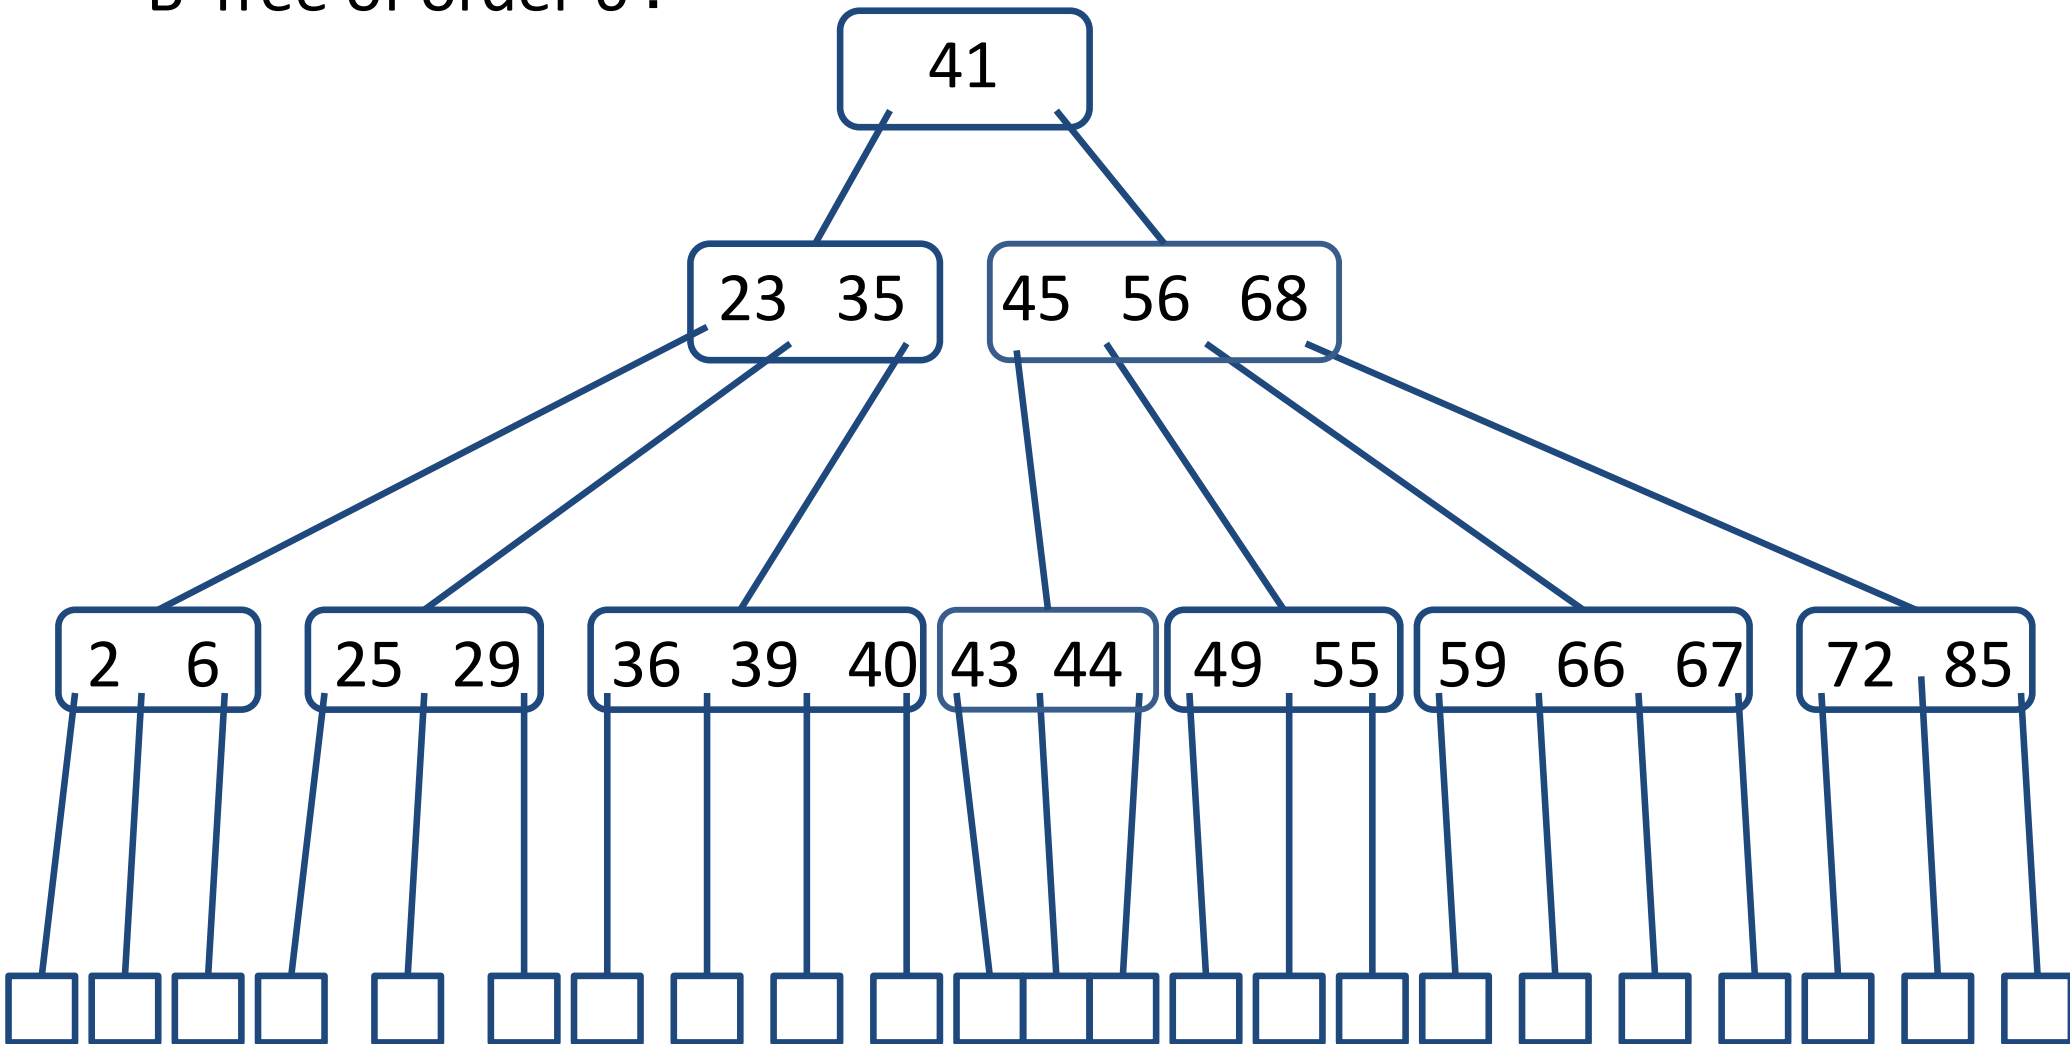
B-Tree of order 6?

min # children = 3



B-Trees

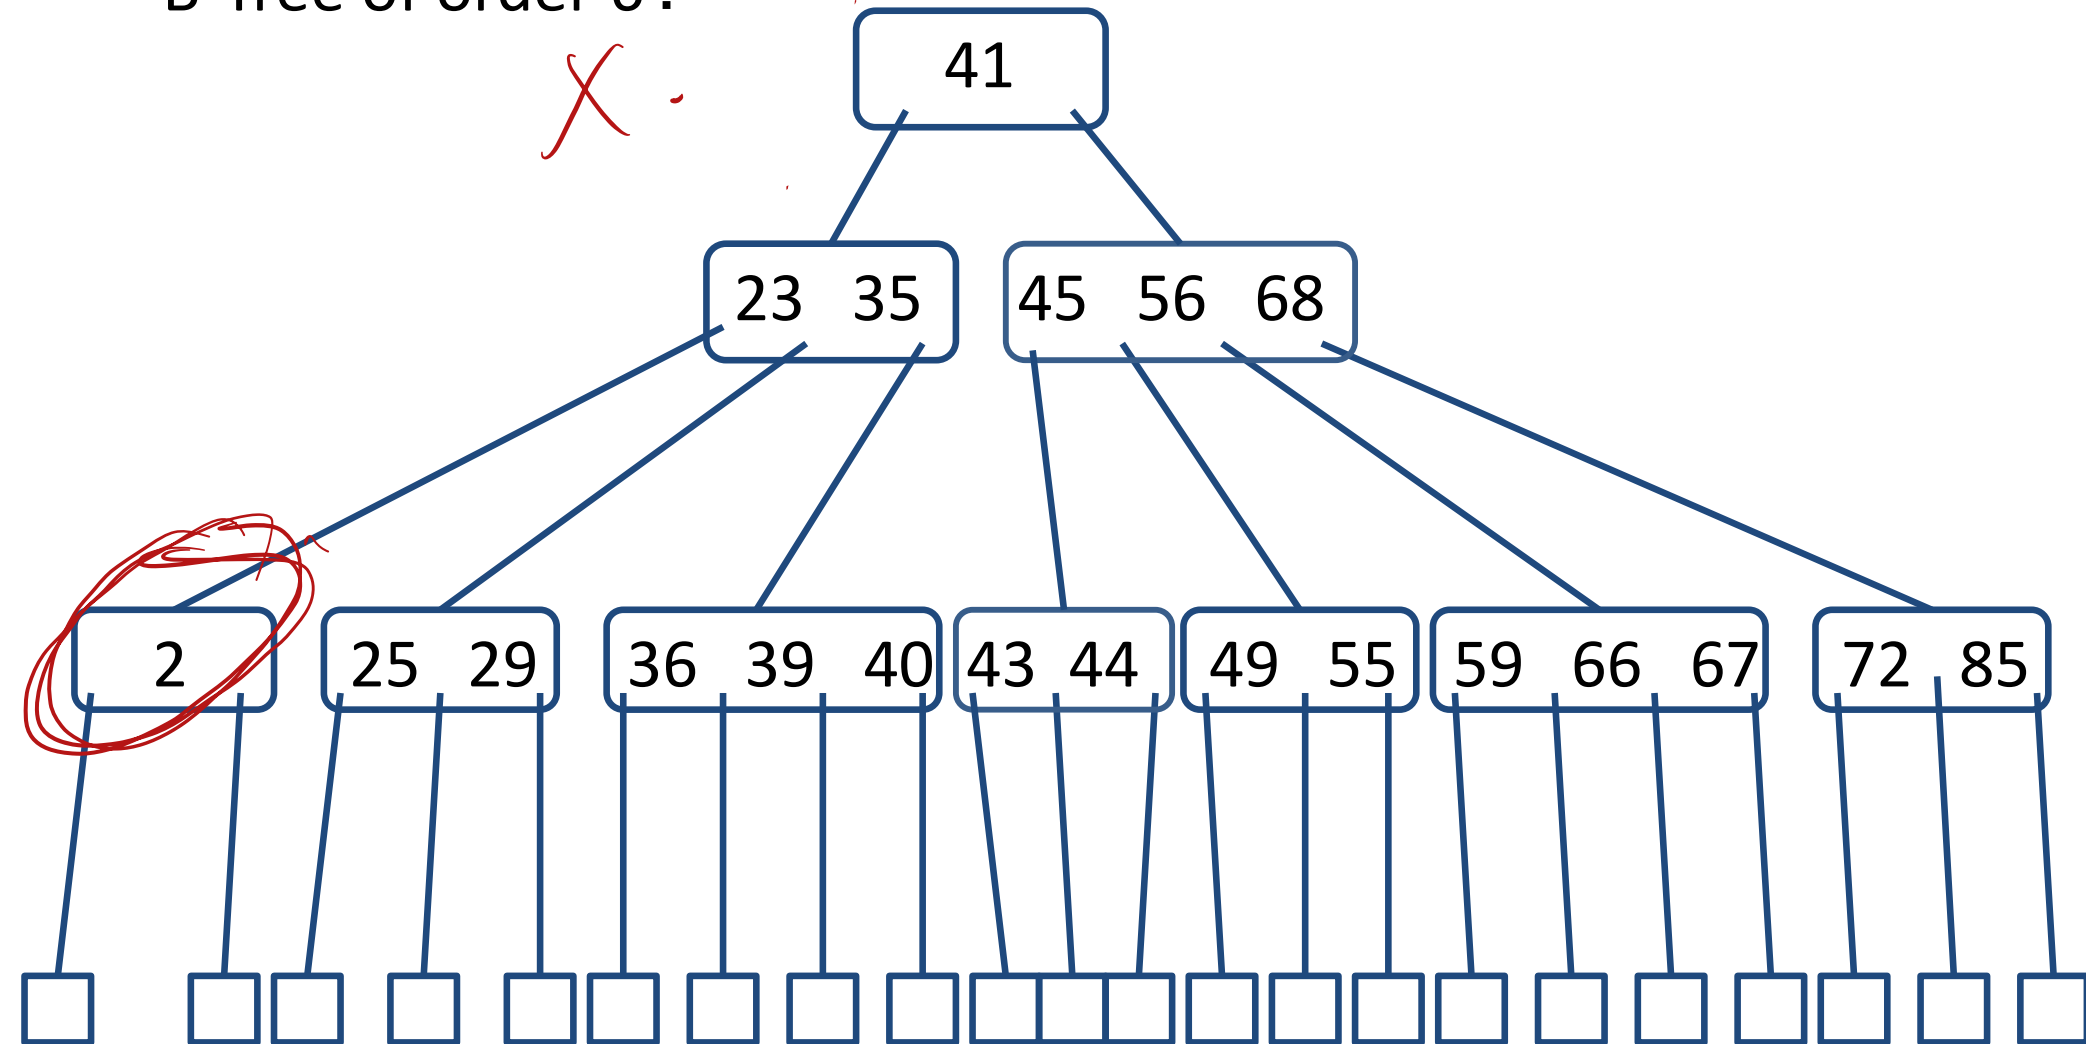
B-Tree of order 6?



B-Trees

B-Tree of order 6?

X.



What is the Maximum Height of a B-Tree?

Height is $O(\log_d n)$

\Rightarrow This may be asked
in the Final.

$$1 + \frac{d}{2} + \left(\frac{d}{2}\right)^2 + \dots + \left(\frac{d}{2}\right)^h = n.$$


$$n = \frac{\left(\frac{d}{2}\right)^h - 1}{\frac{d}{2} - 1} \cdot 1$$

$$h = \left\lceil \log_{\frac{d}{2}} \left(\frac{d}{2} - 1\right) n \right\rceil + 1 \Rightarrow O(h(n)) = O(\log_d n)$$

B-Trees

B-Tree of order 6

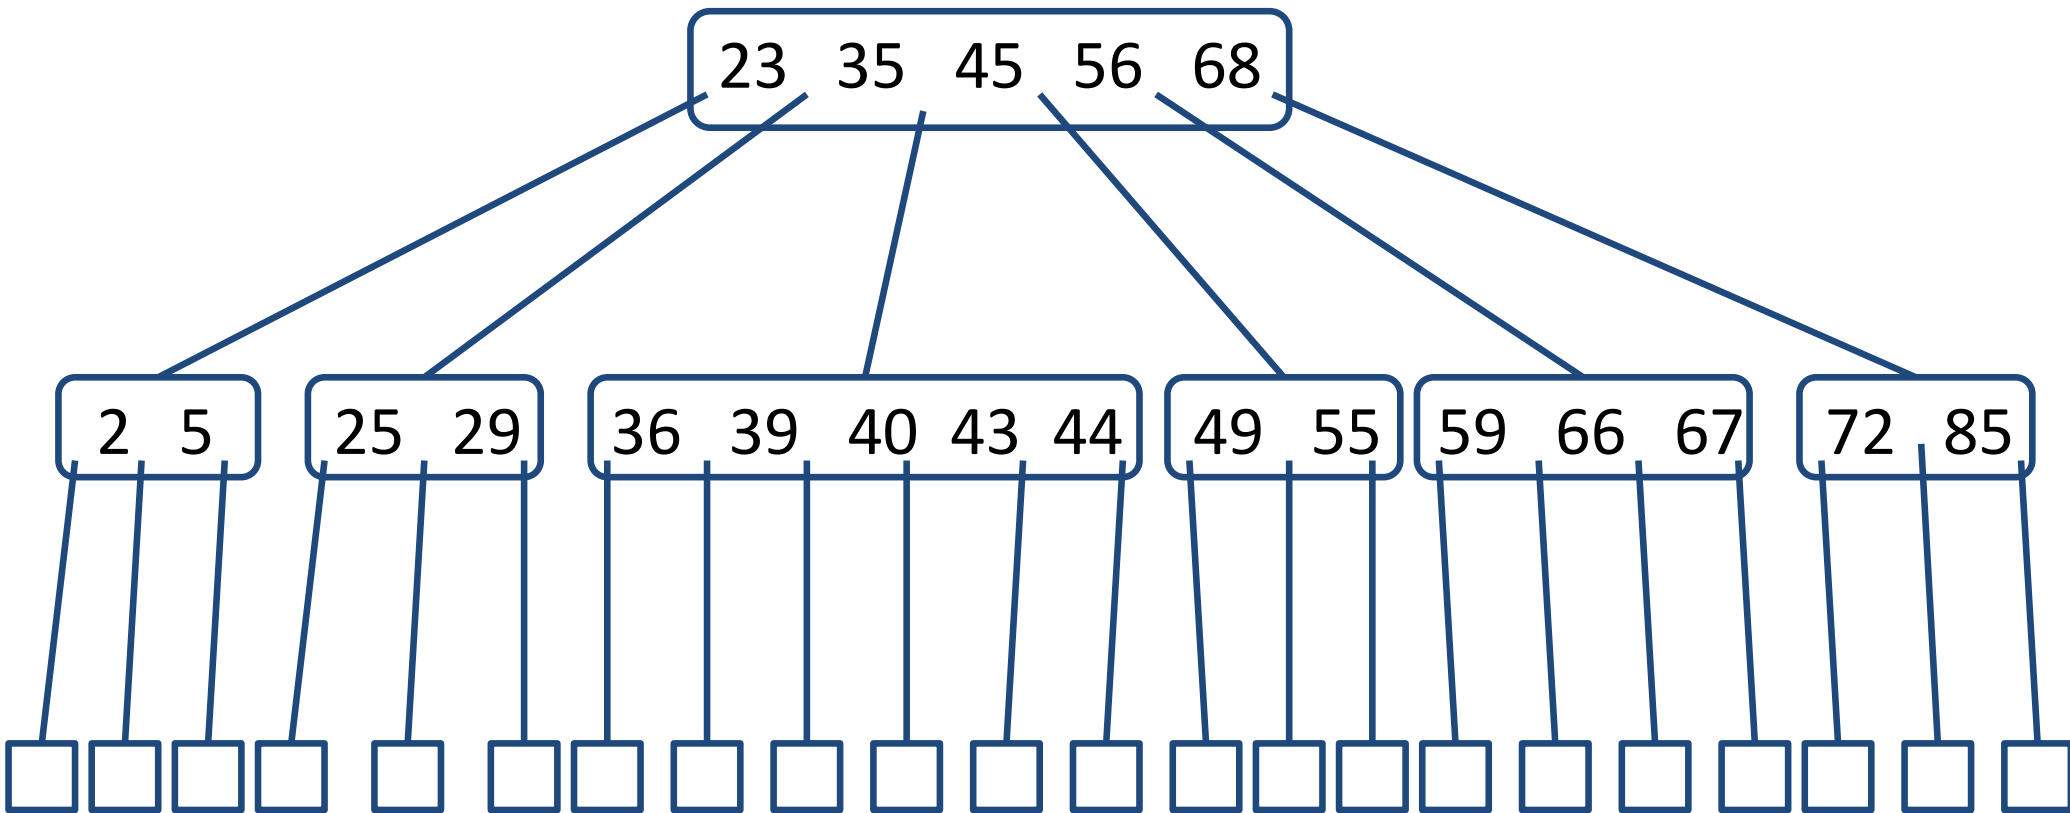
$3 \leq \# \text{ children} \leq 6$



B-Trees

B-Tree of order 6

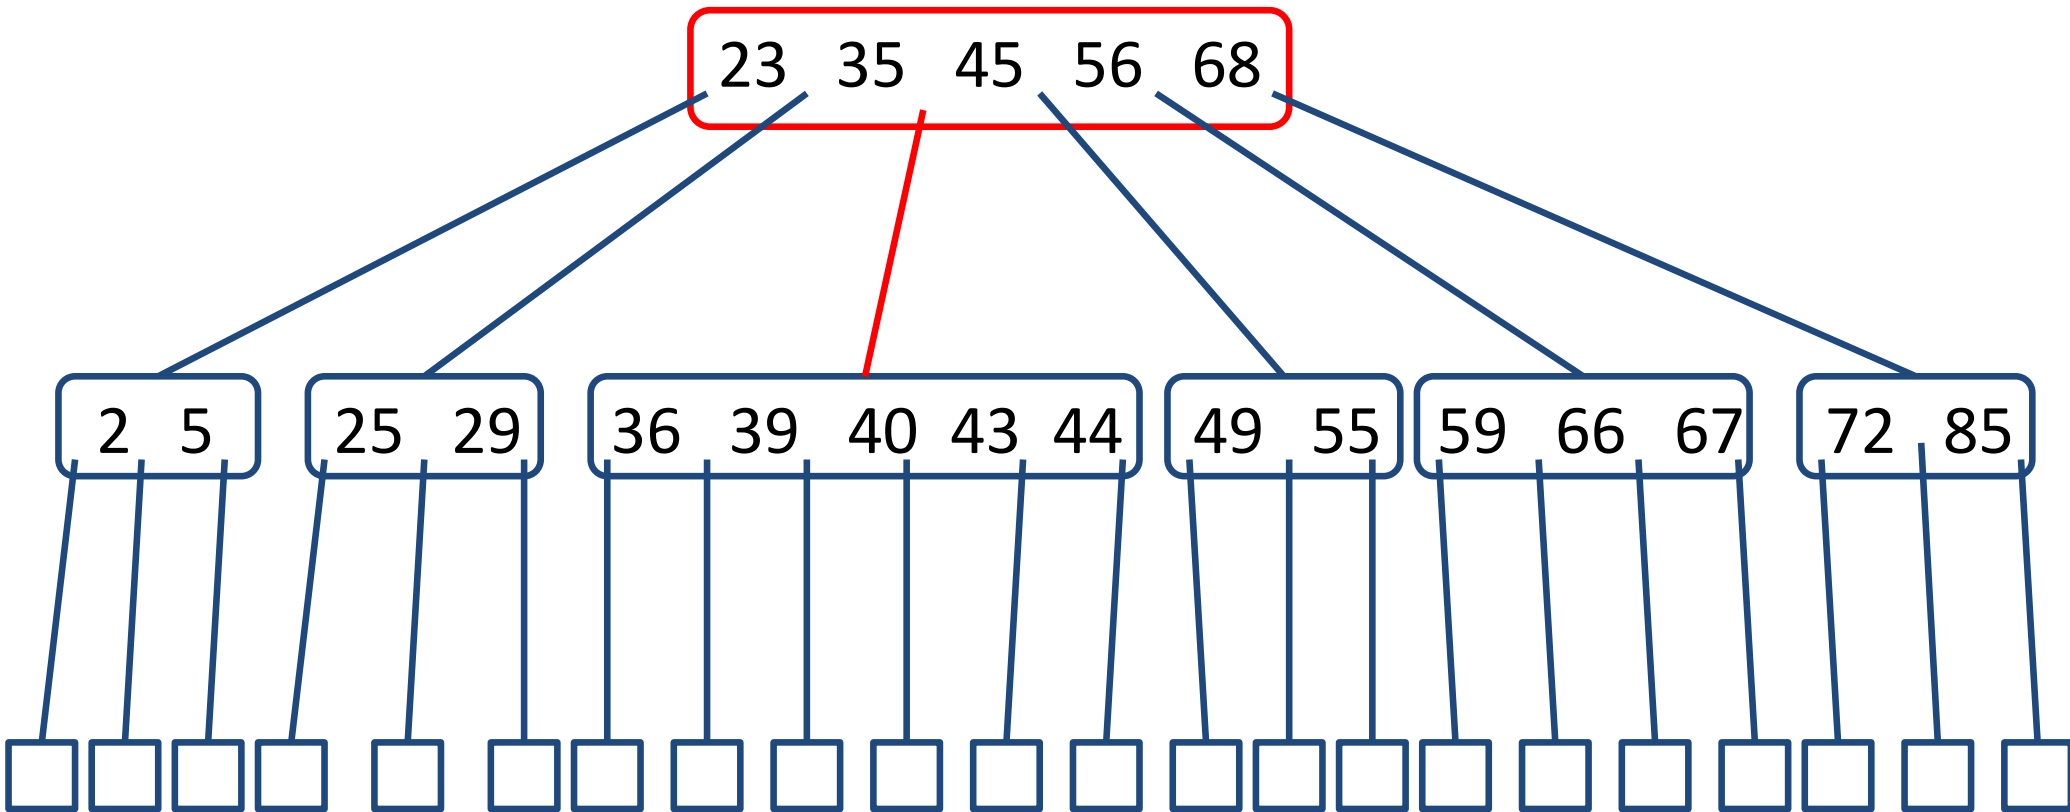
Put 38



B-Trees

B-Tree of order 6

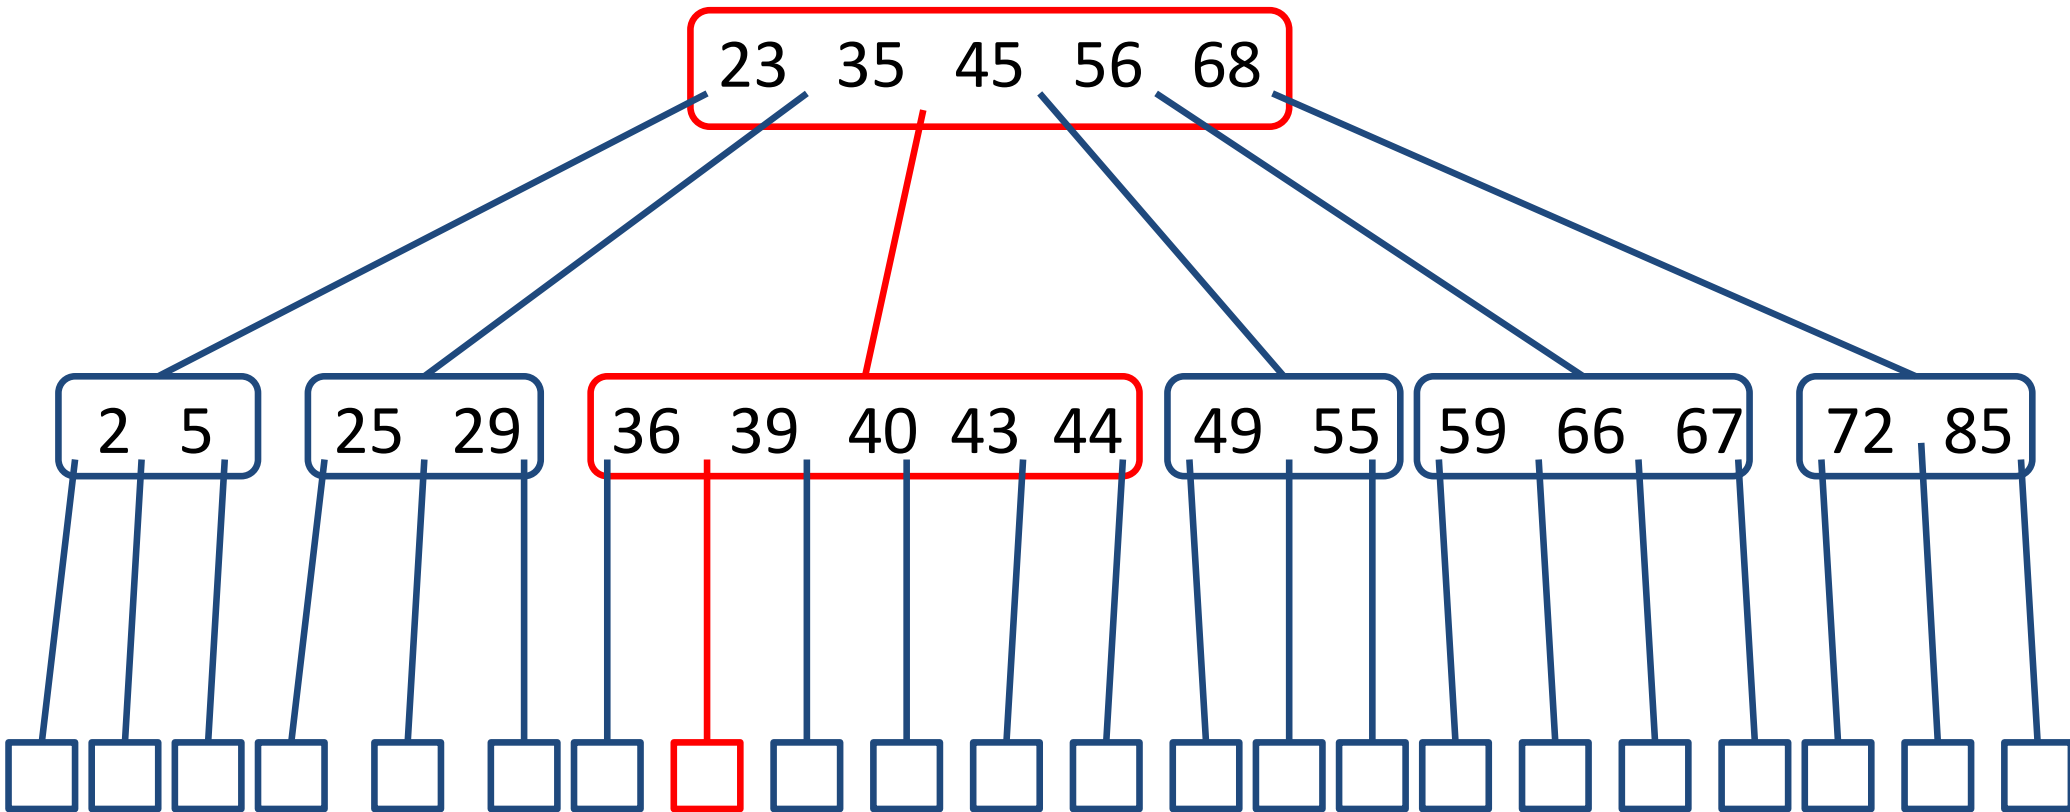
Put 38



B-Trees

B-Tree of order 6

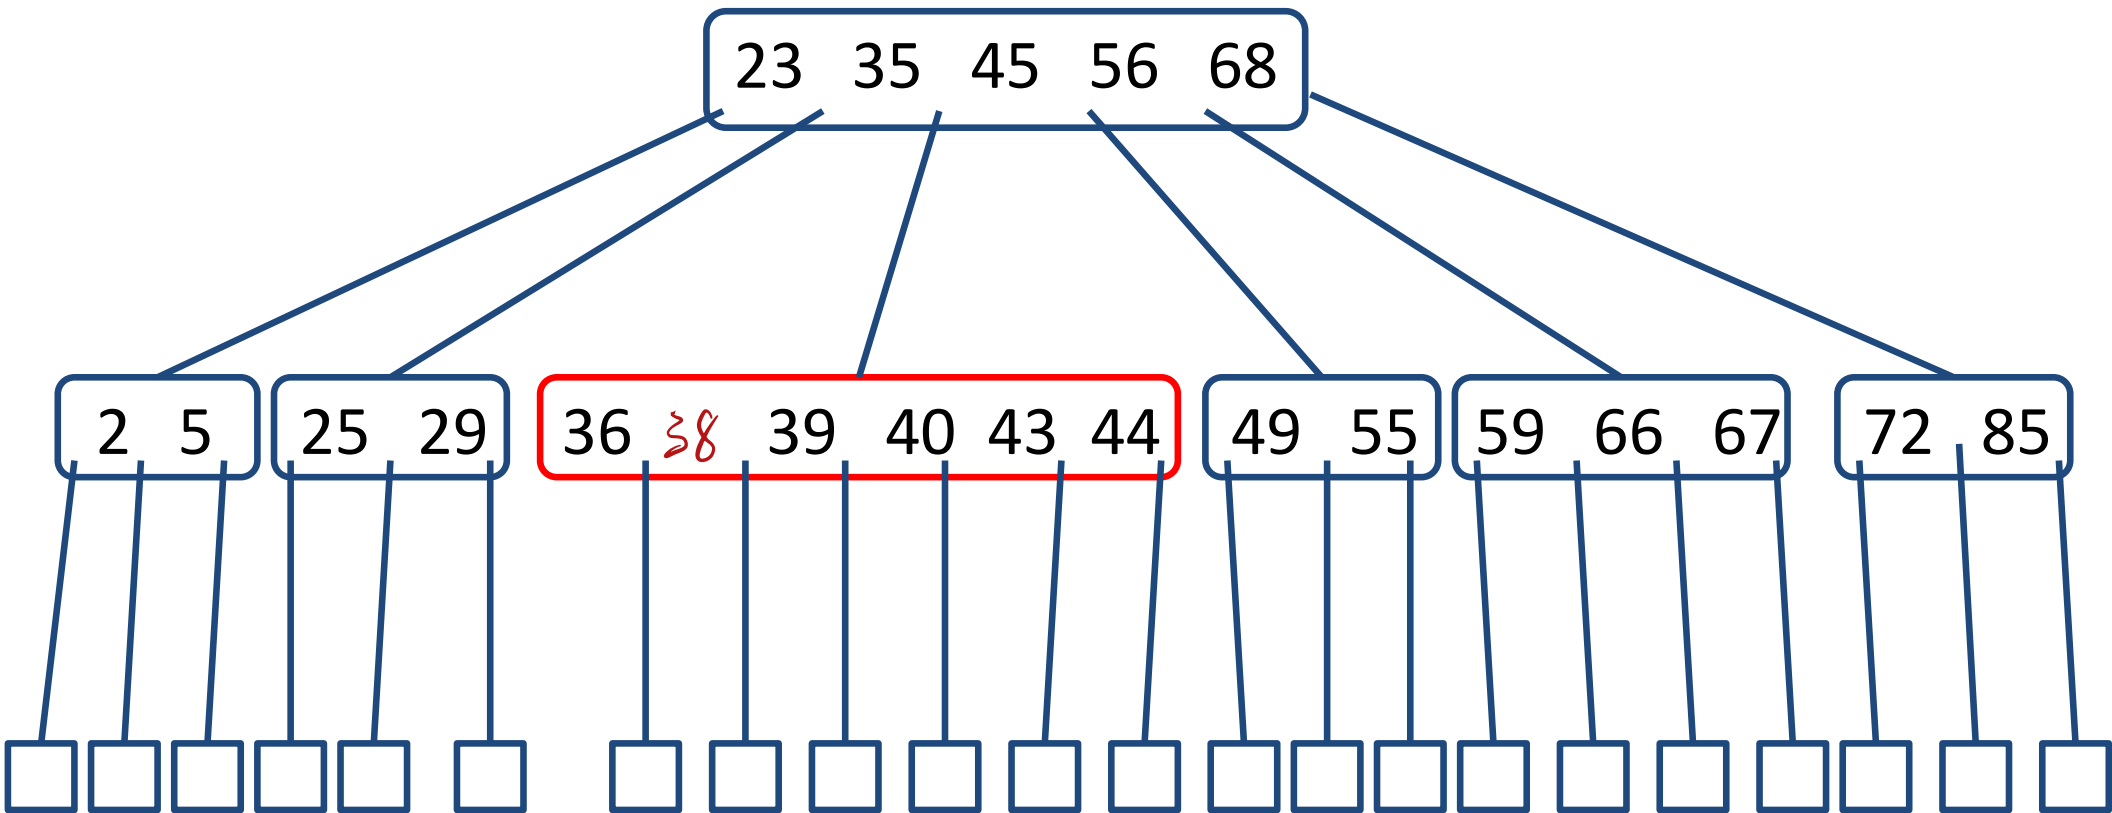
Put 38



B-Trees

B-Tree of order 6


Put 38



B-Trees

B-Tree of order 6


Put 38



B-Trees

B-Tree of order 6

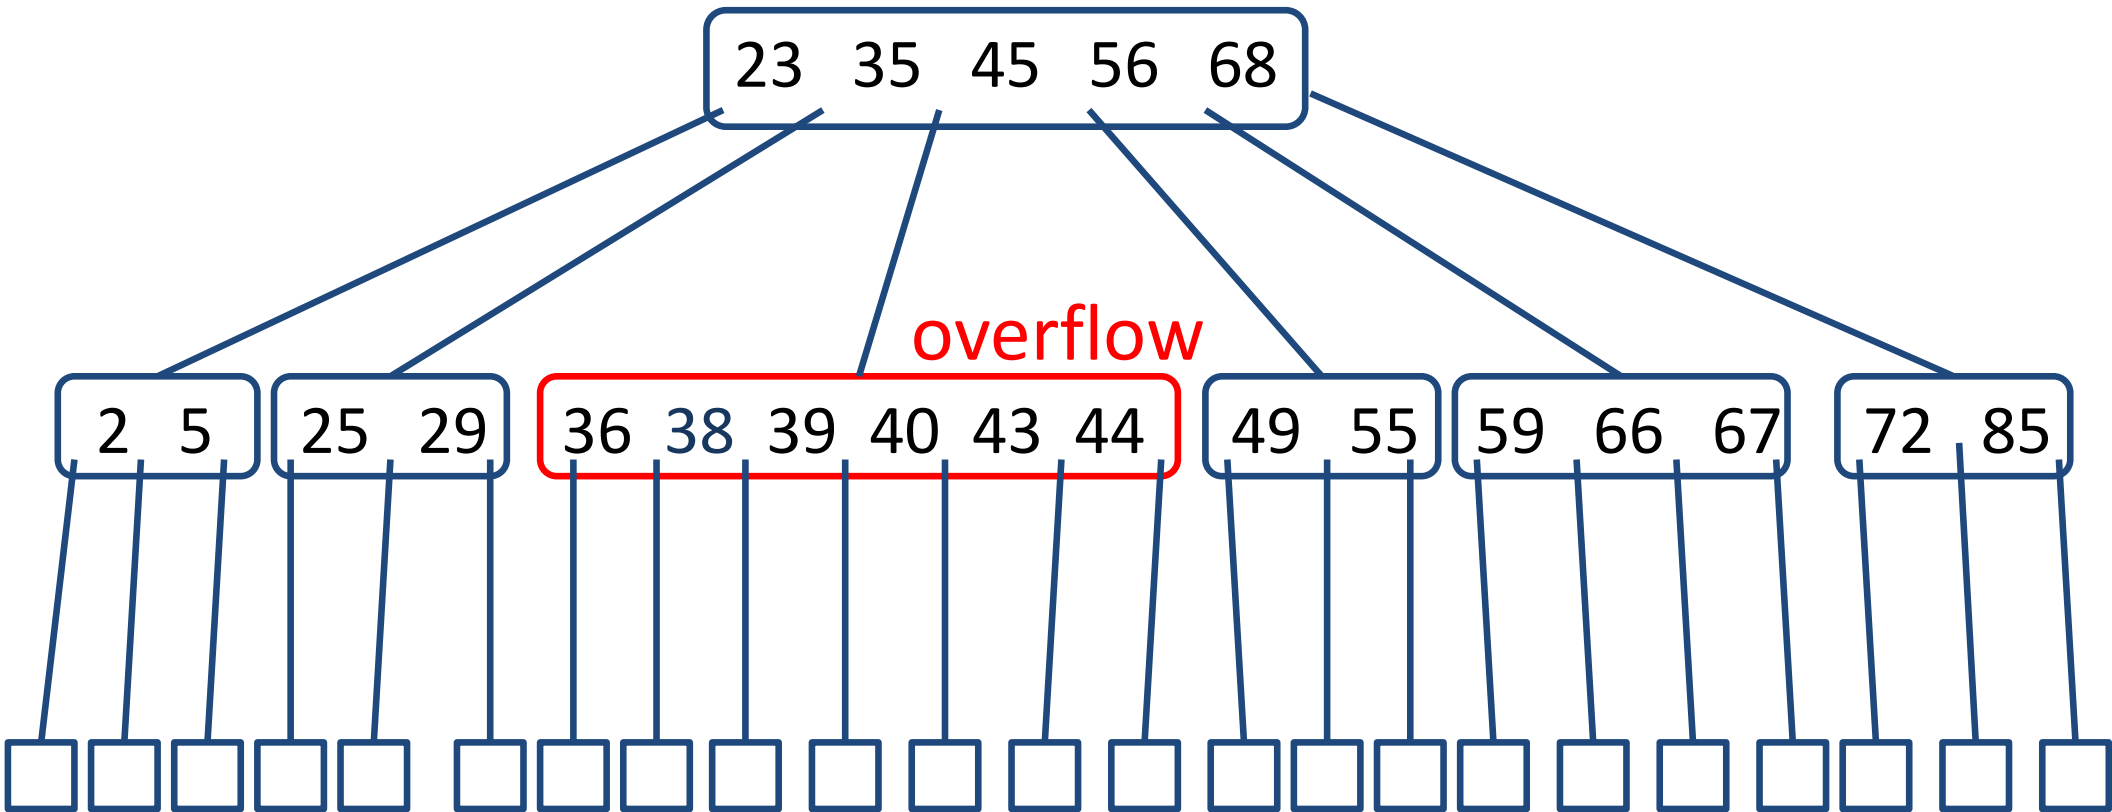
Put 38



B-Trees

B-Tree of order 6

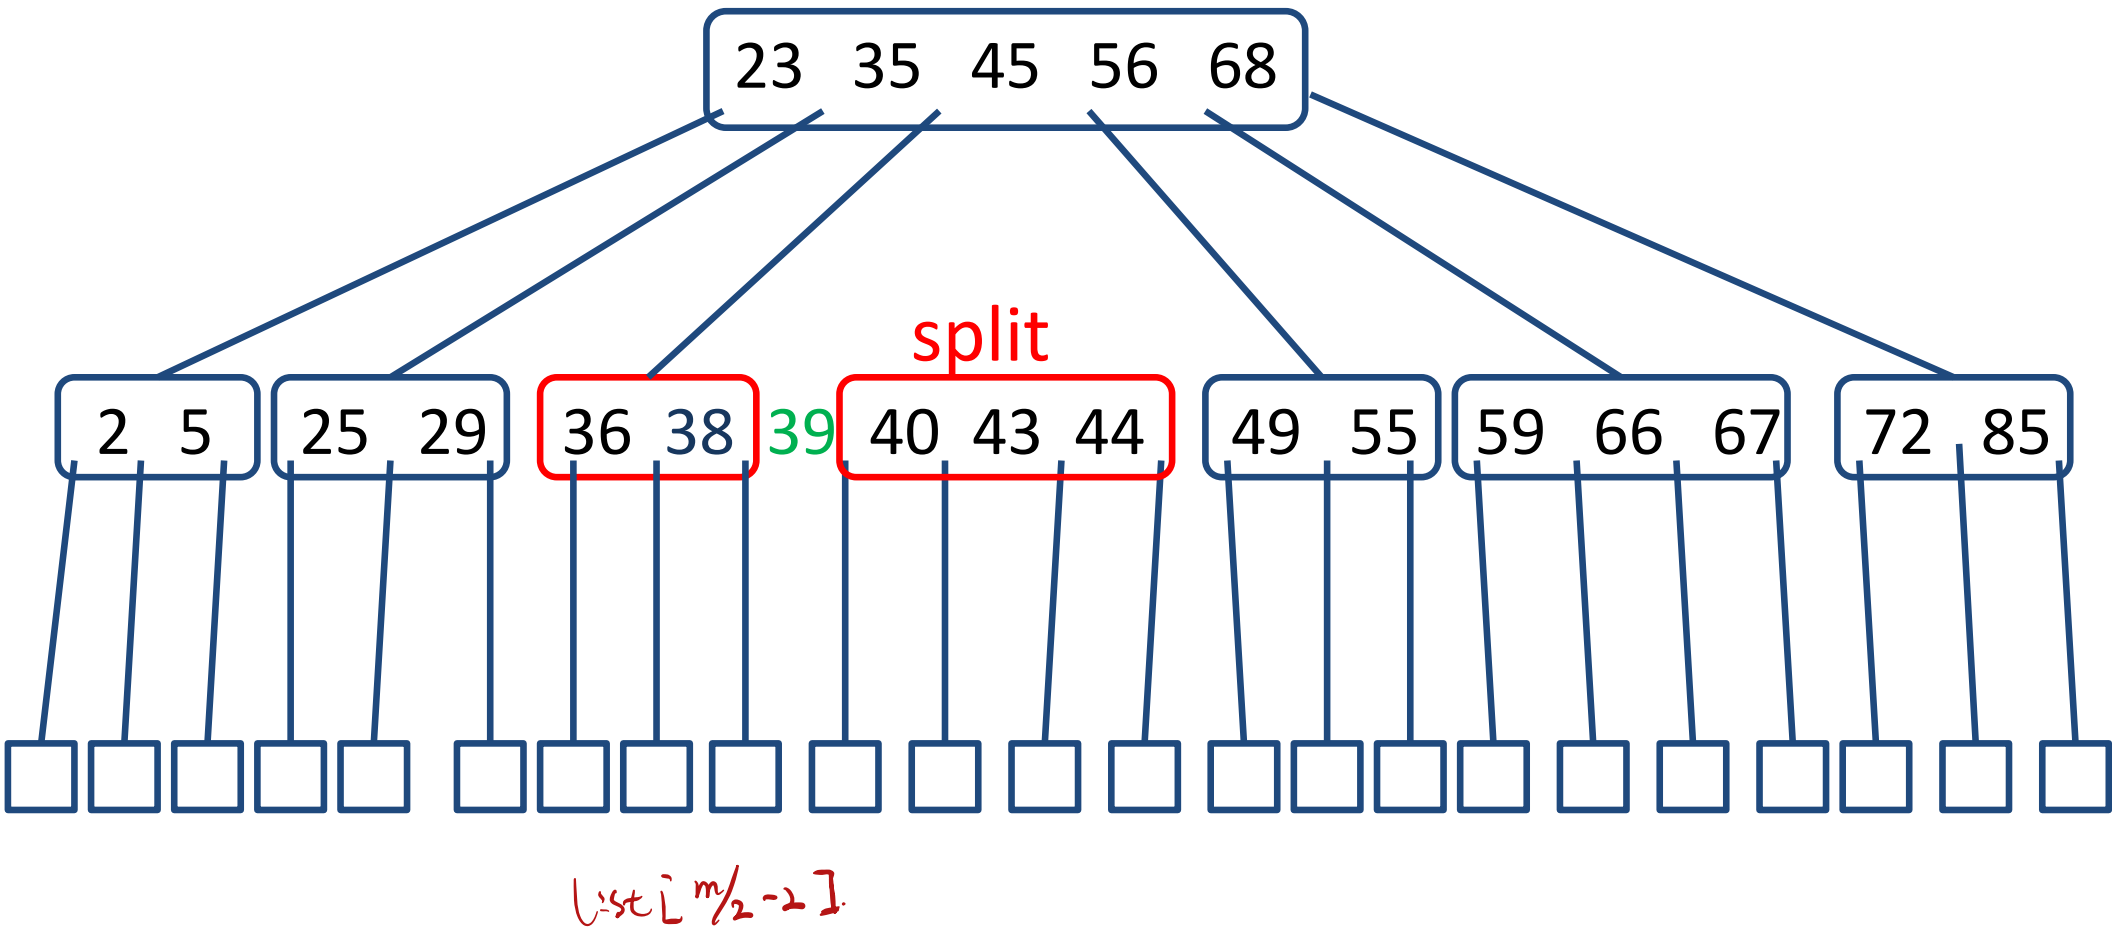
Put 38



B-Trees

B-Tree of order 6


Put 38



B-Trees

B-Tree of order 6

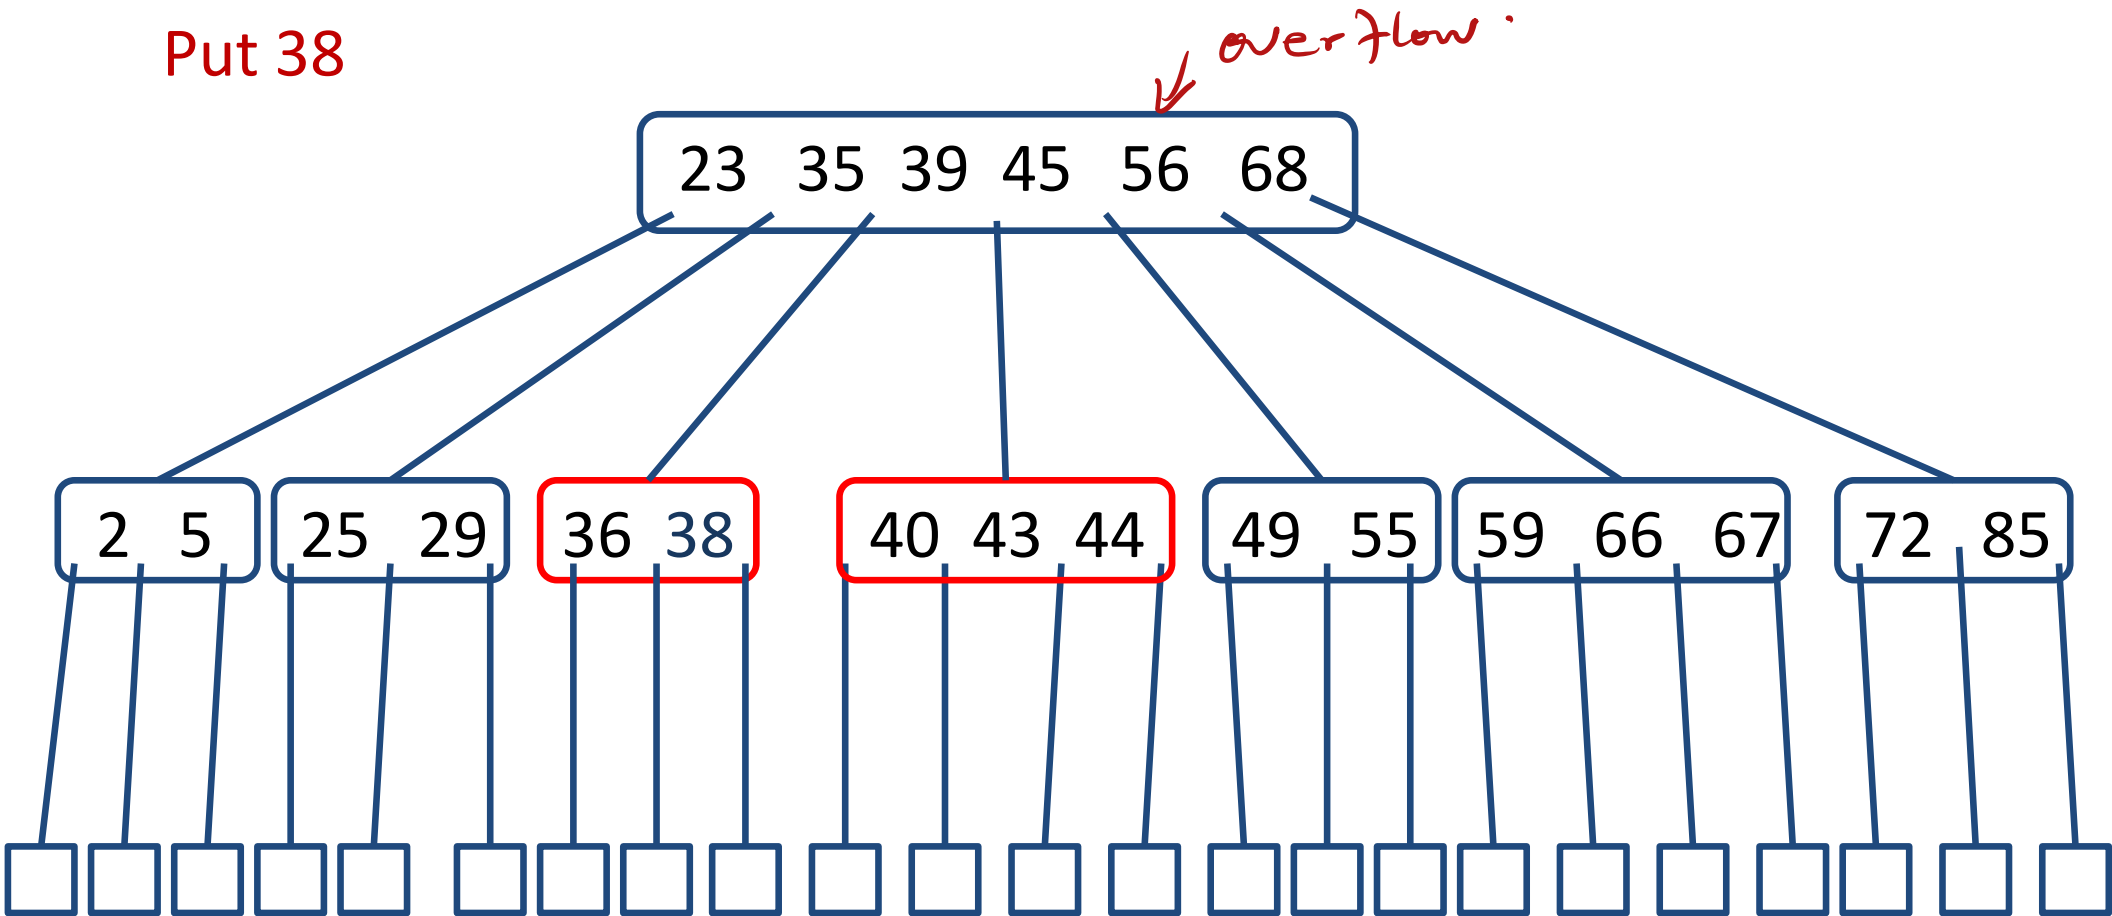
Put 38



B-Trees

B-Tree of order 6

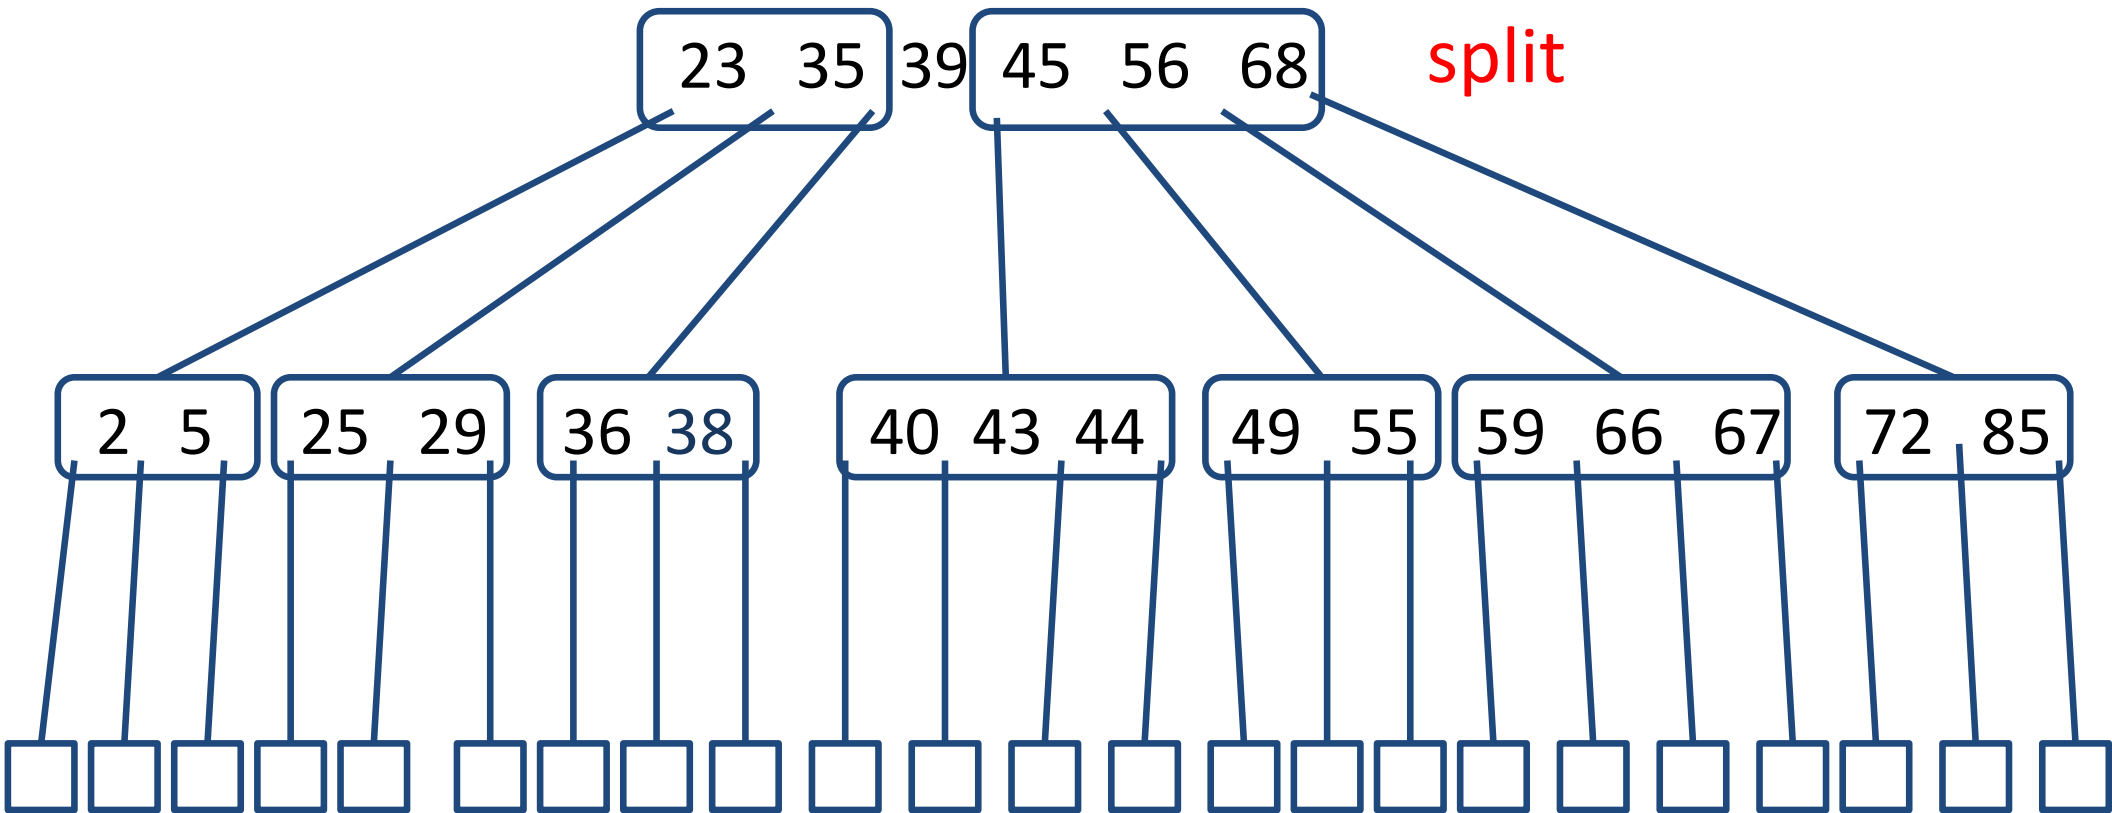
Put 38



B-Trees

B-Tree of order 6

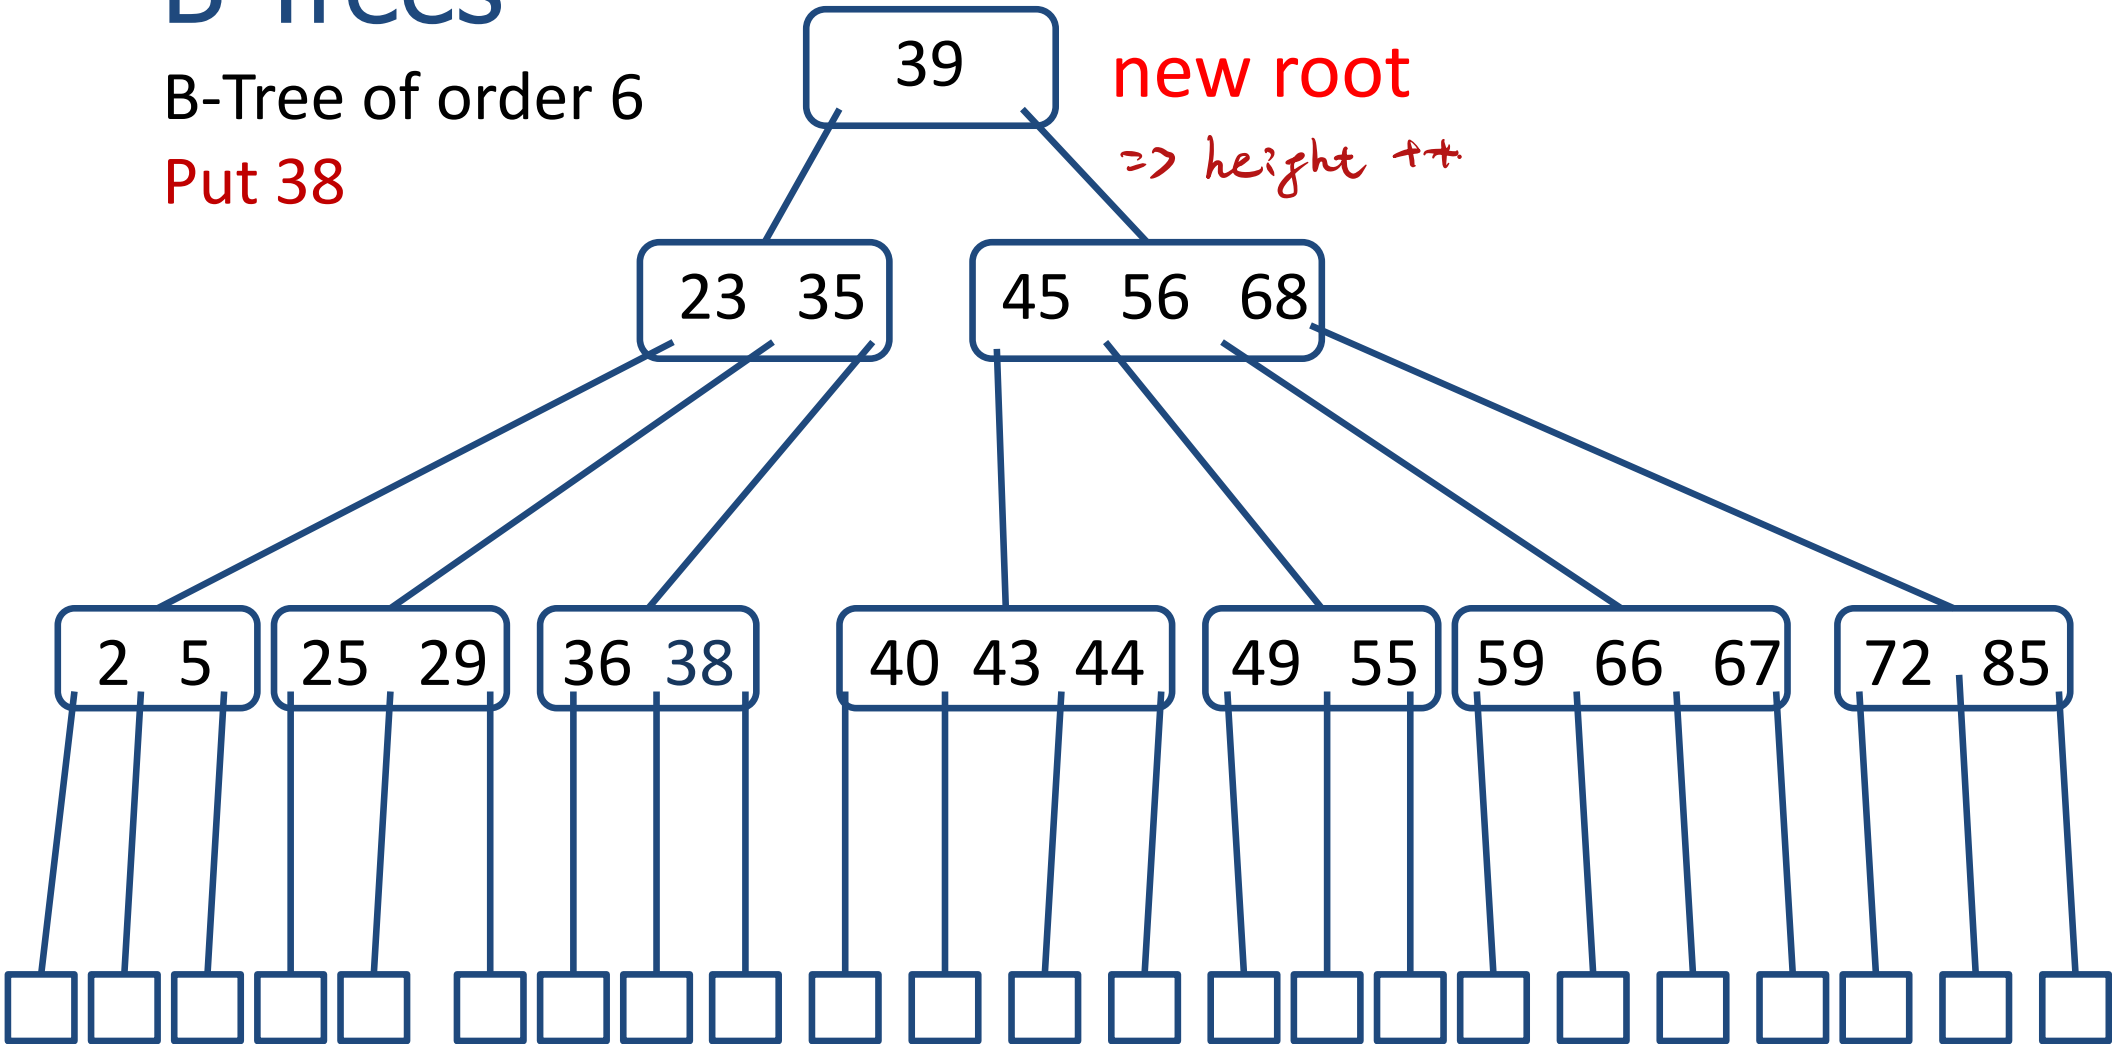
Put 38



B-Trees

B-Tree of order 6

Put 38



Algorithm *put* (r, k, o)

In: Root r of a B-tree, data item (k, o)

Out: {Insert data item (k, o) in the B-tree

Search for k to find the lowest insertion internal node v }

Add the new data item (k, o) at node v

try to fill out the full tree.

while node v overflows do {

if v is the root then

Create a new empty root and set as parent of v

Split v around the **middle** key k' , move k' to parent, and update parent's children

$v \leftarrow$ parent of v

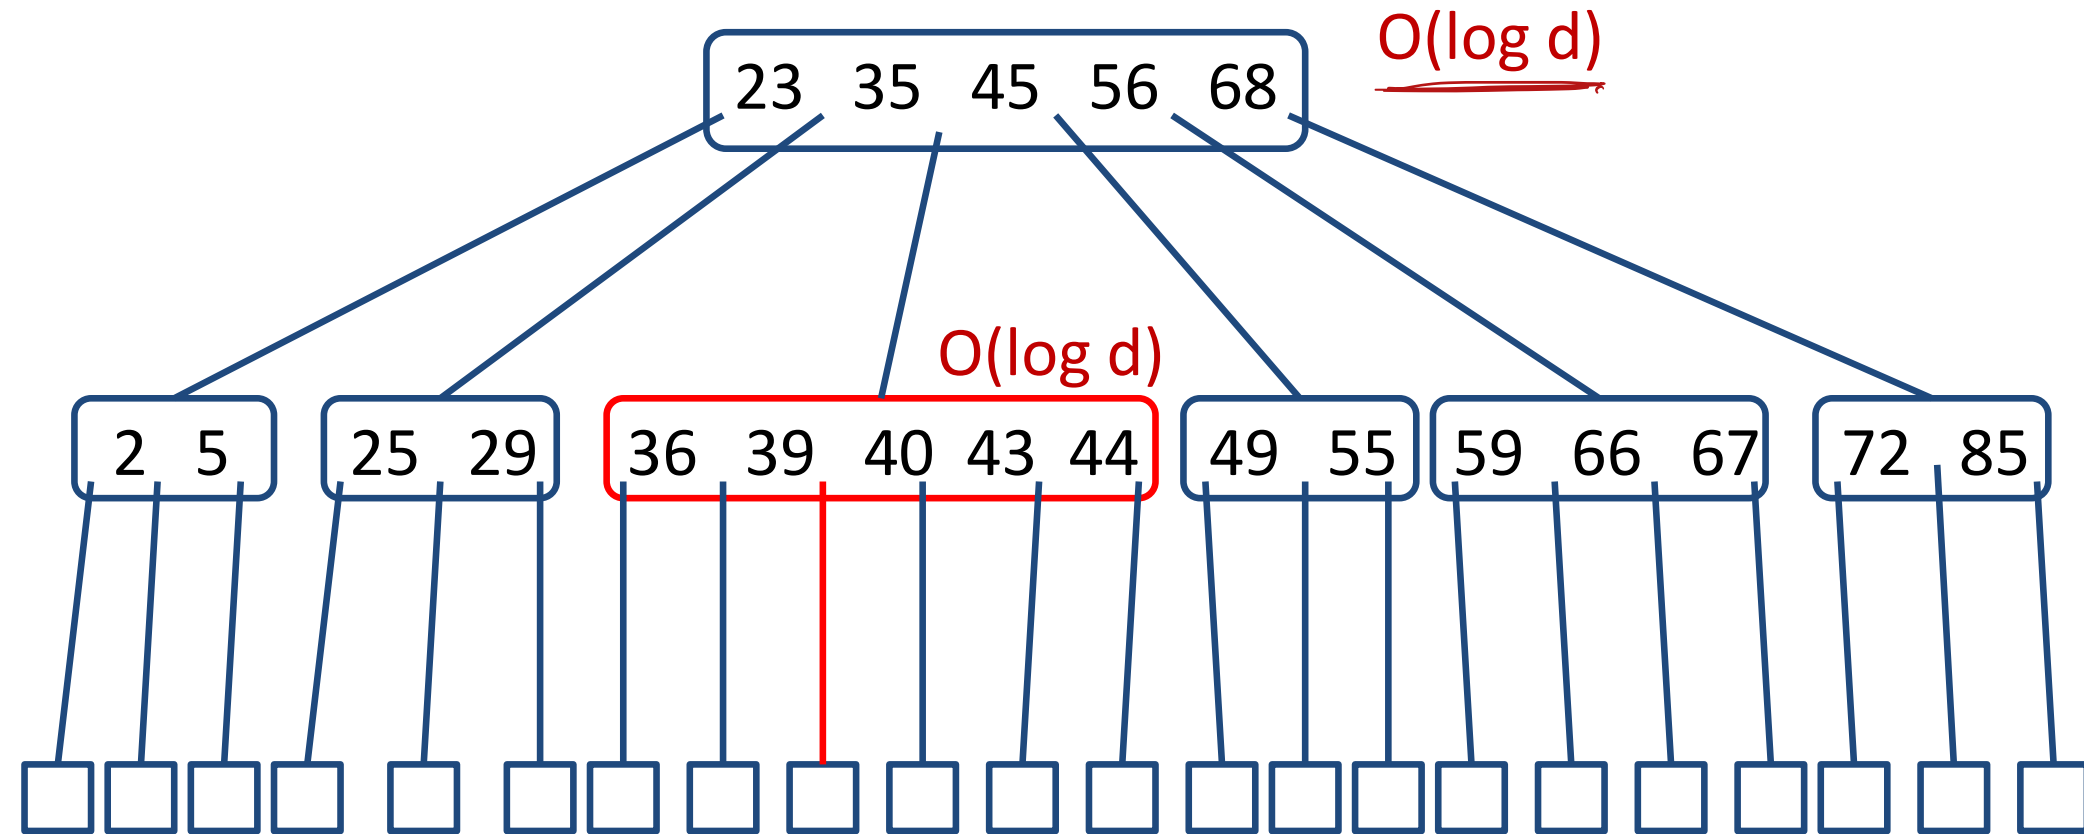
}

move the overflow to upper level.

B-Trees

B-Tree of order 6

B-tree stores both data and keys in internal nodes.
B+ tree stores data at leaf nodes, and keys in internal nodes.
B-tree is another name of B tree.



Algorithm *put* (r, k, o)

In: Root r of a B-tree, data item (k, o)

Out: {Insert data item (k, o) in the B-tree

↓
 $O(\log d \times \log_d n)$ ^{recursive}

Search for k to find the **lowest** insertion **internal** node v

Add the new data item (k, o) at node v

while node v *overflows* **do** {

if v is the root **then**

 Create a new empty root and set as parent of v

 Split v around the **middle** key k' , move k' to parent, and
 update parent's children


$v \leftarrow$ parent of v

}

B-Trees

B-Tree of order 6


Put 38



B-Trees

B-Tree of order 6

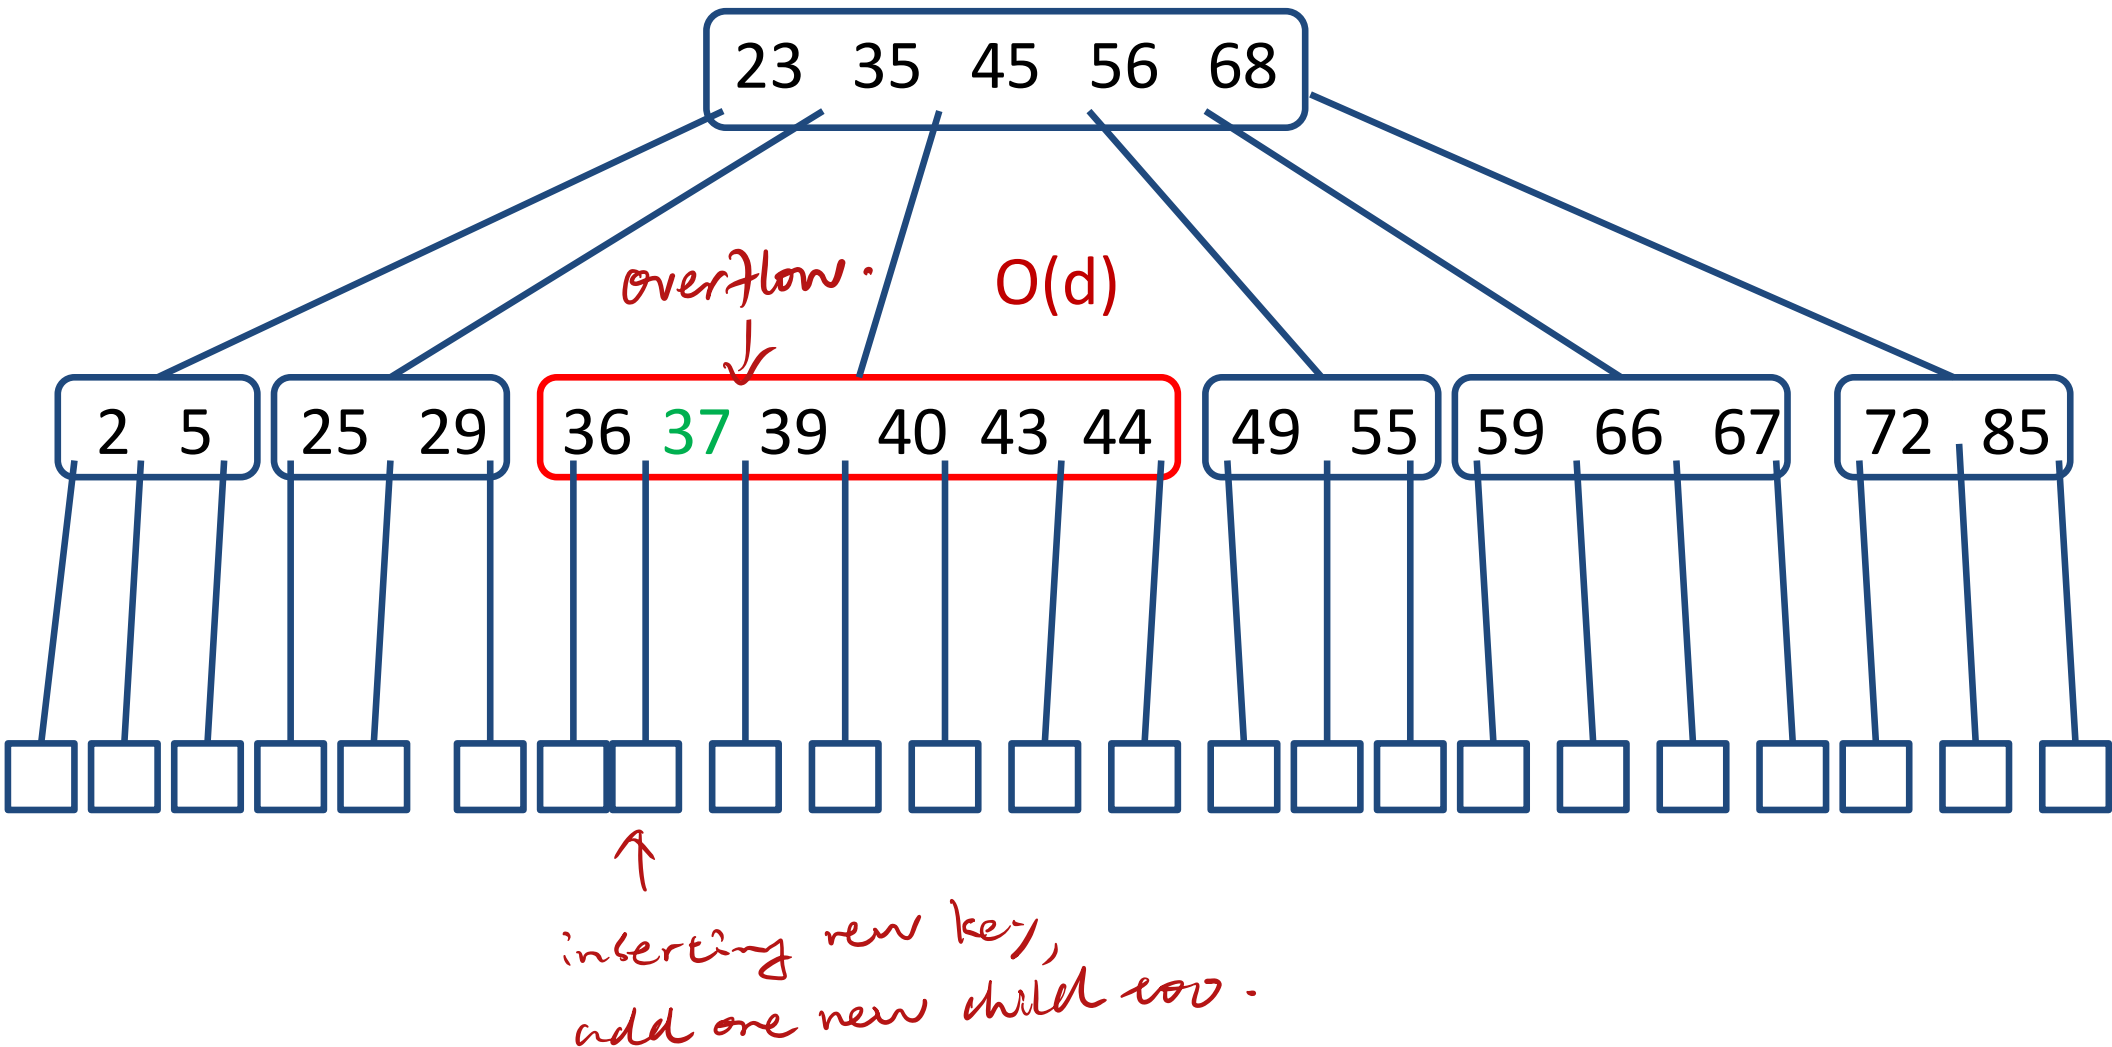
Put 38



B-Trees

B-Tree of order 6

Put 38



Algorithm *put* (r, k, o)

In: Root r of a B-tree, data item (k, o)

Out: {Insert data item (k, o) in the B-tree

$$O\left[\underbrace{\log_2\left(\frac{m}{2}-1\right)}_{\downarrow} \times \underbrace{\log_{\frac{m}{2}}\left(\frac{N}{\frac{m}{2}-1}\right)}_{\uparrow}\right]$$

max time
searching
each level.

max height.

$$O(\log d \times \log_d n)$$

Search for k to find the **lowest** insertion **internal** node v }

Add the new data item (k, o) at node v } $O(d)$

while node v **overflows** **do** {

if v is the root **then** {

Create a new empty root and set as parent of v }.

Split v around the **middle** key k' , move k' to parent, and update parent's children

$v \leftarrow$ parent of v

}

Algorithm *put* (r, k, o)

In: Root r of a B-tree, data item (k, o)

Out: {Insert data item (k, o) in the B-tree

Search for k to find the **lowest** insertion **internal** node v } $O(\log d \times \log_d n)$

Add the new data item (k, o) at node v } $O(d)$ c_i

while node v *overflows* **do** {

if v is the root **then**

 Create a new empty root and set as parent of v

 Split v around the **middle** key k' , move k' to parent, and
 update parent's children

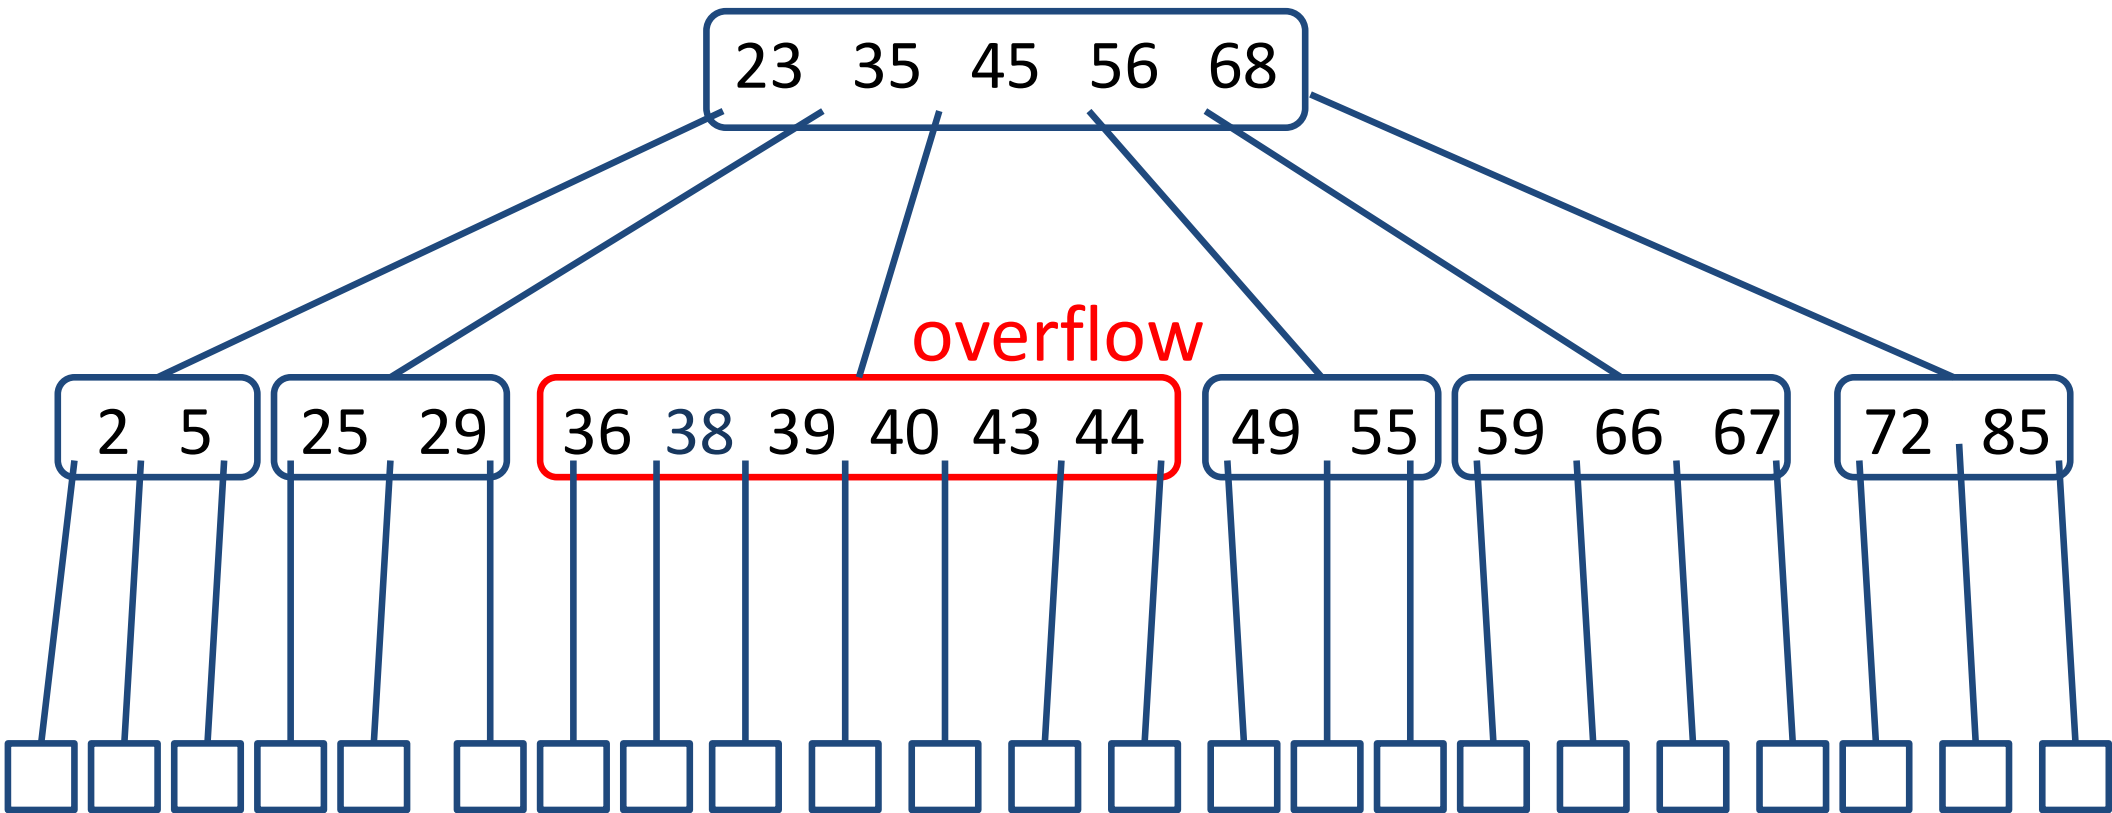
$v \leftarrow$ parent of v

}

B-Trees

B-Tree of order 6


Put 38



B-Trees

B-Tree of order 6

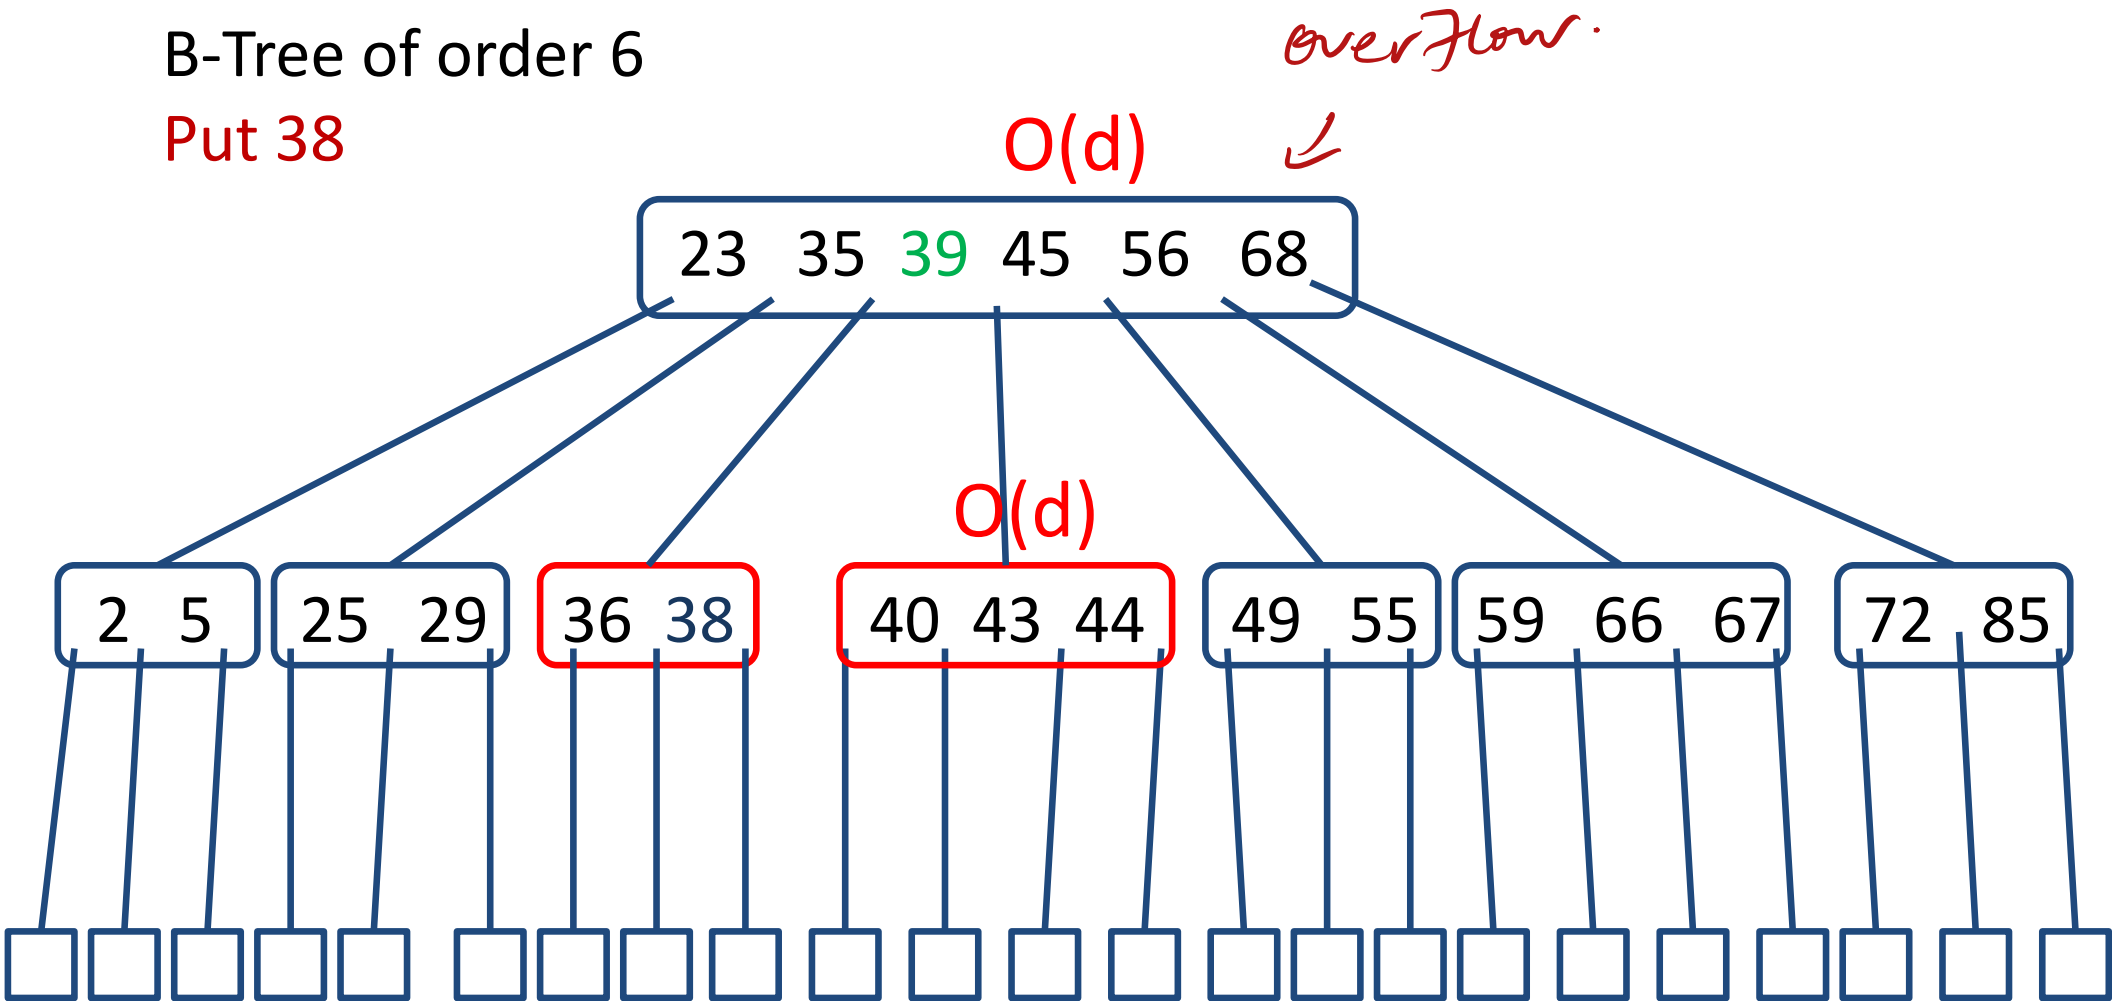
Put 38



B-Trees

B-Tree of order 6

Put 38



Algorithm *put* (r, k, o)

In: Root r of a B-tree, data item (k, o)

Out: {Insert data item (k, o) in the B-tree

Search for k to find the **lowest** insertion **internal** node v } $O(\log d \times \log_d n)$

Add the new data item (k, o) at node v } $O(d)$

while node v *overflows* **do** {

if v is the root **then**

 Create a new empty root and set as parent of v

 Split v around the **middle** key k' , move k' to parent, and
 update parent's children

$v \leftarrow$ parent of v

}

$O(d)$

Algorithm *put* (r, k, o)

In: Root r of a B-tree, data item (k, o)

Out: {Insert data item (k, o) in the B-tree

Search for k to find the **lowest** insertion **internal** node v

$O(\log d \times \log_d n)$
height.

Add the new data item (k, o) at node v

$O(d)$

while node v *overflows* **do** {

if v is the root **then**

 Create a new empty root and set as parent of v

 Split v around the **middle** key k' , move k' to parent, and
 update parent's children

$v \leftarrow$ parent of v

}

$O(d \log_d n)$

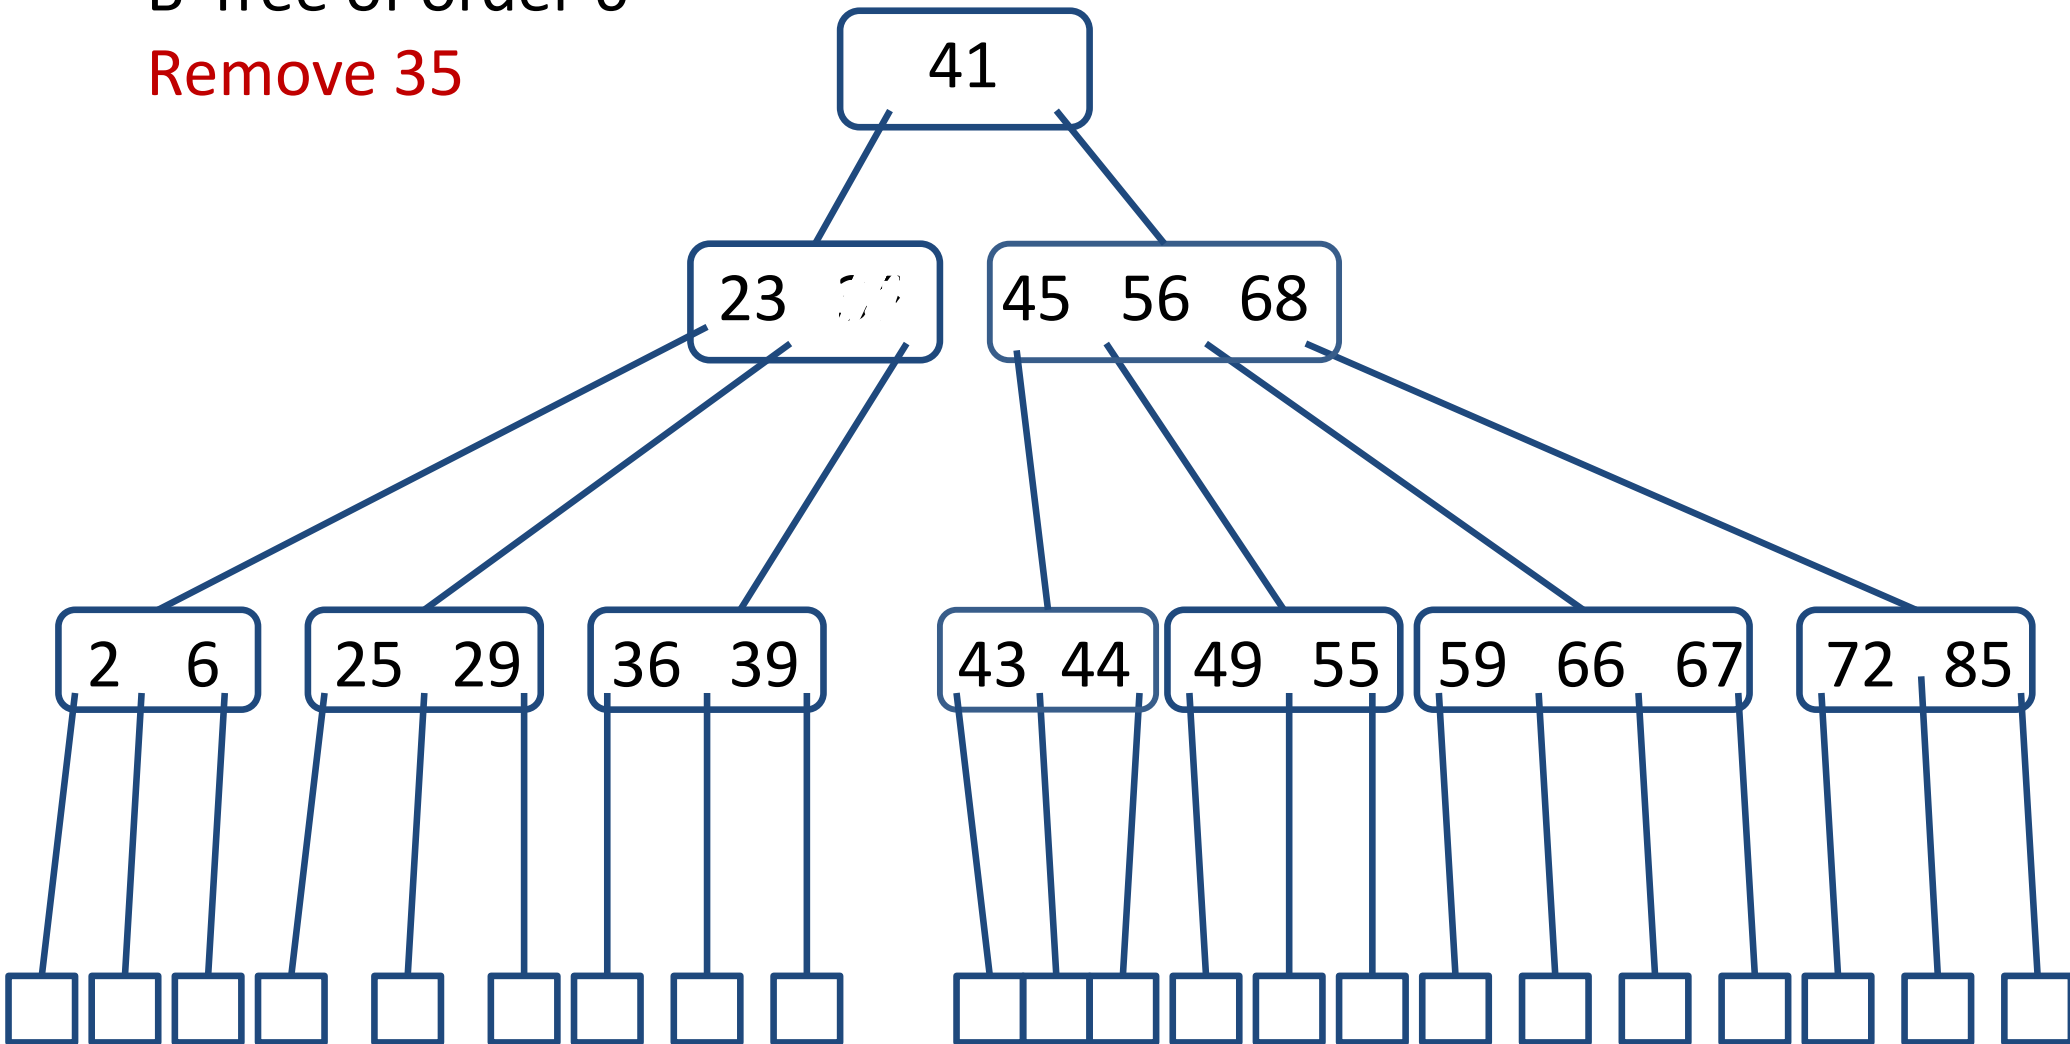
$O(d)$

Time complexity of *put* is $O(d \log_d n)$

B-Trees

B-Tree of order 6

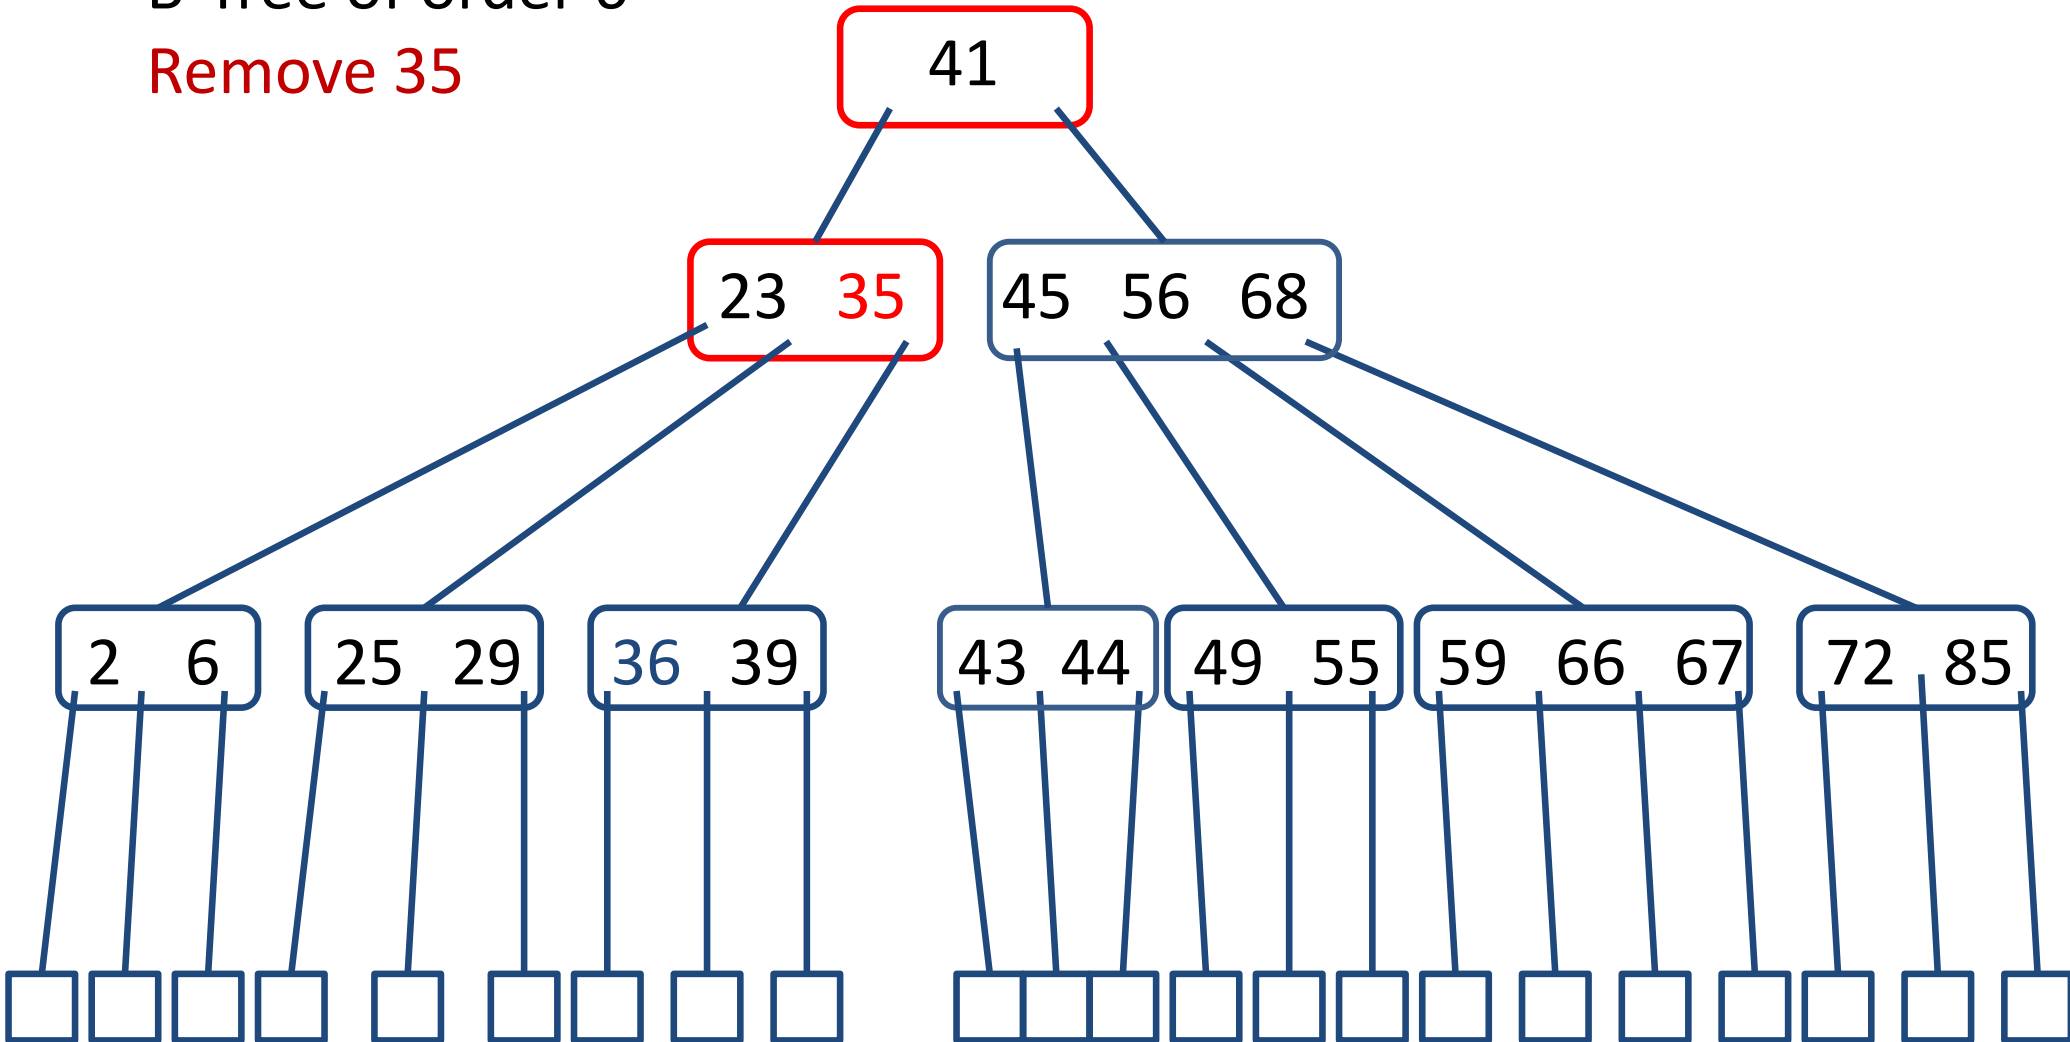
Remove 35



B-Trees

B-Tree of order 6

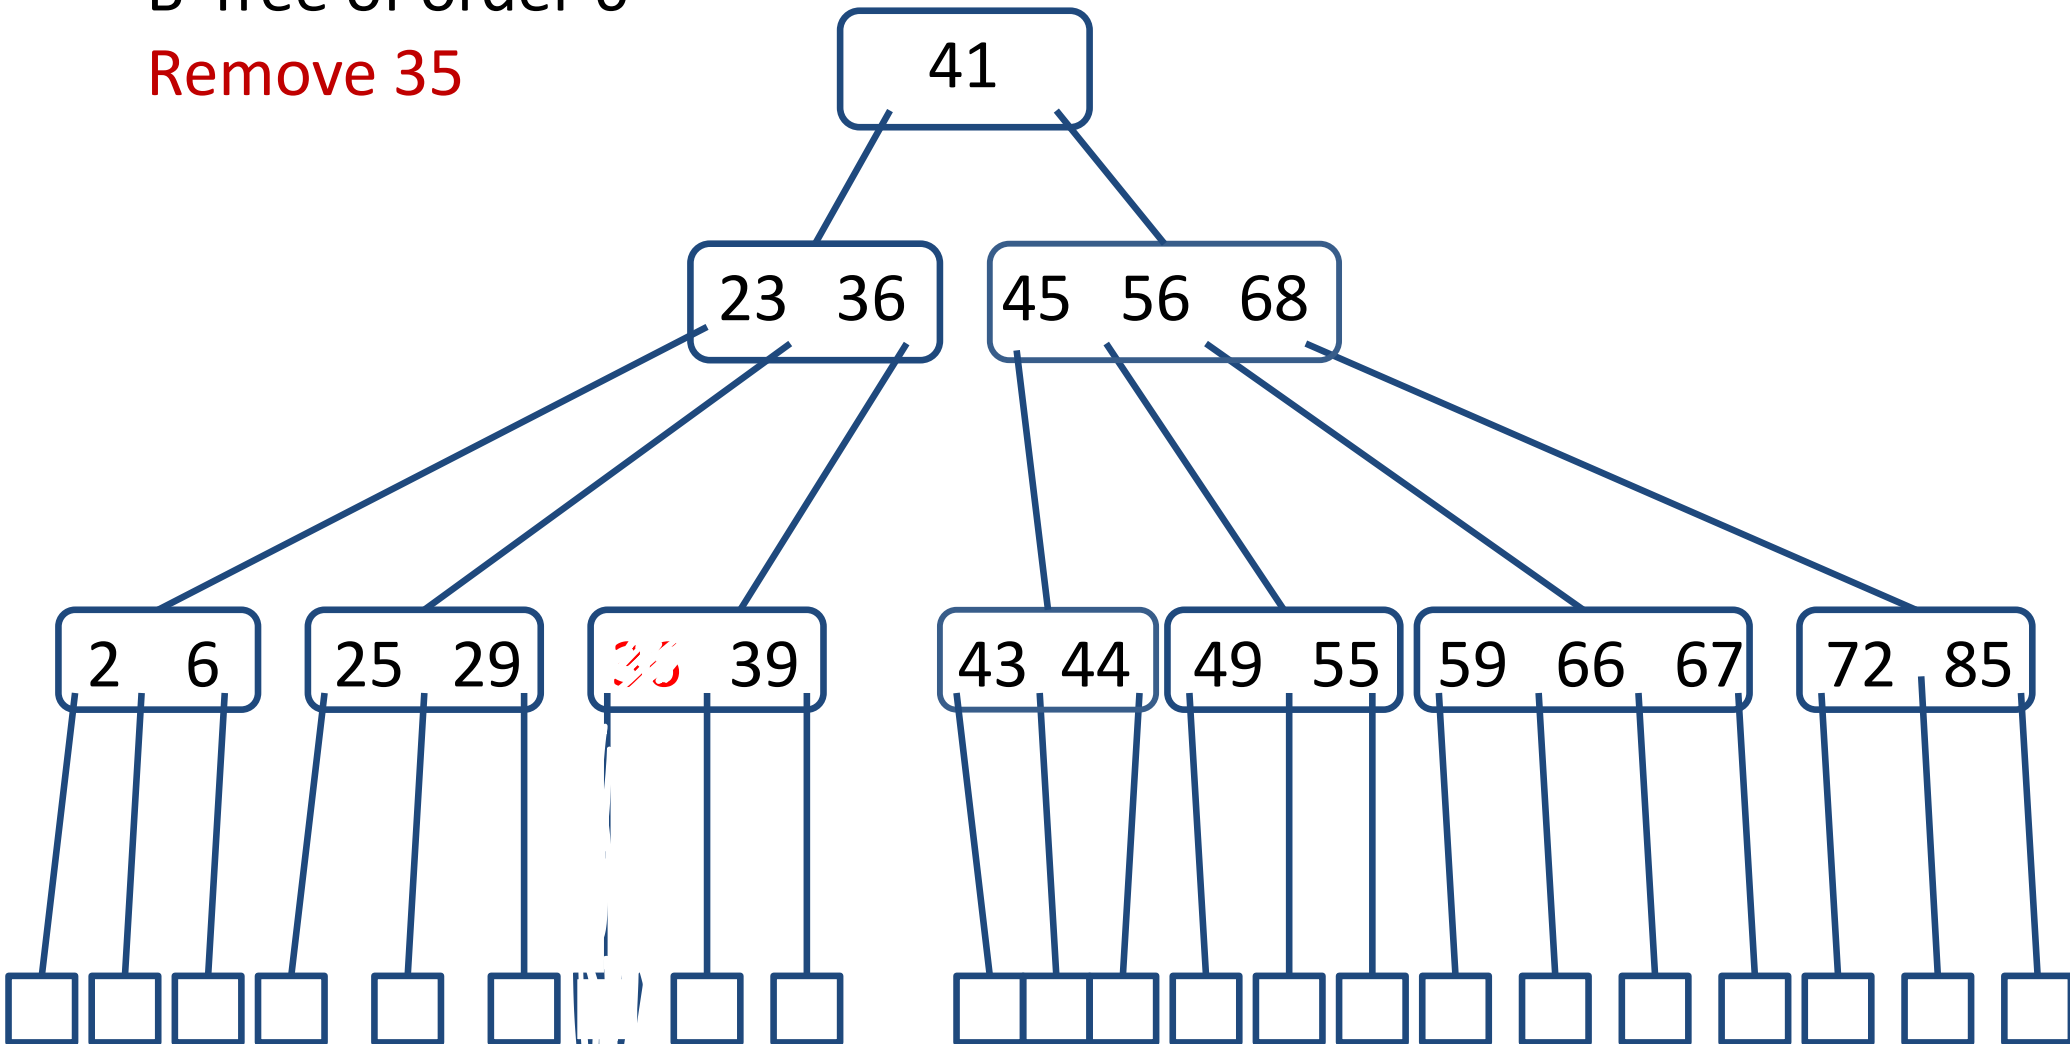
Remove 35



B-Trees

B-Tree of order 6

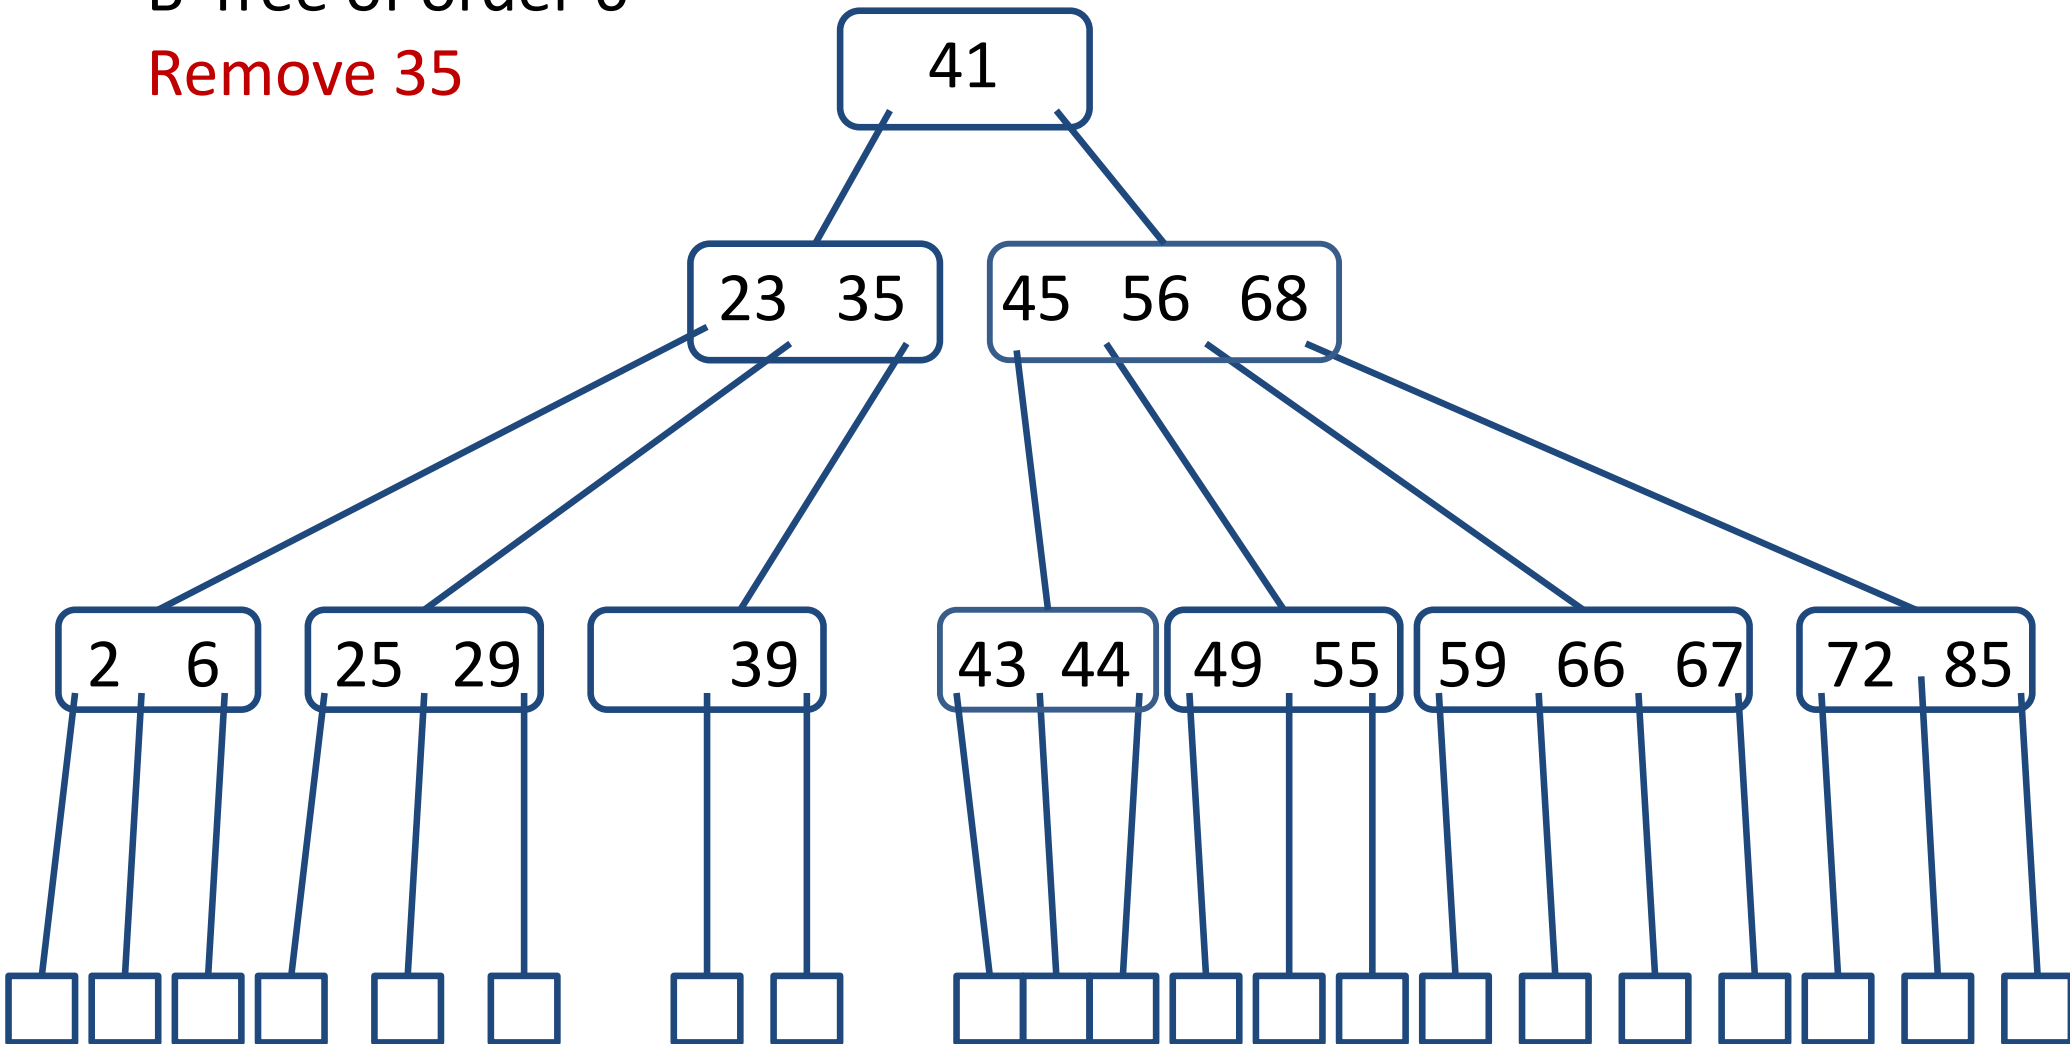
Remove 35



B-Trees

B-Tree of order 6

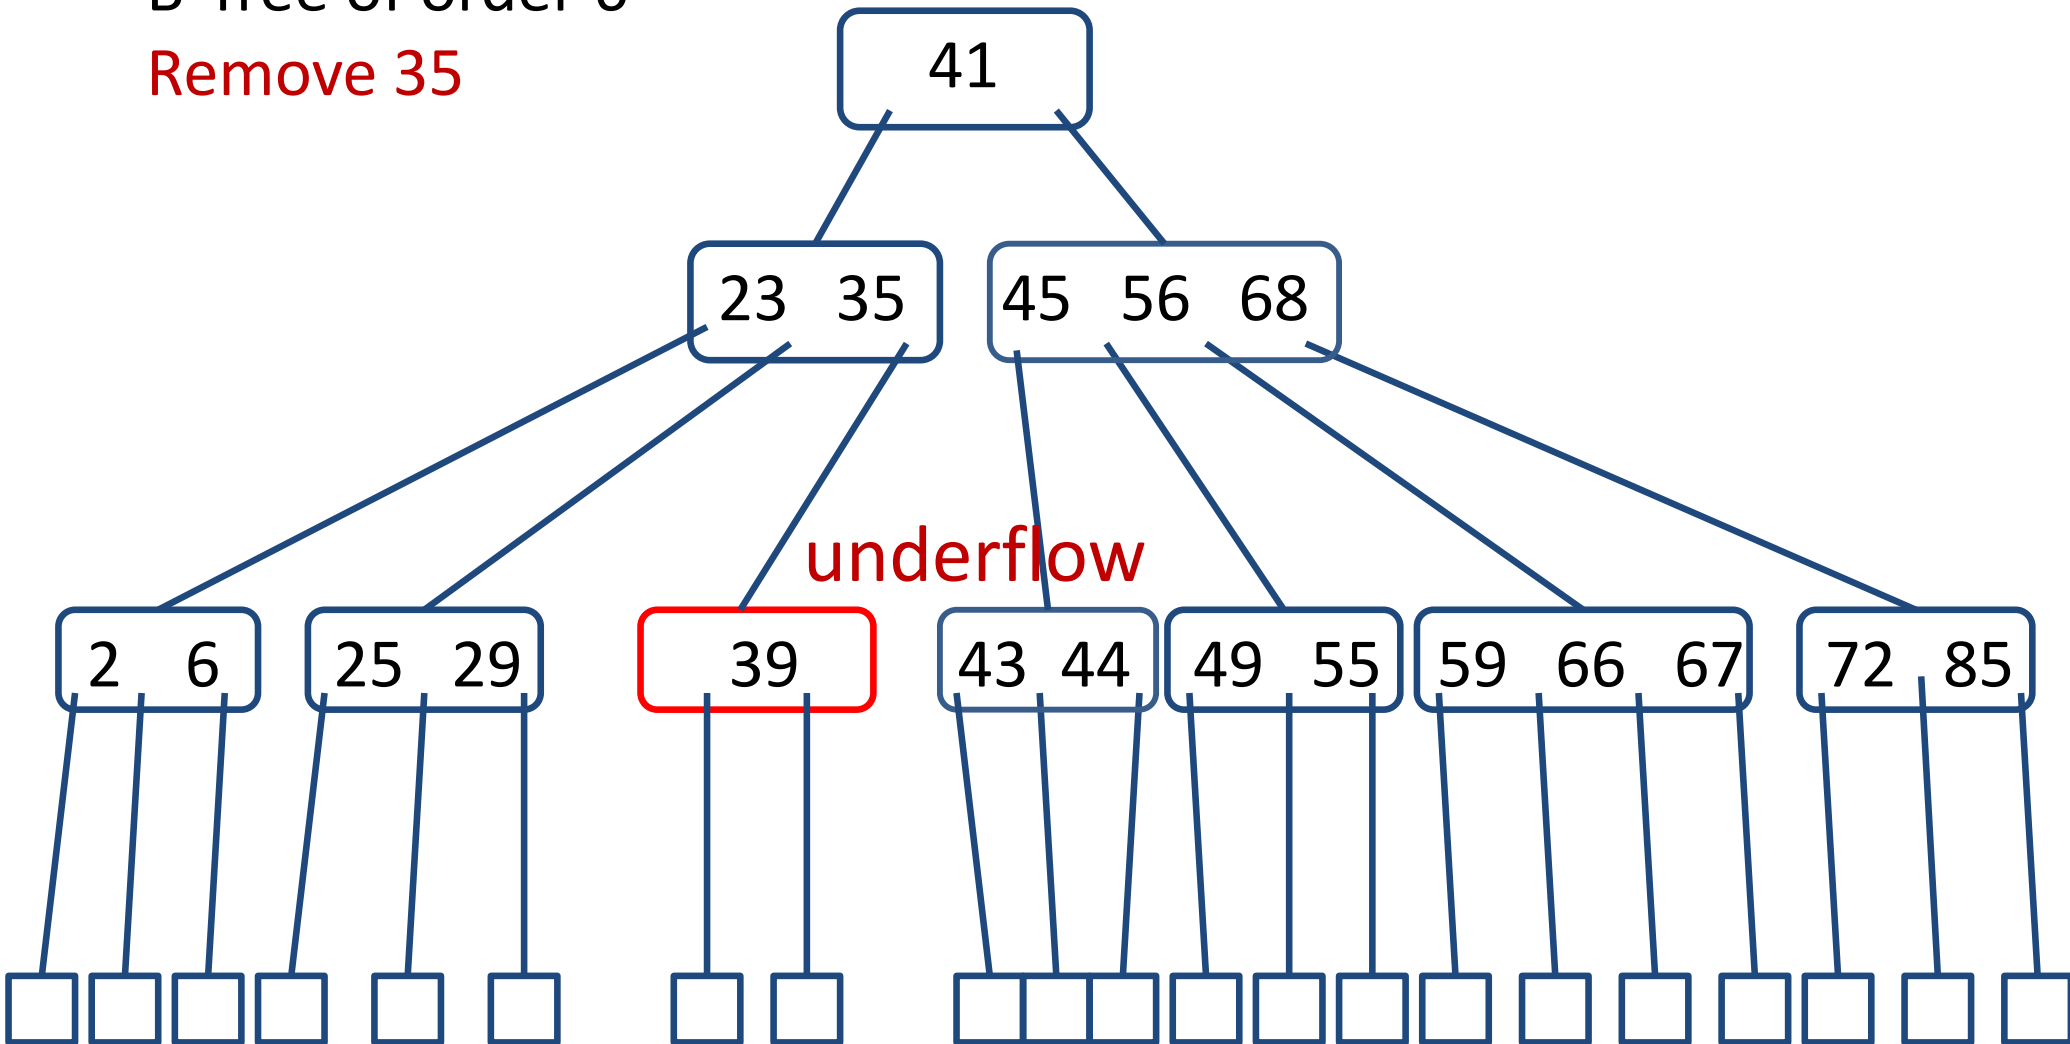
Remove 35



B-Trees

B-Tree of order 6

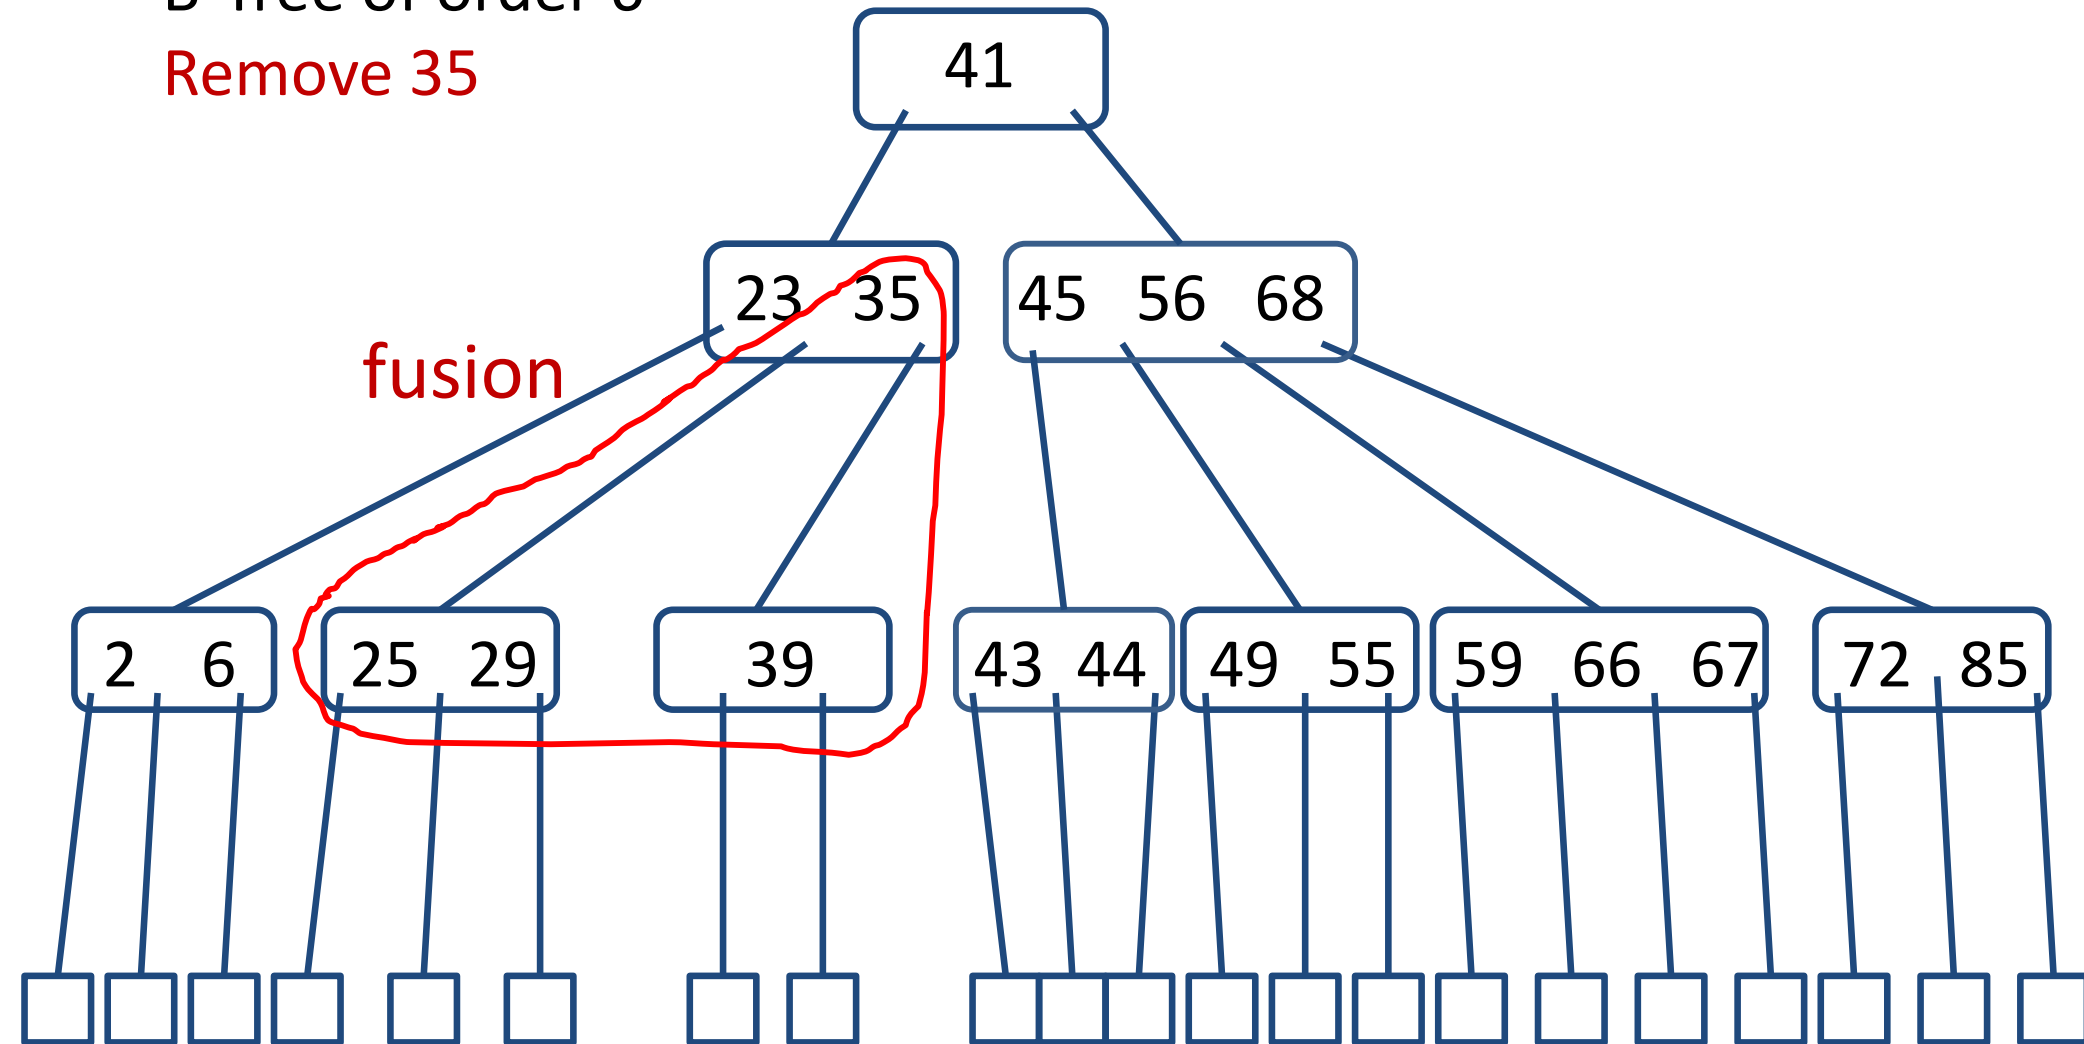
Remove 35



B-Trees

B-Tree of order 6

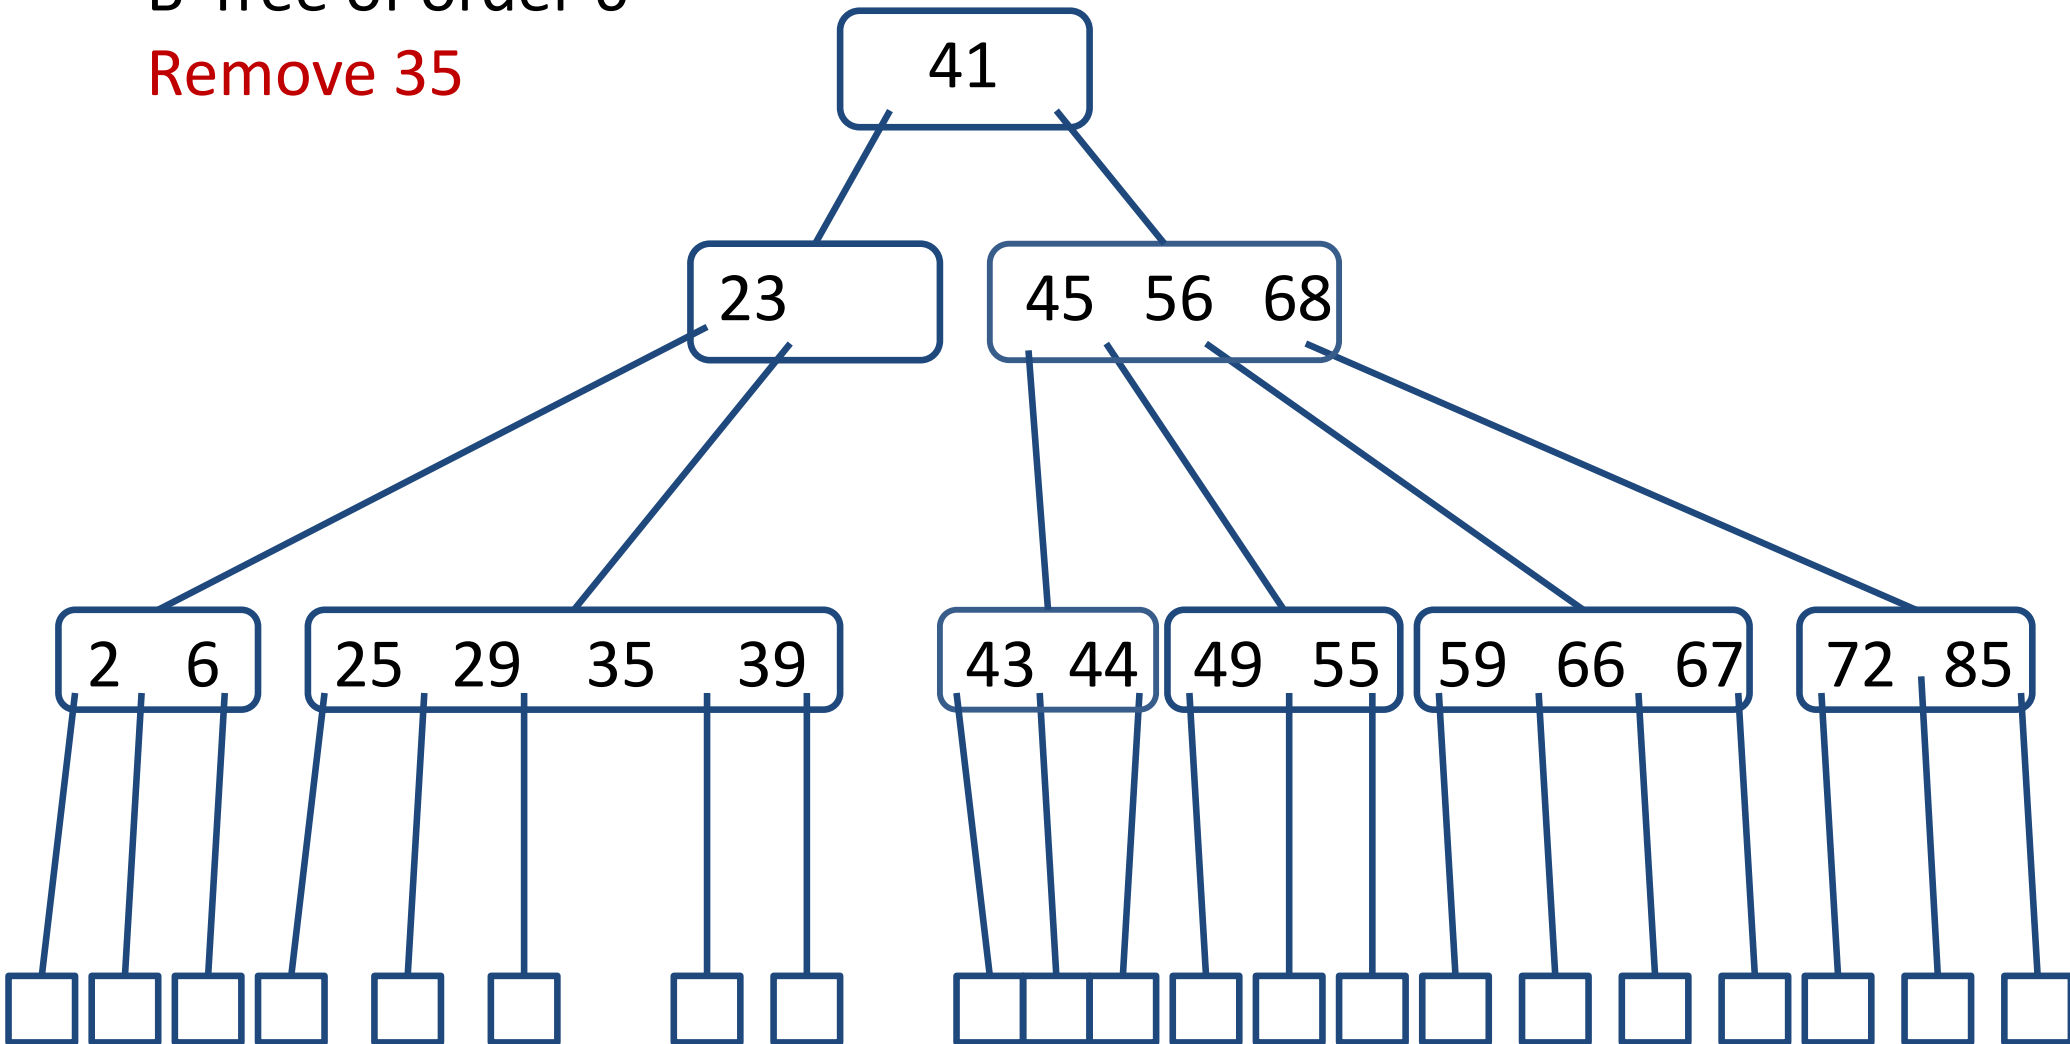
Remove 35



B-Trees

B-Tree of order 6

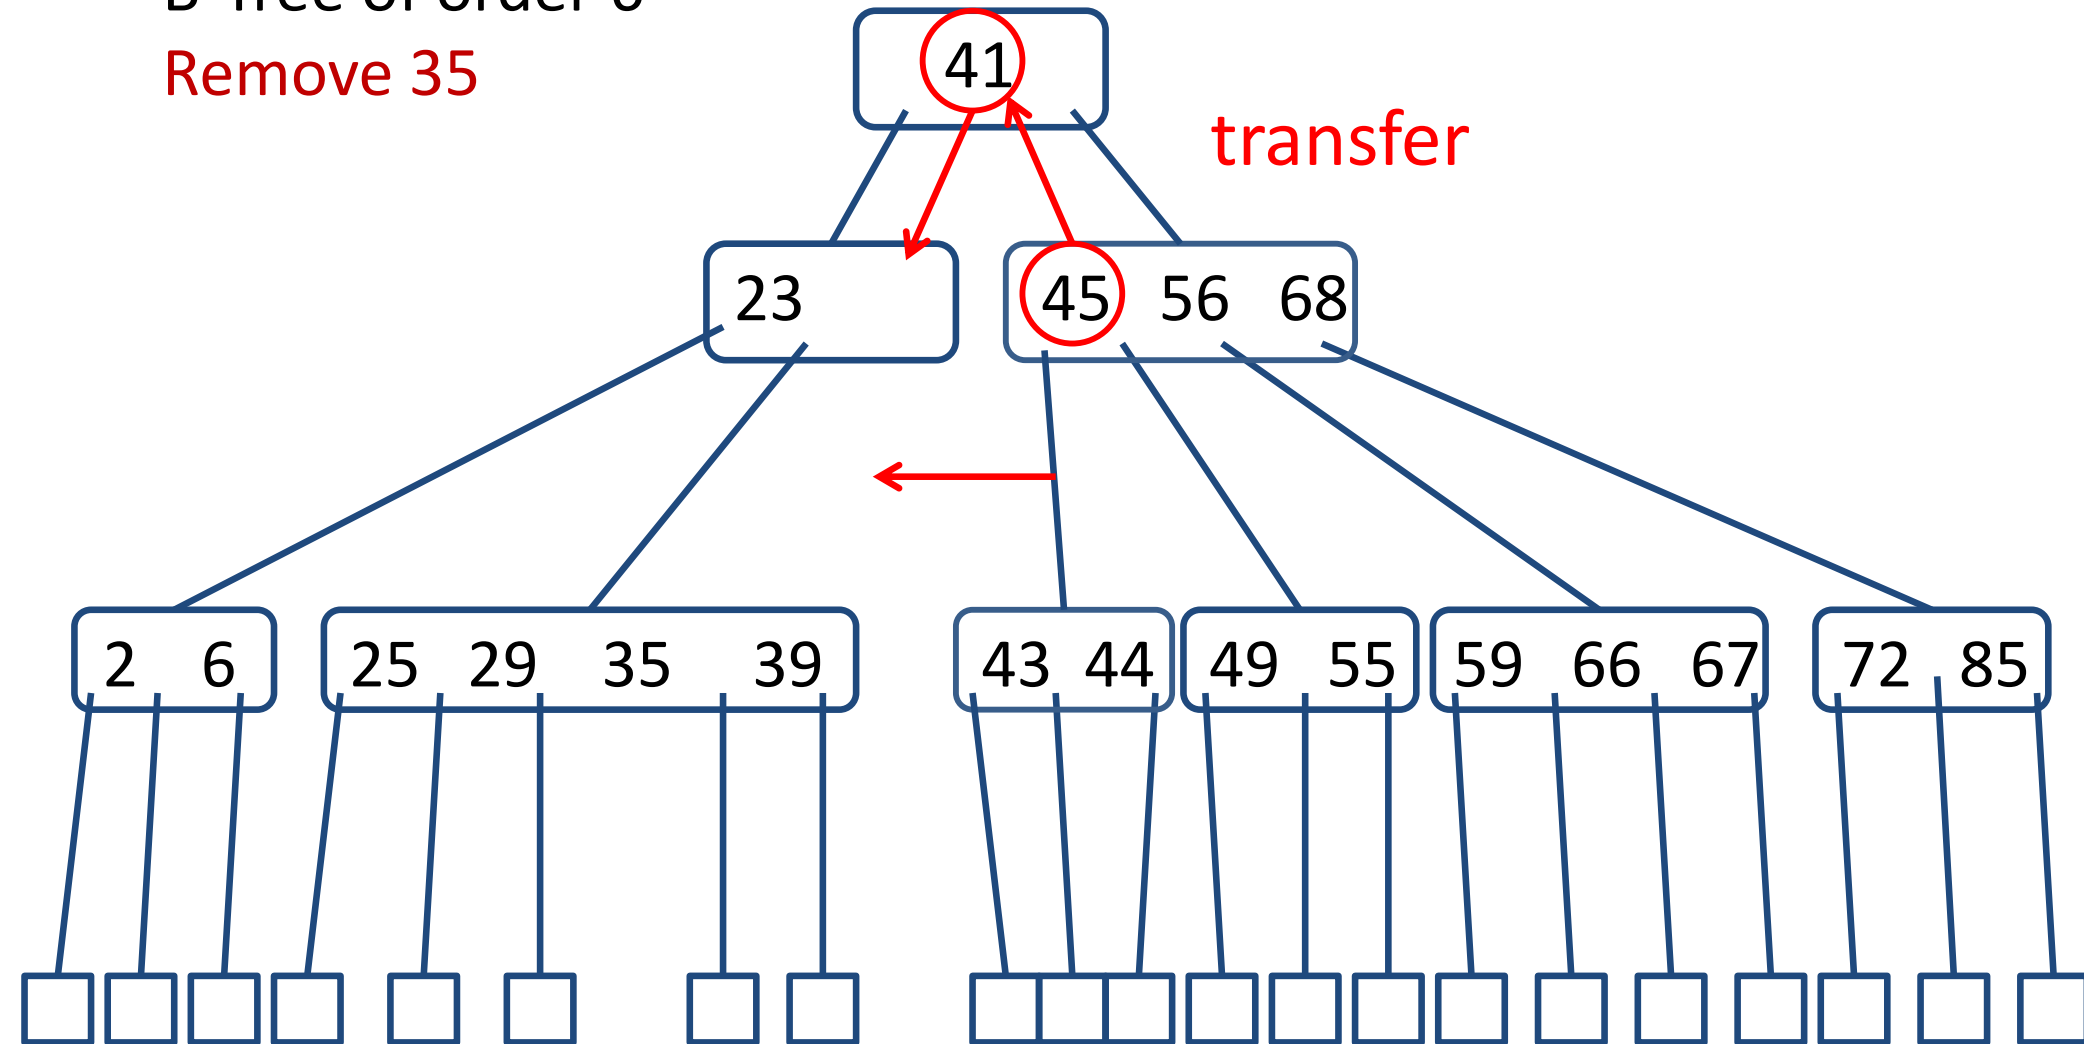
Remove 35



B-Trees

B-Tree of order 6

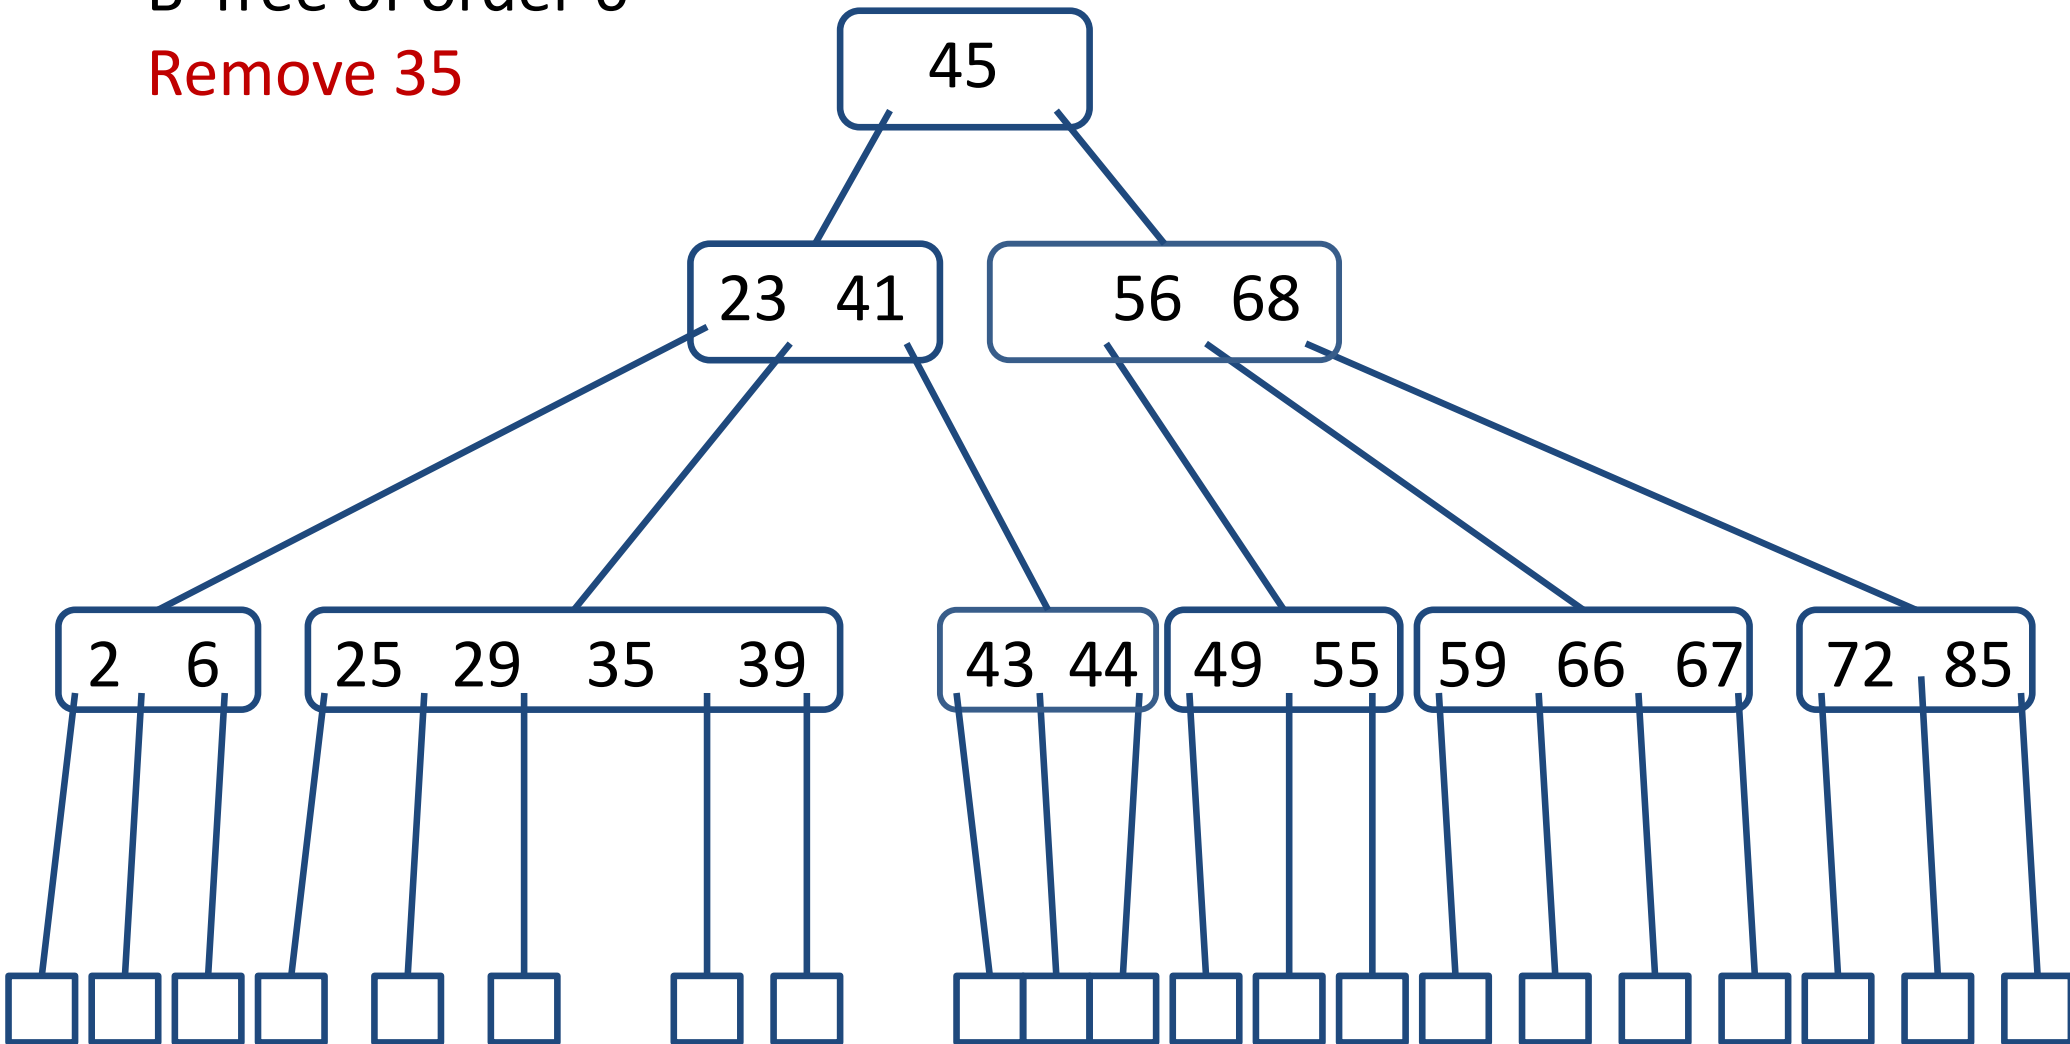
Remove 35



B-Trees

B-Tree of order 6

Remove 35



Algorithm *remove*(r, k)

In: Root r of a B-tree, key k

Out: {remove data item with key k from the tree}

Find the node v storing key k

Remove (k, o) from v replacing it with successor if needed

while node v *underflows* **do** {

if v is the root **then**

 make the first child of v the new root

else if a sibling has more than $\lceil d/2 \rceil$ keys **then**

 perform a transfer operation

else {

 perform a fusion operation

$v \leftarrow$ parent of v

 }

}

Algorithm *remove*(r, k)

In: Root r of a B-tree, key k

Out: {remove data item with key k from the tree}

Find the node v storing key k } $O(\log d \times \log_d n)$

Remove (k, o) from v replacing it with successor if needed

while node v *underflows* **do** {

if v is the root **then**

 make the first child of v the new root

else if a sibling has at least $\lceil d/2 \rceil$ keys **then**

 perform a transfer operation

else {

 perform a fusion operation

$v \leftarrow$ parent of v

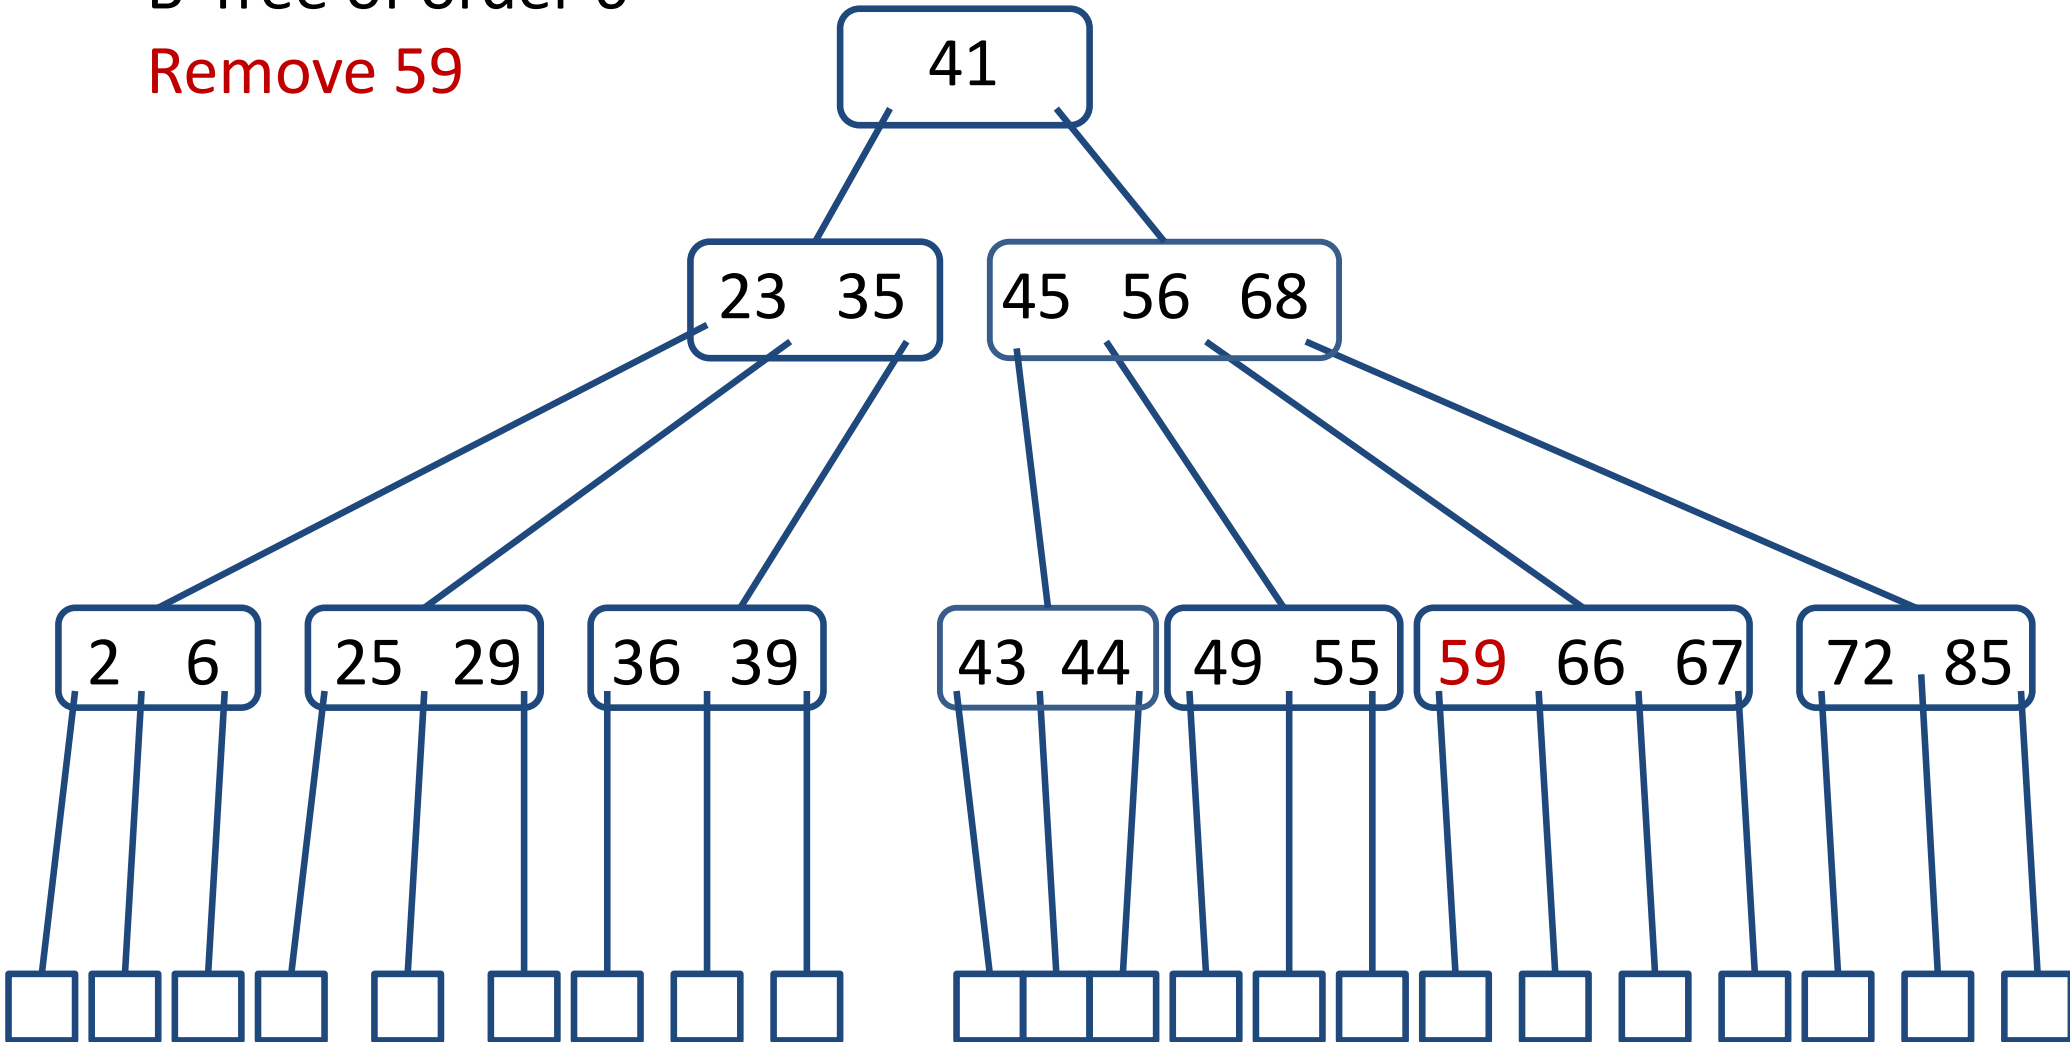
 }

}

B-Trees

B-Tree of order 6

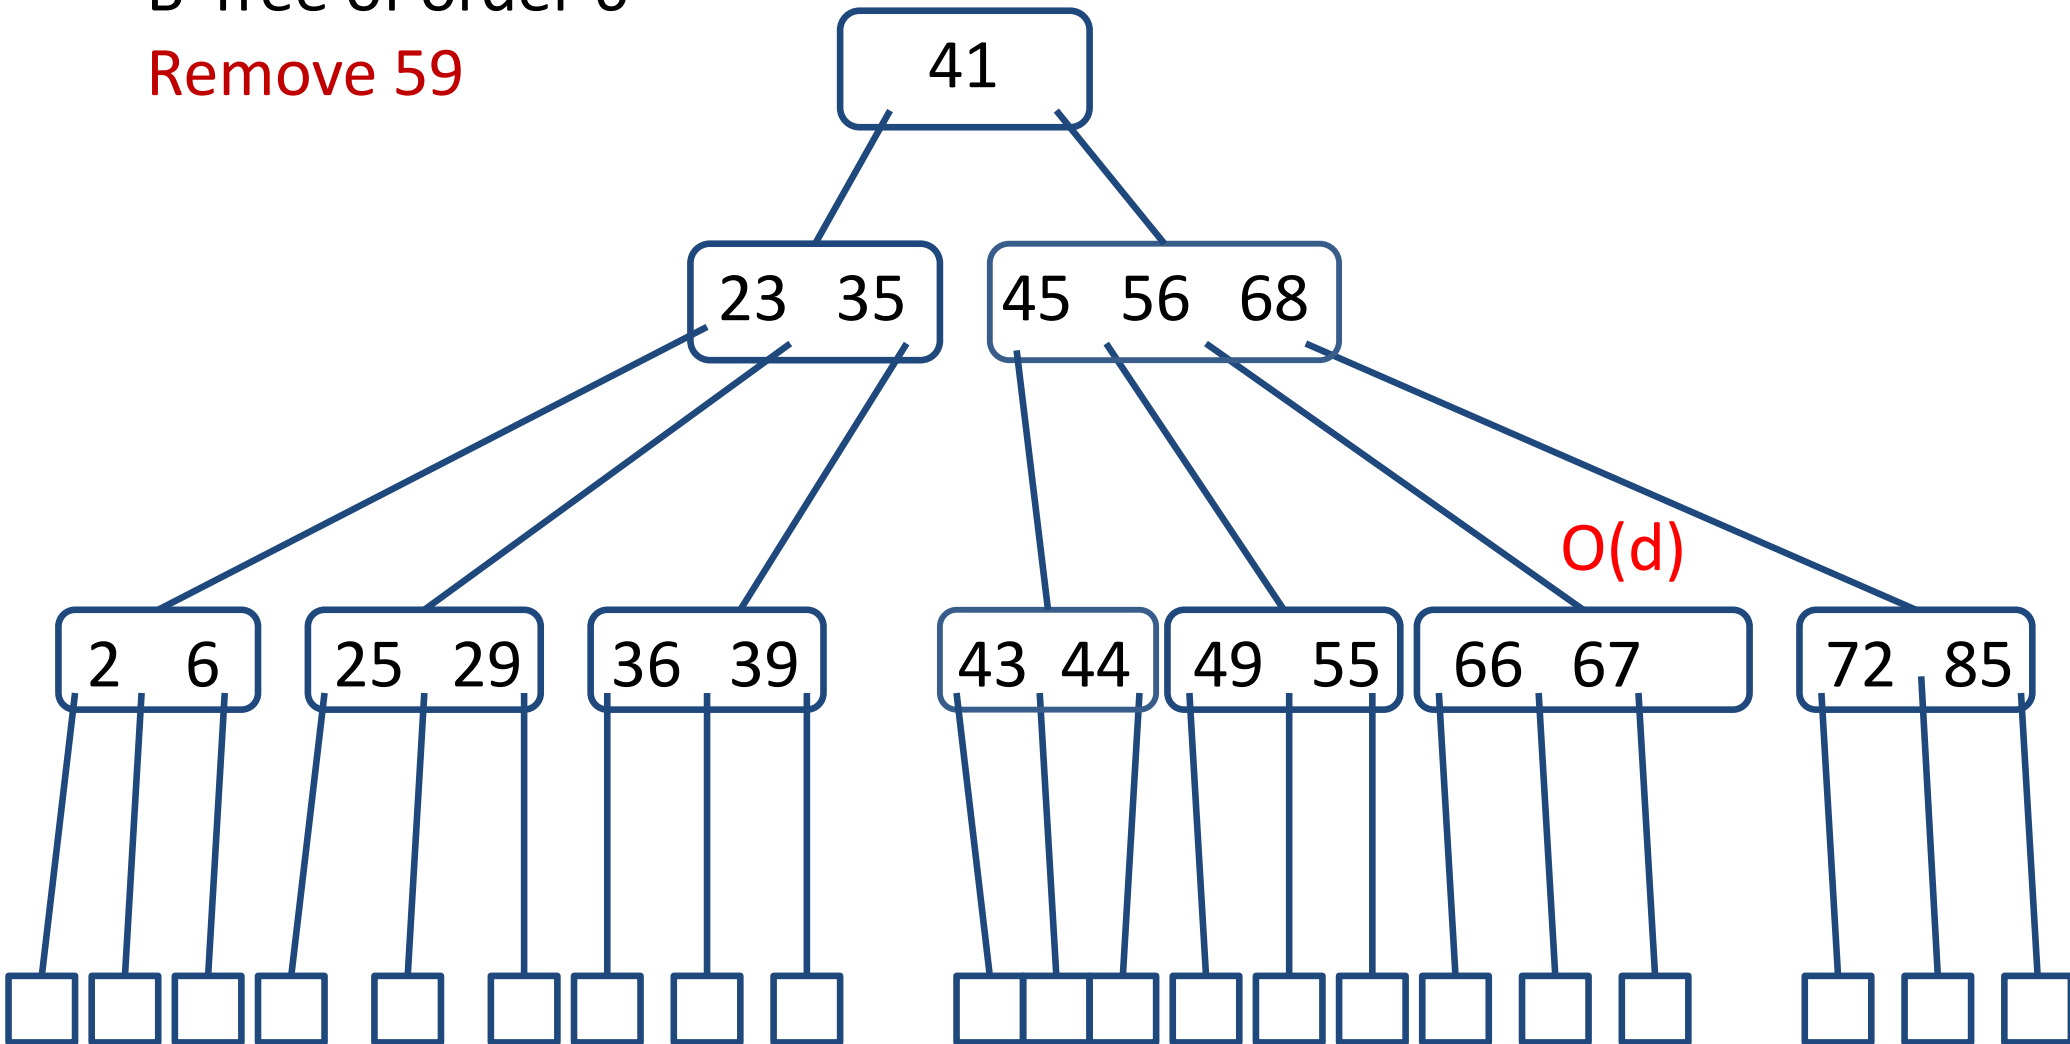
Remove 59



B-Trees

B-Tree of order 6

Remove 59



Algorithm *remove*(r, k)

In: Root r of a B-tree, key k

Out: {remove data item with key k from the tree}

Find the node v storing key k } $O(\log d \times \log_d n)$

Remove (k, o) **from** v replacing it with successor if needed }

while node v underflows **do** { $O(d)$

if v is the root then

make the first child of v the new root $r := v.\text{child}[0]$

else if a sibling has at least $\lceil d/2 \rceil$ keys **then**

 perform a transfer operation }

else {

perform a fusion operation

$v \leftarrow$ parent of v

 }

}

Algorithm *remove*(r, k)

In: Root r of a B-tree, key k

Out: {remove data item with key k from the tree}

Find the node v storing key k } $O(\log d \times \log_d n)$

Remove (k, o) from v **replacing it with successor** if needed }

while node v *underflows* **do** { $O(d + \log d \times \log_d n)$

if v is the root then

 make the first child of v the new root

else if a sibling has at least $\lceil d/2 \rceil$ keys **then**

 perform a transfer operation

else {

 perform a fusion operation

$v \leftarrow$ parent of v

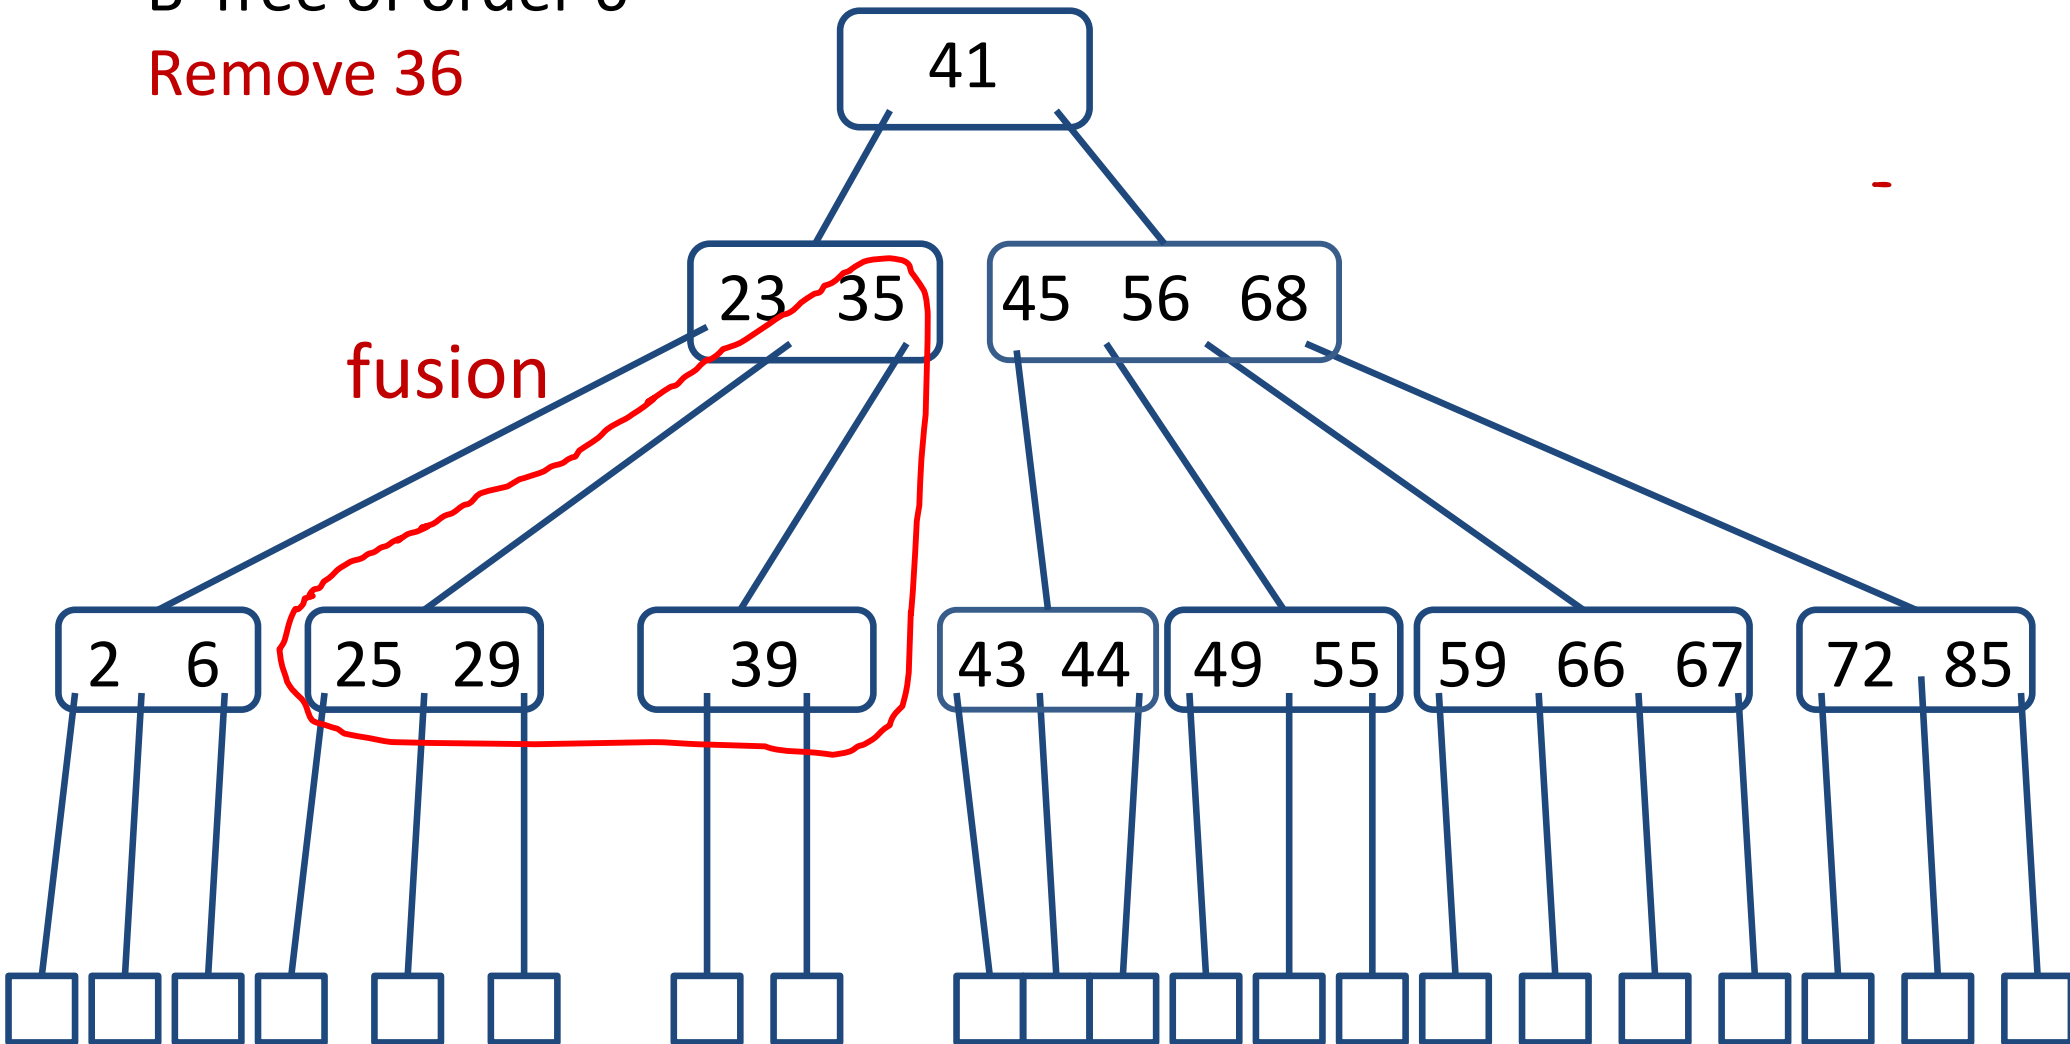
 }

}

B-Trees

B-Tree of order 6

Remove 36

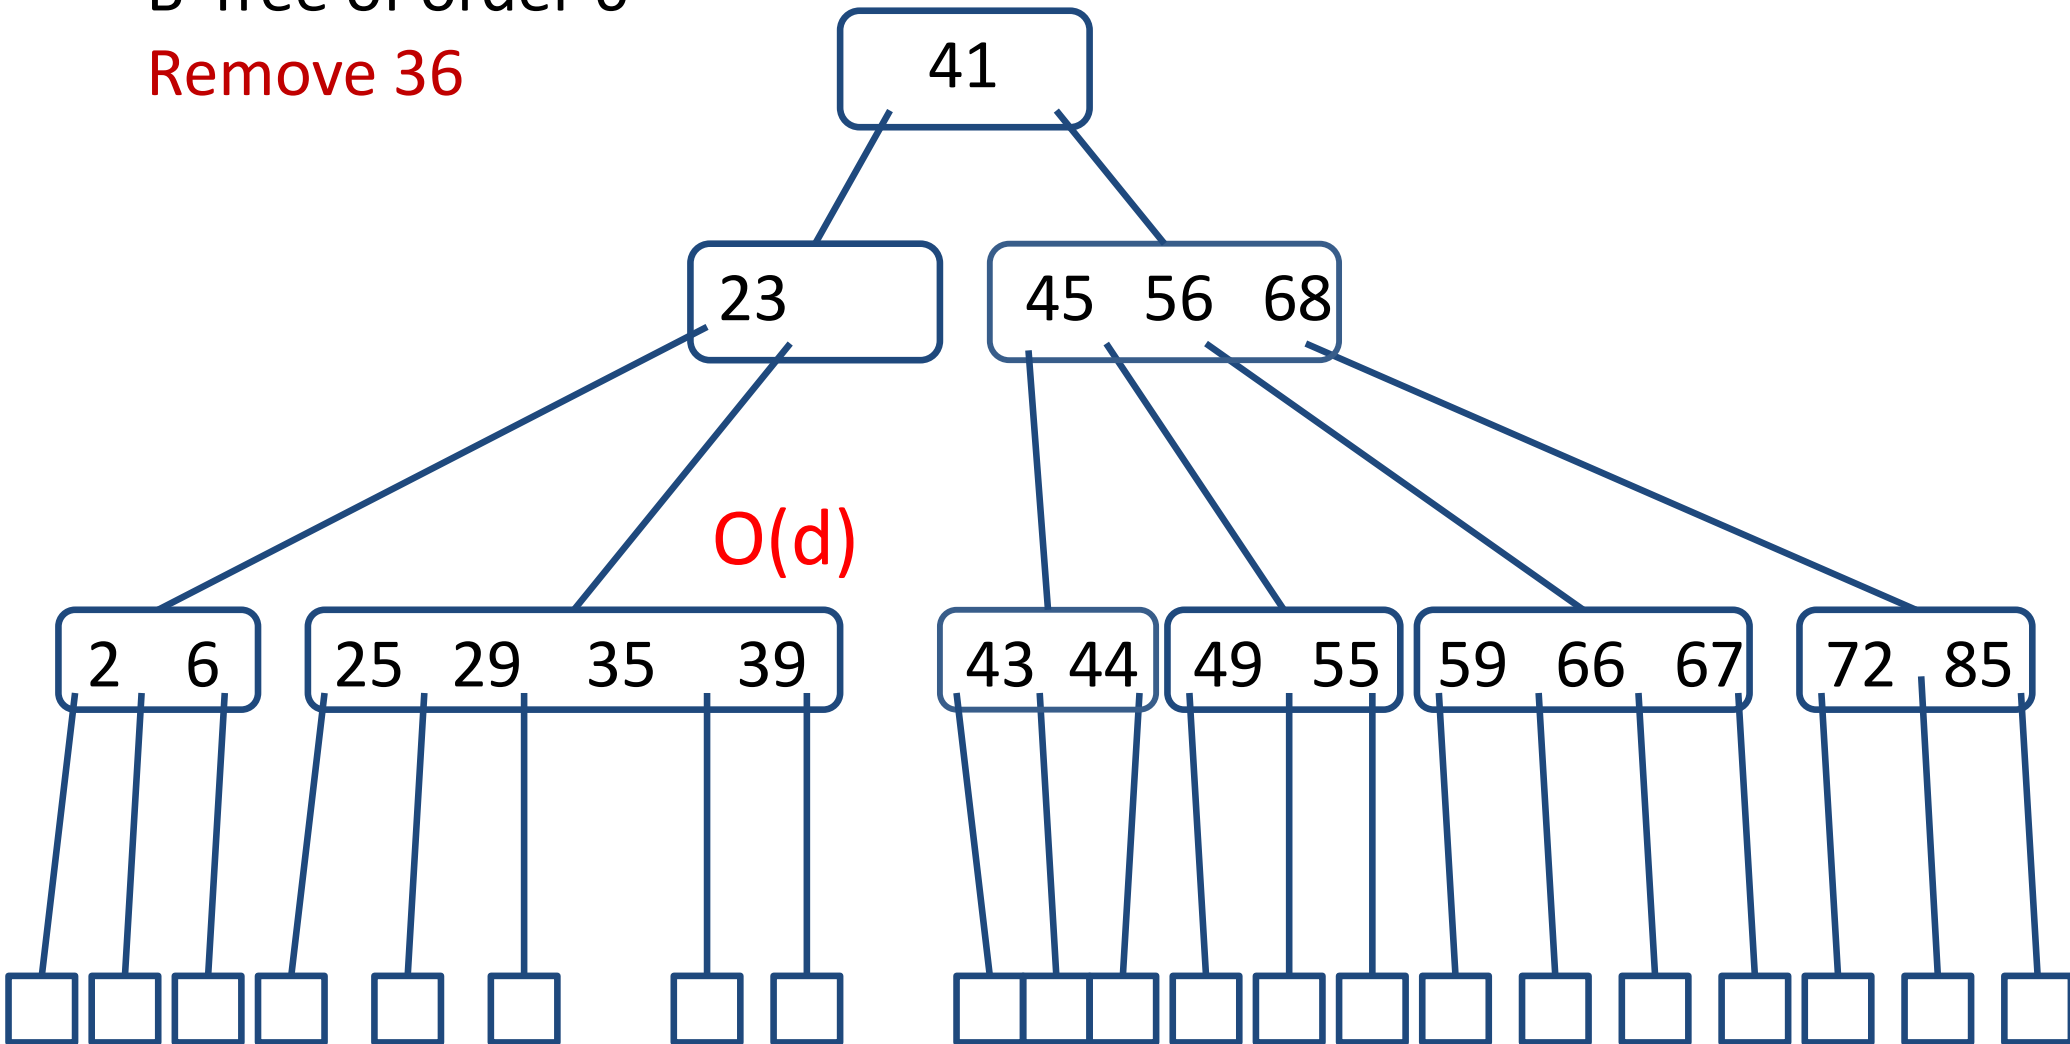


B-Trees

B-Tree of order 6

Remove 36

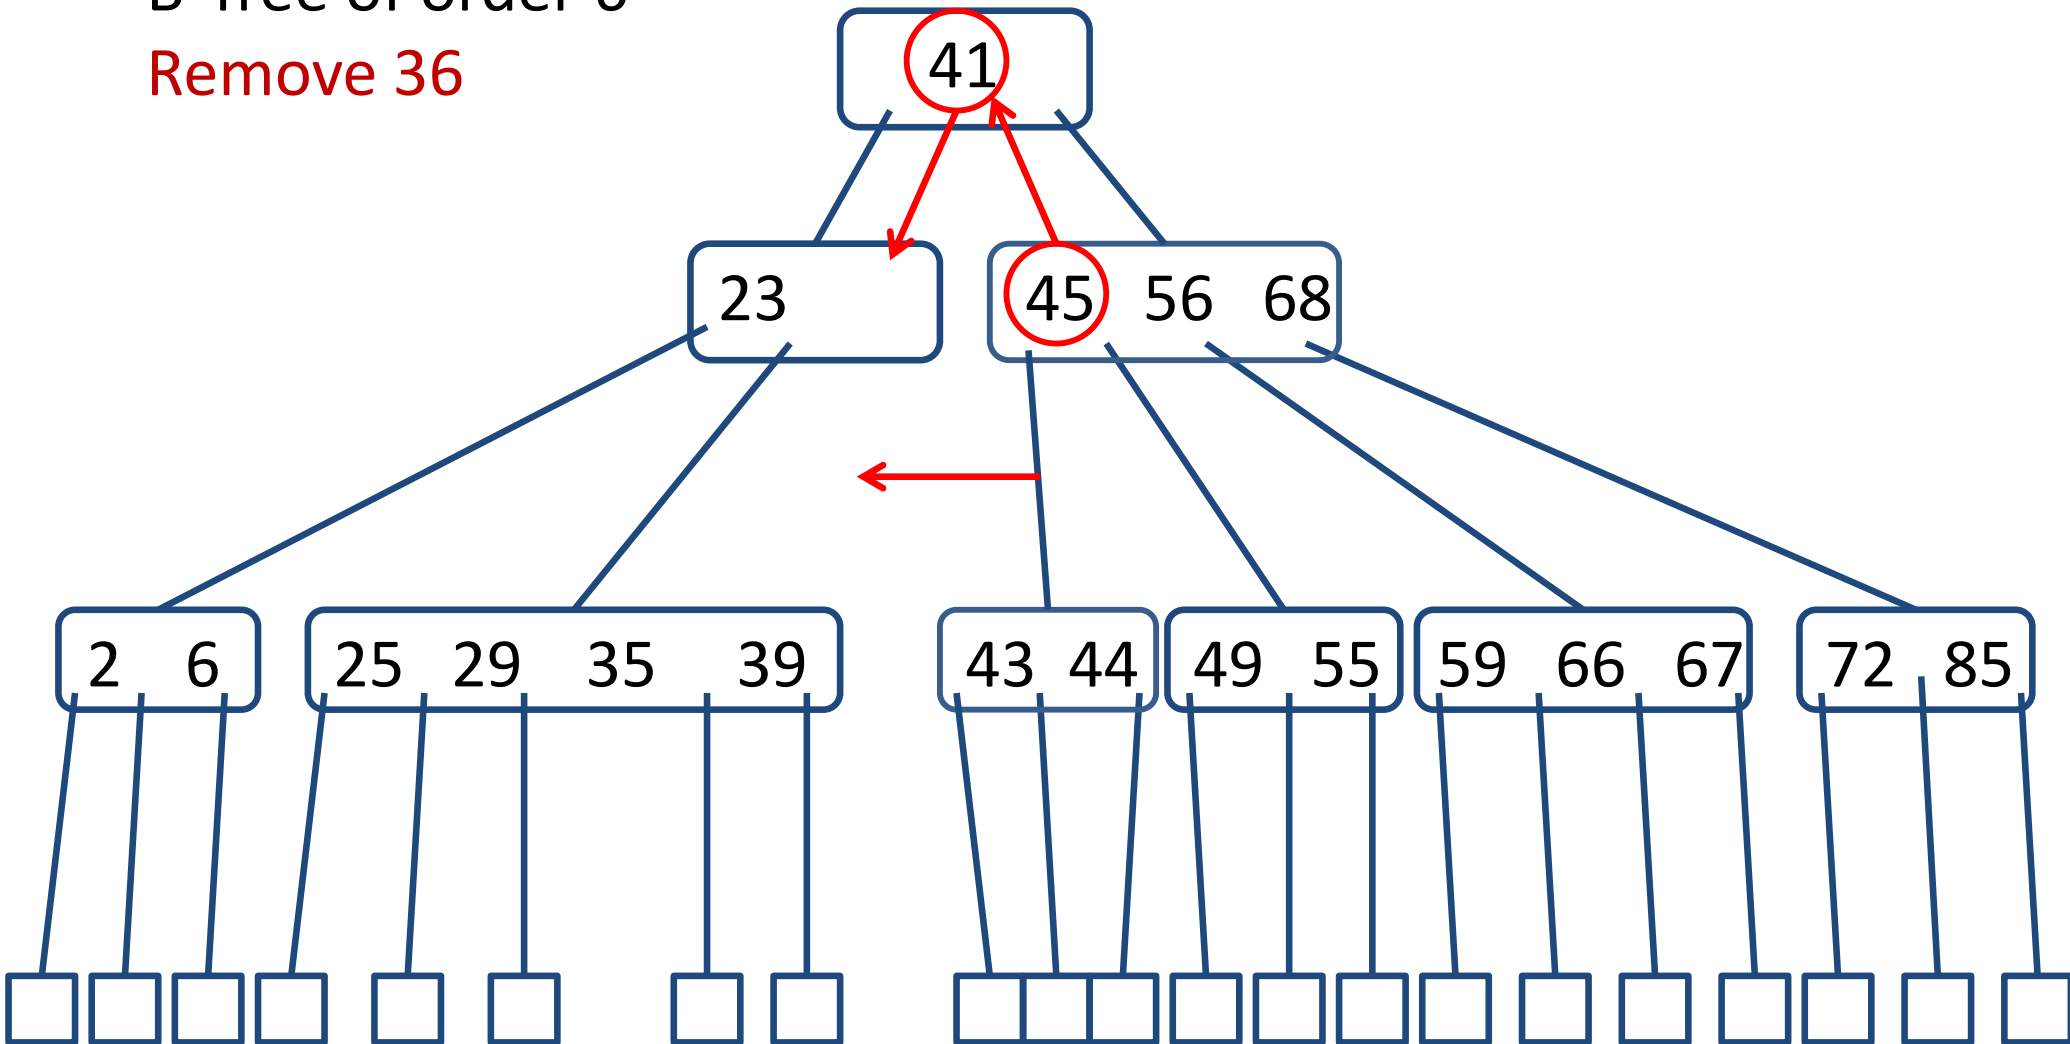
$O(d)$



B-Trees

B-Tree of order 6

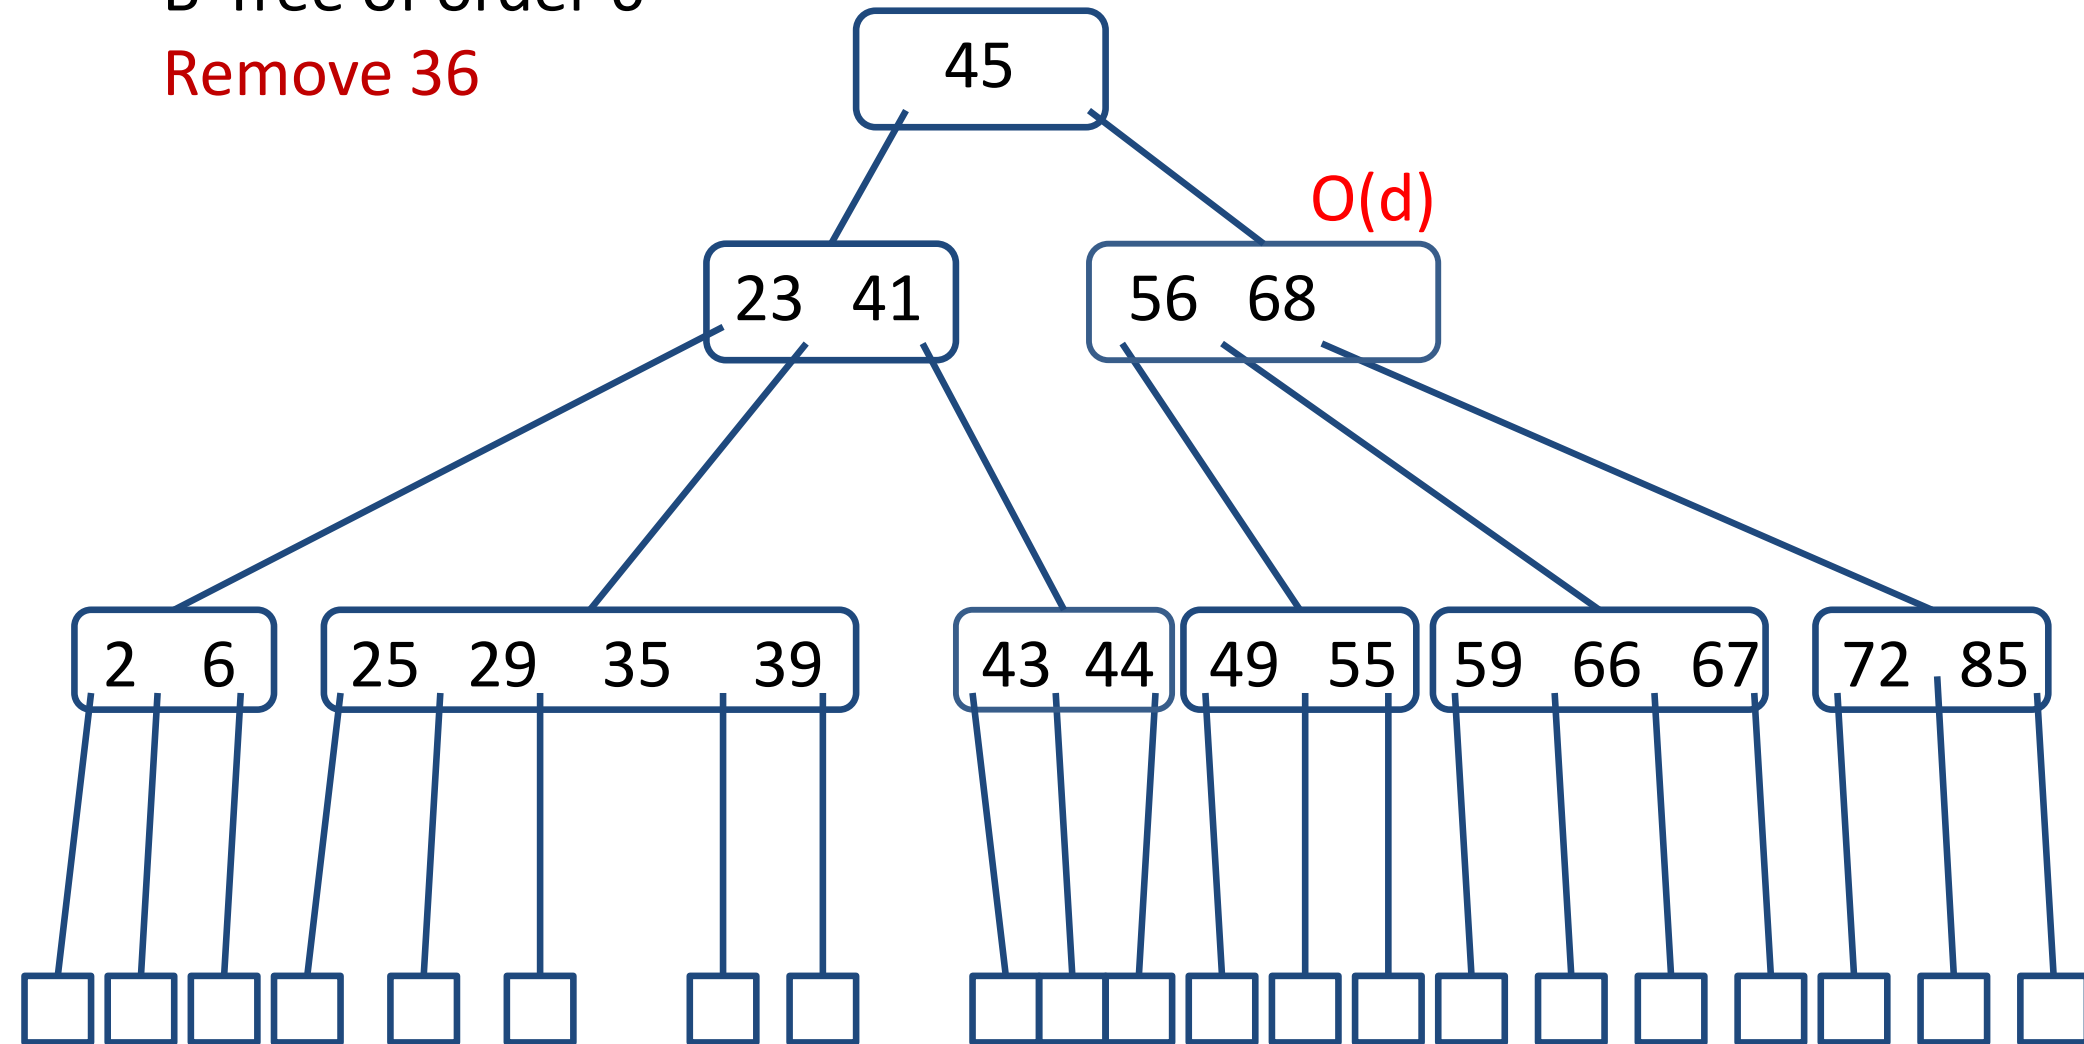
Remove 36



B-Trees

B-Tree of order 6

Remove 36



$O(d)$

Algorithm *remove*(r, k)

In: Root r of a B-tree, key k

Out: {remove data item with key k from the tree}

Find the node v storing key k

} $O(\log d \times \log_d n)$

Remove (k, o) from v replacing it with successor if needed

while node v *underflows* **do** {

$O(d + \log d \times \log_d n)$

if v is the root **then**

 make the first child of v the new root

else if a sibling has at least $\lceil d/2 \rceil$ keys **then**

 perform a transfer operation

else {

 perform a fusion operation

$v \leftarrow$ parent of v

 }

} $O(d)$

}

Algorithm *remove*(r, k) Time complexity $O(d \log_d n)$

In: Root r of a B-tree, key k

Out: {remove data item with key k from the tree}

Find the node v storing key k } $O(\log d \times \log_d n)$

Remove (k, o) from v replacing it with successor if needed }

while node v *underflows* **do** { $O(d + \log d \times \log_d n)$

if v is the root **then**

 make the first child of v the new root

else if a sibling has at least $\lceil d/2 \rceil$ keys **then**

 perform a transfer operation

else {

 perform a fusion operation


$v \leftarrow$ parent of v

 }

}


Collection of 10^9 records

AVL tree



Height: $O(\log_2 n)$
 $\log_2 10^9 \approx 30$


B-tree of degree 10^3



Height: $O(\log_d n)$
 $\log_{1000} 10^9 = 3$

Collection of 10^9 records

AVL tree




Height: $O(\log_2 n)$
 $\log_2 10^9 \approx 30$

Number of comparisons
to find a key ≈ 30

B-tree of degree 10^3

Binary search: $O(\log_2 d)$
 $\log_2 10^3 \approx 10$

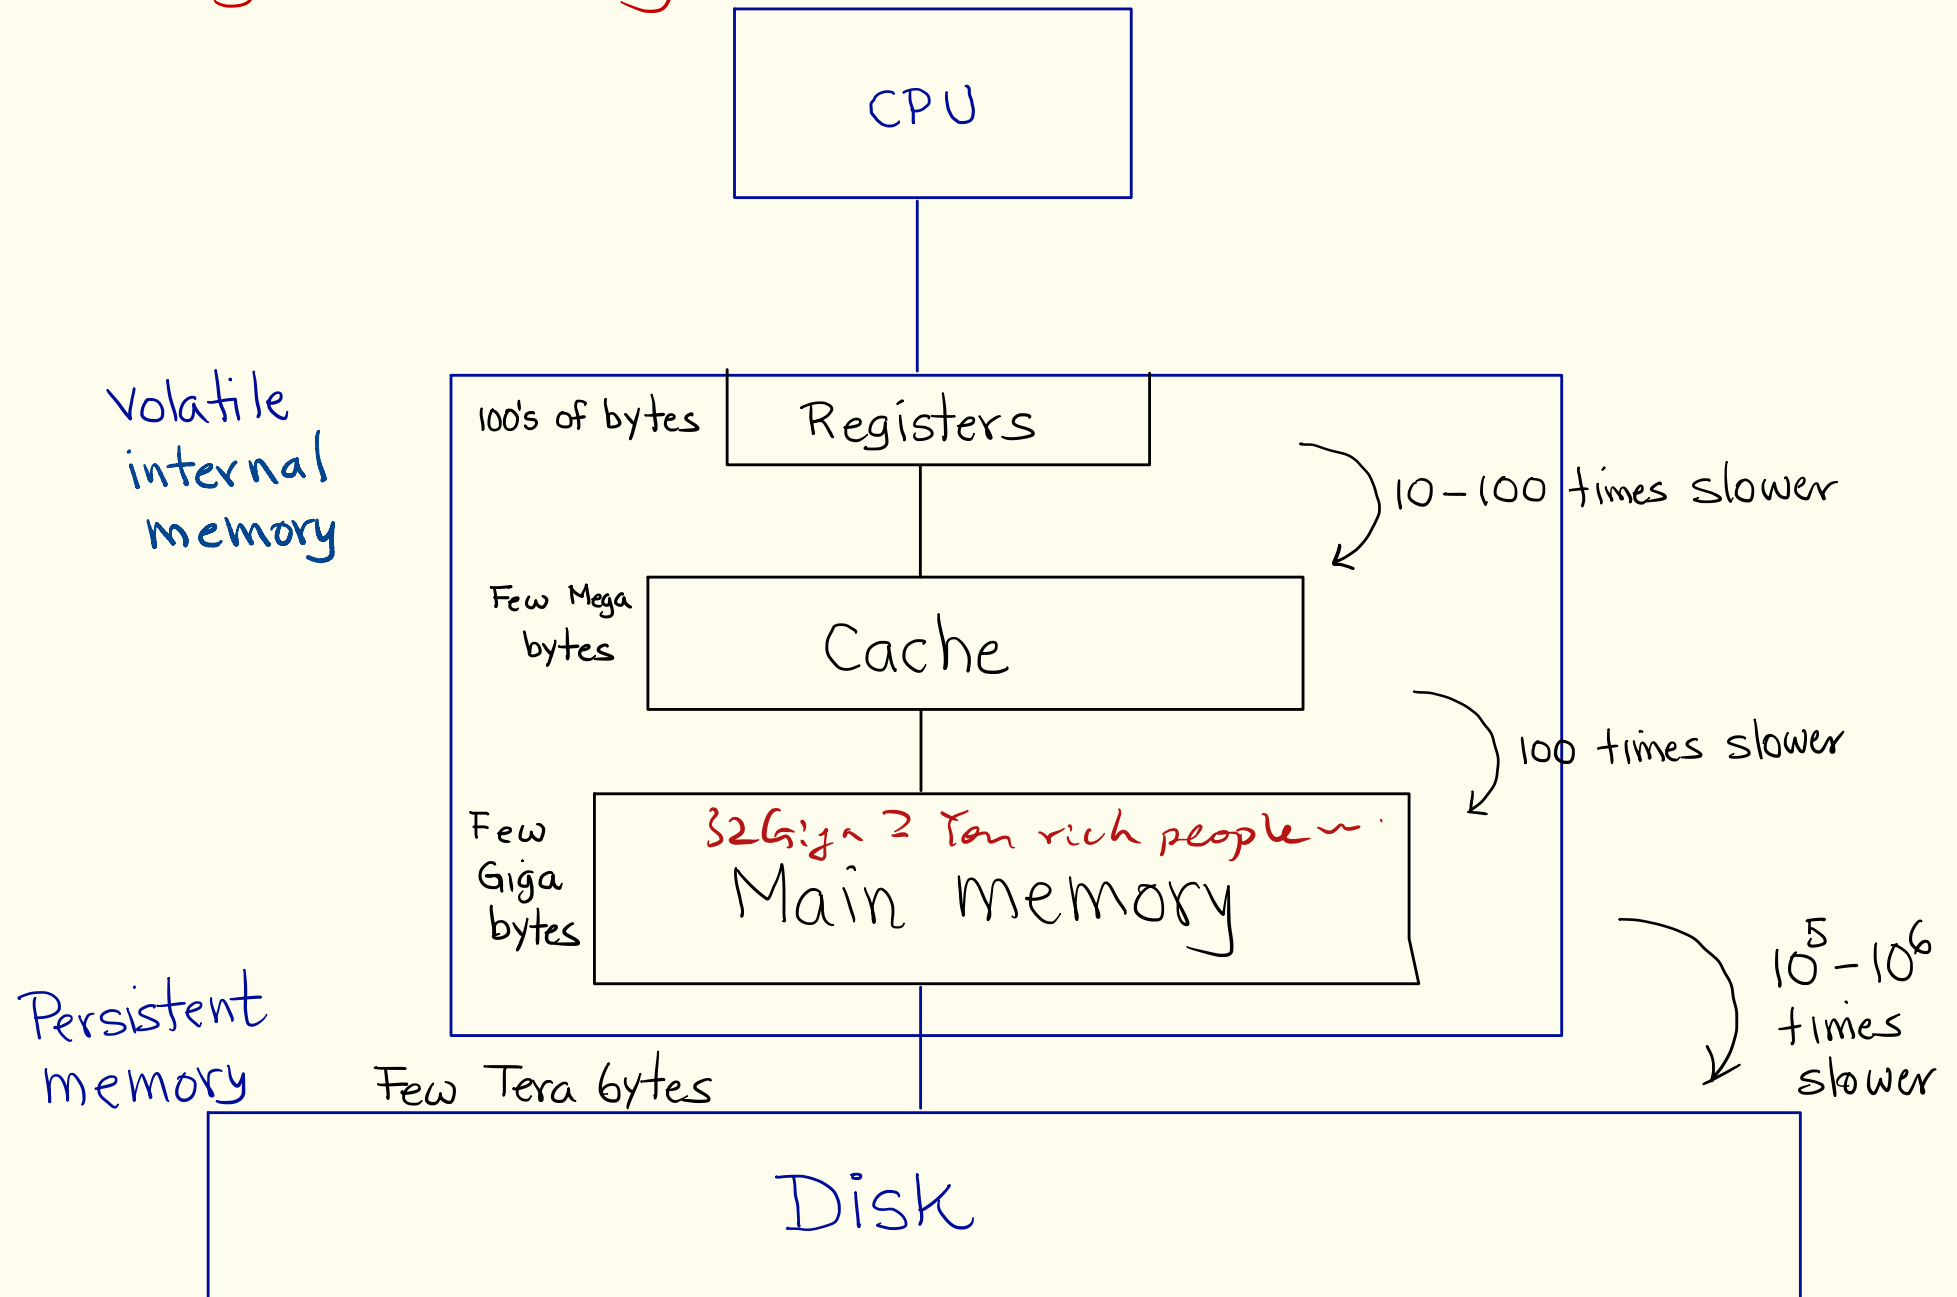


Height: $O(\log_d n)$

$\log_{1000} 10^9 = 3$


$O(\log d \times \log_d n) = \log_{10^3} 10^9$
 ≈ 30

Memory Hierarchy



Collection of 10^9 records

AVL tree




$$\text{Height: } O(\log_2 n)$$
$$\log_2 10^9 \approx 30$$

Number of disk accesses: 30

AVL

B-tree of degree 10^3



$$\text{Height: } O(\log_a n)$$

$$\log_{1000} 10^9 = 3$$

B-tree

3

height = 3

So a B-tree is much better.

Disk Blocks

- Consider the problem of maintaining a large collection of items that does not fit in main memory, such as a typical database.
- In this context, we refer to the external memory is divided into blocks, which we call **disk blocks**.
- The transfer of a block between external memory and primary memory is a **disk transfer** or **I/O**.
- There is a great time difference that exists between main memory accesses and disk accesses
- Thus, we want to minimize the number of disk transfers needed to perform a query or update. We refer to this count as the **I/O complexity** of the algorithm involved.

Memory Hierarchies

- Computers have a hierarchy of different kinds of memories, which vary in terms of their size and distance from the CPU.
- Closest to the CPU are the internal **registers**. Access to such locations is very fast, but there are relatively few such locations.
- At the second level in the hierarchy are the memory **caches**.
- At the third level in the hierarchy is the **internal memory**, which is also known as main memory or core memory.
- Another level in the hierarchy is the **external memory**, which usually consists of disks.

