

Lab 6

October 7, 2021

1 Lab 6

In this lab we discuss statistical modelling and least squares regression.

1.1 Simple Linear Regression

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from scipy import stats
```

```
[2]: # Read .csv data
possum = pd.read_csv("possum.csv")
# The possum dataset consists of morphometric measurements on 46 possums.
possum.head()
```

```
[2]:
```

	sex	age	headL	skullW	totalL	tailL
0	m	8	94.1	60.4	89.0	36.0
1	f	6	92.5	57.6	91.5	36.5
2	f	6	94.0	60.0	95.5	39.0
3	f	6	93.2	57.1	92.0	38.0
4	f	2	91.5	56.3	85.5	36.0

stats.linregress: <https://docs.scipy.org/doc/scipy/reference/generated/scipy.stats.linregress.html>

```
[3]: # Building a simple model with only one term
# Simple linear regression is a linear regression model with a single
#   ↪ explanatory variable.
model = stats.linregress(x = possum['age'], y = possum['headL'])
```

```
[4]: # Finding the slope of the regression line
model.slope
```

```
[4]: 0.5631158455392805
```

```
[5]: # Finding the intercept of the regression line
model.intercept
```

[5]: 90.08288948069242


```
[6]: # Finding the correlation coefficient (R)
model.rvalue
```

[6]: 0.4011016610119052

```
[7]: # Calculating the R-squared
(model.rvalue ** 2)
```


[7]: 0.16088254246650932

```
[8]: # Plotting the fitted line
X = pd.DataFrame({"age":np.linspace(0, 9, 45)})
y_pred = model.intercept + model.slope * X
plt.scatter(possum['age'], possum['headL'], label = 'original data')
plt.plot(X, y_pred, color = 'red', label = 'fitted line')
plt.xlabel("age")
plt.ylabel("headL")
plt.legend()
plt.show()
```




sns.lmplot: <https://seaborn.pydata.org/generated/seaborn.lmplot.html>

```
[9]: # Alternative way to plot the fitted line
sns.lmplot(x = "age", y = "headL", data = possum, ci = None)
plt.xlim(0,10)
plt.show()
```



```
[10]: # Plotting two fitted lines, one for female and another for male possums
sns.lmplot(x = "age", y = "headL", data = possum, hue = "sex", ci = None);
plt.xlim(0,10)
plt.show()
```



1.2 Effect of an Influential Point

```
[11]: # Let's look at the last 5 rows of the original dataset!
# The original dataset has 46 rows and 6 columns.
possum.tail()
```

```
[11]:
```

	sex	age	headL	skullW	totalL	tailL
41	m	3	85.3	54.1	77.0	32.0
42	f	2	90.0	55.5	81.0	32.0
43	m	3	85.1	51.5	76.0	35.5
44	m	3	90.7	55.9	81.0	34.0
45	m	2	91.4	54.4	84.0	35.0

pd.DataFrame.copy: <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.copy.html>

```
[12]: # Now we would like to add an influential point and see how this affects the
↪ regression line.
```

```

# We decide to add a point that behaves like an outlier to the bottom of the
↳ dataframe.
# Let's add a point representing "age" = 9 and "headL" = 86.
# We call the new dataframe "possum_modified".
possum_modified = possum.copy()
possum_modified.loc[46, 'age'] = 9
possum_modified.loc[46, 'headL'] = 86

```

```

[13]: # Now let's look at the last 5 rows of the modified dataframe!
possum_modified.tail()

```

```


[13]:      sex  age  headL  skullW  totalL  tailL
42    f  2.0   90.0    55.5    81.0   32.0
43    m  3.0   85.1    51.5    76.0   35.5
44    m  3.0   90.7    55.9    81.0   34.0
45    m  2.0   91.4    54.4    84.0   35.0
46  NaN  9.0   86.0     NaN     NaN    NaN

```

```

[14]: # The new point that is added to the dataframe is identified in red in the
↳ scatterplot below.
plt.scatter(possum_modified.loc[0:45, 'age'], possum_modified.loc[0:45,
↳ 'headL'], color = 'blue')
plt.scatter(possum_modified.loc[46, 'age'], possum_modified.loc[46, 'headL'],
↳ color = 'red')
plt.xlabel("age")
plt.ylabel("headL")
plt.show()

```



```
[15]: # Building a new model using the modified dataframe
model_modified = stats.linregress(x = possum_modified['age'], y =
↳possum_modified['headL'])
```

```
[16]: # Finding the slope of the regression line
model_modified.slope
```

```
[16]: 0.3585469261900172
```

```
[17]: # Finding the intercept of the regression line
model_modified.intercept
```

```
[17]: 90.71952206590865
```


```
[18]: # Finding the correlation coefficient (R)
model_modified.rvalue
```

```
[18]: 0.2584020565559145
```

```
[19]: # Calculating the R-squared
(model_modified.rvalue ** 2)
```

```
[19]: 0.06677162283232604
```

```
[20]: # Plotting the fitted line
sns.lmplot(x = "age", y = "headL", data = possum_modified, ci = None)
plt.xlim(0,10)
plt.show()
```



[]: