



Western
UNIVERSITY • CANADA

Files and Directories

Winter 2022



File and Directory Types

- What is a file?
 - A container for data
 - Persistent (stays around)
 - Accessible by a name
- In Unix, EVERYTHING is a file

File and Directory Types

- Unix file types
(https://en.wikipedia.org/wiki/Unix_file_types)
 - Regular
 - Directory
 - Device
 - Link
 - Others

File and Directory Types

- Behind the scenes, the name in a directory is translated to a specific location on disk (disk # → cylinder # → track # → sector # → block #). This is known as an "inode number" (index node number).
- This is why it is called a "directory" instead of a "folder". The directory contains an index of files (like the index in a book). The files do not physically "live" inside a folder.

File and Directory Types

- Regular file
 - Text file – Contain "printable" characters
 - Binary file – Contain any ASCII characters from 0 to 255

File and Directory Types

- Directory ("file")
 - Contains the names and inode numbers for all files and directories in this directory
 - We treat it like a folder

File and Directory Types

- Device file
 - Allows I/O from a device (e.g. Soundcard, mouse, etc.)
- Links
 - Hard links
 - Soft links
 - Like a shortcut

File and Directory Types

- Links (continued)
 - Soft links are popular and are often called "symbolic links" or "symlinks"
 - In `-s <target directory> <link name>`

```
[wbeldman@compute ~]$ ln -s /tmp ~/global-temp-directory
[wbeldman@compute ~]$ ls -l ~/global-temp-directory
lrwxrwxrwx 1 wbeldman wbeldman 4 Jan 14 01:50 /home/wbeldman/global-temp-directory -> /tmp
```

File and Directory Types

- How do you find out what type a file is?
Three suggestions:
 - `ls -F`
 - `ls -l`
 - `ls --color=auto`
 - (or all of them: `ls -F -l --color=auto`)

Pathnames and links

- Unix has directories and subdirectories. When referring to a path (e.g. the `cd` command), you will use one of two types
 - Absolute path
 - Relative path

Pathnames and links

- Absolute path
 - Begins with a / (the root)
 - Explicitly uses the entire path from root all the way to the subdirectory
 - E.g. /home/wbeldman

Pathnames and links

- Relative path
 - Dependent on what your current working directory is. E.g.
 - `cd tmp` - means change directory to tmp inside my current working directory. This is not the same as `cd /tmp`

Pathnames and links

- When referring to a location in the system, your command will check in the following order
 - / - The root
 - ~ - The home directory
 - . – The current directory (./ to be more explicit)
 - .. – The parent directory (../ to be more explicit)
 - Otherwise try the current working directory

Pathnames and links

- E.g. If I am in /home/wbeldman, the following are all equivalent
 - /home/wbeldman/cs2211/readme.txt
 - ~/cs2211/readme.txt
 - cs2211/readme.txt
 - ~/cs2211/otherfolder/../readme.txt

Pathnames and links

- Another way to combine shortcuts are like this:
 - `cd ../../../` - Go up three directories in the tree
- When running a command found in the current working directory (e.g. your compiled C program), use the `./` shortcut
 - `./myProgram`

Wildcarding

- We can use special characters to represent a sequence of other characters
- When using a wildcard to match multiple files, this is known as "globbing"
 - * – matches 0 or more characters
 - ? – matches exactly one character
 - [...] – matches any one character in the list

Wildcarding

- E.g.
 - `a*.c*` matches **abc.c** and **abra.cpp**
 - `a?.c` matches **ab.c**, **ax.c**, but not **abc.c**
 - `b[aei]t` matches **bat**, **bet**, or **bit**, but not **but** or **baet**
 - `b[!aei]t` matches **but** but not **bat**, **bet**, **bit**, or **baet**

Wildcarding

- E.g. combining sequences
 - `mv a*[ch] cfiles/` – move all files beginning with a and ending in .c or .h into cfiles
 - `ls [abc]*.*` – list all files beginning with a, b, or c, followed by (possibly) anything, followed by a dot, followed by a single character

Wildcarding

- Wildcards do not traverse directories. It only matches in the current directory
 - E.g. `csnow*c` does NOT match `csnow/codec`
- Wildcards do not match "hidden files".
Hidden files are files that start with a . (dot)
 - E.g. `"cat *profile"` will not find `".bash_profile"`

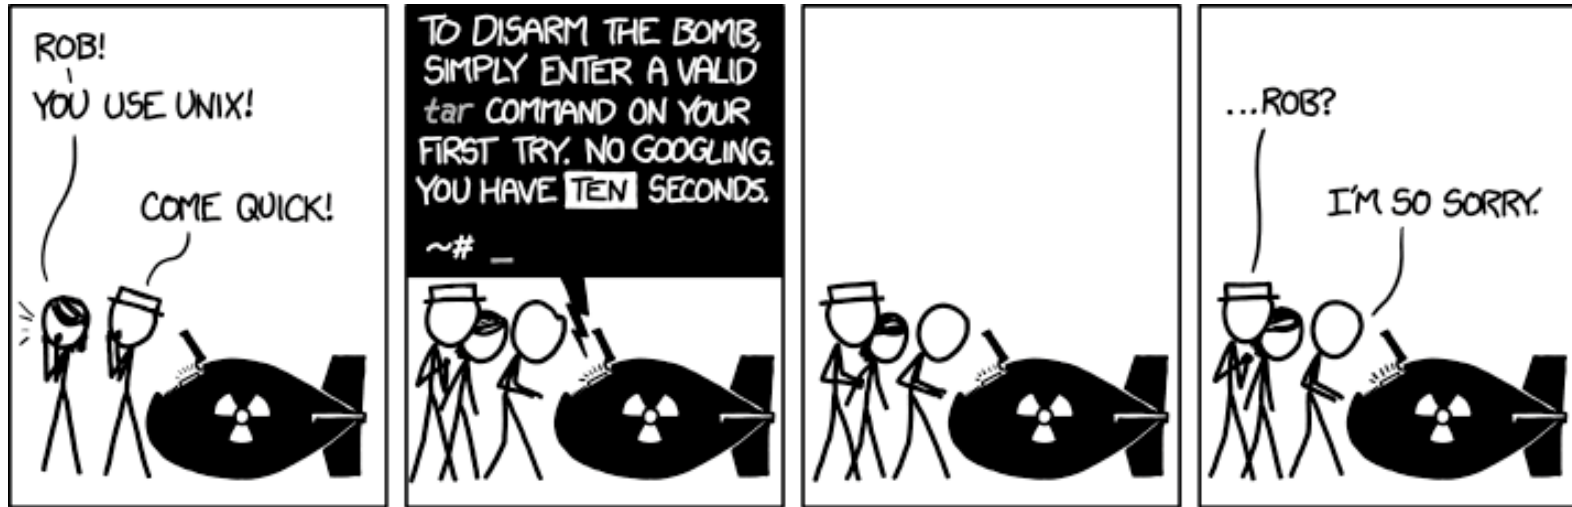
Wildcarding

- [...] allows a range of characters instead of explicitly listing each one
 - E.g. `ls [a-z]*` matches all files beginning with a-z
- The wildcard expansion is done by your shell, NOT by the program you are running
 - E.g. When `ls *` is run, it is actually doing `ls file1.txt file2.txt file3.txt ...`

Wildcarding

- Other advanced examples to think about
 - `ls /bin/*[-_]*`
 - What's the difference between `ls *` and `ls`
 - `mv *.bat *.bit` – This does not rename all .bat files as .bit files as you might expect

Tar



Tar

- An archiving command
- Used to "bundle up" a directory and make a single file out of it, or "unpack" a file into a new directory
- The resulting file by convention has a .tar extension and is called a "tarball"
- Great for backups/snapshots, submitting your code in the assignments ;)

Tar

- To create a tarball out of a directory called Assignment2, use
 - `tar cvf Assignment2.tar Assignment2`
 - `c == create, v == verbose, f == filename`
- To create Assignment2 out of a tarball, use
 - `tar xvf Assignment2.tar`
 - `x == extract, v == verbose, f == filename`

Tar

- You might recognize this as zipping or unzipping a directory
- Technically this is not exactly the same because the tarball is not compressed/uncompressed
- If you supply `z` as in `cvzf/xvzf`, this will create an archive file and compress it. By convention, we use `.tar.gz` as the extension

Finding files

- Use find to recursively locate files in a large directory structure
- A very powerful tool that can (among many things)
 - Match wildcards in file names
 - Match based on file size, permissions, creation time, etc.
 - Execute commands on each file found

Finding files

- To use find, the syntax follows the format:
find <path> <expression>
e.g.
 - find ./ -name "README" – Find all files and directories under the current directory called "README"
 - find /usr/include -name "*.h" – Find all files and directories under /usr/include ending in .h

Finding files

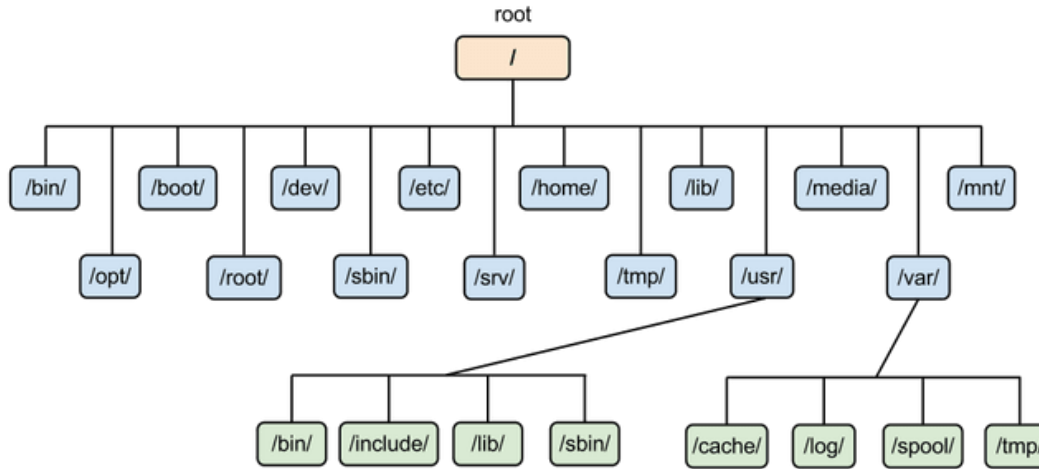
- More examples
 - `find ./ -type f -name "README"` – Find all files (but not directories) under the current directory called "README"
 - `find ./ -type d -name "README"` – Find all directories (but not files) under the current directory called "README"

Finding files

- Use "man find" to read the manual and discover more options

The Unix File System

- A Unix filesystem is organized like an upside down tree



The Unix File System

- The top of the file system is known as the "root" directory and is identified with a single slash (/)
- This is NOT the same as a backslash (\) which is commonly used in Windows
- What kind of files are stored in which directory are done by convention only.

The Unix File System

- The Unix file system can transparently span multiple disks (including network attached disks).
- As a regular user, you don't actually need to know this detail, nor should you care
- You can use the `df` command to see other disks

```
filestore.csd.uwo.ca:/data/cs_homes/wbeldman 38T 3.9T 34T 11% /home/wbeldman
```

The Unix File System

- Some common directories in Unix

- / the root
- /bin **bin**aries (executables)
- /dev **dev**ices (peripherals)
- /devices where the **dev**ices really live
- /etc startup and control files
- /lib **lib**raries (really in /usr)
- /opt **opt**ional software packages
- /proc access to **proc**esses
- /sbin **s**tandalone **bin**aries
- /tmp place for **temp**orary files

- /usr **user** stuff
- /usr/bin **bin**aries again (user)
- /usr/include **include** files for compilers
- /usr/lib **lib**raries of functions etc.
- /usr/local **local** stuff
- /usr/local/bin local **bin**aries
- /usr/local/lib local **lib**raries
- /usr/openwin X11 stuff
- /usr/sbin sysadmin stuff
- /usr/tmp place for more **temp**orary files
- /usr/ucb **ucb** binaries
- /var **var**iable stuff

The Unix File System

- `/bin` – contains small executable programs (binaries). This is where you find the common commands you are used to (e.g. `ls`, `cd`, `mkdir`, etc.)
- `/sbin` – contains small executable programs (binaries) but are only used by the system administrator.
- `/lib` (and `/usr/lib`) – contains binary library files that other programs might call (e.g. `stdio.h`)

The Unix File System

- /dev – Everything in Unix is a file – even devices. This directory contains device files (see slide 8) e.g. disk drives, input devices
- /boot – The OS kernel lives here. The heart of the operating system. If the rest of the Unix system is broken, you at least need /boot to be available so the OS can boot up

The Unix File System

- /etc – Contains system configuration files (regular files). The OS and other programs typically store their configuration information in this directory
- /proc – Virtual files that represent the current state of the kernel. Processes can refer to files here to retrieve information about the system or other processes

The Unix File System

- /mnt – Commonly used to hold sub-directories that are temporarily mounted. This could be something like /mnt/cdrom or /mnt/windows for dual-boot systems
- /usr – Stores programs and files used by end-users. Non default stuff usually goes here. Think of this like C:\Program Files
- /var – Variable data files. Typically log files

The Unix File System

- /home – Contains the home directories for any user with a login to the system (except root). A home directory is the user's personal space
- /root – The home directory of the "root" user

The Unix File System

- /tmp – Contains temporary files and directories. Accessible by everyone. Many systems periodically purge this directory so DON'T store important files here!
- /opt – Some large applications will choose to bundle all their files and directories here instead of /usr, /etc, /bin, etc.

Working with directories

- We covered making and removing directories already. Let's look at moving and copying directories

Working with directories

- A directory is just a file, so you move it the same as you would any other file:
- `mv <directory1> <directory2>` - Moves <directory1> into <directory2>

Working with directories

- Copying is a little different. You have to explicitly tell cp that you want to copy the directory and any files below it.
- Use the -r argument to copy recursively. E.g.
 - `cp -r <directory1> <directory2>` - Copy directory 1 and everything below it into directory 2

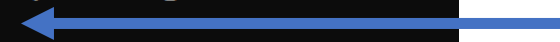
Working with directories

- pushd and popd can be used to maintain a list (a stack) of directories
 - pushd <directory1> - change directory to directory1 and put directory1 on top of the stack
 - popd - change directory to the top of the stack and remove it from the stack
- The current stack is also printed for you

Working with directories

- Use the dirs command to display the current stack

```
[wbeldman@compute ~]$ pushd /tmp
/tmp ~
[wbeldman@compute tmp]$ pushd /
/ /tmp ~
[wbeldman@compute /]$ pushd ~
~ / /tmp ~
[wbeldman@compute ~]$ dirs
~ / /tmp ~
[wbeldman@compute ~]$ popd
/ /tmp ~
[wbeldman@compute /]$ popd
/tmp ~
[wbeldman@compute tmp]$ popd
~
```



File names

- Almost any character can be in a filename except / and the "null" character
- Still, don't use these special characters:
? * [] " ' () & : ; !
- Don't use these as the first character:
- ~
- If you do, you're gonna have a bad time!

File names

- Unix is case sensitive. Upper and lower case are different. A.txt and a.txt are different files
- Unix doesn't use extensions.
 - a
 - a.
 - .a
 - ...
 - a.b.cAre all valid filenames

File names

- Unix is case sensitive. Upper and lower case are different. A.txt and a.txt are different files
- Unix doesn't use extensions.
 - a
 - a.
 - .a
 - ...
 - a.b.care all valid filenames

File names


- Extensions are still useful to the user so in practice they are still used. E.g.
 - .c means a C program
 - .jpg for a JPEG image file
 - .txt means a text file
 - .mp3 for a music file

File names


- Executable files do not have an extension either. Windows usually uses the .exe extension.
- There are limits to the length of the name
 - Typically 255 characters for a file name and 4096 characters for the entire path

File names

- Use the "file" command to gather info on a file



```
[wbeldman@compute Lecture-2]$ file a.out
a.out: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically linked, interpreter /lib64/ld-linux-x86-64.so.2, BuildID[sha1]=14f113374b1fc6f8a009211d78f48bf3cb4f86ad, for GNU/Linux 3.2.0, not stripped
```



```
[wbeldman@compute Lecture-2]$ file pun.c
pun.c: C source, ASCII text
```

File names

- Files beginning with a . (dot) are "hidden" files. `ls` will not list them by default. You have to use "`ls -a`"
- You have a bunch of these in your home directory already. These typically hold personal configuration files rather than storing them in `/etc` for all to see

Unix Quoting

- Use quotations marks to stop the shell from interpreting special characters (e.g. whitespace, *, or ~)
- Here's an example with "

```
[wbeldman@compute ~]$ echo Here are some words
Here are some words
[wbeldman@compute ~]$ echo "Here are some words"
Here are some words
```

Unix Quoting

- Here's an example with ‘

```
[wbeldman@compute ~]$ echo Welcome to $HOME
Welcome to /home/wbeldman
[wbeldman@compute ~]$ echo Welcome to '$HOME'
Welcome to $HOME
```

Unix Quoting

- You can use ``` (backtick) to "insert the results of a command". E.g.

```
[wbeldman@compute Lecture-2]$ file `ls pun*`  
pun:      ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically l  
inked, interpreter /lib64/ld-linux-x86-64.so.2, BuildID[sha1]=d8af98c618c  
59bdae46edec6c27f661848f4e085, for GNU/Linux 3.2.0, not stripped  
pun.c: C source, ASCII text
```

- Note that a backtick (```) and single quote (`'`) are not the same!

Unix Quoting

- The backslash "quotes" or "escapes" the next character (like a newline or another special character). E.g.

```
[wbeldman@compute Lecture-2]$ echo "This is a long line \  
> that continues \  
> over \  
> several lines"  
This is a long line that continues over several lines  
[wbeldman@compute Lecture-2]$ echo \  
*  
[wbeldman@compute Lecture-2]$ echo "Escaping \" character"  
Escaping " character  
[wbeldman@compute Lecture-2]$
```


Miscellaneous

- A few other useful commands to know
 - `head <filename>` - View the first few lines of a file
 - `tail <filename>` - View the last few lines of a file

Miscellaneous

- A few other useful commands to know
 - `alias <mycommand>=<another command>`

```
[wbeldman@compute Lecture-2]$ ll
-bash: ll: command not found
[wbeldman@compute Lecture-2]$ alias ll='ls -l'
[wbeldman@compute Lecture-2]$ ll
total 66
-rwx----- 1 wbeldman wbeldman 24352 Jan 11 23:22 a.out
-rwx----- 1 wbeldman wbeldman 24352 Jan 11 23:49 pun
-rwxr-x--- 1 wbeldman wbeldman  115 Jan 11 23:49 pun.c
```

Miscellaneous

- A few other useful commands to know
 - `which <command>` - Where is a command located. The command `whereis` works similarly. This can help you find out if a command is installed or not.
 - `whatis <command>` - A one-liner description of `<command>` (This is drawn from the man page)

Miscellaneous

- A few other useful commands to know
 - clear – clear all the text off of the screen
 - history – a running history of all the commands you have run
 - touch <filename> - Updates the "update time" on a file. If <filename> does not exist, this is a useful way to create a new empty file

Miscellaneous

- A few other useful commands to know
 - `echo "Some Text"` – Write "Some Text" to the screen. This will be very useful when writing shell scripts
 - `grep "Some Text" <filename>` - Search for "Some Text" inside <filename>



Western
UNIVERSITY • CANADA