

Chapter 12

Pointers and Arrays

Introduction


- C allows us to perform arithmetic—addition and subtraction—on pointers to array elements.
- This leads to an alternative way of processing arrays in which pointers take the place of array subscripts.
- The relationship between pointers and arrays in C is a close one.
- Understanding this relationship is critical for mastering C.

Pointer Arithmetic

- Chapter 11 showed that pointers can point to array elements:

```
int a[10], *p;  
p = &a[0];
```

- A graphical representation:




Pointer Arithmetic

- We can now access `a[0]` through `p`; for example, we can store the value 5 in `a[0]` by writing

`*p = 5;`

- An updated picture:



Pointer Arithmetic

- If p points to an element of an array a , the other elements of a can be accessed by performing *pointer arithmetic* (or *address arithmetic*) on p .
- C supports three (and only three) forms of pointer arithmetic:
 - Adding an integer to a pointer
 - Subtracting an integer from a pointer
 - Subtracting one pointer from another

Adding an Integer to a Pointer


- Adding an integer j to a pointer p yields a pointer to the element j places after the one that p points to.
- More precisely, if p points to the array element $a[i]$, then $p + j$ points to $a[i+j]$.
- Assume that the following declarations are in effect:

```
int a[10], *p, *q, i;
```


Adding an Integer to a Pointer

- Example of pointer addition:


```
p = &a[2];
```



```
q = p + 3;
```



```
p += 6;
```




Subtracting an Integer from a Pointer


- If p points to $a[i]$, then $p - j$ points to $a[i - j]$.

- Example:


```
p = &a[8];
```



```
q = p - 3;
```



```
p -= 6;
```




Subtracting One Pointer from Another

- When one pointer is subtracted from another, the result is the distance (measured in array elements) between the pointers.
- If p points to $a[i]$ and q points to $a[j]$, then $p - q$ is equal to $i - j$.
- Example:

```
p = &a[5];
```

```
q = &a[1];
```



```
i = p - q;    /* i is 4 */
```

```
i = q - p;    /* i is -4 */
```

Subtracting One Pointer from Another

- Operations that cause undefined behavior:
 - Performing arithmetic on a pointer that doesn't point to an array element
 - Subtracting pointers unless both point to elements of the same array

Comparing Pointers

- Pointers can be compared using the relational operators ($<$, $<=$, $>$, $>=$) and the equality operators ($==$ and $!=$).
 - Using relational operators is meaningful only for pointers to elements of the same array.
- The outcome of the comparison depends on the relative positions of the two elements in the array.

- After the assignments

```
p = &a[5];
```

```
q = &a[1];
```

the value of $p <= q$ is 0 and the value of $p >= q$ is 1.

Pointers to Compound Literals (C99)

- It's legal for a pointer to point to an element within an array created by a compound literal:

```
int *p = (int []) {3, 0, 3, 4, 1};
```

- Using a compound literal saves us the trouble of first declaring an array variable and then making `p` point to the first element of that array:

```
int a[] = {3, 0, 3, 4, 1};
```

```
int *p = &a[0];
```


Using Pointers for Array Processing

- Pointer arithmetic allows us to visit the elements of an array by repeatedly incrementing a pointer variable.
- A loop that sums the elements of an array `a`:


```
#define N 10
...
int a[N], sum, *p;
...
sum = 0;
for (p = &a[0]; p < &a[N]; p++)
    sum += *p;
```

Using Pointers for Array Processing


At the end of the first iteration:



At the end of the second iteration:



At the end of the third iteration:



Using Pointers for Array Processing

- The condition `p < &a[N]` in the `for` statement deserves special mention.
- It's legal to apply the address operator to `a[N]`, even though this element doesn't exist.
- Pointer arithmetic may save execution time.
- However, some C compilers produce better code for loops that rely on subscripting.

Combining the * and ++ Operators

- C programmers often combine the * (indirection) and ++ operators.
- A statement that modifies an array element and then advances to the next element:

```
a[i++] = j;
```

- The corresponding pointer version:

```
*p++ = j;
```

- Because the postfix version of ++ takes precedence over *, the compiler sees this as

```
*(p++) = j;
```


Combining the * and ++ Operators

- Possible combinations of * and ++:

<i>Expression</i>	<i>Meaning</i>
*p++ or * (p++)	Value of expression is *p before increment; increment p later
(*p) ++	Value of expression is *p before increment; increment *p later
*++p or * (++p)	Increment p first; value of expression is *p after increment
++*p or ++ (*p)	Increment *p first; value of expression is *p after increment

Combining the * and ++ Operators

- The most common combination of * and ++ is *p++, which is handy in loops.

- Instead of writing

```
for (p = &a[0]; p < &a[N]; p++)  
    sum += *p;
```

to sum the elements of the array a, we could write

```
p = &a[0];  
while (p < &a[N])  
    sum += *p++;
```

Combining the * and ++ Operators

- The * and -- operators mix in the same way as * and ++.
- For an application that combines * and --, let's return to the stack example of Chapter 10.
- The original version of the stack relied on an integer variable named `top` to keep track of the “top-of-stack” position in the `contents` array.
- Let's replace `top` by a pointer variable that points initially to element 0 of the `contents` array:

```
int *top_ptr = &contents[0];
```

Combining the * and ++ Operators

- The new push and pop functions:

```
void push(int i)
{
    if (is_full())
        stack_overflow();
    else
        *top_ptr++ = i;
}

int pop(void)
{
    if (is_empty())
        stack_underflow();
    else
        return *--top_ptr;
}
```

Using an Array Name as a Pointer

- Pointer arithmetic is one way in which arrays and pointers are related.
- Another key relationship:
The name of an array can be used as a pointer to the first element in the array.
- This relationship simplifies pointer arithmetic and makes both arrays and pointers more versatile.

Using an Array Name as a Pointer

- Suppose that `a` is declared as follows:

```
int a[10];
```

- Examples of using `a` as a pointer:

```
*a = 7;    /* stores 7 in a[0] */
```

```
*(a+1) = 12; /* stores 12 in a[1] */
```

- In general, `a + i` is the same as `&a[i]`.
 - Both represent a pointer to element `i` of `a`.
- Also, `*(a+i)` is equivalent to `a[i]`.
 - Both represent element `i` itself.

Using an Array Name as a Pointer

- The fact that an array name can serve as a pointer makes it easier to write loops that step through an array.

- Original loop:

```
for (p = &a[0]; p < &a[N]; p++)  
    sum += *p;
```

- Simplified version:

```
for (p = a; p < a + N; p++)  
    sum += *p;
```

Using an Array Name as a Pointer

- Although an array name can be used as a pointer, it's not possible to assign it a new value.
- Attempting to make it point elsewhere is an error:

```
while (*a != 0)
    a++;                                /* ** WRONG ** */
```

- This is no great loss; we can always copy `a` into a pointer variable, then change the pointer variable:

```
p = a;
while (*p != 0)
    p++;
```


Program: Reversing a Series of Numbers (Revisited)

- The `reverse.c` program of Chapter 8 reads 10 numbers, then writes the numbers in reverse order.
- The original program stores the numbers in an array, with subscripting used to access elements of the array.
- `reverse3.c` is a new version of the program in which subscripting has been replaced with pointer arithmetic.

reverse3.c

```
/* Reverses a series of numbers (pointer version) */  
  
#include <stdio.h>  
  
#define N 10  
  
int main(void)  
{  
    int a[N], *p;  
  
    printf("Enter %d numbers: ", N);  
    for (p = a; p < a + N; p++)  
        scanf("%d", p);  
  
    printf("In reverse order:");  
    for (p = a + N - 1; p >= a; p--)  
        printf(" %d", *p);  
    printf("\n");  
  
    return 0;  
}
```

Array Arguments (Revisited)

- When passed to a function, an array name is treated as a pointer.

- Example:

```
int find_largest(int a[], int n)
{
    int i, max;

    max = a[0];
    for (i = 1; i < n; i++)
        if (a[i] > max)
            max = a[i];
    return max;
}
```

- A call of `find_largest`:

```
largest = find_largest(b, N);
```

This call causes a pointer to the first element of `b` to be assigned to `a`; the array itself isn't copied.

Array Arguments (Revisited)

- The fact that an array argument is treated as a pointer has some important consequences.
- *Consequence 1:* When an ordinary variable is passed to a function, its value is copied; any changes to the corresponding parameter don't affect the variable.
- In contrast, an array used as an argument isn't protected against change.

Array Arguments (Revisited)

- For example, the following function modifies an array by storing zero into each of its elements:

```
void store_zeros(int a[], int n)
{
    int i;

    for (i = 0; i < n; i++)
        a[i] = 0;
}
```

Array Arguments (Revisited)

- To indicate that an array parameter won't be changed, we can include the word `const` in its declaration:

```
int find_largest(const int a[], int n)
{
    ...
}
```

- If `const` is present, the compiler will check that no assignment to an element of `a` appears in the body of `find_largest`.

Array Arguments (Revisited)

- *Consequence 2:* The time required to pass an array to a function doesn't depend on the size of the array.
- There's no penalty for passing a large array, since no copy of the array is made.

Array Arguments (Revisited)

- *Consequence 3:* An array parameter can be declared as a pointer if desired.
- `find_largest` could be defined as follows:

```
int find_largest(int *a, int n)
{
    ...
}
```
- Declaring `a` to be a pointer is equivalent to declaring it to be an array; the compiler treats the declarations as though they were identical.

Array Arguments (Revisited)

- Although declaring a *parameter* to be an array is the same as declaring it to be a pointer, the same isn't true for a *variable*.
- The following declaration causes the compiler to set aside space for 10 integers:

```
int a[10];
```

- The following declaration causes the compiler to allocate space for a pointer variable:

```
int *a;
```

Array Arguments (Revisited)

- In the latter case, `a` is not an array; attempting to use it as an array can have disastrous results.
- For example, the assignment

```
*a = 0;    /** WRONG **/
```

will store 0 where `a` is pointing.
- Since we don't know where `a` is pointing, the effect on the program is undefined.

Array Arguments (Revisited)

- *Consequence 4:* A function with an array parameter can be passed an array “slice”—a sequence of consecutive elements.
- An example that applies `find_largest` to elements 5 through 14 of an array `b`:

```
largest = find_largest(&b[5], 10);
```

Using a Pointer as an Array Name

- C allows us to subscript a pointer as though it were an array name:

```
#define N 10
```

```
...
```

```
int a[N], i, sum = 0, *p = a;
```

```
...
```

```
for (i = 0; i < N; i++)  
    sum += p[i];
```


The compiler treats `p[i]` as `*(p+i)`.

Pointers and Multidimensional Arrays

- Just as pointers can point to elements of one-dimensional arrays, they can also point to elements of multidimensional arrays.
- This section explores common techniques for using pointers to process the elements of multidimensional arrays.

Processing the Elements of a Multidimensional Array

- Chapter 8 showed that C stores two-dimensional arrays in row-major order.
- Layout of an array with r rows:



- If p initially points to the element in row 0, column 0, we can visit every element in the array by incrementing p repeatedly.

Processing the Elements of a Multidimensional Array

- Consider the problem of initializing all elements of the following array to zero:

```
int a[NUM_ROWS][NUM_COLS];
```

- The obvious technique would be to use nested `for` loops:

```
int row, col;
```

```
...
```

```
for (row = 0; row < NUM_ROWS; row++)  
    for (col = 0; col < NUM_COLS; col++)  
        a[row][col] = 0;
```

- If we view `a` as a one-dimensional array of integers, a single loop is sufficient:

```
int *p;
```

```
...
```

```
for (p = &a[0][0]; p <= &a[NUM_ROWS-1][NUM_COLS-1]; p++)  
    *p = 0;
```

Processing the Elements of a Multidimensional Array

- Although treating a two-dimensional array as one-dimensional may seem like cheating, it works with most C compilers.
- Techniques like this one definitely hurt program readability, but—at least with some older compilers—produce a compensating increase in efficiency.
- With many modern compilers, though, there's often little or no speed advantage.

Processing the Rows of a Multidimensional Array

- A pointer variable `p` can also be used for processing the elements in just one *row* of a two-dimensional array.
- To visit the elements of row `i`, we'd initialize `p` to point to element 0 in row `i` in the array `a`:

```
p = &a[i][0];
```

or we could simply write

```
p = a[i];
```

Processing the Rows of a Multidimensional Array

- For any two-dimensional array a , the expression $a[i]$ is a pointer to the first element in row i .
- To see why this works, recall that $a[i]$ is equivalent to $*(a + i)$.
- Thus, $\&a[i][0]$ is the same as $\&(* (a[i] + 0))$, which is equivalent to $\&*a[i]$.
- This is the same as $a[i]$, since the $\&$ and $*$ operators cancel.

Processing the Rows of a Multidimensional Array

- A loop that clears row `i` of the array `a`:

```
int a[NUM_ROWS][NUM_COLS], *p, i;  
...  
for (p = a[i]; p < a[i] + NUM_COLS; p++)  
    *p = 0;
```

- Since `a[i]` is a pointer to row `i` of the array `a`, we can pass `a[i]` to a function that's expecting a one-dimensional array as its argument.
- In other words, a function that's designed to work with one-dimensional arrays will also work with a row belonging to a two-dimensional array.

Processing the Rows of a Multidimensional Array

- Consider `find_largest`, which was originally designed to find the largest element of a one-dimensional array.
- We can just as easily use `find_largest` to determine the largest element in row `i` of the two-dimensional array `a`:

```
largest = find_largest(a[i], NUM_COLS);
```

Processing the Columns of a Multidimensional Array

- Processing the elements in a *column* of a two-dimensional array isn't as easy, because arrays are stored by row, not by column.
- A loop that clears column *i* of the array *a*:

```
int a[NUM_ROWS][NUM_COLS], (*p)[NUM_COLS], i;  
...  
for (p = &a[0]; p < &a[NUM_ROWS]; p++)  
    (*p)[i] = 0;
```

Using the Name of a Multidimensional Array as a Pointer

- The name of *any* array can be used as a pointer, regardless of how many dimensions it has, but some care is required.
- Example:

```
int a[NUM_ROWS][NUM_COLS];
```

a is *not* a pointer to `a[0][0]`; instead, it's a pointer to `a[0]`.
- C regards *a* as a one-dimensional array whose elements are one-dimensional arrays.
- When used as a pointer, *a* has type `int (*)[NUM_COLS]` (pointer to an integer array of length `NUM_COLS`).

Using the Name of a Multidimensional Array as a Pointer

- Knowing that `a` points to `a[0]` is useful for simplifying loops that process the elements of a two-dimensional array.

- Instead of writing

```
for (p = &a[0]; p < &a[NUM_ROWS]; p++)  
    (*p)[i] = 0;
```

to clear column `i` of the array `a`, we can write

```
for (p = a; p < a + NUM_ROWS; p++)  
    (*p)[i] = 0;
```

Using the Name of a Multidimensional Array as a Pointer

- We can “trick” a function into thinking that a multidimensional array is really one-dimensional.
- A first attempt at using using `find_largest` to find the largest element in `a`:

```
largest = find_largest(a, NUM_ROWS * NUM_COLS);  
/* WRONG */
```

This an error, because the type of `a` is `int (*) [NUM_COLS]` but `find_largest` is expecting an argument of type `int *`.

- The correct call:

```
largest = find_largest(a[0], NUM_ROWS * NUM_COLS);
```

`a[0]` points to element 0 in row 0, and it has type `int *` (after conversion by the compiler).

Pointers and Variable-Length Arrays (C99)

- Pointers are allowed to point to elements of variable-length arrays (VLAs).
- An ordinary pointer variable would be used to point to an element of a one-dimensional VLA:

```
void f(int n)
{
    int a[n], *p;
    p = a;
    ...
}
```

Pointers and Variable-Length Arrays (C99)

- When the VLA has more than one dimension, the type of the pointer depends on the length of each dimension except for the first.
- A two-dimensional example:

```
void f(int m, int n)
{
    int a[m][n], (*p)[n];
    p = a;
    ...
}
```

Since the type of `p` depends on `n`, which isn't constant, `p` is said to have a ***variably modified type***.

Pointers and Variable-Length Arrays (C99)

- The validity of an assignment such as `p = a` can't always be determined by the compiler.
- The following code will compile but is correct only if `m` and `n` are equal:

```
int a[m][n], (*p)[m];  
p = a;
```
- If `m` is not equal to `n`, any subsequent use of `p` will cause undefined behavior.

Pointers and Variable-Length Arrays (C99)

- Variably modified types are subject to certain restrictions.
- The most important restriction: the declaration of a variably modified type must be inside the body of a function or in a function prototype.

Pointers and Variable-Length Arrays (C99)

- Pointer arithmetic works with VLAs.

- A two-dimensional VLA:

```
int a[m][n];
```

- A pointer capable of pointing to a row of a:

```
int (*p)[n];
```

- A loop that clears column *i* of a:

```
for (p = a; p < a + m; p++)  
    (*p)[i] = 0;
```