



Western
UNIVERSITY • CANADA

Chapter 3C – Operations on Processes

Spring 2023

Operations on Processes

- Process Creation
- Process Termination
- Examples

Process Creation


- Parent process create children processes, which, in turn create other processes, forming a tree of processes
- Generally, process identified and managed via a process identifier (pid)

Process Creation

- Resource sharing options
 - Parent and children share *mem, data, etc ----* all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate

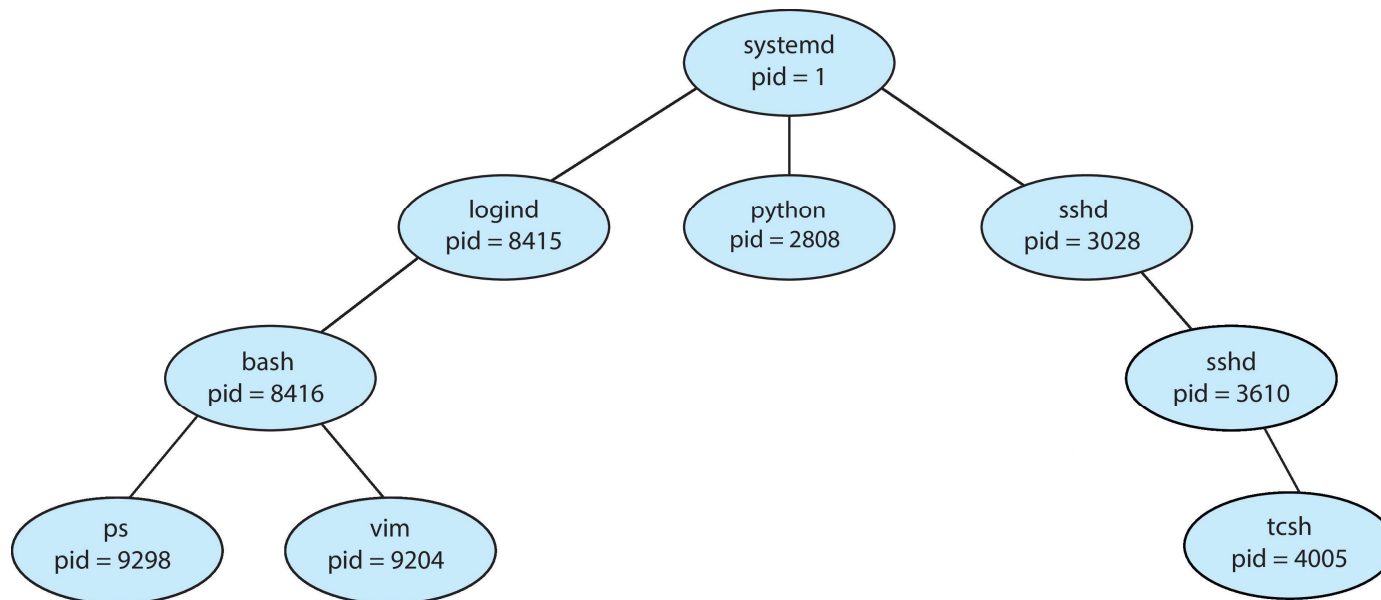
Process Creation

- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - `fork()` system call creates new process
 - `exec()` system call used after a `fork()` to replace the process' memory space with a new program
 - Parent process calls `wait()` waiting for the child to terminate



Process Creation

- A tree of processes in Linux



Process Termination

- Process executes last statement and then asks the operating system to delete it using the `exit()` system call.
- Returns status data from child to parent (via `wait()`)
- Process' resources are deallocated by operating system

Process Termination

- Parent may terminate the execution of children processes using the abort() system call. Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting, and the operating system does not allow a child to continue if its parent terminates

Process Termination

- Some operating systems do not allow child to exist if its parent has terminated. If a process terminates, then all its children must also be terminated.
- **Cascading termination:** All children, grandchildren, etc., are terminated.
- The termination is initiated by the operating system.
- The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the pid of the terminated process
- `pid = wait(&status);`
- If no parent waiting (did not invoke `wait()`) process is a **zombie**
- If parent terminated without invoking `wait()`, process is an **orphan**

Process Termination

- `pid = wait(&status);`
is shorthand for:
`pid = waitpid(-1, &wstatus, 0);`
- `-1` – Wait for any child process
- `wstatus` – Load the exit status into this integer
 - This can be `NULL` if we don't need the exit status
- `0` – wait but depending on how the child returned. We will always just use 0
one spec child.
↓
- `waitpid(pid, NULL, 0);` – wait until child `pid` returns
- `while (wait(NULL) > 0);` – wait for all child processes

Examples

- For all system and library calls, consult the man page (you may need to specify section 2 or 3)
- ```
$ man man
...
1 Executable programs or shell commands
2 System calls (functions provided by the kernel)
3 Library calls (functions within program libraries)
...
```
- `$ man wait` ← The wait command built into bash
- `$ man 2 wait` ← The wait system call

# Examples

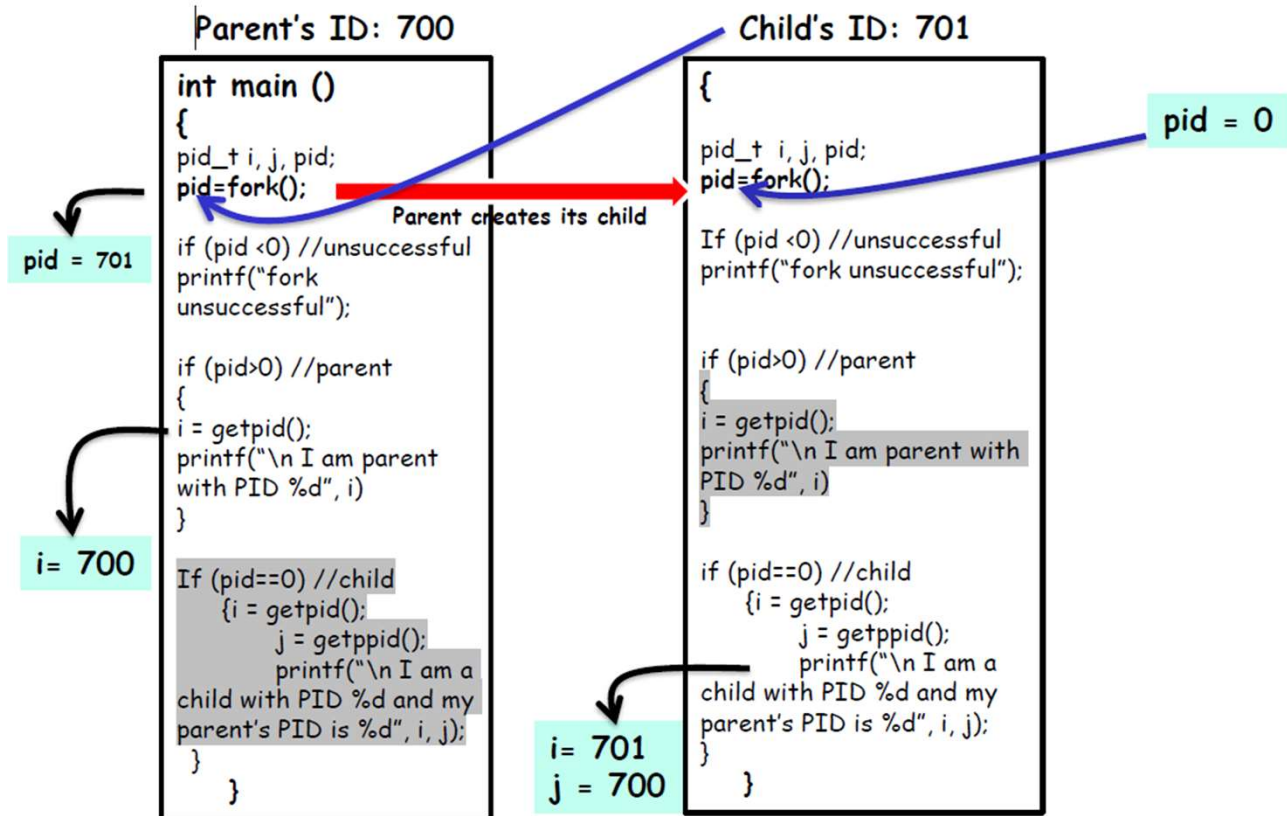
- The Unix system call for process creation is `fork()`
- The fork system call creates a child process that is a duplicate of the parent.
  - Child inherits state from parent process
    - Same program instructions, variables have the same values, same position in the code.
  - Parent and child have separate copies of that state
    - They are stored in separate locations in memory
    - Important: updating the value of a variable inside the child will NOT update that variable in the parent, and vice-versa.

# Examples

- If `fork()` succeeds it returns the child PID to the parent and returns 0 to the child
- If `fork()` fails, it returns -1 to the parent (no child is created)
- `pid_t` data type represents process identifiers
- Other calls:
  - `pid_t getpid()` – returns the PID of calling process. Call is always successful.
  - `pid_t getppid()` – returns the PID of parent process. Call is always successful.

# Examples

## fork() example



# Examples

- ```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main(){
    pid_t pid;
    int i;
    pid=fork();
    if( pid> 0 ) { /* parent */
        for( i=0; i< 10; i++ )
            printf("\t\t\tPARENT%d\n", i);
    }
    else { /* child */
        for( i=0; i< 10; i++ )
            printf( "CHILD %d\n", i);
    }
    return 0;
}
```

- What is the possible output?

- ```
PARENT 0
PARENT 1
PARENT 2
PARENT 3
PARENT 4
PARENT 5
PARENT 6
PARENT 7
PARENT 8
PARENT 9
CHILD 0
CHILD 1
CHILD 2
CHILD 3
CHILD 4
CHILD 5
CHILD 6
CHILD 7
CHILD 8
CHILD 9
```



# Examples

- ```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main(){
    pid_t pid;
    int i;
    pid=fork();
    if( pid> 0 ) { /* parent */
        for( i=0; i< 10; i++ )
            printf("\t\t\tPARENT%d\n", i);
    }
    else { /* child */
        for( i=0; i< 10; i++ )
            printf( "CHILD %d\n", i);
    }
    return 0;
}
```

- What is the possible output?

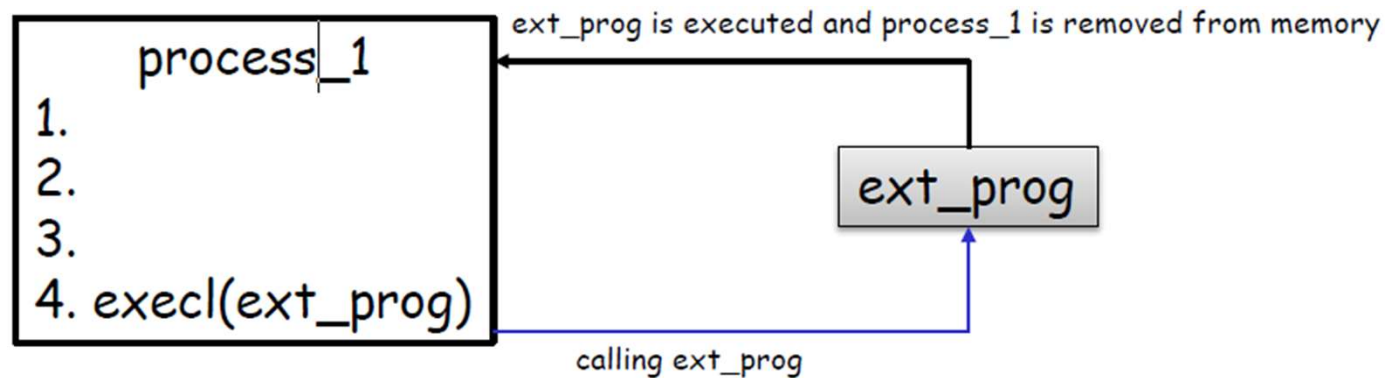
- ```
PARENT 0
PARENT 1
PARENT 2
PARENT 3
PARENT 4
PARENT 5
PARENT 6
CHILD 0
CHILD 1
CHILD 2
PARENT 7
PARENT 8
PARENT 9
CHILD 3
CHILD 4
CHILD 5
CHILD 6
CHILD 7
CHILD 8
CHILD 9
```

# Examples

- Output is **nondeterministic** - Cannot determine output by looking at code
- Processes get a share of the CPU to give another process a turn
  - The switching between the parent and child depends on many factors: machine load, process scheduling

# Examples

- The system call `exec1()` replace a process (the caller process) with a new loaded program
- `exec1()` loads a binary file into memory (destroying the memory image of the program calling it)
- On success, `exec1()` **never** returns; on failure, `exec1()` returns -1



# Examples

- Program A

- ```
int i= 5;
printf("%d\n",i);
execl("B", "", NULL);
printf("%d\n",i);
```

- Program B

- ```
main() {
 printf("hello\n");
}
```

- What is the possible output?

- ```
5
hello
```

- Why not?

- ```
5
hello
5
```

# Examples

- `fork()` and `execl()`

- ```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main(){
    pid_t pid;
    pid = fork();
    if (pid > 0){
        wait(NULL);
        printf("Child Complete");
    }
    else{
        if (pid == 0) {
            execl("B", "", NULL);
            printf("\n You'll never see this line..");
        }
    }
}
```

this is a child process.

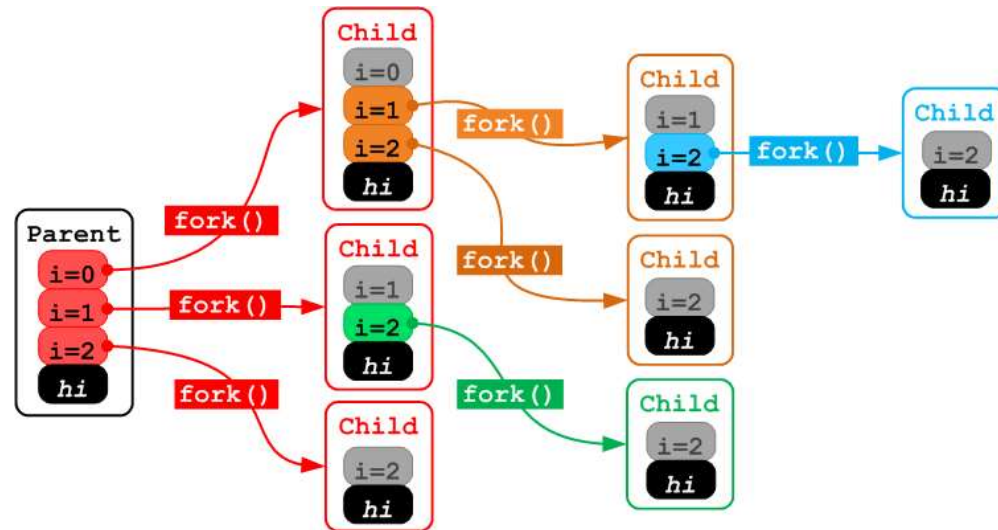
Examples

- How many processes are created by this program?

- ```
#include <stdio.h>
#include <unistd.h>
void main() {
 int i;
 for (i=0;i<3;i++){
 fork();
 }
 printf("hi\n");
}
```


# Examples

- <https://stackoverflow.com/questions/26793402/visually-what-happens-to-fork-in-a-for-loop>
- $2^n$  processes
- $2^n - 1$  children



# Examples

- Forks and files
  - Every process has a **process file descriptor table**
  - Each entry in process file descriptor table represents **stdin**, **stdout**, **stderr**, and **file pointer**.
  - Assume that there was something like this in a program
    - ```
int in_file;  
in_file= open("list.txt", O_RDONLY);
```



Examples

- Open a file before a fork
 - The child process gets a copy of its parent's process file descriptor table.
 - The child and parent share a file pointer because the open came before the fork.
- Open a file after a fork
 - Assume that parent and child each open a file after the fork
 - They get their own entries in the file descriptor table
 - This implies that the file position information is different
- Suppose that hello.txt consists of Hello, world!. What is the output of...
 - `open()` before `fork()`
 - `open()` after `fork()`

Examples

- ```
#include <stdio.h>
#include <unistd.h> //fork()
#include <fcntl.h> //open() (unbuffered open)
#include <sys/wait.h> //wait()
int main(){
 int fd; char c; pid_t pid;
 fd=open("hello.txt", O_RDONLY);
 pid=fork(); //Open a file before a fork
 if (pid> 0){
 read(fd, &c, 1);
 printf("fd: %d, parent: c = %c\n", fd, c);
 wait(NULL);
 }
 else if (pid== 0){
 read(fd, &c, 1);
 printf("fd: %d, child: c = %c\n", fd, c);
 return 0;
 }
 fclose(fd);
}
```

- Output

- fd: 3, parent: c = H  
fd: 3, child: c = e  
or  
fd: 3, child: c = H  
fd: 3, parent: c = e

# Examples

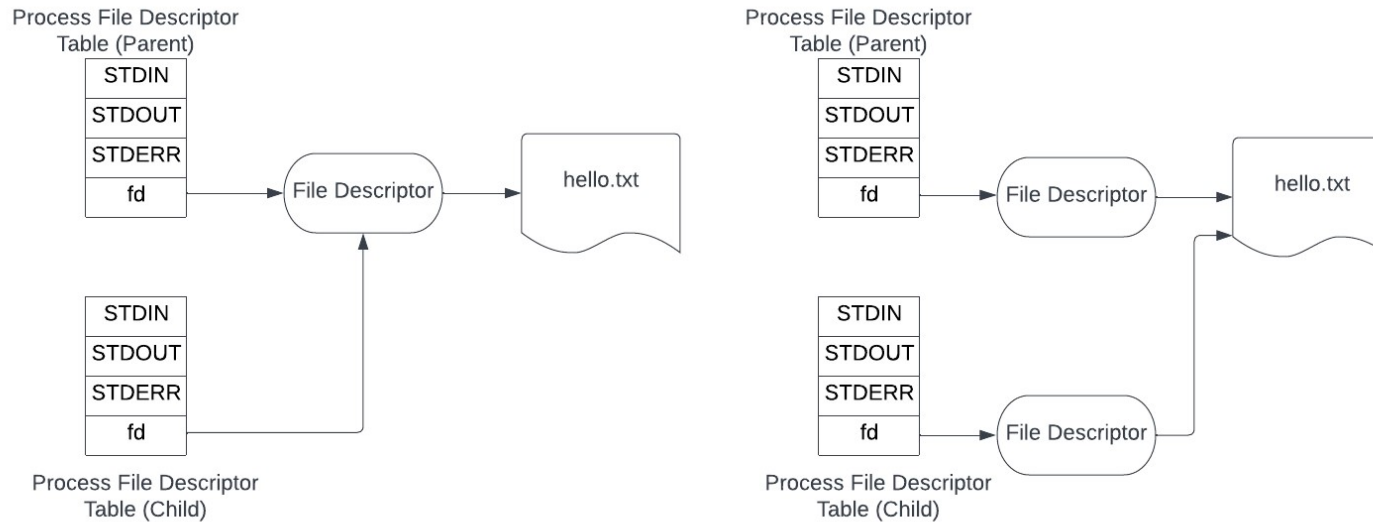
- ```
#include <stdio.h>
#include <unistd.h> //fork()
#include <fcntl.h>  //open() (unbuffered open)
#include <sys/wait.h> //wait()
int main(){
    int fd; char c; pid_t pid;
    pid=fork(); //Open a file after a fork
    fd=open("hello.txt", O_RDONLY);
    if (pid> 0){
        read(fd, &c, 1);
        printf("fd: %d, parent: c = %c\n", fd, c);
        wait(NULL);
    }
    else if (pid== 0){
        read(fd, &c, 1);
        printf("fd: %d, child: c = %c\n", fd, c);
        return 0;
    }
    fclose(fd);
}
```

- Output

- fd: 3, parent: c = H
fd: 3, child: c = H
or
fd: 3, child: c = H
fd: 3, parent: c = H

Examples

- It is probably better to open files **after** a fork so the parent and child do not confuse each other





Western
UNIVERSITY • CANADA