



Western
UNIVERSITY • CANADA

Chapter 8 - Deadlocks

Spring 2023

Overview


- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock

System Model


- Any system has a set of resources denoted $R_1, R_2, \dots R_n$
 - These could be CPU cycles, memory space, I/O devices, mutex locks, semaphores
- Each resource R_i has 1 or more instances
 - E.g. 4 CPUs, 2 Network interfaces, Semaphore value set to 5, etc.
- Each process or thread that requests a resources follows these steps:
 - 1) Request
 - 2) Use
 - 3) Release

System Model

- For example, consider a traffic jam



(a) Deadlock possible



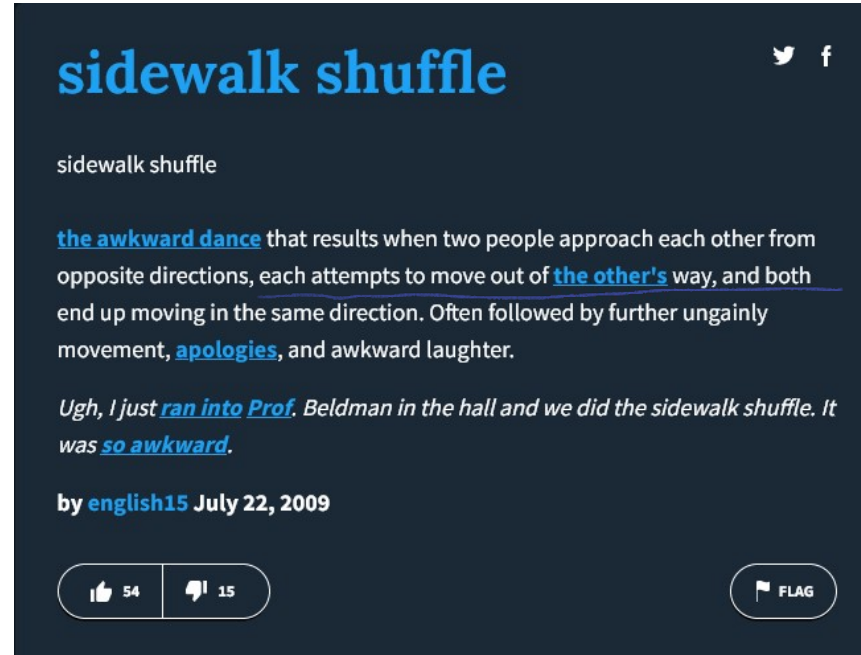
(b) Deadlock

System Model

- For example, return to the Dining Philosopher's problem
 - Suppose all five philosophers pick up their left chopstick. None of the philosophers can pick up their right chopstick. Some potential solutions:
 - Allow only 4 philosophers or less at the table
 - Allow a philosopher to pick up both chopsticks if and only if both are available
 - Odd numbered philosophers pick up their left chopstick, then the right. Even numbered philosophers pick up the right chopstick, then the left.

System Model

- For example, consider the "Sidewalk Shuffle". This is a special case known as livelock. There is no blocking but there is no progress either.



System Model

- For example, consider two threads with two semaphores

- T1:

- wait(s1)
 - wait(s2)

- T2:

- wait(s2)
 - wait(s1)

when T1 wait for s1, T2 would wait for s2,
so it is a deadlock.

T1: wait(s1)
=> wait for s1 to finish
=> turn to T2
T2: wait(s2)
=>

Deadlock Characterization


- A deadlock occurs under all the following four conditions ^{12/28} simultaneously
- **Mutual Exclusion** – only one thread at a time can use a resource
- **Hold and Wait** – a thread holding at least one resource is waiting to acquire additional resources held by other threads
- **No preemption** – a resource can be released only voluntarily by the thread holding it, after that thread has completed its task
- **Circular wait** – Thread T_0 waits on a resource held by T_1 , T_1 waits on a resource held by T_2 , ..., T_n waits on a resource held by T_0

Deadlock Characterization

- Identifying circular wait with resource allocation graphs
 - Consider two types of vertices
 - Threads – denoted $T = \{T_1, T_2, \dots, T_n\}$
 - Resources – denoted $R = \{R_1, R_2, \dots, R_m\}$
 - Two types of edges
 - Request edge – A directed edge $T_i \rightarrow R_j$ ("wants to hold")
 - Assignment edge – A directed edge $R_j \rightarrow T_i$ ("is holding")


Deadlock Characterization

- Identifying circular wait with resource allocation graphs
 - One instance of R_1
 - Two instances of R_2
 - One instance of R_3
 - Three instance of R_4
 - T_1 holds one instance of R_2 and is waiting for an instance of R_1
 - T_2 holds one instance of R_1 , one instance of R_2 , and is waiting for an instance of R_3
 - T_3 holds one instance of R_3




Deadlock Characterization

- Identifying circular wait with resource allocation graphs



Deadlock Characterization

- Identifying circular wait with resource allocation graphs
- If a graph does not contain cycles, then we do not have deadlock
- If a graph does contain cycles, we may have deadlock



not 100% having deadlock.
we have to check the case.

Handling Deadlocks



- Four options
 - **Ignore the deadlock** – Force application developers to ensure deadlock does not occur. For reasons of efficiency, this is the most common approach by most operating systems.
 - **Deadlock Prevention** – Ensure one of the four characteristics of deadlock does not occur. This typically underutilizes system resources.
 - **Deadlock Avoidance** – Declaring all resources that may be required before starting. Force threads to wait if there is a potential for overlap. Very inefficient.
 - **Deadlock Detection** – Periodically check to see if there is a deadlock and recover appropriately. Common in relational database management systems.

Deadlock Prevention

- Ensure one of the four characteristics of deadlock does not occur
 - Mutual Exclusion – Some resources are just ^{本質上} intrinsically unshareable. Preventing mutual exclusion is just not feasible
 - Hold and Wait – Force a thread to ^{可能} hold all its resources before it can start. Most resources will be held for no reason. Popular resources will create starvation. *prevent the circular wait.*
 - No preemption – If a resource cannot be secured, voluntarily release all existing resources and restart. Only feasible if state can be saved such as CPU registers and database transactions
 - Circular wait – Enforce an increasing order on all resources for all threads. It's up to the application developer to obey the order. E.g. R_1 then R_2 . Never R_2 then R_1


Deadlock Avoidance

- Resource Allocation Graphs – Do not allow any thread to enter the graph if it will create a cycle
- E.g. T_2 must wait until T_1 finishes because introducing it could create a cycle




Deadlock Avoidance

- Deadlock Trajectory Diagram – Consider two Processes that require access to two resources





Deadlock Avoidance

- Deadlock Trajectory Diagram – Consider two Processes that require access to two resources



Deadlock Detection

- Collapse a resource allocation graph into a "wait-for" graph. Periodically invoke an algorithm to search for a cycle in the graph

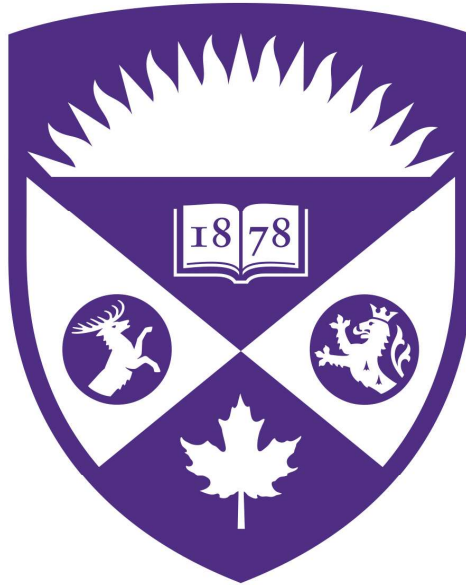


Deadlock Detection

- How often to run the algorithm depends on a few factors
 - How often is a deadlock likely to occur?
 - When it does occur, how many threads will need to be aborted or rolled back?
- Running the algorithm uses precious CPU cycles, so we don't want to run it too often
- However, if the algorithm is not invoked often enough, deadlock will occur too frequently. The process to abort or roll back threads may take some time

Recovery from Deadlock

- Abort or rollback all deadlocked threads?
- Abort or rollback one thread at a time until deadlock stops?
- Should aim to choose the minimum cost
 - Priority of the thread(s)?
 - Computation time completed? Computation time remaining?
 - Number of resources held? Number of remaining resources needed?
 - Fewest number of thread(s) to clear the deadlock
- Starvation could occur if the same thread is always the victim



Western
UNIVERSITY • CANADA