



Western
UNIVERSITY • CANADA

Chapter 9B – Memory Management (Paging)

Spring 2023

Overview

- Basic Method
- Hardware Support
- Protection
- Shared Pages
- Structure
- Swapping

Basic Method

- With paging, process space is no longer contiguous
- Used by most modern operating systems
- Partition the physical memory into fixed-size blocks called **frames**
- Partition the logical memory into same-sized fixed-size blocks called **pages**
 - There can be more pages than there are frames. This is virtual memory and will be covered in the next chapter.
- Use a **Page Table** for each process to translate logical addresses to physical addresses

Basic Method

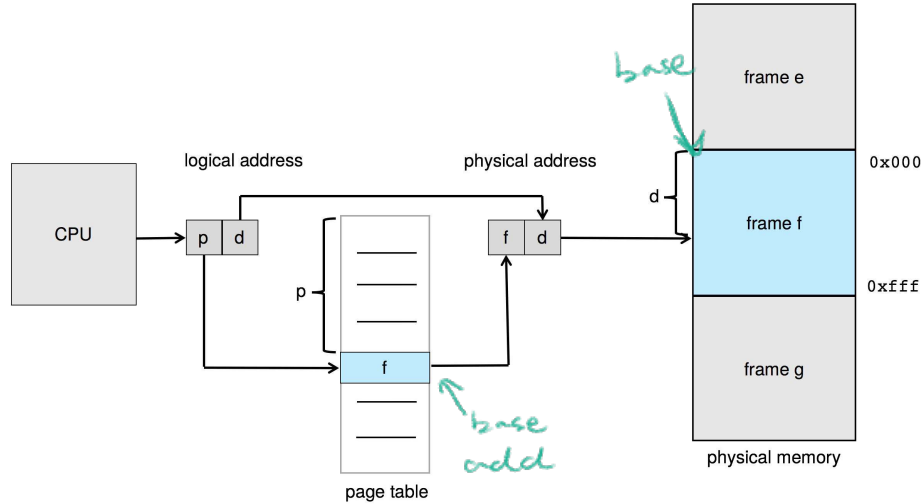
- Processes use pages 0 to N so it thinks the memory is contiguous
- However, the physical location of each frame in memory is non-contiguous
- Internal fragmentation is still possible *← not all frames are used.
but at least all processes could*
- External fragmentation is avoided *← find a place to go.*

Basic Method

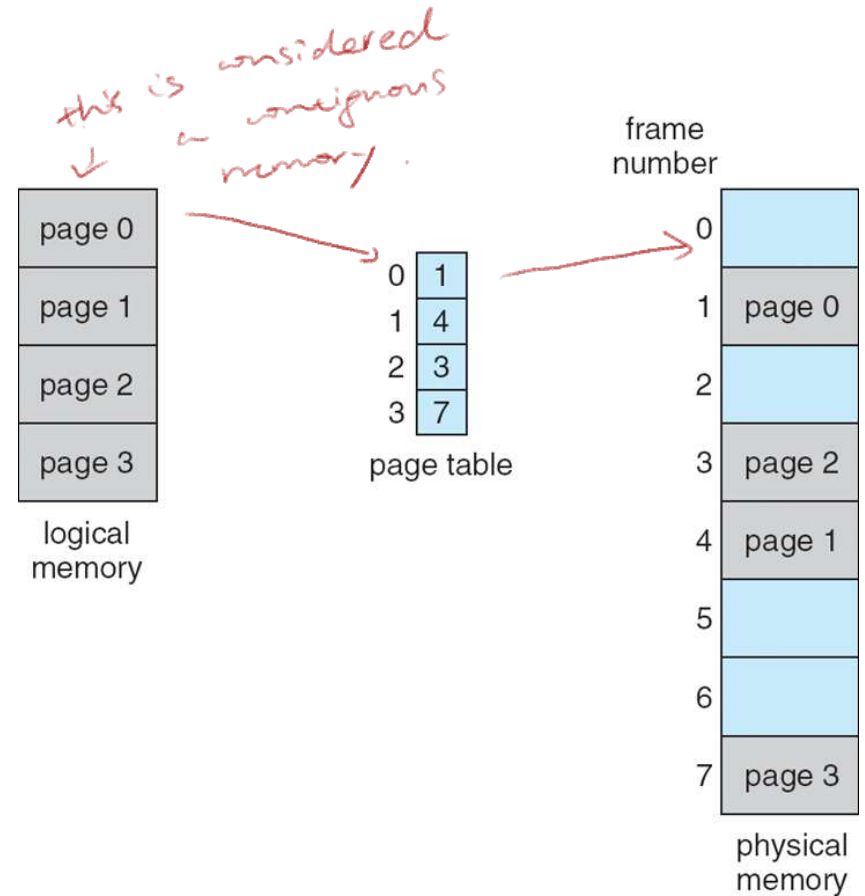
- The logical address generated by the CPU is divided into a bitmap of length m . That's 2^m logical addresses with:
 - **Page number (p)** - used as an index into a page table which contains base address of each page in physical memory. It is of length $m-n$ or 2^{m-n} bits.
 - **Page offset (d)** – combined with base address to define the physical memory address that is sent to the memory unit. It is length n or 2^n bits. This is the size of each page.

page number	page offset
p	d
$m - n$	n

Basic Method



$f + d = \text{address}$



Basic Method

- Example page numbers and page offsets
 - 3 bits – 2^1 | 2^2 (2 **pages**, 4 **page offsets** each)

000	Page 0	Location 0
001	Page 0	Location 1
010	Page 0	Location 2
011	Page 0	Location 3
100	Page 1	Location 0
101	Page 1	Location 1
110	Page 1	Location 2
111	Page 1	Location 3

Basic Method

- Example page numbers and page offsets
 - 4 bits – $2^3 \mid 2^1$ (8 **pages**, 2 **page offsets** each)

0000	Page 0	Location 0
0001	Page 0	Location 1
0010	Page 1	Location 0
0011	Page 1	Location 1
0100	Page 2	Location 0
0101	Page 2	Location 1
0110	Page 3	Location 0
0111	Page 3	Location 1

1000	Page 4	Location 0
1001	Page 4	Location 1
1010	Page 5	Location 0
1011	Page 5	Location 1
1100	Page 6	Location 0
1101	Page 6	Location 1
1110	Page 7	Location 0
1111	Page 7	Location 1

Basic Method

- Example page numbers and page offsets
- 4 bits – 2^2 | 2^2 (4 **pages**, 4 **page offsets** each)

0000	Page 0	Location 0
0001	Page 0	Location 1
0010	Page 0	Location 2
0011	Page 0	Location 3
0100	Page 1	Location 0
0101	Page 1	Location 1
0110	Page 1	Location 2
0111	Page 1	Location 3

1000	Page 2	Location 0
1001	Page 2	Location 1
1010	Page 2	Location 2
1011	Page 2	Location 3
1100	Page 3	Location 0
1101	Page 3	Location 1
1110	Page 3	Location 2
1111	Page 3	Location 3


Basic Method

- Mapping pages to frames

- 5 bits – $2^3 \mid 2^2$ (8 **pages**, 4 **page offsets** each) *length of each page.*

- Page 0, offset 0 (**a**) is at physical address 20
 - Page 0, offset 3 (**d**) is at physical address 23
 - Page 1, offset 0 (**e**) is at physical address 24
 - Page 3, offset 2 (**n**) is at physical address 9

- (Note: In this diagram, if a physical address is handled in 1-byte increments, there are 32 bytes of memory, a page is 4 bytes, and there are 8 pages/frames)



Basic Method

- Internal fragmentation
 - Consider a page size of 2048 bytes and process size is 72,766 bytes
 - 35 pages + 1,086 bytes
 - The last page is using 1,086 bytes with 962 bytes left over
 - A process could use n pages + 1 byte so n+1 frames in the worst case. Almost the entire frame is useless. On average, $\frac{1}{2}$ of the last frame is unused
 - Smaller page sizes minimizes the waste but adds overhead to the operating system
 - Larger page sizes make disk I/O to memory more efficient but adds more internal fragmentation

Basic Method

- Internal fragmentation
 - Most operating systems support two page sizes
 - 4KB or 8KB for regular processes
 - 2MB or 4MB or higher for processes that use a lot of memory

Basic Method


- External fragmentation
 - We have none. Any free frame can be allocated to a process as needed.

Basic Method

- Free frames
 - The operating system maintains a system-wide frame table
 - One entry per frame
 - Each entry has a flag indicating whether the frame is free or not
 - If it is not free, which process is using it

Basic Method

- Free frames




Hardware Support

- Due to the size of a process' page table, the table is kept in memory
 - The address of the page table is stored in the **Page-table base register (PTBR)**
 - The size of the page table is stored in the **Page-table length register (PTLR)**
 - Every instruction that accesses memory will trigger two memory accesses:
 - One to reference the page table itself and find the physical address
 - One to access the found address in memory
 - Solution: Cache the page table (Most systems have multiple levels)
 - This is called the **Translation Look-aside Buffer (TLB)**

Hardware Support

- Some TLBs hold address-space identifiers (ASIDs) so processes can only use their own entries
 - Otherwise, the TLB will need to be reloaded at each context switch
- Typically small (e.g. 64 – 1024 entries). Some entries can be made permanent.
- If the page is found in the TLB (a **hit**), use the frame
- If the page is not found in the TLB (a **miss**)
 - Consult the page table in memory instead
 - Load it into the TLB so we can avoid a miss in the future (Use an algorithm like Least-recently used, round-robin, or random to replace existing entries)

Hardware Support



Hardware Support

- Effective memory-access time
 - The percentage of times the page number of interest is found in the TLB is the **hit ratio**
 - If it is not found in the TLB, we need to consult the page table (1 access)
 - Then we can use the address (1 access)

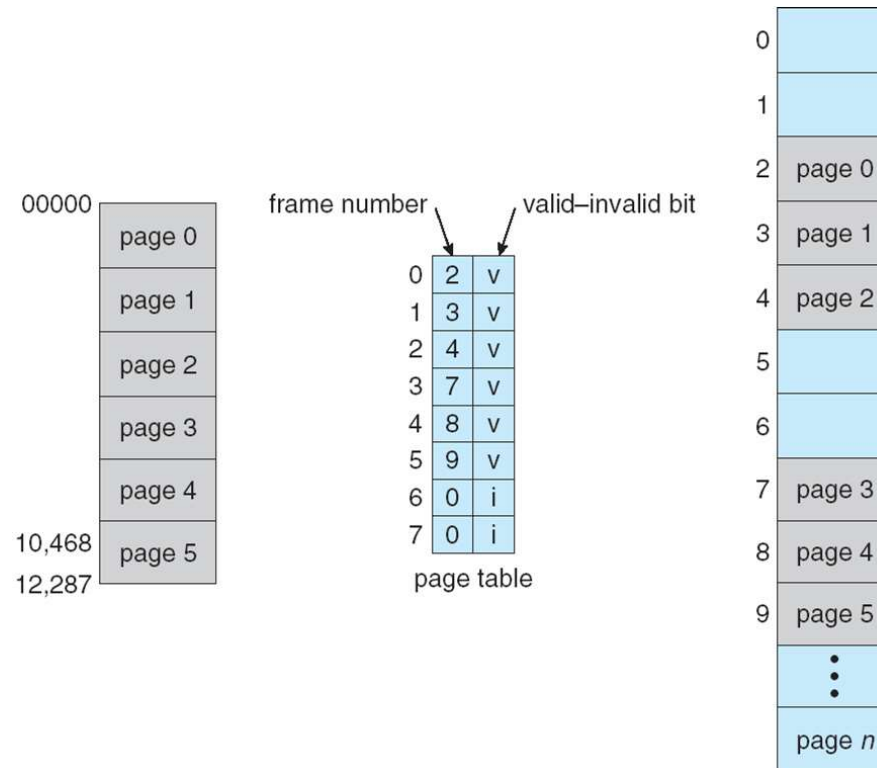
$$\text{Effective Access Time} = (\text{Hit}) \bullet \text{AccessTime} + (1 - \text{Hit}) \bullet 2(\text{AccessTime})$$

- Suppose a hit ratio of 80% and a memory access time of 10 nanoseconds, then effective access time is $0.80 \times 10 + 0.2 \times 20 = 12$ nanoseconds
- A hit ratio of 99% is 10.1 nanoseconds

Protection


- Each entry in the page table can hold a protection bit to indicate whether the frame is read-only or read-write.
- To determine whether the frame is in the process' logical address space
 - Use a page-table length register
 - Each entry in the page table can hold a valid-invalid bit for all frames in the address space

Protection



Shared Pages


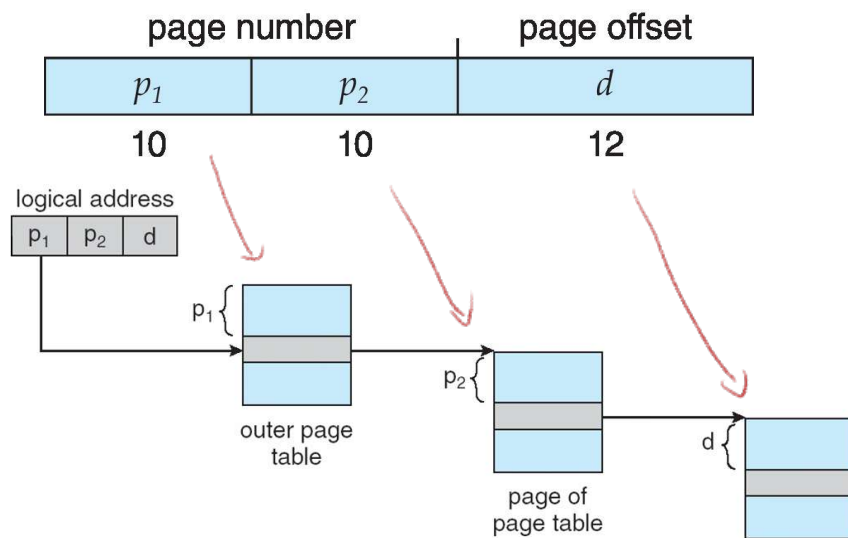
- Paging makes sharing address space easy
- These pages should be read-only
- Commonly used for DLLs



Structure

- Page tables can be very large. We don't want to put this in one contiguous block

- Solution: Hierarchical Page Tables (multilevel)



Structure

- Here are some examples using 64 bits
- More page tables means more memory accesses, so this strategy only goes so far

outer page	inner page	offset
p_1	p_2	d
42	10	12


2nd outer page	outer page	inner page	offset
p_1	p_2	p_3	d
32	10	10	12

Structure

- **Hashed Page Tables**
 - Common in address spaces larger than 32 bits
 - Each entry in the page table holds
 - A hash of the page number
 - The frame
 - A pointer to the next hash
 - A logical address is hashed and then the matching hash is searched in the list

Structure

- Hashed Page Tables

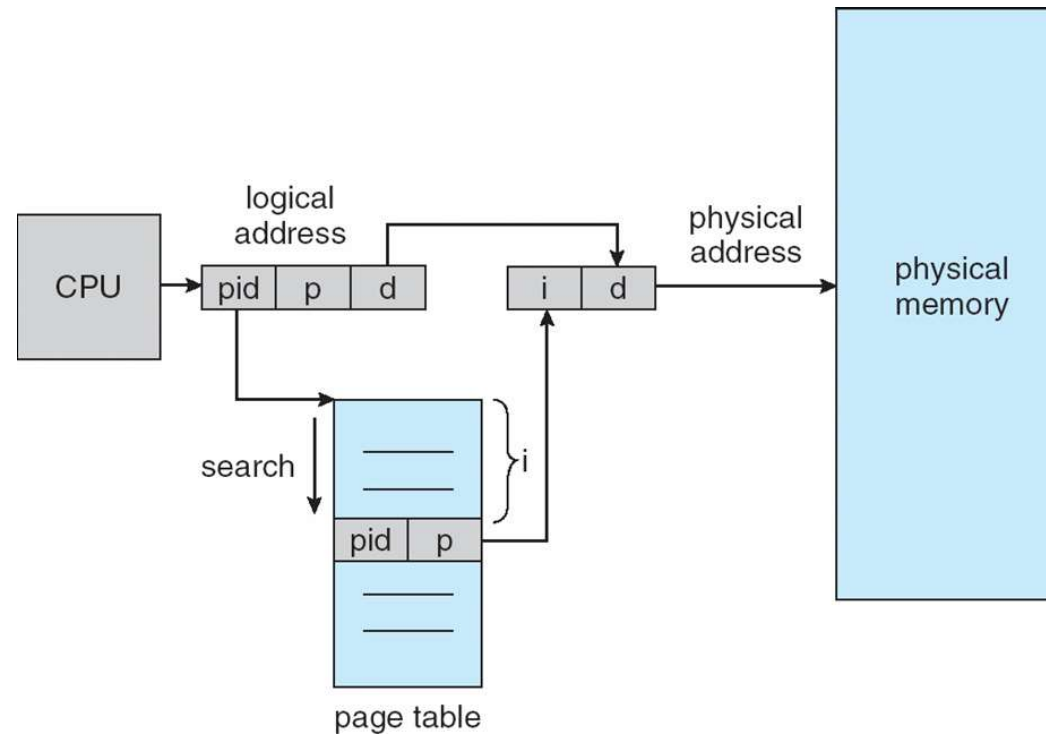


Structure

- **Inverted Page Tables**
 - Instead of one page table per process, one page table for all processes
 - One entry in the table for each frame in the system
 - This removes the need to maintain many page tables
 - This adds to the amount of time it takes to search the page table
 - Since pages to frames is one-to-one, shared memory will be regularly mapped and re-mapped

Structure

- Inverted Page Tables




Structure

- In reality, modern operating systems will employ a mixture of these structures


Swapping


- A process can be temporarily **swapped** to a backing store (usually disk) and then brought back in
- Total memory space for processes can exceed the physical memory (oversubscribed)



Swapping

- The backing store needs to be reasonably fast
- Swapping entire processes is generally no longer used
- Modern operating systems may still swap pages
- Swapping pages due to low memory is an "emergency" situation and should be avoided
 - Terminate unnecessary processes
 - Get more memory





Western
UNIVERSITY • CANADA