

A decorative graphic on the left side of the slide, consisting of a network of white lines and circles on a teal background, resembling a circuit board or a neural network.

WEEK 10

TRANSACTION – CONCURRENCY CONTROL

STUDENT OBJECTIVES

- Upon completion of this video, you should be able to:
 - Identify the 3 main problems with interleaving schedules
 - Distinguish between a serial and a non serial schedule
 - Determine if a schedule is serializable

CONCURRENCY CONTROL

- We want to allow multiple transactions to get at all the data and resources at the same time.
- Objective is to ensure serializability of transactions in multi-user database environment (interleaving of transactions where the resulting state is the same as one the states that would occur if the transactions were done serially in some order).
- **Schedule:** A schedule S of n transactions T_1, T_2, \dots, T_N is an ordering of the operations of the transactions. Each operation of an individual transaction T_i in S must appear in the same order that it appears in T_i . The scheduler interleaves the execution of database operations to ensure serializability

these states does not interleave.

SAMPLE SCHEDULES:

Schedule 1:

| Time | Transaction | Step | Stored Value |
|------|-------------|------------------|--------------|
| 1 | T1 | Read Bal | 35 |
| 2 | T1 | Bal = $35 + 100$ | |
| 3 | T1 | Write Bal | 135 |
| 4 | T1 | Commit | |
| 5 | T2 | Read Bal | 135 |
| 6 | T2 | Bal = $135 - 30$ | |
| 7 | T2 | Write Bal | 105 |
| 8 | T2 | Commit | |

Schedule 2:

| Time | Transaction | Step | Stored Value |
|------|-------------|-----------------|--------------|
| 1 | T2 | Read Bal | 35 |
| 2 | T2 | Bal = $35 - 30$ | |
| 3 | T2 | Write Bal | 5 |
| 4 | T2 | Commit | |
| 5 | T1 | Read Bal | 5 |
| 6 | T1 | Bal = $5 + 100$ | |
| 7 | T1 | Write Bal | 105 |
| 8 | T1 | Commit | |

Schedule 3:

| Time | Transaction | Step | Stored Value |
|------|-------------|------------------|--------------|
| 1 | T1 | Read Bal | 35 |
| 2 | T2 | Read Bal | 35 |
| 3 | T1 | Bal = $35 + 100$ | |
| 4 | T2 | Bal = $35 - 30$ | |
| 5 | T1 | Write Bal | 135 |
| 6 | T2 | Write Bal | 5 |
| 7 | T1 | Commit | |
| 8 | T2 | Commit | |

**NOTICE THE PROBLEM
IN SCHEDULE 3!**

5 ← this is incorrect.

↑
this is an invalid
state.

CONFLICT

- 2 operations in a schedule *conflict* if
 - 1) they belong to different transactions,
 - 2) they access the same data item X
 - 3) at least one of the transactions issues a WRITE(X).
- 3 Main Problems:
 - Lost Updates
 - Uncommitted Data
 - Inconsistent Retrievals

EXAMPLE: LOST UPDATE

- T1: Homer puts in 100 dollars to the savings account ($Bal = Bal + 100$)
- T2: Marge takes out 30 dollars from the savings account ($Bal = Bal - 30$)

Correct Schedule

| Time | Transaction | Step | Stored Value |
|------|-------------|------------------|--------------|
| 1 | T1 | Read Bal | 35 |
| 2 | T1 | $Bal = 35 + 100$ | |
| 3 | T1 | Write Bal | 135 |
| 4 | T1 | Commit | |
| 5 | T2 | Read Bal | 135 |
| 6 | T2 | $Bal = 135 - 30$ | |
| 7 | T2 | Write Bal | 105 |
| 8 | T2 | Commit | |

↑ correct

Incorrect Schedule:

| Time | Transaction | Step | Stored Value |
|------|-------------|------------------|-------------------|
| 1 | T1 | Read Bal | 35 |
| 2 | T2 | Read Bal | 35 |
| 3 | T1 | $Bal = 35 + 100$ | |
| 4 | T2 | $Bal = 35 - 30$ | |
| 5 | T1 | Write Bal | 135 ← Lost Update |
| 6 | T2 | Write Bal | 5 |
| 7 | T1 | Commit | |
| 8 | T2 | Commit | |

this one override 135.

UNCOMMITTED DATA

Something has happened here, maybe the ATM rollers broke before the money was sucked in

When 2 transactions T1 and T2 are executed, T1 is rolled back. T2 reads uncommitted data (violates the Isolation property)
Transaction starts reading before the rollback is completed

Correct Schedule

| Time | Transaction | Step | Stored Value |
|------|-------------|---------------------|--------------|
| 1 | T1 | Read Bal | 35 |
| 2 | T1 | Bal = 35+100 | |
| 3 | T1 | Write Bal | 135 |
| 4 | T1 | Rollback | 35 |
| 5 | T2 | Read Bal | 35 |
| 6 | T2 | Bal = 35-30 | |
| 7 | T2 | Write Bal | 5 |
| 8 | T2 | Commit | |

Incorrect Schedule:

| Time | Transaction | Step | Stored Value |
|------|-------------|--------------|------------------------|
| 1 | T1 | Read Bal | 35 |
| 2 | T1 | Bal = 35+100 | |
| 3 | T1 | Write Bal | 135 |
| 4 | T2 | Read Bal | 135 ← Uncommitted Data |
| 5 | T2 | Bal = 135-30 | too late. |
| 6 | T1 | Rollback | 35 |
| 7 | T2 | Write Bal | 105 |
| 8 | T2 | Commit | |

INCONSISTENT RETRIEVALS

- Occurs when a transaction reads some data before they are changed and other data after they are changed.

- For example if T1 calculates the sum of all inventory:

SELECT SUM(Quantity_on_hand) FROM inventory

at the same time as T2 is update inventory for some items, the total at the end will be wrong.

SCHEDULER

- Want to ensure that once a transaction T is committed, it should never be necessary to roll back T . (A recoverable schedule)
- Recoverable schedules may have cascading rollbacks that allows uncommitted transactions to be rolled back because it read an item from a transaction that aborted (try to avoid this because it is time consuming!)

- **Scheduler:** A scheduler sets the order that concurrent transactions are executed. It interweaves the operations to ensure serialization.
- A schedule can use a number of methods, we will look at 3: locking, time stamping and optimistic.
 - It MUST preserve the order of operations within the original transactions
 - It makes some system component schedules
 - It can't really make a planned schedule because it doesn't know what transactions the users are going to submit and in what order.

QUESTION: If the scheduler just had to ensure serialization, what would it always do to ensure this?

ANSWER: Never ever let them mix, always finish one transaction before doing the next transaction

keep + they not overlap.

SERIALIZABILITY

- Suppose 2 users (e.g. airline clerks) submit the 2 transactions: T1 and T2 at approximately the same time: 3 resolutions:

- fine* {
- Execute all the transaction of T1 in sequence followed by T2
 - Execute all the transactions of T2 in sequence followed by T1
 - Some interleaving of operations may be allowed

*↑
But this one may cause problem.*

EXAMPLE

- Say Account A has 1000 and Account B has 2000
 - T1: Transfer \$50 from account A to account B
 - T2: Transfer 10% of account A to account B
- NOTE: Even if we do the transactions as atomic units we get 2 different answers depending on the order we execute them!!
- **THAT IS OKAY**, either order: T1 then T2 OR T2 then T1 produces a serial schedule

Because these actions are serial

- **Serializability Theory:** determines which schedules are "correct" and those that are not and tries to allow only correct schedules
- **Serial:** A schedule is serial if transactions are executed consecutively with out interleaving. Every serial schedule is considered correct
- **Non-serial:** A schedule is non-serial if it allows interleaving of transactions

QUESTION: Why not make all schedules Serial?

ANSWER: Too slow, we **MUST** allow some interleaving

- A schedule of transactions is serializable if it is equivalent to some serial schedule.

QUESTION: Is result equivalence (a non-serial schedule that produces the same final database state as a serial schedule) enough?

*It may be wrong
for other data items*

ANSWER: No! It might have just been by accident!

Example:

If X starts at 100, everything is fine with both schedules but not with other numbers!

Example 1

| Schedule 1 | Schedule 2 |
|--------------|--------------|
| Read X | Read X |
| $X = X + 10$ | $X = X - 6$ |
| Write x | Write X |
| Read X | Read X |
| $X = X - 5$ | $X = X + 11$ |
| Write X | Write X |

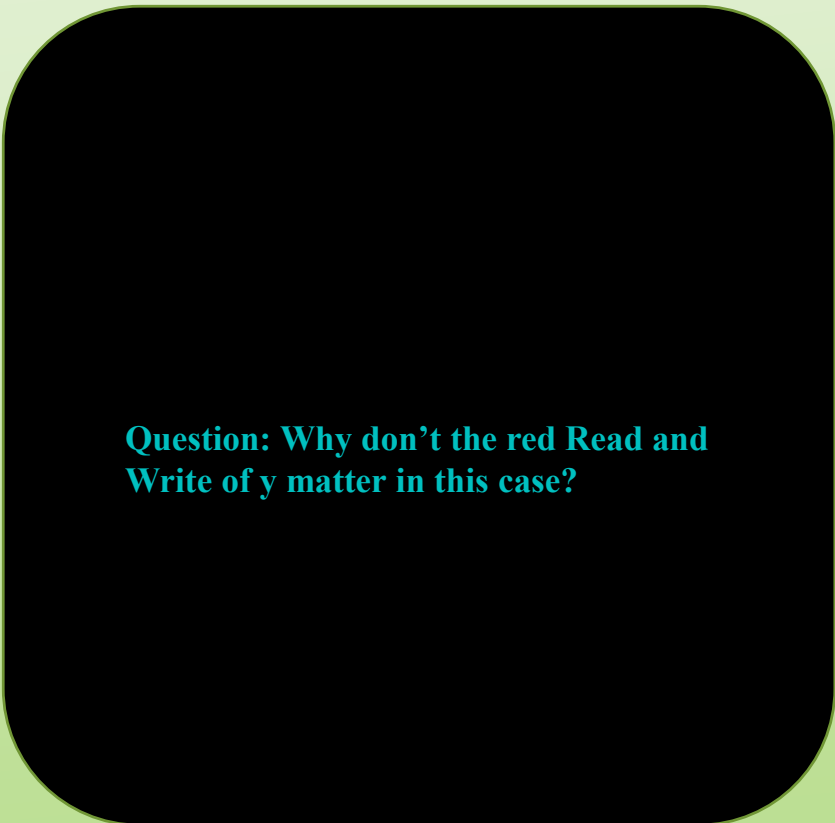
Example 2

| Schedule 1 | Schedule 2 |
|--------------|---------------|
| Read X | Read X |
| $X = X + 10$ | $X = X * 1.1$ |
| Write X | Write X |

CONFLICT EQUIVALENCE

- **Conflict equivalence**: if the order of any 2 conflicting operations (operations from different transactions, using the same data and one of them is a write operation) is the same in both schedules
- **Testing for Conflict Serializability** (using a serialization graph)
 - for each transaction T_i participating in schedule S create a node labeled T_i in the graph
 - for each case in S where T_j executes a read (X) after a write(X) executed by T_i create an edge ($T_i \rightarrow T_j$)
 - for each case in S where T_j executes write(X) after T_i executes a read(X) create an edge ($T_i \rightarrow T_j$)
 - for each case in S where T_j executes a write (X) after T_i executes a write(X) create an edge ($T_i \rightarrow T_j$)
 - The schedule is serializable if and only if the precedence graph has no cycles

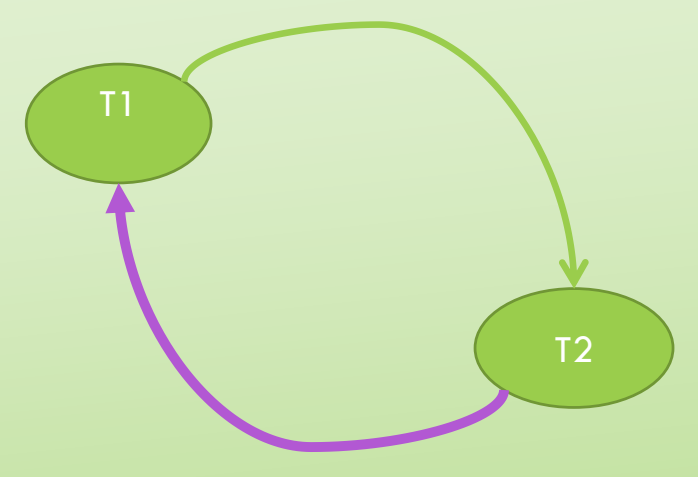
EXAMPLE 1:

| Time | Transaction | Graph |
|------|-------------|---|
| T1 | Read x |  <p>Question: Why don't the red Read and Write of y matter in this case?</p> |
| T1 | $x = x - n$ | |
| T1 | Write x | |
| T1 | Read y | |
| T1 | $y = y + n$ | |
| T1 | Write y | |
| T2 | Read x | |
| T2 | $x = x + M$ | |
| T2 | Write x | |



these two
are in same
transaction.

EXAMPLE 2:

| Time | Transaction | Graph |
|------|-------------|---|
| T1 | Read x |  <pre>graph LR; T1((T1)) -- green --> T2((T2)); T2 -- purple --> T1</pre> |
| T1 | $x = x - n$ | |
| T2 | Read x | |
| T2 | $x = x + M$ | |
| T1 | Write x | |
| T1 | Read y | |
| T2 | Write x | |
| T1 | $y = y + n$ | |
| T1 | Write y | |

[illegible]

EXAMPLE 4:

| Time | Transaction | Graph |
|------|-------------|---|
| T2 | Read z | <pre>graph TD; T2((T2)) -- purple --> T1((T1)); T3((T3)) -- orange --> T2((T2)); T2((T2)) -- blue --> T3((T3)); T3((T3)) -- green --> T2((T2));</pre> |
| T2 | Read y | |
| T2 | Write y | |
| T3 | Read y | |
| T3 | Read z | |
| T1 | Read x | |
| T1 | Write x | |
| T3 | Write y | |
| T3 | Write z | |
| T2 | Read x | |
| T1 | Read y | |
| T1 | Write y | |
| T2 | Write y | |

CS319

19

- NOTE: In general, the above test for every schedule is too time consuming, so instead of testing for serializability, we use protocols (sets of rules) that ensure serializability without having to test the schedule, such as 2-phase locking.