


SVD/LSA as a Neural Network

Singular Value Decomposition

$$M_{m \times n} \approx \dot{M}_{m \times n} = U_{m \times p} D_{p \times p} V^T_{p \times n}$$



Rows of V represent documents

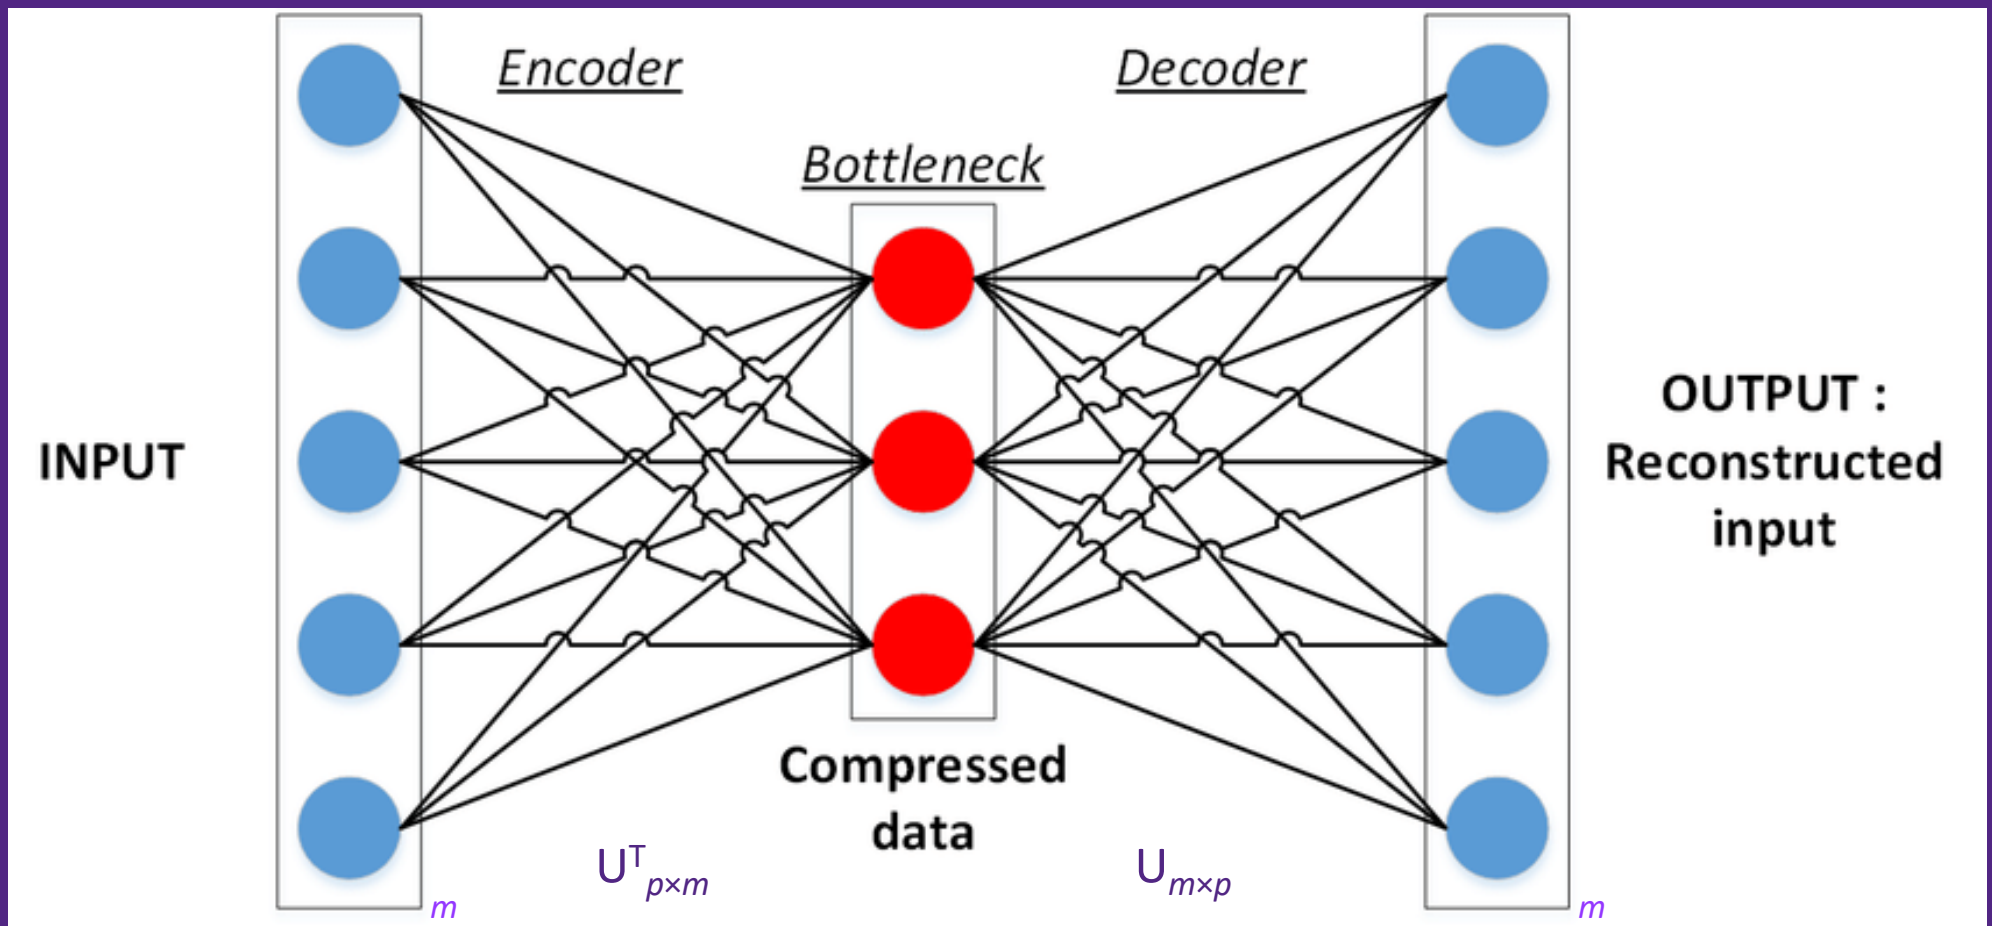
- $(\Sigma_{p \times p}^{-1} U_{p \times m}^T) \dot{M}_{m \times n} = V_{p \times n}^T$
- Each element of column j of V^T is a weighted sum of a column of \dot{M}
- In NN terms:
 - one-layer network
 - m inputs, p outputs
 - fully-connected
 - linear transfer function (no ReLU or anything)
 - weights are $(\Sigma_{p \times p}^{-1} U_{p \times m}^T)$

Rows of V represent documents

- $U_{m \times p} \Sigma_{p \times p} V_{p \times n}^T = \dot{M}_{m \times n}$
- Each element of column j of \dot{M} is a weighted sum of a column of V^T
- In NN terms:
 - one-layer network
 - p inputs, m outputs
 - fully-connected
 - linear transfer function (no ReLU or anything)
 - weights are $(U_{m \times p} \Sigma_{p \times p})$ (I'll call them $\tilde{U}_{m \times p}$)

“LSA” Autoencoder

$$\dot{M}_{:,j} = \tilde{U}_{m \times p} \tilde{U}_{p \times m}^T \dot{M}_{:,j}$$



Question: - Are there any issues? (using m neurons in hidden units) A) Yes B) No

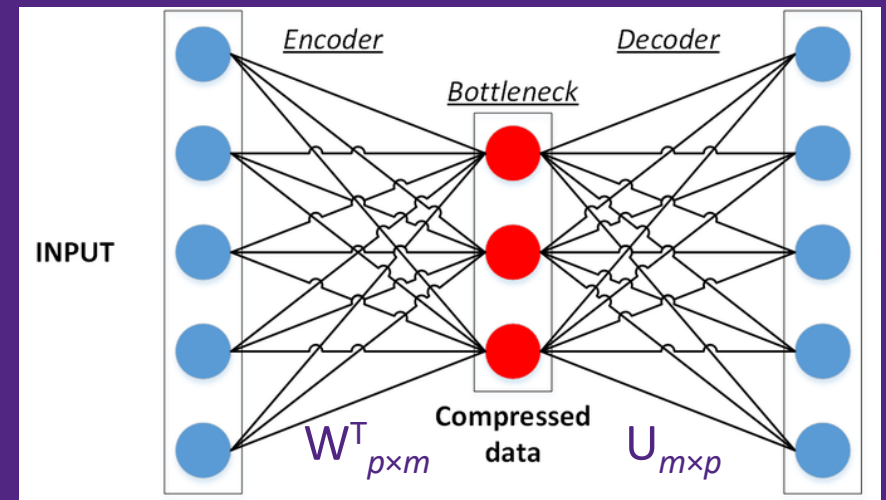
SVD/LSA Versus Autoencoder

- “auto”-encoder meaning “self”-encoder
 - Learns smaller representation for each input (column, document) vector
- Learned weights $\tilde{U}_{m \times p}$ give row (word) representations
- The entries in \tilde{U} will not exactly match the ones you would get from SVD depending on how you train, but they span the same space.
- They’re not “ordered” in terms of importance.
- Full discussion here (optional):
 - <https://arxiv.org/pdf/1804.10253.pdf>

word2vec

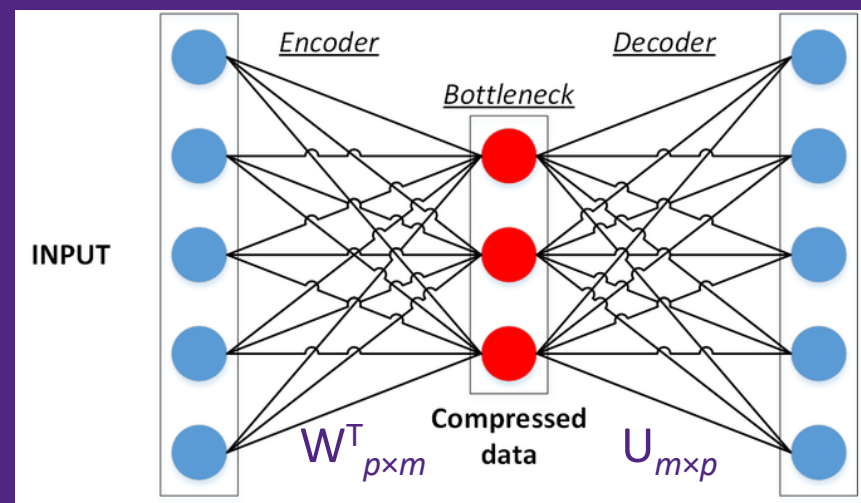
- *Simplified from Mikolov, Tomas; et al. “Efficient Estimation of Word Representations in Vector Space”. arXiv:1301.3781 [Optional]*
- NN-based word representation learner
 - m input units (one per word)
 - p hidden units (dimension of representation)
 - m output units (one per word)
- **Output not same as input**

$$\zeta(U_{m \times p} W_{p \times m}^T \mathbf{x}_{m \times 1}) = \mathbf{y}_{m \times 1}$$



word2vec

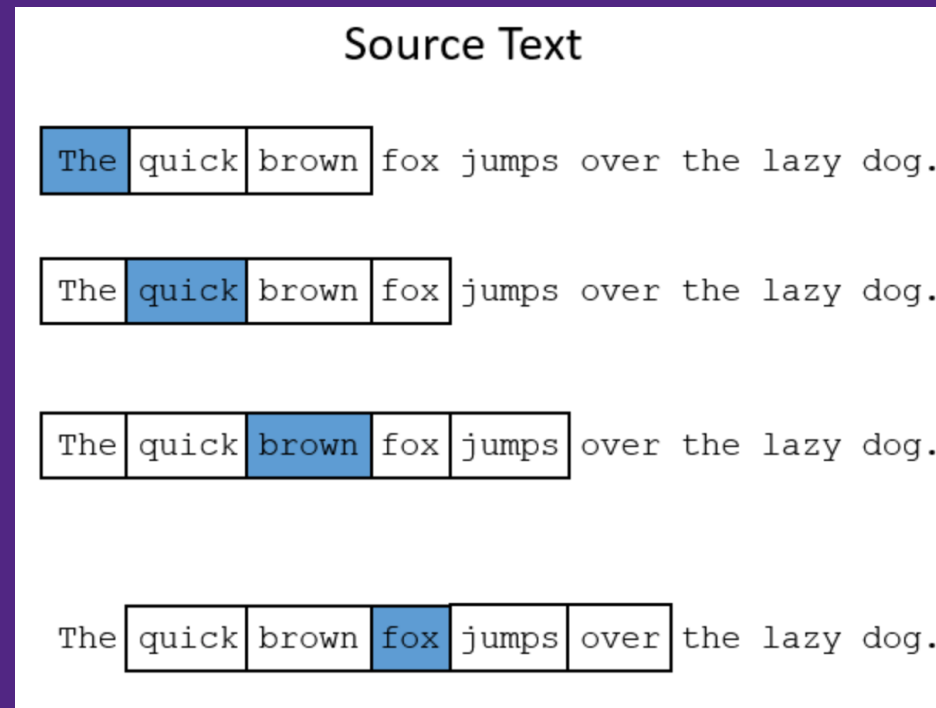
- $\zeta(U_{m \times p} W^T_{p \times m} \mathbf{x}_{m \times 1}) = \mathbf{y}_{m \times 1}$
- \mathbf{x} and \mathbf{y} both represent a word or collection of words
- Matrix $U_{m \times p}$ that we learn will hold our word representations



Choosing x and y :

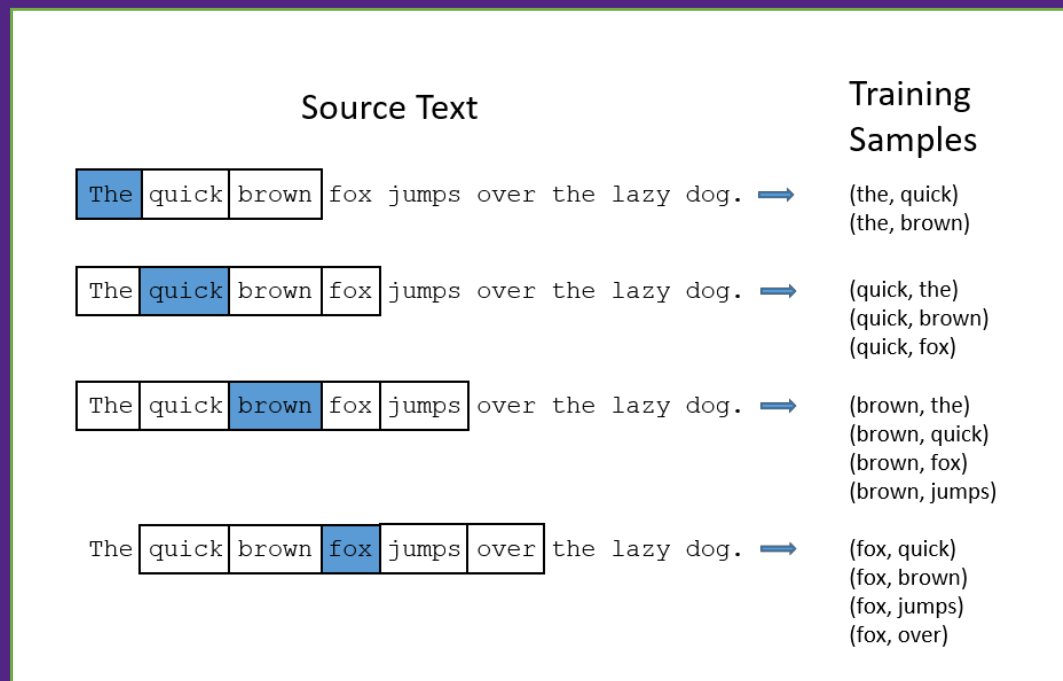
“Continuous Bag of Words”

- CBoW
 - Vector y (target) is one-hot encoding of a word
 - Vector x (input) is average of one-hot encodings of “context words”
 - Input-output pair created for every position of a “sliding window”




Choosing x and y : “Skip-Gram”

- Vector x (input) is one-hot encoding of a word
- Vector y (target) is one-hot encoding of a “nearby” word
- Multiple training vectors per sliding window position



“Analogy Task”

- “France is to Paris as Germany is to _____?”
- $\text{Paris} - \text{France} + \text{Germany} = ???$



Structured Representations from Neural Nets – In General

word2vec idea

- Purpose of word2vec was never to “predict an output”
- What matters is the representation created as a “side-effect” of training

Intermediate Representations

- Same intermediate representation \rightarrow same output
- Advantageous for NN to map different inputs that should give same output to similar representation


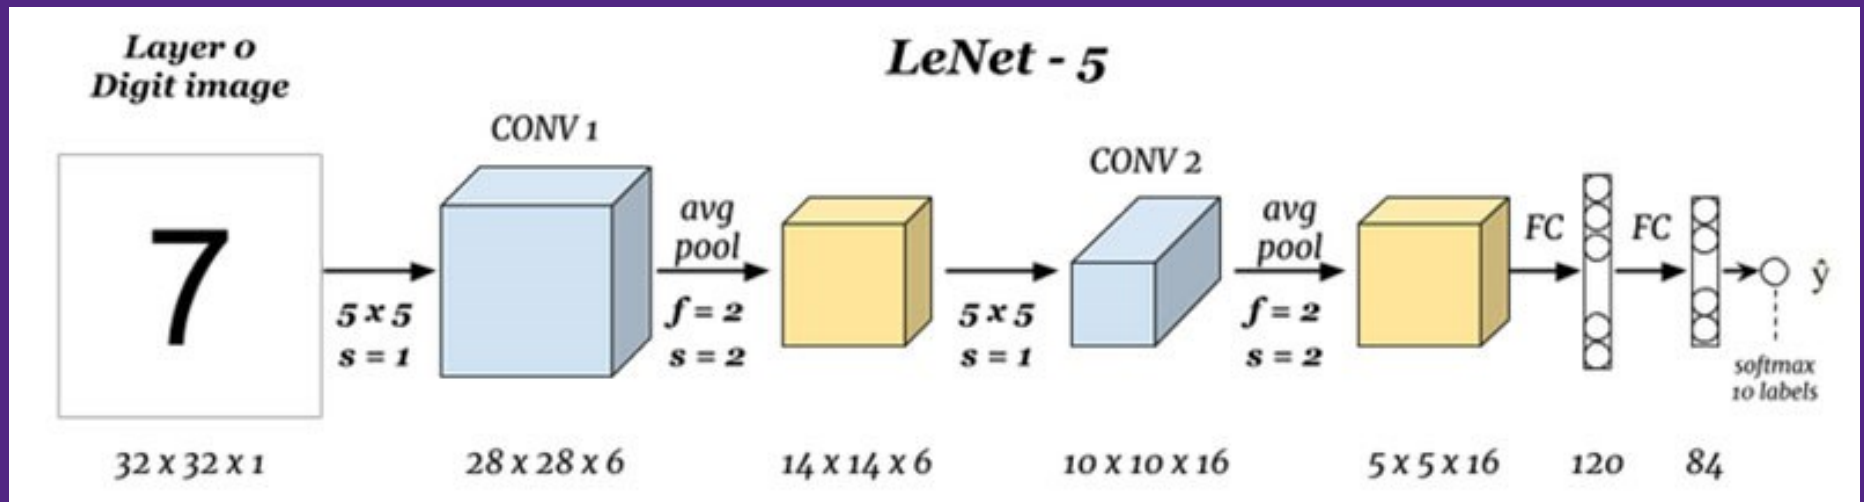


Image Recognition



1024 inputs

Designing Your Own Representations

- Is there a labelling task that can define a useful notion of similarity for you?
- Even if you can't learn that task "100%" you might still learn structure!
 - Recall: Skip Gram model trains on (brown, the), (brown, quick), (brown, fox), (brown, jumps) – impossible to produce a single "right answer" for input "brown"!
- Train a NN that has at least one layer that "compresses" the input.
- Output of that hidden layer, or the hidden layer weights, become your representation

Summary

- We can learn representations that capture relational structure
- LSA/SVD Does this
- LSA/SVD is similar to Autoencoding neural networks
- Other learning tasks and architectures learn different structures
 - CBoW and Skip Grams for word2vec
 - Classification for images
- You can design your own!