Count-based Language Models

CS 4417B

The University of Western Ontario

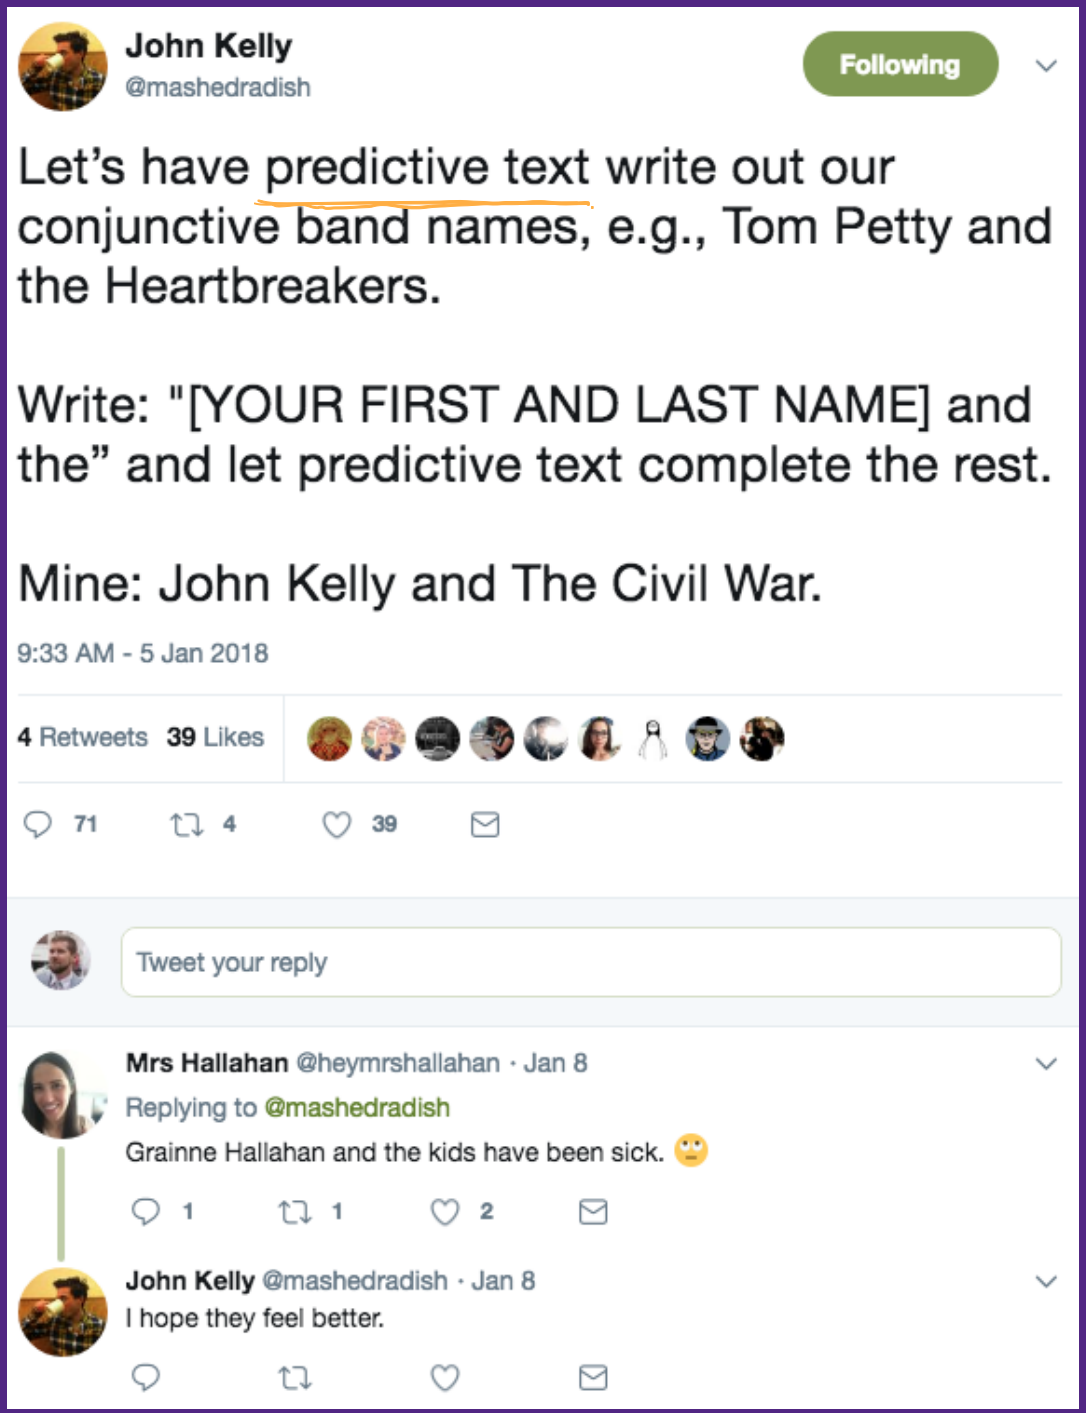
John Kelly
@mashedradish
Following
Let's have predictive text write out our conjunctive band names, e.g., Tom Petty and the Heartbreakers.
Write: "[YOUR FIRST AND LAST NAME] and the" and let predictive text complete the rest.
Mine: John Kelly and The Civil War.
9:33 AM - 5 Jan 2018
4 Retweets 39 Likes
71 4 39
Tweet your reply
Mrs Hallahan @heymrshallahan · Jan 8
Replying to @mashedradish
Grainne Hallahan and the kids have been sick.
1 1 2
John Kelly @mashedradish · Jan 8
I hope they feel better.

Statistical Language Models

[BCC Ch. 1.3.4]

give prediction based on the probability learnt from previous contexts

- Attempt to capture probabilities
  - Of observing a term or sequence of terms
  - Usually given some context
- Captures the sequential structure of a language
  - Grammar/word choice
    - "There are" more likely than "Their are"
  - Idioms
    - "Raining cats and dogs" more likely than "Raining dogs and cats"
  - Topics
    - "peas and carrots" more likely than "peas and briefcases"

corpus needed

in corpus, this one has higher possibility.

use statistics to infer probability.

Add'l reading: [https://web.stanford.edu/~jurafsky/slp3/4.pdf]

Uses

- Predictive text
- Language generation
- Grammar checking
- Evaluating machine-generated text
  - 他 向 记者 介绍了 主要 内容
    He to reporters introduced main content

    - "He, to reporters, introduced the main content."
    - "He introduced reporters to the main contents."
    - "He briefed to reporters the main contents."
    - "He briefed reporters on the main contents."

- Internals of language models are often useful representations for words and documents

huge amount of data required.

Distribution over terms

- Suppose we choose a position uniformly randomly from a corpus.
  - What is the probability that the term at that position is "the"?
  - What is the probability that the term at that position is "coconut"?

- "Term model" provides estimate of P(*t*).

$\sum P(t) = 1$

- If there are *m* terms in our vocabulary, how many numbers do we need to store to describe this distribution?

Little less than m ⇒ $(m-1)$; actually.

since $P(x) = 1 - \sum_{t=0}^{t=m-1} P(t)$

Conditional Models

- Suppose we choose a position uniformly randomly from a corpus, excluding the last position...
- ...and we're told the word (e.g. "the")
  - What is the probability that the next word is "coconut?"
  - What is the probability that the next word is "the?"

- "(Conditional) bigram model" provides estimate of $P(t_2 | t_1)$ next one's possibility. giving the current word. e.g. "coconut", getting
- If there are *m* terms in our vocabulary, how many numbers do we need to store this distribution?

$m \times (m-1)$

$\Rightarrow$ roughly m^2

More complex models

- $P(t_1 | t_2, t_3, t_4)$

- How many parameters? (iClicker)

A) m-1
B) m*(m-1)
C) m*m*(m-1)
D) m*m*m*(m-1)
E) m*m*m*(m-1)

More complex models

- $P(t_1 | t_2, t_3, t_4)$
- How many parameters?
- Suppose a 10000-word vocabulary (modest.)
  - There are 10^{16} 4-grams
  - English Wikipedia has about $10^{9.5}$ words
  - Best-case scenario, 99.9999% of 4-grams never occur

Language Model Sparsity

- Often, we want a language model to *generalize* beyond the corpus it was learned from
  - E.g. for evaluating sentences/text in that language, and to be able to create new text

- Good sparsity: *n*-grams that don't make sense have probability 0. ("he you bird now")
- Bad sparsity: Plausible *n*-grams that happen not to be present in the corpus get probability zero.

- More data reduces bad sparsity.

wrong case!

Simplifying Assumptions

- Suppose we choose to represent $P(t_1,t_2)$ like so:

given.

- $P(t_1,t_2) = P(t_1)\,P(t_2)$

- When would this model give exactly the right probability? (What property of language?)

this one is the probability for t_1 and t_2 are independent to each other

- How many parameters?

$P(t_1) \Rightarrow (m-1)$

$P(t_2) \Rightarrow (m-1)$] $\Rightarrow 2\times(m-1)$.

Independence Model

- $P(t_1,t_2,t_3,\ldots,t_k) = P(t_1)\ P(t_2)\ P(t_3)\ \ldots\ P(t_k)$

- This is an *independence assumption.*
  - Knowing t_1 tells us nothing about what $P(t_2)$ will be, etc.

- Note under this assumption $P(t_1,t_2) = P(t_2,t_1)$
  - All order information is lost; much like Bag of Words

(Markov) Bigram Model

$$P(\langle s \rangle, t_1, t_2, \ldots, t_k, \langle /s \rangle)$$
$$= P(t_1 | \langle s \rangle)\, P(t_2 | t_1) \ldots P(t_{k-1} | t_k)\, P(\langle /s \rangle | t_k)$$

- <s>, </s> are special tokens for start and end of sentence
  - This kind of model is only used on whole sentences.
- Probability of next term only depends on term immediately before
- Assigns probabilities to sequences of arbitrary length
- Number of parameters fixed: $m \times (m - 1)$

Trigram Model

$$P(\langle s\rangle,\langle s\rangle,t_1,t_2,\ldots,t_k,\langle/s\rangle)$$
$$= P(t_1|\langle s\rangle,\langle s\rangle)\ P(t_2|t_1,\langle s\rangle)\ \ldots\ P(t_k|t_{k-1},t_{k-2})\ P(\langle/s\rangle|t_k)$$

- Probability of next term only depends previous two terms

- 4-gram, 5-gram, ... are similar.

Estimating models from data

- $P(t_k | t_1,t_2,t_3,\ldots,t_{k-1}) = P(t_1,t_2,t_3,\ldots,t_k) \,/\, P(t_1,t_2,t_3,\ldots,t_{k-1})$
- $P(t_k | t_1,t_2,t_3,\ldots,t_{k-1}) \approx C(t_1,t_2,t_3,\ldots,t_k) \,/\, C(t_1,t_2,t_3,\ldots,t_{k-1})$

Smoothing

- Zero probabilities cause problems.
  - In our Markov/bigram/trigram/ngram models, if just one of the probabilities is zero, the whole sequence is given probability zero
- “Smoothing” modifies our probability estimates to avoid zero probabilities.

Laplace Smoothing

- So far, all probability estimates we have seen are derived from counts (of terms, bigrams, trigrams, etc.)
- Laplace smoothing adds 1 to each count before normalizing appropriately.
- Avoids zero probabilities, but doesn’t work that well.

Back-off

- If we encounter a never-before-seen 3-gram, maybe we have seen the 2-gram. (Or the term.)
- E.g. never seen “mustard ice cream”
  but have seen “ice cream”
- “Back-off” methods default to simpler models if the more complicated ones do not have reliable estimates.
  - Note we can’t just “substitute” simpler model or things won’t normalize correctly. “Katz back-off” is one strategy
  - “Stupid back-off” ignores this problem; works well on large corpora. Brants, T., Popat, A. C., Xu, P., Och, F. J., and Dean, J. (2007). Large language models in machine translation. In *EMNLP/CoNLL 2007*.

Context-based Smoothing

- Suppose want P(T|reading) but “reading” was not in our corpus.
- Which of our known words are more likely to appear after an unknown word?
  - “San Francisco”
  - “safety glasses”
  - “drinking glasses”
  - “water glasses”
  - …
- Kneser-Ney smoothing boosts these “likely continuations”

Applications

- Text generation
- Filtering output from other systems
- Intelligent “spell checking”
- Corpus membership classification

Babbling

- Given a distribution over words, it's possible to draw a word according to that distribution.
  - If P(coconut = 0.75) and P(pear = 0.25), "coconut" will appear about 3 times out of 4, and pear about 1 out of 4.

- "Babbling" is drawing a sequence of words, always conditioning on the most recently generated word(s) to produce the next one.

- Gives a sense of what structure is captured by the model.

1 gram	–To him swallowed confess hear both. Which. Of save on trail for are ay device and rote life have –Hill he late speaks; or! a more to leg less first you enter
2 gram	–Why dost stand forth thy canopy, forsooth; he is this palpable hit the King Henry. Live king. Follow. –What means, sir. I confess she? then all sorts, he is trim, captain.
3 gram	–Fly, and will rid me these news of price. Therefore the sadness of parting, as they say, 'tis done. –This shall forbid it should be branded, if renown made it empty.
4 gram	–King Henry. What! I will go seek the traitor Gloucester. Exeunt some of the watch. A great banquet serv'd in; –It cannot be but so.

Figure 4.3 Eight sentences randomly generated from four *N*-grams computed from Shakespeare's works. All characters were mapped to lower-case and punctuation marks were treated as words. Output is hand-corrected for capitalization to improve readability.

Additional reading: [https://web.stanford.edu/~jurafsky/slp3/4.pdf]

Filtering

- Statistical machine translation methods often produce several plausible translations
  - We can get the translation model to “babble” likely sentences
  - Choose the one with the highest probability under the language model
- Speech-to-text also produces statistical models over possible sentences

- Language modeling allows both of these to be improved with huge amounts of “unlabeled” data

Intelligent “spell checking”

Class will end in five minuets.

Corpus membership classification

- “Generative” model for classification.
- E.g., classify passage as Stephen King or Shakespeare:
  - Build language model for Stephen King
  - Build language model for Willy Shakes
  - Ask probability of passage belonging to each, see which one gives higher probability.

Summary

- Language models that predict next word given past context – (sometimes called “causal” language models)
- Challenge to naïve approaches: Data sparsity
- Applications
  - Generation – create new text
  - Evaluation – assess whether text is plausible
  - Classification – assess relative likelihood of authorship
  - Many more

End of the session